

Oracle Linux System Administration

Activity Guide

D74508GC20

Edition 2.0

February 2014

D85362

ORACLE®

Copyright © 2014, Oracle and/or its affiliates. All rights reserved.

Disclaimer

This document contains proprietary information and is protected by copyright and other intellectual property laws. You may copy and print this document solely for your own use in an Oracle training course. The document may not be modified or altered in any way. Except where your use constitutes "fair use" under copyright law, you may not use, share, download, upload, copy, print, display, perform, reproduce, publish, license, post, transmit, or distribute this document in whole or in part without the express authorization of Oracle.

The information contained in this document is subject to change without notice. If you find any problems in the document, please report them in writing to: Oracle University, 500 Oracle Parkway, Redwood Shores, California 94065 USA. This document is not warranted to be error-free.

Restricted Rights Notice

If this documentation is delivered to the United States Government or anyone using the documentation on behalf of the United States Government, the following notice is applicable:

U.S. GOVERNMENT RIGHTS

The U.S. Government's rights to use, modify, reproduce, release, perform, display, or disclose these training materials are restricted by the terms of the applicable Oracle license agreement and/or the applicable U.S. Government contract.

Trademark Notice

Oracle and Java are registered trademarks of Oracle and/or its affiliates. Other names may be trademarks of their respective owners.

Author

Craig McBride

Technical Contributors and Reviewers

Avi Miller, Elena Zannoni, Lenz Grimmer, Sergio Leunissen, Waseem Daher, Wim Coekaerts, Al Flournoy, Harald Van Breederode, Joel Goodman, Manish Kapur, Soeren Binner, Jeremy Smyth, Yasar Akthar, Javier Saiz, Ozgur Yuksel, Antoinette O'Sullivan, Frank Allan, Gavin Bowe, Gino Kawalski, Jeff Suchomel, Rob Swank, Ron Hardin, Michele Dady, Matt Taylor

This book was published using: Oracle Tutor

Table of Contents

Practices for Lesson 1: Course Introduction.....	1-1
Practices for Lesson 1: Course 1	1-3
Practice 1-1: Exploring the dom0 Environment	1-4
Practice 1-2: Starting, Stopping, and Listing Guests.....	1-10
Practice 1-3: Connecting to a Guest VM	1-12
Practice 1-4: Logging Off Your Student PC.....	1-16
Practices for Lesson 2: Introduction to Oracle Linux.....	2-1
Practices for Lesson 2: Introduction to Oracle Linux.....	2-3
Practice 2-1: Quiz – Introduction to Oracle Linux	2-4
Solution 2-1: Quiz – Introduction to Oracle Linux	2-6
Practice 2-2: Viewing Kernel Information	2-7
Practices for Lesson 3: Installing Oracle Linux	3-1
Practices for Lesson 3: Installing Oracle Linux	3-3
Practice 3-1: Installing Oracle Linux.....	3-4
Practice 3-2: Using FirstBoot	3-36
Practice 3-3: Logging In to Oracle Linux and Shutting Down	3-43
Practice 3-4: Re-Creating the host03 VM Guest	3-45
Practices for Lesson 4: Linux Boot Process	4-1
Practices for Lesson 4: Linux Boot Process	4-3
Practice 4-1: Exploring the GRUB Boot Loader	4-4
Practice 4-2: Booting Different Kernels	4-12
Practice 4-3: Using the GRUB Menu.....	4-14
Practice 4-4: Changing the Default Run Level	4-20
Practice 4-5: Exploring and Configuring init Services.....	4-23
Practice 4-6: Exploring and Configuring Upstart Jobs.....	4-28
Practices for Lesson 5: System Configuration.....	5-1
Practices for Lesson 5: System Configuration	5-3
Practice 5-1: Exploring /etc/sysconfig	5-4
Practice 5-2: Exploring /proc	5-9
Practice 5-3: Exploring the sysfs File System	5-13
Practice 5-4: Using sysctl.....	5-16
Practices for Lesson 6: Package Management	6-1
Practices for Lesson 6: Package Management.....	6-3
Practice 6-1: Using the rpm Utility	6-4
Practice 6-2: Creating a Local Yum Repository	6-8
Practice 6-3: Using the yum Utility	6-13
Practice 6-4: Unbreakable Linux Network (ULN).....	6-17
Practices for Lesson 7: Ksplice	7-1
Practices for Lesson 7: Ksplice	7-3
Practice 7-1: Viewing Ksplice Offline Client Packages from the Unbreakable Linux Network (ULN)	7-4
Practice 7-2: Using sftp to Upload Ksplice Packages.....	7-7
Practice 7-3: Installing the Ksplice Offline Client and Kernel Updates	7-8

Practices for Lesson 8: Automating Tasks	8-1
Practices for Lesson 8: Automating Tasks	8-3
Practice 8-1: Automating Tasks	8-4
Practices for Lesson 9: Kernel Module Configuration.....	9-1
Practices for Lesson 9: Kernel Module Configuration	9-3
Practice 9-1: Using Loadable Kernel Modules	9-4
Practices for Lesson 10: User and Group Administration	10-1
Practices for Lesson 10: User and Group Administration.....	10-3
Practice 10-1: User Account Administration.....	10-4
Practice 10-2: Group Account Administration	10-12
Practice 10-3: User Private Groups.....	10-13
Practice 10-4: Password Aging	10-16
Practice 10-5: Using the User Manager GUI	10-18
Practice 10-6: Restricting the Use of the su Command.....	10-25
Practice 10-7: Allowing the Use of the sudo Command	10-27
Practices for Lesson 11: File Systems.....	11-1
Practices for Lesson 11: File Systems	11-3
Practice 11-1: Listing the Current Disk Partitions	11-4
Practice 11-2: Partitioning a Storage Device	11-8
Practice 11-3: Creating ext File Systems	11-13
Practice 11-4: Increasing Swap Space	11-16
Practices for Lesson 12: Storage Administration	12-1
Practices for Lesson 12: Storage Administration.....	12-3
Practice 12-1: Creating Linux LVM Partitions.....	12-4
Practice 12-2: Creating a Logical Volume	12-7
Practice 12-3: Creating a File System and Mounting a Logical Volume.....	12-11
Practice 12-4: Backing Up Volume Group Metadata.....	12-13
Practice 12-5: Creating a Logical Volume Snapshot.....	12-17
Practice 12-6: Increasing the Capacity of a Logical Volume	12-20
Practice 12-7: Restoring Volume Group Metadata.....	12-25
Practice 12-8: Creating a RAID Device	12-26
Practices for Lesson 13: Network Configuration	13-1
Practices for Lesson 13: Network Configuration	13-3
Practice 13-1: Configuring the eth1 Network Interface	13-4
Practice 13-2: Using NetworkManager.....	13-8
Practice 13-3: Using the system-config-network Utility	13-13
Practice 13-4: Accessing the Public Yum Repository.....	13-17
Practices for Lesson 14: File Sharing	14-1
Practices for Lesson 14: File Sharing.....	14-3
Practice 14-1: Configuring an NFS Server and an NFS Client.....	14-4
Practice 14-2: Using Automounter	14-9
Practice 14-3: Configuring an FTP Server	14-11
Practice 14-4: Downloading a File from an FTP Server	14-13
Practice 14-5: Configuring xinetd	14-16

Practices for Lesson 15: OpenSSH	15-1
Practices for Lesson 15: OpenSSH.....	15-3
Practice 15-1: Connecting to a Remote Server by Using ssh	15-4
Practice 15-2: Configuring OpenSSH to Connect Without a Password.....	15-7
Practices for Lesson 16: Pluggable Authentication Modules (PAM).....	16-1
Practices for Lesson 16: Pluggable Authentication Modules (PAM).....	16-3
Practice 16-1: Configuring PAM for a Single Login Session	16-4
Practice 16-2: Configuring PAM to Prevent Non-root Login	16-8
Practices for Lesson 17: Security Administration	17-1
Practices for Lesson 17: Security Administration	17-3
Practice 17-1: Configuring a chroot Jail	17-4
Practice 17-2: Configuring a chroot Jail for ftp Users	17-7
Practice 17-3: Configuring iptables	17-13
Practice 17-4: Configuring a TCP Wrapper.....	17-20
Practices for Lesson 18: Oracle on Oracle	18-1
Practices for Lesson 18: Oracle on Oracle.....	18-3
Practice 18-1: Using sftp to Upload oracle-rdbms-server Package	18-4
Practice 18-2: Installing and Running Oracle RDBMS Pre-Install	18-5
Practice 18-3: Preparing Disks for ASM Use	18-9
Practice 18-4: Installing and Configuring ASMLib	18-14
Practices for Lesson 19: System Monitoring.....	19-1
Practices for Lesson 19: System Monitoring	19-3
Practice 19-1: Using sosreport to Collect System Information	19-4
Practice 19-2: Using Standard Linux Performance Monitoring Tools	19-7
Practice 19-3: Installing and Using OSWatcher	19-30
Practice 19-4: Using OSWatcher Black Box Analyzer.....	19-37
Practices for Lesson 20: System Logging	20-1
Practices for Lesson 20: System Logging	20-3
Practice 20-1: Viewing the System Log File Configuration.....	20-4
Practice 20-2: Using rsyslog Templates.....	20-9
Practice 20-3: Using logwatch.....	20-11
Practices for Lesson 21: Troubleshooting	21-1
Practices for Lesson 21: Troubleshooting	21-3
Practice 21-1: Transferring Utilities from dom0	21-4
Practice 21-2: System Boots into Single-User Mode.....	21-6
Solution 21-2: System Boots into Single-User Mode.....	21-8
Practice 21-3: Status Commands Fail.....	21-10
Solution 21-3: Status Commands Fail.....	21-12
Practice 21-4: cron Job Fails to Run	21-13
Solution 21-4: cron Job Fails to Run	21-15
Practice 21-5: User Cannot Log In	21-17
Solution 21-5: User Cannot Log In	21-19
Practice 21-6: File System Troubleshooting.....	21-21
Solution 21-6: File System Troubleshooting.....	21-23
Practice 21-7: Logical Volume Space is Exhausted	21-25
Solution 21-7: Logical Volume Space is Exhausted	21-30

Practice 21-8: Network Connectivity Problem	21-32
Solution 21-8: Network Connectivity Problem	21-34
Practice 21-9: NFS Permission Problem	21-36
Solution 21-9: NFS Permission Problem	21-42
Practice 21-10: Remote Access Problem.....	21-46
Solution 21-10: Remote Access Problem.....	21-48
Practice 21-11: Log File Is Not Getting Updated.....	21-54
Solution 21-11: Log File Is Not Getting Updated	21-57
Appendix A: Source Code for Problem-Causing Executables.....	21-59
Appendices: Remote Access Options	22-1
Appendix	22-3
Appendix A: Using an NX Client to Connect to dom0	22-4
Appendix B: Using an NX Player to Connect to dom0	22-8
Appendix C: Using VNC (TightVNC) to Connect Directly to VM Guests	22-14

Practices for Lesson 1: Course Introduction

Chapter 1

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 1: Course 1

Practices Overview

In these practices, you become familiar with the Oracle VM Server environment installed on your classroom PC. You log on to your classroom PC and connect to the virtual machines used for the hands-on practices.

Assumptions

- Your instructor has assigned a student PC to you.
- The student PC is running OVM 2.2.1.
- The GNOME desktop is installed on **dom0**.
- Three guests (virtual machines) are created: **host01**, **host02**, and **host03**.
- Guest VMs **host01** and **host02** are running Oracle Linux 6.5.
- You install Oracle Linux 6.5 on **host03** in Practice 3.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.


Practice 1-1: Exploring the dom0 Environment

Overview

In this practice, you explore **dom0** configuration and directory structure.

Tasks

1. Open a terminal window.
- Begin this task from the **dom0** GNOME virtual desktop window as shown below.




Double-click the **Terminal** icon on the GNOME desktop.

- A terminal window opens.

2. Become the `root` user.

- Enter the commands from an open terminal window as shown below.



Become `root` by using the `su -` command. The `root` password is `oracle`. Use the `whoami` command to print the user identity to confirm that you are `root`:

```
$ su -  
Password: oracle  
# whoami  
root
```

3. Determine the operating system running on `dom0`.

Use the `uname -a` command to display the operating system version. The output shown is a sample and may not represent the host name and kernel version running on your system.

```
# uname -a  
Linux EDDDR9P1 2.6.18-128.2.1.4.25.el5xen #1 SMP Tue Mar 23  
12:43:27 EDT 2010 i686 i686 i386 GNU/Linux
```

- In this example, the operating system is Linux.
- The Linux kernel is 2.6.18-128.2.1.4.25.el5xen.
- The host name is EDDDR9P1 (your host name is different).

4. Determine the network configuration of **dom0**.

Use the `ifconfig -a` command to display the network configuration. The output shown is a sample and may not represent the network configuration of your system. Only partial output is shown.

```
# ifconfig -a
...
eth0      Link encap:Ethernet
          inet addr:10.150.30.72
...
lo        Link encap:Local Loopback
          inet addr:127.0.0.1
...
vif...    Link encap:Ethernet
...
virbr0    Link encap:Ethernet
          inet addr:192.0.2.1
...
virbr1    Link encap:Ethernet
          inet addr:192.168.1.1
...
xenbr0    Link encap:Ethernet
          inet addr:10.150.30.72
...
```

- In this example, the network interface for **dom0** is `eth0` and is assigned an IP address of `10.150.30.72`. Connection to the outside world is provided from this interface.
- The `lo` interface is a software loopback interface that identifies `localhost`. It is always assigned an IP address of `127.0.0.1`.
- The `virbr0` interface is a `xen` bridge interface used by the VM guests. It is assigned an IP address of `192.0.2.1`.
- The `virbr1` interface is a second `xen` bridge interface used by the VM guests. It is assigned an IP address of `192.168.1.1`.
- As you will see later, each of the guest virtual machines is on the `192.0.2` subnet and the `192.168.1` subnet. Communication to the classroom network is through `xenbr0`.
- You also notice `vif<#>.0` and `vif<#>.1` entries. These are the virtual interfaces that are tied to the VM/domU IDs. You can get the VM/domU IDs from the `xm list` command, which you run later in this practice.

5. Explore the /OVS directory structure on **dom0**.

- The output shown is a sample and may not represent your system. Only partial output is shown.
- a. Explore the top level of the /OVS directory:

```
# ls -l /OVS
lrwxrwxrwx  /OVS -> /var/ovs/mount/...
# cd /OVS
# ls -l
drwxrwxrwx  iso_pool
drwxrwxrwx  lost+found
drwxrwxrwx  publish_pool
drwxrwxrwx  running_pool
drwxrwxrwx  seed_pool
drwxrwxrwx  sharedDisk
```

- Note that /OVS is a symbolic link to the /var/ovs/mount/... directory.
- There are six directories in the /OVS directory.

b. Explore the /OVS/running_pool directory:

```
# cd /OVS/running_pool
# ls -l
drwxr-xr-x  host01
drwxr-xr-x  host02
drwxr-xr-x  host03
```

- The files needed to create the VMs are in separate directories in the /OVS/running_pool directory.
- This example shows that three VMs currently exist: **host01**, **host02**, and **host03**.

c. Explore the **host03** VM directory.

```
# cd /OVS/running_pool/host03
# ls -l
-rw-r--r--  system.img
-rw-r--r--  u01.img
-rw-r--r--  u02.img
-rw-r--r--  vm.cfg
```


- The **system.img** file is the operating system virtual disk.
- The **u01.img** and **u02.img** files are additional virtual disks used by the storage administration practices.
- The **vm.cfg** file is the configuration file for the virtual machine. This file is read when the virtual machine is created.

d. View the `vm.cfg` file.

```
# cat vm.cfg
name = 'host03'
builder = 'hvm'
memory = 2048
boot = 'cd'
disk = [ 'file:/OVS/running_pool/host03/system.img,hda,w',
         'file:/OVS/running_pool/host03/u01.img,hdb,w',
         'file:/OVS/running_pool/host03/u02.img,hdd,w',
         'file:/OVS/seed_pool/OracleLinux-R6-U5-Server-x86_64-
dvd.iso,hdc:cdrom,r']
vif = [ 'mac=00:16:3e:00:01:03, bridge=virbr0',
        'mac=00:16:3e:00:02:03, bridge=virbr1']
device_model = '/usr/lib/xen/bin/qemu-dm'
kernel = '/usr/lib/xen/boot/hvmloader'
vnc = 1
vncunused=1
vcpus = 1
timer_mode = 0
apic = 1
acpi = 1
pae = 1
serial = 'pty'
on_reboot = 'restart'
on_crash = 'restart'
usb = 1
usbdevice = 'tablet'
```

- Notice there are three virtual disks represented by the three `.img` files.
- Notice the Oracle Linux 6.5 `dvd.iso` is mounted on a virtual CDROM device.
- Notice there are two virtual network interfaces. The interface on the `virbr0` bridge is `eth0`, and the interface on the `virbr1` bridge is `eth1`.

The following displays the configuration. All three VMs (**host01**, **host02**, and **host03**) have two network interfaces and three disks.



Practice 1-2: Starting, Stopping, and Listing Guests

Overview

In this practice, you use `xm` commands to list, create, and shut down virtual machines.

Assumptions

- You are logged on to **dom0**.
- You have a terminal window open.
- You are the `root` user.

Tasks

1. List all currently active guests, as well as **dom0** itself.

Use the `xm list` command. The output shown is a sample and will not represent your system exactly.

Name	ID	Mem	VCPUs	State	Time(s)
Domain-0	0	1024	2	r-----	281.1
host01	3	2048	1	-b----	157.6
host02	1	2048	1	-b----	159.0
host03	2	2048	1	-b----	13.2

- You should have three guests (**host01**, **host02**, and **host03**) listed in the preceding output.

2. Shut down the specified VMs.

Use the `xm shutdown -w <VM name>` command to shut down the VMs. The `-w` option tells the system to wait until all services in the domain shut down cleanly.

xm shutdown -w host01
Domain host01 terminated
All domains terminated
xm shutdown -w host02
Domain host02 terminated
All domains terminated
xm shutdown -w host03
Domain host03 terminated
All domains terminated
xm list
Name ID Mem VCPUs State Time(s)
Domain-0 0 1024 2 r----- 289.6

3. Start the VMs.

Use the `xm create <config_file>` command to start the **host01** and **host02** VMs. The `<config_file>` is named `vm.cfg` and is located in the `/OVS/running_pool/<VM_name>` directory.

```
# cd /OVS/running_pool/host01
# xm create vm.cfg
Using config file "./vm.cfg".
Started domain host01 (id=4)
# xm list
Name           ID   Mem  VCPUs      State      Time(s)
Domain-0        0    1024       2      r-----  304.5
host01         4    2048       1      -b----  18.7
# cd /OVS/running_pool/host02
# xm create vm.cfg
Using config file "./vm.cfg".
Started domain host02 (id=5)
# xm list
Name           ID   Mem  VCPUs      State      Time(s)
Domain-0        0    1024       2      r-----  318.4
host01         4    2048       1      -b----  24.7
host02         5    2048       1      -b----  13.7
```

- Do not start the **host03** VM.

Practice 1-3: Connecting to a Guest VM

Overview

In this practice, you connect to a virtual machine guest.

Assumptions

- You are logged on to **dom0**.
- You have a terminal window open.
- You are the `root` user.
- Both **host01** and **host02** VMs are running.

Tasks

1. Connect to **host01** guest by using **vncviewer**.

- Determine the vnc port number for **host01** by running the `xm list -l host01 | grep location` command.


```
# xm list -l host01 | grep location
          (location 0.0.0.0:5900)
          (location 2)
```

- The sample shown indicates that the port number is **5900**. This may not be true in your case.


- b. Run the `vncviewer&` command:

```
# vncviewer&
```

- The **VNC Viewer: Connection Details** dialog box is displayed:




- c. Enter the `localhost :<port_number>` command, substituting the port number displayed from the previous `xm list` command. For example, if the port number is 5900, enter `localhost :5900` and click **OK**.



- You could also enter `localhost : 0` (omit the "590" from the port number).

- d. The window, which is shown as follows, is displayed. Log in as `root` with password `oracle` (all lowercase).

- If the screen is blank, press Enter to display the login prompt.



- e. Use the `hostname` command to confirm that you are logged on to the **host01** VM.

```
# hostname
host01.example.com
```

- f. Log out by entering either the `logout` command or the `exit` command.

```
# logout
```

- g. Close the VNC window by clicking the X at the top-right corner of the window.

2. Connect to **host01** guest by using `ssh`.

- a. Alternatively, use the `ssh` command to connect to the VM guest.

- You are presented with the RSA key fingerprint question the first time you connect. Answer **yes** to the question and then enter the `root` password `oracle`.
- RSA and the `ssh` utility are covered in Lesson 15, “OpenSSH.”

```
# ssh host01
The authenticity of host 'host01 (192.0.2.101)' can't be
established. RSA key fingerprint is ...
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'host01,192.0.2.101' (RSA) to the
list of known hosts.
root@host01's password: oracle
```

- Most of the practices instruct you to log in using `vncviewer` rather than `ssh`.

- b. To confirm that you have connected to **host01**, enter the `hostname` command:

```
# hostname  
host01.example.com
```

- c. Terminate the `ssh` connection by entering the `exit` command. To verify that you are back to **dom0**, enter the `hostname` command:

```
# exit  
Connection to host01 closed.  
# hostname  
EDDR9P1
```

- In this example, the host name is **EDDR9P1**. Your host name is different.


Practice 1-4: Logging Off Your Student PC

Overview

In this practice, you learn how to log off your system.

Tasks

1. Log off your student PC.
 - a. Open the **System** menu on the GNOME desktop.



- b. Select **Log Out vncuser** from the System menu.

Do not perform this next step. This is information only.

- c. Click the **Log Out** button to log out.



- Alternatively, you could log off by clicking the red X in the upper-right corner of the window. You are then prompted to click Terminate to exit the session and close all running programs.

Practices for Lesson 2: Introduction to Oracle Linux

Chapter 2

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 2: Introduction to Oracle Linux

Practices Overview

In these practices, you test your knowledge of Linux and of Oracle's contributions and commitment to Linux. You also view the different kernel packages that are installed with Oracle Linux.

Practice 2-1: Quiz – Introduction to Oracle Linux

Overview

In this quiz, you answer questions about the history of Linux in general, as well as Oracle's commitment and contributions to Linux.

Quiz Questions

Choose the best answer from those provided for each multiple choice or True/False question.

1. Which of the following statements are true? (Choose all that apply).
 - a. Linux is a UNIX-like operating system named for its developer, Linus Torvalds.
 - b. Richard Stallman released version 0.01 of the Linux kernel in September 1991.
 - c. The goal of the GNU Project to create a free, UNIX-like operating system was realized in 1992 with the addition of the Hurd kernel.
 - d. Distributors of GPL software must provide the binaries at no cost, but they can charge for the source code.
2. Thousands of programmers all over the world participate in the development of the Linux operating system.
 - a. True
 - b. False
3. Linus Torvalds releases a new version of the mainline kernel approximately:
 - a. Every month
 - b. Every three months
 - c. Every six months
 - d. Every year
4. The revision control system for the Linux kernel source code is called:
 - a. SCCS
 - b. RCS
 - c. BitKeeper
 - d. Git
5. A Linux distribution is a collection of software built on top of the Linux kernel and offered as a complete package.
 - a. True
 - b. False

6. Which of the following statements about Oracle's commitment to Linux are true? (Choose all that apply).
 - a. Oracle runs its business on Linux.
 - b. Oracle has a global Linux support organization.
 - c. Oracle has a dedicated Linux kernel development team.
 - d. Oracle's goal is to deliver the best performing, most modern, and most reliable Linux operating system for the enterprise.
7. Oracle Linux is fully compatible with which of the following Linux distributions?
 - a. Debian
 - b. Ubuntu
 - c. SUSE
 - d. RHEL
8. The Unbreakable Enterprise Kernel is optimized for Oracle applications and is the only kernel offered with Oracle Linux.
 - a. True
 - b. False

Solution 2-1: Quiz – Introduction to Oracle Linux

Quiz Solutions

1. Which of the following statements are true? (Choose all that apply).
 - a. Linux is a UNIX-like operating system named for its developer, Linus Torvalds.
2. Thousands of programmers all over the world participate in the development of the Linux operating system.
 - a. True
3. Linus Torvalds releases a new version of the mainline kernel approximately:
 - b. Every three months
4. The revision control system for the Linux kernel source code is called:
 - d. Git
5. A Linux distribution is a collection of software built on top of the Linux kernel and offered as a complete package.
 - a. True
6. Which of the following statements about Oracle's commitment to Linux are true? (Choose all that apply).
 - a. Oracle runs its business on Linux.
 - b. Oracle has a global Linux support organization.
 - c. Oracle has a dedicated Linux kernel development team.
 - d. Oracle's goal is to deliver the best performing, most modern, and most reliable Linux operating system for the enterprise.
7. Oracle Linux is fully compatible with which of the following Linux distributions?
 - d. RHEL
8. The Unbreakable Enterprise Kernel is optimized for Oracle applications and is the only kernel offered with Oracle Linux.
 - b. False. It is true that the Unbreakable Enterprise Kernel is optimized for Oracle applications. But Oracle Linux also includes a Red Hat-compatible kernel for customers requiring strict RHEL compatibility.

Practice 2-2: Viewing Kernel Information

Overview


In this practice, you connect to the **host01** VM, view the installed kernel packages, and determine which kernel is running. You also view the text file used by Oracle support teams to easily identify the active kernel. Note that the displayed sample output might not represent your own system.

Assumptions

- You are logged on to **dom0**.
- VM **host01** is running.


Tasks

1. Open a terminal window.
- Begin this task from the GNOME virtual desktop window:



Double-click the **Terminal** icon on the GNOME desktop.

- A terminal window opens.
2. Become the `root` user.
 - Enter the commands from an open terminal window:



Become `root` by using the `su -` command. The `root` password is `oracle`. Use the `whoami` command to print the user identity to confirm that you are `root`:

```
$ su -
Password: oracle
# whoami
root
```

3. Connect to **host01** guest as `root` user (password is `oracle`) by using `ssh`.

Use the `ssh` command to connect to the VM guest **host01** as `root`. The password is `oracle`. Confirm that you have connected to **host01** by entering the `hostname` command:

```
# ssh host01
root@host01's password: oracle
# hostname
host01.example.com
```

4. List the kernel release that is currently running on your system.

Use the `uname -r` command to print the kernel release:

```
# uname -r
3.8.13-16.2.1.el6uek.x86_64
```

- Note that the Unbreakable Enterprise Kernel is running by default.

5. List the kernel packages that are installed on your system.

Use the `rpm -qa` command to query the kernel packages installed on your system:

```
# rpm -qa | grep "^(kernel"
kernel-2.6.32-431.el6.x86_64
kernel-uek-3.8.13-16.2.1.el6uek.x86_64
...
```

- Note that the following two kernel packages are installed (associated firmware packages, headers and doc packages are also installed):
 - Unbreakable Enterprise Kernel package (`kernel-uek-3.8.13-16.2.1.el6uek.x86_64`)
 - Red Hat-compatible kernel package (`kernel-2.6.32-431.el6.x86_64`)

6. View the `/etc/*release` files on your system.

Change to the `/etc/` directory and view the “release” files:

```
# cd /etc
# ls -l *release
...
-rw-r--r--. oracle-release
-rw-r--r--. redhat-release
lrwxrwxrwx. system-release -> oracle-release
# cat oracle-release
Oracle Linux Server release 6.5
# cat redhat-release
Red Hat Enterprise Linux Server release 6.5 (Santiago)
```

- Note that `/etc/system-release` is a symbolic link to `/etc/oracle-release`.

7. Log out from **host01**.

```
# logout
Connection to host01 closed.
```

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 3: Installing Oracle Linux

Chapter 3

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 3: Installing Oracle Linux

Practices Overview

In these practices, you:

- Install Oracle Linux on the **host03** virtual machine
- Run FirstBoot on **host03**
- Log in to **host03** and perform shutdown operations

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practice 3-1: Installing Oracle Linux

Overview

In this practice, you install Oracle Linux on the **host03** virtual machine (VM).

Assumptions

- You are logged on as the `root` user on **dom0**.
- You have a terminal window open.

Tasks

1. Create the **host03** VM.

Change to the `/OVS/running_pool/host03` directory and use the `xm create` command as follows:

```
# cd /OVS/running_pool/host03  
# xm create vm.cfg  
Using config file "./vm.cfg".  
Started domain host03 (id=6)
```

2. Connect to the **host03** guest by using **vncviewer**.

Determine the vnc port number for **host03** by running the `xm list -l host03 | grep location` command.

```
# xm list -l host03 | grep location  
    (location 0.0.0.0:5903)  
    (location 3)
```

- The sample shown indicates that the port number is **5903**. Your port number might be different.


Run the `vncviewer&` command.

```
# vncviewer&
```

- The VNC Viewer: Connection Details dialog box is displayed.




Enter `localhost:<port_number>`, substituting the port number displayed from the previous `xm list -l host03 | grep location` command. For example, if the port number is 5903, enter `localhost:5903` and click **OK**.



- The Oracle Linux boot menu appears.


3. Press Enter to select **Install or upgrade an existing system** from the Oracle Linux boot menu.




- The default option of **Install or upgrade an existing system** is selected automatically after 60 seconds. Therefore, this boot menu may or may not appear, depending on how quickly you connected to **host03** using **vncviewer**.

4. The next window to appear is the Media Test window.

Use the Tab key to select the **Skip** button, and then press Enter.




5. The Logo window appears.



Scroll down if necessary and click the **Next** button.

6. The Language Selection window appears.




- a. Select the appropriate language.

If taking this class in the United States, select **English (English)**.


- b. Scroll down if necessary and click **Next**.

7. The Keyboard Selection window appears.




- Select the appropriate keyboard for your system.
If taking this class in the United States, select **U.S. English**.
- Scroll down if necessary and click **Next**.

8. The Storage Devices Selection window appears.




- a. Ensure that **Basic Storage Devices** is selected.
- b. Scroll down if necessary and click **Next**.

9. The Storage Devices Warning dialog box appears.



Click the **Yes, discard any data** button.


10. The Set Hostname window appears.



- Enter host03.example.com as the Hostname.
- Scroll down if necessary and click the **Configure Network** button.

11. Configure Network.

- The Network Connections window appears.



- a. Select **System eth0** from the list and click **Edit**. Then click the **IPv4 Settings** tab to display the following screen.




- b. Make the following changes:
- 1) Select **Connect automatically**.
 - 2) Change **Method=Manual**.
 - 3) Click **Add** to add the following:
 - a) **Address=192.0.2.103**
 - b) **Netmask=24**
 - c) **Gateway=192.0.2.1**
 - 4) Accept the defaults for all other settings.
 - c. Click **Apply**.

- d. The Network Connections window appears, select **System eth0** and click **Edit**.
- 1) Click the **IPv4 Settings** tab.
 - The screen should appear as follows.




- 2) Click **Cancel** to close the Editing System eth0 window.
- 3) Click **Close** to close the Network Connections window.
- 4) Scroll down if necessary and click **Next**.

12. The Time Zone Selection window appears.



- a. Select the appropriate time zone for your location.
The time zone can be chosen by either selecting from the drop-down list or clicking the map.
- b. Select the **System clock uses UTC** check box (this is the default).
- c. Scroll down if necessary and click **Next**.


13. The Set Root Password window appears.



- a. Enter the password **oracle** for the **root** user.
- b. Press Tab or click the **Confirm** field and re-enter the same password **oracle**.
- c. Scroll down if necessary and click **Next**.
- d. A Weak Password box appears. Click **Use Anyway**.

14. The Disk Partitioning Setup window appears.

- You select an option that creates a default partition layout.
- You then delete the default partitions and create a custom layout.
- An easier way is to select **Create Custom Layout** from the following window. However, it is beneficial to see the default partition layout.




- Select **Use All Space**.
- Scroll down if necessary.
- Do not select the **Encrypt system** check box.
- Select the **Review and modify partitioning layout** check box.




- Click **Next**.

15. The Storage Devices window appears.



- a. Three virtual disk images appear. These were pre-created for you:
 - 12 GB disk image (**system.img**) for the operating system
 - 5 GB disk image (**u01.img**) for the storage administration practices
 - 5 GB disk image (**u02.img**) for the storage administration practices
- b. Select the 12 GB disk (**xen-vbd-768** from the list in the Storage Devices window) and click the **right arrow** to move the disk to the right column.
- c. Ensure that **Boot Loader** is selected for this 12 GB disk.


- After you select the designated disk, the Storage Devices window should look like the following screenshot:



- d. Scroll down if necessary and click **Next**.


16. Delete all the default partitions.

- The Default Partition Layout window appears:



- Notice that the default layout uses Logical Volume Manager (LVM).
 - One volume group (**vg_host03**) is created.
 - Two logical volumes (**lv_root** and **lv_swap**) are created.
- Select the first entry, **vg_host03**, scroll down if needed, and click **Delete**.
- A **Confirm Delete** box appears. Click **Delete** to confirm.


- After you delete a volume group, all associated logical volumes are deleted. The window should look like the following:



- Select the first virtual hard drive, **xvda**, and click **Delete**.
- A **Confirm Delete** box is again displayed. Click **Delete** to confirm.


17. Create Custom Layout.

- After deleting all the default partitions, the window should look like the following:



- Scroll down if necessary and click **Create**.

- The Create Storage dialog box appears:




- b. Select **Standard Partition** and click **Create**.

- The Add Partition box appears.

- c. Add a standard partition with the following characteristics:
- Mount Point: /
 - File System Type: **ext4**
 - Allowable Drives: **xvda**
 - Size (MB): **6000**
 - Additional Size Options: **Fixed size**


The completed Add Partition dialog box should look like the following:



- d. Click **OK**.

- e. Click **Create** to add a second Standard Partition with the following characteristics:
- Mount Point: **/boot**
 - File System Type: **ext4**
 - Allowable Drives: **xvda**
 - Size (MB): **500**
 - Additional Size Options: **Fixed size**

The completed Add Partition dialog box should look like the following:




Note that the recommended size of /boot went from 100 MB to 500 MB between Oracle Linux 5 and Oracle Linux 6. The `initramfs` files are much larger in Oracle Linux 6 and multiple `initramfs` files can easily fill a small /boot partition.

- f. Click **OK**.

- g. Click **Create** to add a third Standard Partition with the following characteristics:
- Mount Point: **/home**
 - File System Type: **ext4**
 - Allowable Drives: **xvda**
 - Size (MB): **2000**
 - Additional Size Options: **Fixed size**


The completed Add Partition dialog box should look like the following:



- h. Click **OK**.


- i. Click **Create** to add a fourth Standard (swap) Partition with the following characteristics:
- Mount Point: **no selection**
 - File System Type: **swap**
 - Allowable Drives: **xvda**
 - Size (MB): **no selection**
 - Additional Size Options: **Fill to maximum allowable size**

The completed Add Partition dialog box should look like the following:




- j. Click **OK**.

- Confirm that the final custom layout looks like the following:



- Scroll down if necessary and click **Next**.

- A Format Warnings box appears.




- I. Click **Format**.

- A **Writing storage configuration to disk** dialog box appears.



- m. Click **Write changes to disk**.

18. The Boot Loader Selection window appears.




- Ensure that the **Install boot loader on /dev/xvda** check box is selected (this is the default).
- Click **Change device** to display the Boot loader device dialog box:



- Do not make any changes. Click **Cancel** to install the boot loader on the MBR.
- Scroll down if necessary and click **Next**.


19. Select software to install.

- The Software Package Selection window appears:




- Select **Desktop** from the list of software sets.
- Scroll down if necessary and click **Customize Now**.
- Accept all the other default selections.
- Click **Next**.

- The Customize Package Selection window appears.




- Select **Servers** from the left pane, and then select **System administration tools** from the right pane (as shown in the screenshot).
- Scroll down if necessary and click **Optional packages**.

- The Packages in System administration tools screen appears as follows:



- Scroll down if necessary and select the **oracleasm-support** package (as shown):
- Click **Close**.
- Click **Next**.
 - The Software Install screen is displayed, showing the progress as the selected software packages are installed.

- After the installation is complete, the following window appears:



j. Click **Reboot**.

- After you reboot your system, your vnc session closes.

Practice 3-2: Using FirstBoot

Overview

In this practice, you complete FirstBoot.

Assumptions

- You have completed the installation of Oracle Linux.
- You have rebooted your system.


Tasks

1. Connect to the **host03** guest by using **vncviewer**.

- a. Run the **vncviewer&** command.


```
# vncviewer&
```

- The VNC Viewer: Connection Details dialog box is displayed.
- b. Enter `localhost :<port_number>`, substituting the correct port number for the **host03** guest. For example, if the port number is 5903, enter `localhost:5903` and click **OK** (as shown in the following screenshot):




2. Complete FirstBoot.

- The **Welcome** window appears.




- Scroll down if necessary and click **Forward**.

- The **License Information** window appears.




- Scroll down if necessary and click **Yes, I agree to the License Agreement**.
- Click **Forward**.

- The **Set Up Software Updates** window appears.




- The Unbreakable Linux Network (ULN) is discussed in Lesson 6, "Package Management."
- d. Select **No, I prefer to register at a later time.**
- e. Scroll down if necessary and click **Forward**.
 - A FirstBoot window displays reasons you should connect to ULN.
- f. Click **No thanks, I'll connect later.**
- g. Scroll down if necessary and click **Forward**.

- The **Finish Updates Setup** window appears.




- You are notified that your system is not set up for software updates.
- h. Scroll down if necessary and click **Forward**.

- The Create User window appears.



- i. FirstBoot requires you to create an initial user to log in to your system
- j. Create a user with the following attributes:
 - Username: oracle
 - Full Name: Oracle Student
 - Password: oracle
 - Confirm Password: oracle
- k. Scroll down if necessary and click **Forward**.
- l. Click **Yes** to use the chosen password.

- The **Date and Time** window appears.



- Set the current date and time.
- Scroll down if necessary and click **Forward**.
 - Kdump is the last component to configure with FirstBoot.
 - However, do not enable kdump at this time.
- Scroll down if necessary and click **Finish**.

The GNOME desktop login window appears after completing FirstBoot.


Practice 3-3: Logging In to Oracle Linux and Shutting Down

Overview

In this practice, you log in as the `oracle` user. Then you shut down the **host03** VM.

Tasks


1. Log in to Oracle Linux from the GNOME desktop.
 - a. Select Oracle Student.
 - The following window appears:




- b. Enter the password `oracle` for Oracle Student.
- c. Click **Log In**.

- The GNOME desktop appears.

2. Shut down Oracle Linux.
 - a. Open the **System** menu:



- b. Select **Shut Down** from the System menu.
 - The following dialog box appears:



- c. Click **Shut Down**.
 - The VM **host03** is shut down and the vnc session closes.

Practice 3-4: Re-Creating the host03 VM Guest

Overview

In this practice, you re-create the **host03** VM.

Task

Re-create the **host03** VM:

- From **dom0**, run the `xm list` command as follows to verify **host03** is not running.

```
# xm list
Name           ID   Mem  VCPUs      State      Time(s)
Domain-0       0    1024     2        r-----  758.9
host01         4    2048     1        -b----  37.4
host02         5    2048     1        -b----  37.3
```

- Only **dom0** (Domain-0) and two guests (**host01** and **host02**) are running.

- Use the `cd` command to change to the `/OVS/running_pool/host03` directory:

```
# cd /OVS/running_pool/host03
```

- Use the `xm create` command as follows:

```
# xm create vm.cfg
Using config file "./vm.cfg".
Started domain host03 (id=8)
```

- Confirm that **host03** VM is running by using the `xm list` command as follows:

```
# xm list
Name           ID   Mem  VCPUs      State      Time(s)
Domain-0       0    1024     2        r-----  774.6
host01         4    2048     1        -b----  37.5
host02         5    2048     1        -b----  37.4
host03         8    2048     1        -b----  3.3
```

- This output shows **host03** is now running.

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 4: Linux Boot Process

Chapter 4

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 4: Linux Boot Process

Practices Overview

In these practices, you perform the following tasks:

- Explore the GRUB boot loader.
- Edit the GRUB configuration file to boot different kernels.
- Use the GRUB menu to modify kernel boot parameters.
- Change the default run level.
- Explore the `init` script directories.
- Configure a service to start and stop at a given run level.
- Explore and configuring Upstart jobs.

In Lesson 14 titled “File Sharing,” you configure `xinetd` to start `vsftpd` on demand.

Practice 4-1: Exploring the GRUB Boot Loader

Overview

In this practice, you relate boot loader options selected during the installation of Oracle Linux to actual disk partitions and directories. You explore the `/boot` directory, the GRUB configuration file, and kernel boot parameters. The displayed sample output may not exactly match what you see on your system. In some cases, only partial output is shown.

Assumptions

- You are logged on to **dom0**.
- You have a terminal window open.
- You are the `root` user.
- You completed Practice 3 (OS install and FirstBoot configuration).
- VM **host03** is running.

If you were unable to complete the OS install and FirstBoot configuration on **host03** in Practice 3, substitute **host02** for **host03** in this practice and all future practices in this course.

Tasks

1. Connect to **host03** guest using **vncviewer**.

- a. Run the `xm list -l host03 | grep location` command to determine the vnc port number for **host03**.


```
# xm list -l host03 | grep location
        (location 0.0.0.0:5903)
        (location 3)
```

- The sample shown indicates that the port number is **5903**. This may not be true in your case.

- b. Run the `vncviewer&` command.


```
# vncviewer&
```

- The VNC Viewer: Connection Details dialog box is displayed.
- c. Enter `localhost:<port_number>`, substituting the port number displayed from the previous `xm list` command. For example, if the port number is 5903, enter `localhost:5903`. Then click **OK**.



- After you connect, the GNOME login window appears.


2. Log in from the GNOME desktop.
 - a. Select **Oracle Student**.
 - The following window appears:



- b. Enter the password `oracle` for Oracle Student.
- c. Click **Log In**.

- After you log in, the GNOME desktop appears.

3. Open a terminal window and become the `root` user.
 - a. Right-click the desktop to display the pop-up menu:




- b. Click **Open in Terminal** to display a terminal window.
 - c. In the terminal window, enter the `su -` command. Then enter the `root` password `oracle` to become `root`.
- The password is displayed in the following example but does not appear on your screen.

```
$ su -
Password: oracle
# whoami
root
```


- The `whoami` command confirms that you are the `root` user.

4. The following screens serve as a reminder of the partition layout that you created when installing Oracle Linux.
- The Storage Devices that you selected during the installation process are shown here:




- You indicated that the GRUB boot loader should be installed on the 12 GB virtual drive.

- You configured the boot loader during the installation process as follows:



- You accepted the defaults, which resulted in installing the GRUB boot loader on the MBR on /dev/xvda.
- Recall from the lecture portion of the lesson that only a subset of the GRUB boot loader is written to the MBR. The remainder of the GRUB boot loader is written to the /boot partition.

- The partition layout that you created during the installation process is shown here:



Use the `df -h` command to confirm the mounted partitions.

```
# df -h
Filesystem      Size  Used  Avail Use% Mounted on
/dev/xvda2      5.7G  3.0G  2.5G  55% /
tmpfs           1002M   80K  1002M   1% /dev/shm
/dev/xvda1      477M   55M   397M  13% /boot
/dev/xvda3      1.9G   3.9M   1.8G   1% /home
/dev/sr0         3.7G   3.7G     0 100% /media/OL6.5 ...
```

- Notice that `/boot` is a separate partition mounted on `/dev/xvda1`.

5. Explore the `/boot` partition. (Your system may be different from the following sample output.)

- a. Use the `ls` command to view the `/boot/grub` directory.

```
# ls /boot/grub
device.map      grub.conf          minix_stage1_5    stage2
e2fs_stage1_5  iso9660_stage1_5  reiserfs_stage1_5
ufs2_stage1_5  fat_stage1_5     jfs_stage1_5
splash.xpm.gz   vstafs_stage1_5  ffs_stage1_5
stage1          xfs_stage1_5
```

- GRUB works in stages: stage 1, stage 1_5, and stage 2.
- Notice the various “stage” files in the `/boot/grub` directory.
- Notice the splash screen file, `splash.xpm.gz`, displayed during the boot process.
- Notice the GRUB configuration file, `grub.conf`, in this directory.

- b. Use the `cat` command to view the GRUB configuration file `/boot/grub/grub.conf`.

```
# cat /boot/grub/grub.conf
# grub.conf generated by anaconda
...
#boot=/dev/sda
default=0
timeout=5
splashimage=(hd0,0)/grub/splash.xpm.gz
hiddenmenu
title Oracle Linux Server Unbreakable Enterprise Kernel (3.8.13-
16.2.1.el6uek.x86_64)
    root (hd0,0)
    kernel /vmlinuz-3.8.13-16...
    initrd /initramfs-3.8.13-16...
title Oracle Linux Server Red Hat Compatible Kernel (2.6.32-
431.el6.x86_64)
    root (hd0,0)
    kernel /vmlinuz-2.6.32-431...
    initrd /initramfs-2.6.32-431...
```

- Each **title** line specifies a bootable kernel.
- With Oracle Linux, two bootable kernels are included:
 - UEK (3.8.13-16.2.1.el6uek.x86_64)
 - Red Hat Compatible Kernel (2.6.32-431.el6.x86_64)
- Each title has associated **root**, **kernel**, and **initrd** directives.
- With **default=0**, the kernel associated with the first title is booted.

- c. Use the `grep` command to list the bootable kernels.

```
# grep title /boot/grub/grub.conf
title Oracle Linux Server Unbreakable Enterprise Kernel (3.8.13-
16.2.1.el6uek.x86_64)
title Oracle Linux Server Red Hat Compatible Kernel (2.6.32-
431.el6.x86_64)
```

- d. Use the `ls` command to view the kernel files.

```
# ls /boot/vmlinuz*
vmlinuz-2.6.32-431.el6.x86_64
vmlinuz-3.8.13-16.2.1.el6uek.x86_64
```

- e. Use the `ls` command to view the initial RAM disk files.

```
# ls /boot/init*
initramfs-2.6.32-431.el6.x86_64.img
initramfs-3.8.13-16.2.1.el6uek.x86_64.img
```

6. View the kernel boot parameters.

- a. Use the `grep` command to view the kernel boot parameters.

```
# grep kernel /boot/grub/grub.conf
...
kernel /vmlinuz-3.8.13-16.2.1.el6uek.x86_64 ro
root=UUID=... rd_NO_LUKS rd_NO_LVM LANG=en_US.UTF-8 rd_NO_MD
SYSFONT=latarcyrheb-sun16 KEYBOARDTYPE=pc KEYTABLE=us rd_NO_DM
rhgb quiet
...
```

- Kernel boot parameters are appended to the “kernel” entry in `grub.conf`.
- These parameters are written to the `/proc/cmdline` file and viewable after boot.
- The UUID value is not displayed.

- b. Use the `cat` command to view the `/proc/cmdline` file.

```
# cat /proc/cmdline
ro root=UUID=... rd_NO_LUKS rd_NO_LVM LANG=en_US.UTF-8 rd_NO_MD
SYSFONT=latarcyrheb-sun16 KEYBOARDTYPE=pc KEYTABLE=us rd_NO_DM
rhgb quiet
```

- Notice that the content is the same as the kernel boot parameters in `grub.conf`.

Practice 4-2: Booting Different Kernels

Overview

In this practice, you edit the GRUB configuration file to boot different kernels.

Assumptions

You are the `root` user on **host03** VM.

Tasks

1. Boot the Red Hat Compatible Kernel.

a. Use the `uname -r` command to display the current running kernel:

```
# uname -r
3.8.13-16.2.1.el6uek.x86_64
```

- The current running kernel is the UEK.

b. Use the `vi` editor to change the **default** entry in `grub.conf` from 0 to 1, and change the **timeout** entry from 5 to 30.

```
# vi /boot/grub/grub.conf
...
default=0 (old entry)
timeout=5 (old entry)

default=1 (new entry)
timeout=30 (new entry)
```

- The **timeout** entry is increased to 30 seconds primarily for the next practice.

c. Use the `reboot` command to reboot your system:

```
# reboot
...
```

- After you reboot your system, your VNC session closes.

d. Connect to **host03** by using VNC.

1) Run the `vncviewer&` command.

```
# vncviewer&
```

- The VNC Viewer: Connection Details dialog box is displayed.

2) Enter the command `localhost:<port_number>`, substituting the correct port number for the **host03** guest. For example, if the port number is 5903, enter the following and click **OK**:

```
localhost:5903
```

e. Select **Oracle Student** from the GNOME login window, password is **oracle**.

f. Right-click the GNOME desktop and select **Open in Terminal** from the pop-up menu.

- g. In the terminal window, become the `root` user by entering the `su -` command followed by the `root` password `oracle`.

```
$ su -  
Password: oracle  
# whoami  
root
```

- h. Use the `uname -r` command to display the new running kernel:

```
# uname -r  
2.6.32-431.el6.x86_64
```

- With `default=1`, the kernel associated with the second title in `grub.conf` is booted.
- The second bootable kernel is the Red Hat Compatible Kernel.

2. Boot the UEK.

- Use the `vi` editor to change the `default` entry in `/boot/grub/grub.conf` from 1 back to 0.
- Use the `reboot` command to reboot your system.
 - After you reboot your system, your VNC session closes.
- Connect to `host03` by using VNC.
 - Run the command `vncviewer&` from `dom0`.
 - Enter `localhost:<port_number>`, substituting the correct port number for the `host03` guest (for example, `localhost:5903`).
- Select `Oracle Student` from the GNOME login window. The password is `oracle`.
- Right-click the GNOME desktop and open a terminal window.
- Become the `root` user by entering the `su -` command. The password is `oracle`.
- Use the `uname -r` command to ensure that you are running the UEK:

```
# uname -r  
3.8.13-16.2.1.el6uek.x86_64
```

- With `default=0`, the kernel associated with the first title in `grub.conf` is booted.
- The first bootable kernel is the UEK.

Practice 4-3: Using the GRUB Menu

Overview

In this practice, you access the GRUB menu and modify the kernel boot parameter to boot into single-user mode.

Assumptions

You are the `root` user on **host03** VM.

Tasks

1. Display the current run level.

Use the `who -r` command (or `runlevel` command) to display the current run level:

```
# who -r
    run-level 5  2013-12-06 10:25
# runlevel
N 5
```

- The current run level in the example is 5.

2. Reboot your system and display the GRUB menu.


- a. Reboot your VM by using the `reboot` command:

```
# reboot
```

- After you reboot your system, your VNC session closes.
- b. Run the command `vncviewer&` from **dom0**.
- c. Enter `localhost:<port_number>`, substituting the correct port number for the **host03** guest (for example, `localhost:5903`).


- d. Press the **Esc** key before the timeout expires to display the GRUB menu.

- The GRUB menu is shown here:




- Each entry in the GRUB menu is associated with a **title** entry in the `grub.conf` file.

3. Edit a bootable kernel entry to boot into single-user mode.
 - a. Use the up and down arrow keys if necessary to highlight the first entry (UEK).
 - b. Press the **e** key to edit the entry.
 - The associated **root**, **kernel**, and **initrd** directives for the selected title are shown:



- c. Use the up-arrow and down-arrow keys if necessary to highlight the **kernel** entry (as shown in the preceding screenshot).
- d. Press the **e** key to edit the kernel entry.


- The end of the **kernel** entry is shown as follows:



```
VNC: Xen-host03
[ Minimal BASH-like line editing is supported. For the first word, TAB lists possible command completions. Anywhere else TAB lists the possible completions of a device/filename. ESC at any time cancels. ENTER at any time accepts your changes.]
<GRUB_CMDLINE_LINUX_DEFAULT="quiet" |
```

- Use the left and right arrow keys to view the entire kernel entry.
- After viewing the entire entry, use the right arrow key to move the cursor to the end of the line.

- g. Add the word `single` to the end of the kernel line:



```
VNC: Xen-host03
[ Minimal BASH-like line editing is supported. For the first word, TAB lists possible command completions. Anywhere else TAB lists the possible completions of a device/filename. ESC at any time cancels. ENTER at any time accepts your changes. ]

<ABLE=us rd_NO_DM rhgb quiet single
```

- h. After adding the word `single` as shown in the preceding screenshot, press the **Enter** key.
i. Press the **b** key to boot.
- Note that the system boots and you are automatically logged in as the `root` user.

- j. Use the `who -r` command to determine the current run level:

```
# who -r
run-level S 2013-12-06 10:37
```

- The current run level in the example is `S`, or single-user mode.

4. Bring the system from single-user mode to the previous run level.
- Press **Ctrl + D** to bring the system back to run level 5.
 - The GNOME login window appears.
 - Select `Oracle Student` in the GNOME login window. The password is `oracle`.
 - Right-click the GNOME desktop and open a terminal window.

- d. Become the root user by entering the su – command. The password is oracle.

```
$ su -  
Password: oracle  
# whoami  
root
```

- e. Use the who -r command to determine the current run level:

```
# who -r  
run-level 5 2013-12-06 10:50 last=S
```

- The run level in the example is 5.
- The previous run level was single-user mode.

5. Change the **timeout** entry in /boot/grub/grub.conf back to 5.

Use the vi editor to change the **timeout** entry from 30 to 5.

```
# vi /boot/grub/grub.conf  
...  
timeout=30 (old entry)  
  
timeout=5 (new entry)
```

Practice 4-4: Changing the Default Run Level

Overview

In this practice, you change the default run level.

Assumptions

You are the `root` user on **host03** VM.

Tasks

1. Edit `/etc/inittab` and change the default run level to 3.
 - a. Use the `vi` editor to change the **default** run level from 5 to 3:

```
# vi /etc/inittab
...
id:5:initdefault: (old entry)
id:3:initdefault: (new entry)
```

- b. Use the `reboot` command to reboot your system:

```
# reboot
...
```

- After you reboot your system, your VNC session closes.
- c. Connect to **host03** by using VNC.

- 1) Run the `vncviewer&` command.

```
# vncviewer&
```


- The VNC Viewer: Connection Details dialog box is displayed.

- 2) Enter the command `localhost:<port_number>`, substituting the correct port number for the **host03** guest. For example, if the port number is 5903, enter the following and click **OK**:

```
localhost:5903
```

- The GNOME login does not appear. Instead, a text-based login window appears.
- Booting to run level 3 does not start the X Window services.

- The text-based login window appears:



```
VNC: Xen-host03
Oracle Linux Server release 6.5
Kernel 3.8.13-16.2.1.el6uek.x86_64 on an x86_64

host03 login: _
```

- Log in as `root`. The password is `oracle`.
- Use the `who -r` command to display the current run level:

```
# who -r
run-level 3 2013-12-06 10:52
```

- The default run level in the example is now 3.
2. Edit `/etc/inittab` and change the default run level from 3 back to 5.
- Use the `vi` editor to change the **default** run level from 3 back to 5:

```
# vi /etc/inittab
...
id:3:initdefault: (old entry)
id:5:initdefault: (new entry)
```

- Use the `reboot` command to reboot your system:

```
# reboot
...
```

- After you reboot your system, your VNC session closes.
- Connect to **host03** by using VNC.

- Run the `vncviewer&` command.

```
# vncviewer&
```

- The VNC Viewer: Connection Details dialog box is displayed.

- 2) Enter the command `localhost:<port_number>`, substituting the correct port number for the **host03** guest. For example, if the port number is 5903, enter the following and click **OK**:

```
localhost:5903
```

- The GNOME login does appear in run level 5 because the X Window services are started.
- d. Select **Oracle Student** from the GNOME login window. The password is **oracle**.
- e. Right-click the GNOME desktop and select **Open in Terminal** from the pop-up menu.
- f. In the terminal window, become the **root** user by entering the `su -` command followed by the **root** password **oracle**.

```
$ su -
Password: oracle
# whoami
root
```

- g. Use the `who -r` command to display the current run level:

```
# who -r
run-level 5 2013-12-06 10:54
```

- The default run level in the example is now 5.

Practice 4-5: Exploring and Configuring init Services

Overview

In this practice, you explore the `init` script directories. You also configure a service to start and stop at a specific run level.

Assumptions

You are the `root` user on **host03** VM.

Tasks

1. Explore the `/etc/rc.d` directory.

- a. Use the `ls -l` command to display the contents of the directory.

```
# ls -l /etc/rc.d
drwxr-xr-x.    init.d
-rwxr-xr-x.    rc
drwxr-xr-x.    rc0.d
drwxr-xr-x.    rc1.d
drwxr-xr-x.    rc2.d
drwxr-xr-x.    rc3.d
drwxr-xr-x.    rc4.d
drwxr-xr-x.    rc5.d
drwxr-xr-x.    rc6.d
-rwxr-xr-x.    rc.local
-rwxr-xr-x.    rc.sysinit
```

- `/etc/rc.d/rc.sysinit`: Runs once at boot time
- `/etc/rc.d/rc`: Starts and stops services when the run level changes
- `/etc/rc.d/rc.local`: The last script that the `init` program runs

- b. Use the `ls` command to display the contents of the `/etc/rc.d/init.d` directory.

```
# ls /etc/rc.d/init.d
abrt-ccpp  firstboot  messagebus  quota_nld  snmpd
abrtfd     functions   netconsole  rdisc      snmptrapd
...
```

- The directory contains scripts that start and stop services.
- Note that `/etc/init.d` is a symbolic link to `/etc/rc.d/init.d`.

- c. Use the `ls -l` command to display the contents of the `/etc/rc.d/rc5.d` directory.

```
# ls -l /etc/rc.d/rc5.d
lrwxrwxrwx.    K01smartd -> ../init.d/smard
lrwxrwxrwx.    K02oddjobd -> ../init.d/oddjobd
...
lrwxrwxrwx.    K95firstboot -> ../init.d/firstboot
lrwxrwxrwx.    K99rngd -> ../init.d/rngd
lrwxrwxrwx.    S01sysstat -> ../init.d/sysstat
lrwxrwxrwx.    S02lvm2-monitor -> ../init.d/lvm2-monitor
...
lrwxrwxrwx.    S99local -> ../rc.local
```

- The directory contains symbolic links to scripts in the `/etc/rc.d/init.d` directory.
- This allows services to be stopped (K files) or started (S files) when entering a particular run level (run level 5 in this example).
- Note that `/etc/rc5.d` is a symbolic link to `/etc/rc.d/rc5.d`.

- d. Use the `ls -l` command to display the services that are stopped and started when entering run level 3.

```
# ls -l /etc/rc.d/rc3.d
...
```

- Note that `/etc/rc3.d` is a symbolic link to `/etc/rc.d/rc3.d`.

2. Start the `httpd` service when entering run level 3 or run level 5.

- a. Use the `find` command to display all `http` files in the `/etc/rc.d` directory.

```
# find /etc/rc.d -name "*http*"
/etc/rc.d/rc2.d/K15httpd
/etc/rc.d/rc0.d/K15httpd
/etc/rc.d/rc3.d/K15httpd
/etc/rc.d/rc6.d/K15httpd
/etc/rc.d/init.d/httpd
/etc/rc.d/rc5.d/K15httpd
/etc/rc.d/rc1.d/K15httpd
/etc/rc.d/rc4.d/K15httpd
```

- Currently, the `httpd` service is configured to stop when entering each run level.

- b. Use the `service` command to verify that the `httpd` service is stopped.

```
# service httpd status
httpd is stopped
```

- c. Use the `chkconfig` command to query run level information for the `httpd` service.

```
# chkconfig --list httpd
httpd 0:off 1:off 2:off 3:off 4:off 5:off 6:off
```

- This confirms the `httpd` service is not configured to run at any run level.

- d. Use the `chkconfig` command to configure the `httpd` service to start at run levels 3 and 5.

```
# chkconfig --level 35 httpd on
# chkconfig --list httpd
httpd 0:off 1:off 2:off 3:on 4:off 5:on 6:off
```

- Now the service is configured to start at run levels 3 and 5.

- e. Use the `find` command to display all `http` files in the `/etc/rc.d` directory.

```
# find /etc/rc.d -name "*http*"
/etc/rc.d/rc2.d/K15httpd
/etc/rc.d/rc0.d/K15httpd
/etc/rc.d/rc3.d/S85httpd
/etc/rc.d/rc6.d/K15httpd
/etc/rc.d/init.d/httpd
/etc/rc.d/rc5.d/S85httpd
/etc/rc.d/rc1.d/K15httpd
/etc/rc.d/rc4.d/K15httpd
```

- Now a “start” script exists in the run level 3 directory, `/etc/rc.d/rc3.d/S85httpd` and in the run level 5 directory, `/etc/rc.d/rc5.d/S85httpd`.

- f. Use the `service` command to check the status of the `httpd` service.

```
# service httpd status
httpd is stopped
```

- The `chkconfig` command only configures a service to start or stop at a given runlevel.
- The `chkconfig` command does not actually start or stop a service.

- g. Use the `reboot` command to reboot your system.

```
# reboot
...
```

- After you reboot your system, your VNC session closes.

- h. Connect to **host03** by using VNC, and log in as Oracle Student (with password `oracle`).

- From **dom0**, run the `vncviewer&` command.

```
# vncviewer&
```

- The VNC Viewer: Connection Details dialog box is displayed.

- Enter the command `localhost:<port_number>`, substituting the correct port number for the **host03** guest. For example, if the port number is 5903, enter the following and click **OK**:

```
localhost:5903
```

- Select Oracle Student in the GNOME login window (password is `oracle`).

- i. Open a terminal window and become the `root` user.
 - 1) Right-click the GNOME desktop and select **Open in Terminal** from the pop-up menu.
 - 2) In the terminal window, become the `root` user by entering the `su -` command followed by the `root` password `oracle`.

```
$ su -  
Password: oracle  
# whoami  
root
```

- j. Use the `service` command to check the status of the **httpd** service.

```
# service httpd status  
httpd (pid 1405) is running...
```

- The `httpd` service is now running.

3. Stop the **httpd** service when entering run level 3 or run level 5.

- a. Use the `chkconfig` command to configure the **httpd** service so that it does not run at run levels 3 and 5.

```
# chkconfig --level 35 httpd off  
# chkconfig --list httpd  
httpd 0:off 1:off 2:off 3:off 4:off 5:off 6:off
```

- Now the service is configured so that it does not start at run levels 3 and 5.

- b. Use the `find` command to display all **http** files in the `/etc/rc.d` directory.

```
# find /etc/rc.d -name "*http*"  
/etc/rc.d/rc2.d/K15httpd  
/etc/rc.d/rc0.d/K15httpd  
/etc/rc.d/rc3.d/K15httpd  
/etc/rc.d/rc6.d/K15httpd  
/etc/rc.d/init.d/httpd  
/etc/rc.d/rc5.d/K15httpd  
/etc/rc.d/rc1.d/K15httpd  
/etc/rc.d/rc4.d/K15httpd
```

- Now the “start” scripts in the run level 3 directory, `/etc/rc.d/rc3.d/`, and in the run level 5 directory, `/etc/rc.d/rc5.d/`, are gone.

- c. Use the `service` command to check the status of the **httpd** service.

```
# service httpd status  
httpd (pid 1405) is running...
```

- The `chkconfig` command only configures a service. It does not start or stop the service.

- d. Use the `service` command to stop the `httpd` service.

```
# service httpd stop
Stopping httpd:                                [  OK  ]
# service httpd status
httpd is stopped
```

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practice 4-6: Exploring and Configuring Upstart Jobs

Overview

In this practice, you:

- View the files contained in the Upstart package
- View Upstart jobs
- Create a test Upstart job
- Start, stop, and view the status of the test job

Assumptions

You are the `root` user on **host03** VM.

Tasks

1. Query the Upstart package.

- a. Use the `rpm -q upstart` command to view the installed Upstart package.
 - The `rpm` command is discussed in Lesson 6, “Package Management.”

```
# rpm -q upstart
upstart-0.6.5-12.el6.x86_64
```

- b. Use the `rpm -ql upstart` command view the files contained in the Upstart package.

```
# rpm -ql upstart
/etc/dbus-1/system.d/Upstart.conf
/etc/init
/etc/init/init-system dbus.conf
/sbin/halt
/sbin/init
/sbin/initctl
/sbin/poweroff
/sbin/reboot
/sbin/reload
/sbin/restart
/sbin/runlevel
/sbin/shutdown
/sbin/start
/sbin/status
/sbin/stop
/sbin/telinit
...
```

- Notice that the `/etc/init` directory is provided by the Upstart package.
- Notice that some of the key system utilities such as `halt`, `init`, `reboot`, `shutdown`, and others are provided by the Upstart package.

- Notice that the `initctl` utility and the `start`, `stop`, and `status` commands are provided by the Upstart package.
 - Also notice that the package provides the `man` pages for these utilities.
2. View Upstart jobs.
- Upstart jobs are located in the `/etc/init` directory.
 - Use the `cd` command to change to the `/etc/init` directory. Use the `ls` command to list the contents of the directory.

```
# cd /etc/init
# ls
control-alt-delete.conf    rcs-sulogin.conf
init-system-dbus.conf      readahead-collector.conf
kexec-disable.conf         readahead.conf
plymouth-shutdown.conf     readahead-disable-services.conf
prefdm.conf                 serial.conf
quit-plymouth.conf          splash-manager.conf
rc.conf                     start-ttys.conf
rcS.conf                    tty.conf
rcS-emergency.conf
```

- Notice all files have a `.conf` extension.

- b. Use the `initctl list` command to list Upstart jobs.

```
# initctl list
rc stop/waiting
tty (/dev/tty3) start/running, process ...
tty (/dev/tty2) start/running, process ...
tty (/dev/tty6) start/running, process ...
tty (/dev/tty5) start/running, process ...
tty (/dev/tty4) start/running, process ...
plymouth-shutdown stop/waiting
control-alt-delete stop/waiting
rcS-emergency stop/waiting
...
```

- Notice there is a one-to-one association for each job defined in `/etc/init` with the exception of the `tty` job.
- The number of active `tty` consoles is set by the `ACTIVE_CONSOLES` directive in the `/etc/sysconfig/init` file.

3. Create a test job file.

- a. Use the `vi` editor to create `/etc/init/test.conf` with the following entries:

```
# vi test.conf
author "<your_name>"
description "Test Upstart Configuration"
```

- Both author and description are valid stanzas in Upstart job configuration files.

- Use the `initctl list` command to list Upstart jobs.

```
# initctl list
rc stop/waiting
test stop/waiting
...
```

- Notice the test job is now included in the output of `initctl list`.

- Use the vi editor to edit `/etc/init/test.conf` and add the following entries (in bold):

```
# vi test.conf
author "<your_name>"
description "Test Upstart Configuration"
script
  sleep 1000
end script
```

- Start, stop, and view the status of the test job.

- Use the `initctl start test` command to start the Upstart test job.

```
# initctl start test
test start/running, process ...
```

- Use the `initctl status test` command to view the status of the Upstart test job.

```
# initctl status test
test start/running, process ...
```

- Use the `initctl stop test` command to stop the Upstart test job.

```
# initctl stop test
test stop/waiting
```

- Use the `status test` command to view the status of the Upstart test job.

```
# status test
test stop/waiting
```

- Use the `start test` command to start the Upstart test job.

```
# start test
test start/running, process ...
```

- Use the `stop test` command to stop the Upstart test job.

```
# stop test
test stop/waiting
```

- The `start`, `stop`, and `status` can be run directly without using `initctl`.

5. View Upstart logging information and set logging priority.

- a. Use the `tail` command to view the latest entries written to the `/var/log/messages`, file.

```
# tail /var/log/messages
...
<date_time> host03 init: test main process (...) killed ...
```

- Notice that by default, only the stopping of Upstart jobs is logged.

- b. Use the `initctl log-priority` command to display the logging priority.

```
# initctl log-priority
message
```

- Notice the default logging priority is `message`.

- c. Use the `initctl log-priority debug` command to change the logging priority to `debug`.

```
# initctl log-priority debug
```

- d. Use the `start test` command to start the Upstart `test` job.

```
# start test
test start/running, process ...
```

- e. Use the `tail` command to view the latest entries written to the `/var/log/messages`, file.

```
# tail /var/log/messages
...
<date_time> host03 init: Connection from private client
<date_time> host03 init: test goal changed from stop to start
<date_time> host03 init: test state changed from waiting to
starting
<date_time> host03 init: Handling starting event
<date_time> host03 init: test state changed from starting to
pre-start
<date_time> host03 init: test state changed from pre-start to
spawned
<date_time> host03 init: test main process ...
<date_time> host03 init: test state changed from spawned to
post-start
<date_time> host03 init: test state changed from post-start to
running
<date_time> host03 init: Handling starting event
```

- Notice with logging priority set to `debug`, state changes (waiting, starting, pre-start, spawned, post-start, and running) are logged.

- f. Use the `stop test` command to stop the Upstart `test` job.

```
# stop test  
test stop/waiting
```

- g. Use the `initctl log-priority message` command to change the logging priority to message.

```
# initctl log-priority message
```

6. Start and stop the `test` job based on an event.

- a. Use the `vi` editor to edit `/etc/init/test.conf` and add the following entries (in bold):

```
# vi test.conf  
author "<your_name>"  
description "Test Upstart Configuration"  
start on event1  
stop on event2  
script  
    sleep 1000  
end script
```

- b. Use the `status test` command to display the status of the Upstart `test` job.

```
# status test  
test stop/waiting
```

- c. Use the `initctl emit event1` command to cause the `test` job to start.

```
# initctl emit event1
```

- d. Use the `status test` command to display the status of the Upstart `test` job.

```
# status test  
test start/running, process ...
```

- Notice the `test` job started on the `event1` event.

- e. Use the `initctl emit event2` command to cause the `test` service to stop.

```
# initctl emit event2
```

- f. Use the `status test` command to display the status of the Upstart `test` job.

```
# status test  
test stop/waiting
```

- Notice the `test` job stopped on the `event2` event.

- g. Do not enter the following; this is information only. You can also start and/or stop a job based on the starting/stopping of another service.

- The following entry in `test.conf` starts the `test` job when `<other-service>` is started:

```
start on started <other-service>
```

- The following entry in test.conf stops the test job when <other-service> is stopped:
stop on stopping <other-service>
- h. Do not enter the following; this is information only. You can also start and/or stop a job based on the runlevel.
 - The following entry in test.conf starts the test job on run levels 3 and 5:
start on runlevel [35]
 - The following entry in test.conf stops the test job on run level not equal to 3 or 5:
stop on runlevel [|35]
- i. Use the rm command to remove the /etc/init/test.conf file.

```
# rm /etc/init/test.conf  
rm: remove regular file '/etc/init/test.conf'? y
```

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 5: System Configuration

Chapter 5

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 5: System Configuration

Practices Overview

In these practices, you modify the system configuration files, view and modify kernel settings, and view hardware device and device driver attributes.

Practice 5-1: Exploring /etc/sysconfig

Overview

In this practice, you perform the following tasks:

- Explore the /etc/sysconfig directory.
- Explore selected files in this directory.
- See that many initialization scripts derive values from files in the directory.
- View the documentation file that describes the entries in the directory.
- Make changes to the content of files in the directory.
- Observe the effect of the changes.

Assumptions

You are the `root` user on the **host03** VM.

Tasks

1. Explore the /etc/sysconfig directory.
 - a. Use the `ls` command to display the contents of the /etc/sysconfig directory.

```
# ls /etc/sysconfig
atd          iptables-config   nfs           selinux
auditd      iptables.old     nspluginwrapper smartmontools
autoconfig   iptables        ntpd          snmpd
...
...
```

- System configuration settings are stored in these files.
- The values from many of these files are initialized by their respective initialization script when the service is started.

- b. Use the `grep` command to display specific `/etc/sysconfig` entries in initialization scripts.
- The following commands search for the existence of files in the `/etc/sysconfig` directory contained in initialization scripts in the `/etc/init.d` directory.
 - The purpose of these commands is to show that initialization scripts for various Linux services get configuration information and/or additional command-line arguments from files in the `/etc/sysconfig` directory.

```
# grep /etc/sysconfig/sshd /etc/init.d/sshd
[ -f /etc/sysconfig/sshd ] && . /etc/sysconfig/sshd
# grep /etc/sysconfig/nfs /etc/init.d/nfs
[ -f /etc/sysconfig/nfs ] && . /etc/sysconfig/nfs
# grep /etc/sysconfig/atd /etc/init.d/atd
config=/etc/sysconfig/atd
# grep /etc/sysconfig/init /etc/init.d/functions
if [ -f /etc/sysconfig/init ]; then
    . /etc/sysconfig/init
...
# grep /etc/sysconfig/network /etc/init.d/*
/etc/init.d/dnsmasq:.. /etc/sysconfig/network
/etc/init.d/netconsole:.. /etc/sysconfig/network
...
```

- The `/etc/init.d/sshd` script starts and stops the OpenSSH server daemon.
 - The configuration file for the `sshd` service is `/etc/sysconfig/sshd`.
- The `/etc/init.d/nfs` script starts and stops the NFS services.
 - The configuration file for the NFS services is `/etc/sysconfig/nfs`.
- The `/etc/init.d/atd` script starts and stops the `atd` daemon.
 - The `/etc/sysconfig/atd` file specifies additional command-line arguments for `atd`.
- The `/etc/init.d/functions` script contains functions used by most or all scripts in the `/etc/init.d` directory.
 - The `functions` script reads configuration information from the `/etc/sysconfig/init` file, which controls how the system appears and functions during the boot process.
- The `/etc/sysconfig/network` file contains network configuration information which is read by several scripts in the `/etc/init.d` directory.

2. Explore the `/usr/share/doc/init*/sysconfig.txt` file.

- a. Use the `cd` command to change to the `/usr/share/doc/init*` directory.

```
# cd /usr/share/doc/init*
# pwd
/usr/share/doc/initscripts-9.03.40
# ls
changes.ipv6          ipv6-6to4.howto      README-init
sysconfig.txt          COPYING            ipv6-tunnel.howto
static-routes-ipv6    sysvinitfiles
```

- b. Use the `less` command to display the contents of the `sysconfig.txt` file.

```
# less sysconfig.txt
...
```

- c. Search the contents of the `sysconfig.txt` file for `/etc/sysconfig`.

- While viewing the file using the `less` command, use the forward slash (/) followed by `etc/sysconfig` to search for this string.
- Press the `n` key (lowercase for “next”) to display the next instance of the string.

- d. Continue to view the `/etc/sysconfig` entries in this file.

- e. At the `/etc/sysconfig/network` entry, notice some of the variables initialized in this file. Here are some examples:

```
NETWORKING=yes|no
HOSTNAME=<fqdn by default, but whatever hostname you want>
GATEWAY=<gateway IP>
```

- f. Press the `q` key to quit the “`less`” command.

3. Change the system time zone.

- a. Use the `cat` command to display the contents of the `/etc/sysconfig/clock` file.

```
# cat /etc/sysconfig/clock
# The time zone of the system...
# This file is only for evaluation by system-config-date...
ZONE="America/Denver"
```

- Time zone settings are stored in this file.
- You can change the time zone and date/time settings after installation by editing this file or by running the `system-config-date` command.

- b. Change the time zone by running the `system-config-date` command.

```
# system-config-date
```

- Running the `system-config-date` command displays the following window.



- This window has both a **Date and Time** tab and a **Time Zone** tab.
- Change the **Time Zone** to any different zone (for example, Detroit) and click **OK**.
 - Use the `cat` command to display the contents of the `/etc/sysconfig/clock` file again.
- ```
cat /etc/sysconfig/clock
...
ZONE="America/Detroit"
```
- Note that the `/etc/sysconfig/clock` file is updated with the new time zone.
- Run the `system-config-date` command again and change the time zone back to the appropriate zone.
  - Display the contents of the `/etc/sysconfig/clock` file to confirm that the change is correct.

- Alternatively, you can edit the `/etc/sysconfig/clock` file, link the `/usr/share/zoneinfo` file to `/etc/localtime`, and then run the `system-config-date` command to confirm the change. See the following page for details:  
<http://www.redhat.com/advice/tips/timezone.html>.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 5-2: Exploring /proc

---

### Overview

In this practice, you explore the proc file system (directory), view various files and directories that represent the current state of the kernel, and change the value of current settings.

### Assumptions

You are the root user on the **host03** VM.

### Tasks

- Explore the proc file system (directory).

- Use the ls command to display the contents of the /proc directory.

```
ls /proc
1 1305 1612 1840 2012 3 879 kallsyms self
10 1310 1618 1841 2013 30 883 kcore slabinfo
1052 1328 1659 1843 2014 353 9 keys softirqs
...
...
```

- Directories with numerical names are named after a program's process ID.

- Use the ls command to display entries without numerical names.

```
ls -ld /proc/*[a-z]* | less
dr-xr-xr-x. acpi
-r--r--r--. buddyinfo
dr-xr-xr-x. bus
-r--r--r--. cgroups
-r--r--r--. cmdline
-r--r--r--. cpuinfo
...
...
```

- Notice that some entries are files and some entries are directories.

- View the meminfo file to display information about RAM.

```
less /proc/meminfo
MemTotal: 2050816 kB
MemFree: 1037256 kB
...
...
```

- d. View the `cpuinfo` file to display information about the processor(s) used by your system.

```
less /proc/cpuinfo
processor : 0
vendor_id : GenuineIntel
cpu_family : 6
model : 23
...
...
```

- e. View the `devices` file to display information about the various character and block devices currently configured.

```
less /proc/devices
Character devices:
 1 mem
 4 /dev/vc/0
 4 tty
...
Block devices:
 1 ramdisk
 2 fd
...
```

- f. View the `/proc/fs/ext4` directory to display the ext4 file systems currently configured.

```
ls /proc/fs/ext4
xvda1 xvda2 xvda3
```

## 2. Explore process directories in the `proc` file system.

- a. Use the `ls` command to display only the directories in `/proc` with numerical names.

```
ls -d /proc/*[0-9]*
1 1149 13 1464 1525 1698 1819 1846 19 2014 ...
10 1150 1304 1472 1527 17 1822 1848 1903 2015 ...
1052 1167 1305 1483 1542 1707 1825 1849 1905 2054 ...
...
```

- b. Show that one process directory exists for each process running on your system

```
ls -d /proc/*[0-9]* | wc -l
128
ps -e | wc -l
129
```

- The `ps` command produces one more entry due to the column headings.

```
ps -e | less
PID TTY TIME CMD
 1 ? 00:00:00 init
...
...
```

- View the contents of the process directory for PID=1.

```
cd /proc/1
ls -l
dr-xr-xr-x. attr
-r-----. auxv
-r--r--r--. cgroup
--w-----. clear_refs
-r--r--r--. cmdline
-rw-r--r--. comm
-rw-r--r--. coredump_filter
-r--r--r--. cpuset
lrwxrwxrwx. cwd -> /
-r-----. environ
lrwxrwxrwx. exe -> /sbin/init
dr-x-----. fd
...
...
```

- Notice that some entries are files, some entries are directories, and some entries are symbolic links.

- Use the `less` command to display the status of PID=1.

```
less status
Name: init
State: S (sleeping)
Tgid: 1
Pid: 1
PPid: 0
...
...
```

- Change the values of kernel settings.

- Use the `cat` command to check whether IP Forwarding is enabled.

```
cat /proc/sys/net/ipv4/ip_forward
0
...
```

- With the value set to 0, IP Forwarding is disabled.

- b. Use the `echo` command to enable IP Forwarding (for example, to set up a Linux router or gateway).

```
echo 1 > /proc/sys/net/ipv4/ip_forward
cat /proc/sys/net/ipv4/ip_forward
1
```

- c. Use the `echo` command to define the local port range used by TCP and UDP traffic.

```
echo "32768 61000" > /proc/sys/net/ipv4/ip_local_port_range
```

- Notice that the first local port and the last local port allowed are surrounded by quotation marks.

## Practice 5-3: Exploring the `sysfs` File System

### Overview

In this practice, you explore the `sysfs` file system. You view the virtual block devices and virtual interfaces and determine which power states are supported on your system.

### Assumptions

You are the `root` user on the **host03** VM.

### Tasks

1. Explore the `/sys` directory


- a. Use the `ls` command to display the contents of the `/sys` directory.

```
ls -l /sys
drwxr-xr-x. block
drwxr-xr-x. bus
drwxr-xr-x. class
drwxr-xr-x. dev
drwxr-xr-x. devices
drwxr-xr-x. firmware
...
```

- b. Use the `ls` command to display the virtual disk block (**xvd**) devices on your system.

```
ls -l /sys/block | grep xvd
lrwxrwxrwx. xvda -> ../devices/vbd-768/block/xvda
lrwxrwxrwx. xvdb -> ../devices/vbd-832/block/xvdb
lrwxrwxrwx. xvdd -> ../devices/vbd-5696/block/xvdd
```

- Recall these entries from the installation:



- c. Traverse the /sys/bus directory and display the virtual interface (**vif**) devices.

```
cd /sys/bus/xen/devices
ls vif*
vif-0:
devtype driver modalias net nodename power subsystem
uevent
vif-1:
devtype driver modalias net nodename power subsystem
uevent
```

- d. Display the operational state, MAC address, and MTU of **eth0**.

```
cd vif-0/net/eth0
pwd
/sys/bus/xen/devices/vif-0/net/eth0
ls
...
cat operstate
up
cat address
00:16:3e:00:01:03
cat mtu
1500
```

- e. Determine which power states are supported.

```
cd /sys/power
ls
disk image_size pm_test pm_trace resume state
cat state
mem disk
```

- mem means Suspend-to-RAM.
- disk means Suspend-to-Disk.

## Practice 5-4: Using `sysctl`

### Overview

In this practice, you use the `sysctl` utility and view the `sysctl` configuration file.

### Assumptions

You are the root user on the **host03** VM.

### Tasks

1. Use the `sysctl` utility.
  - a. Use the `sysctl` command to disable IP Forwarding.

```
cat /proc/sys/net/ipv4/ip_forward
1
sysctl -w net.ipv4.ip_forward=0
net.ipv4.ip_forward = 0
cat /proc/sys/net/ipv4/ip_forward
0
```

- b. Use the `sysctl` command to display the current kernel settings.

```
sysctl -a
...
debug.exception-trace = 1
...
dev.raid.speed_limit_min = 1000
dev.scsi.logging_level = 0
...
fs.inode-nr = 48928 370
fs.inode-state = 48928 370 0 0 0 0 0
...
kernel.sched_child_runs_first = 0
kernel.sched_latency_ns = 6000000
...
vm.page-cluster = 3
vm.panic_on_oom = 0
...
net.ipv4.route.gc_thresh = -1
net.ipv4.route.gc_timeout = 300
...
```

```
net.ipv6.neigh.default.mcast_solicit = 3
net.ipv6.neigh.default.proxy_delay = 80
...
net.netfilter.nf_log.0 = NONE
net.netfilter.nf_log.1 = NONE
...
```

2. Use the `less` command to view the contents of the `/etc/sysctl.conf` file.

```
less /etc/sysctl.conf
Kernel sysctl configuration file for Red Hat Linux
For binary values, 0 is disabled, 1 is enabled. See sysctl(8) and
sysctl.conf(5) for more details.
Controls IP packet forwarding
net.ipv4.ip_forward = 0
Controls source route verification
net.ipv4.conf.default.rp_filter = 1
...
```

- Changes that are made by using both `echo` and `sysctl` are lost when the system is rebooted.
- To preserve custom settings, add them to the `/etc/sysctl.conf` file.
- Values that are added to this file take effect each time the system boots.

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## **Practices for Lesson 6: Package Management**

**Chapter 6**

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practices for Lesson 6: Package Management

---

### Practices Overview

In these practices, you do the following:

- Use the `rpm` utility to perform software package management.
- Create a local `yum` repository and use the `yum` utility.
- Become familiar with the Oracle Unbreakable Linux Network (ULN).

You access the Oracle Public Yum Server in Practices for Lesson 13, after the network is completely configured.

## Practice 6-1: Using the rpm Utility

### Overview

In this practice, you use the `rpm` utility to query installed packages, install packages, upgrade packages, and remove packages.

### Assumptions

- You are the `root` user on the **host03** VM.
- Sample output is provided. Package names and versions shown might differ from actual output.

### Tasks

1. Query packages by using `rpm`.
  - a. Query all installed packages.

```
rpm -qa
...
rng-tools-2-13.el6.x86_64
media-player-info-6-1.el6.noarch
gnome-python2-gnome-2.28.0-3.el6.x86_64
libshout-2.2.2-5.1.el6.x86_64
```

- b. Query whether a specific package (**bash**, for example) is installed.

```
rpm -qa bash
bash-4.1.2-15.el6_4.x86_64
```

- c. Get detailed information about the **bash** package.

```
rpm -qi bash
Name : bash
Version : 4.1.2
Release : 15.el6_4
Install Date: Mon 09 Dec 2013
Group : System Environment/Shells
...
```

- d. List the files in the **bash** package.

```
rpm -ql bash
/bin/bash
/bin/sh
/etc/skel/.bash_logout
/etc/skel/.bash_profile
...
```

- e. Perform a reverse search. That is, determine what package the /etc/sysconfig file belongs to.

```
rpm -qf /etc/sysconfig
filesystem-2.4.30-3.el6.x86_64
```

- f. List the configuration files associated with the cups package.

```
rpm -qc cups
/etc/cups/classes.conf
/etc/cups/client.conf
/etc/cups/cupsd.conf
/etc/cups/lpoptions
...
```

2. Install packages by using rpm.

- a. Use the df command to determine the mount point for Oracle Linux installation media.

```
df -h
Filesystem ... Mounted on
...
/dev/sr0 ... /media/OL6.5 x86_64 Disc 1 20131125
```

- In this example, the Oracle Linux installation media is mounted on /media/OL6.5 x86\_64 Disc 1 20131125.
- If you connected to **host03** using the ssh command rather than using vncviewer as instructed, you might not have this mount point.
- Perform the following optional task only if you do not have this mount point; otherwise, proceed to task 2b.

Optional task (skip this optional task if task 2a successfully returned the mount point for the Oracle Linux installation media). Use the mkdir command to create the /media/"OL6.5 x86\_64 Disc 1 20131125" directory (quotes are required because the directory name contains white space). Ignore the error message if the OL\* directory already exists. Use the mount command to mount the Oracle Linux installation media:

```
mkdir /media/"OL6.5 x86_64 Disc 1 20131125"
mount -o ro /dev/sr0 /media/"OL6.5 x86_64 Disc 1 20131125"
```

- b. Use the cd command to change to the /media/OL\* directory.

```
cd /media/OL*
```

- c. Use the ls command to list the contents of the directory.

```
ls
EFI Packages ResilientStorage
EULA README-en RPM-GPG-KEY
...
```

- Notice the **Packages** subdirectory.

- d. Use the `cd` command to change to the **Packages** directory, and then list the contents of the directory.

```
cd Packages
ls
389-ds-base-1.2.11.15-29.el6.x86_64.rpm
...
zsh-4.3.10-7.el6.x86_64.rpm
```

- e. Verify that the **zsh** package (Z-Shell) is not already installed.

```
zsh
-bash: zsh: command not found
rpm -qa zsh
```

- In this example, the **zsh** package is not installed.

- f. Install the **zsh** package by using `rpm`.

```
rpm -Uvh zsh-4.3.10-7.el6.x86_64.rpm
...
Preparing... ##### [100%]
1:zsh ##### [100%]
```

- g. Verify that the **zsh** package is now installed.

```
rpm -qa zsh
zsh-4.3.10-7.el6.x86_64
```

- h. Run the `zsh` command, and then display the process ID of `zsh`.

```
zsh
ps
PID
7887 pts/0 00:00:00 su
7895 pts/0 00:00:00 bash
16690 pts/0 00:00:00 zsh
16697 pts/0 00:00:00 ps
```

- i. Use the `exit` command to log out of `zsh`.

```
exit
ps
PID
7887 pts/0 00:00:00 su
7895 pts/0 00:00:00 bash
16697 pts/0 00:00:00 ps
```

3. Remove packages by using rpm.

- a. Remove the **zsh** package.

```
rpm -e zsh
```

- b. Verify that the **zsh** package has been removed.

```
rpm -qa zsh
zsh
-bash: /bin/zsh: No such file or directory
```

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 6-2: Creating a Local Yum Repository

---

### Overview

In this practice, you install the `createrepo` utility (if necessary), create a Local Yum repository, and create a `repo` file to enable the local repository.

### Assumptions

- You are the `root` user on the **host03** VM.
- The Oracle Linux `dvd.iso` image is mounted on `/media`.

### Tasks

1. Disable the Public Yum repositories.

- a. Use the `vi` editor to edit the `/etc/yum.repos.d/public-yum-ol6.repo` file and set all “`enabled=1`” and “`enabled=$uek|uekr3`” to “`enabled=0`”.
  - You enable the Public Yum repository in a practice for Lesson 13, after the network is completely configured.

```
vi /etc/yum.repos.d/public-yum-ol6.repo
[public_ol6_latest]
...
enabled=1 (old value)
enabled=0 (new value)

...
[public_ol6_UEK_latest]
...
enabled=$uek (old value)
enabled=0 (new value)

...
[public_ol6_UEKR3_latest]
...
enabled=$uekr3 (old value)
enabled=0 (new value)

...
```

- b. Run the `yum clean all` command to clean up the `yum` cache.

```
yum clean all
Loaded plugins: refresh-packagekit
Cleaning repos:
Cleaning up Everything
```

2. Ensure that the Oracle Linux `dvd.iso` image is mounted on `/media`.

Use the `df` command to display the mounted file systems.

```
df -h
Filesystem Size Used Avail Use% Mounted on
...
/dev/sr0 3.7G 3.7G 0 100% /media/OL6.5_x86_64
Disc...
```

- Notice the OL6.5 media is mounted on `/media`.

3. Ensure that the `createrepo` utility is installed. If it is not, install it now.

- a. Use the `rpm` command to check whether the `createrepo` package is installed.

```
rpm -qa createrepo
```

- In this example, the package is not installed.

- b. Change to the `/media/OL*/Packages` directory to install `createrepo`.

```
cd /media/OL*/Packages
ls creat*
createrepo-0.9.9-18.0.1.el6.noarch.rpm
```

- c. Use the `rpm` command to install the package.

```
rpm -Uvh createrepo-0.9.9-18.0.1.el6.noarch.rpm
...
error: Failed dependencies:
 python-deltarpm is needed by createrepo...
```

- The package does not install because a dependent package is not installed.

- d. Attempt to use `yum` to install because `yum` resolves dependencies.

```
yum install createrepo
Loaded plugins: refresh-packagekit
Setting up Install Process
No package createrepo available
Error: Nothing to do
```

- The `yum` command fails because no repository is available.

- e. The only alternative is to use `rpm` and install the dependent package first, and then install `createrepo`.

```
rpm -Uvh python-deltarpm-3.5-0.5.20090913git.el6.x86_64.rpm
error: Failed dependencies:
 deltarpm = 3.5-0.5.20090913git.el6 is needed by python-
deltarpm-3.5-0.5.20090913git.el6.x86_64
```

- f. This package failed to install due to another dependency. Use the `rpm` command to install this `deltarpm` package. Then install `python-deltarpm` then `createrepo`.

```
rpm -Uvh deltarpm-3.5-0.5.20090913git.el6.x86_64.rpm
...
Preparing... ## [100%]
1:deltarpm ## [100%]
rpm -Uvh python-deltarpm-3.5-0.5.20090913git.el6.x86_64.rpm
...
Preparing... ## [100%]
1:python-deltarpm ## [100%]
rpm -Uvh createrepo-0.9.9-18.0.1.el6.noarch.rpm
...
Preparing... ## [100%]
1:createrepo ## [100%]
```

- g. Verify that `createrepo` is installed and in the search path.

```
which createrepo
/usr/bin/createrepo
```

#### 4. Create the local repository.

- a. Change to the `/media` directory.

```
cd /media
ls
OL6.5 x86_64 Disc 1 20131125
```

- b. Use the `createrepo` command to create a repository of the current directory.

- This command takes several minutes to complete.

```
createrepo .
Spawning worker 0 with 26278 pkgs
Workers Finished
Gathering worker results
Saving Primary metadata
Saving file lists metadata
Saving other metadata
Generating sqlite DBs
Sqlite DBs complete
```

- c. View the results of the `createrepo` command.

```
pwd
/media
ls -l
drwxr-xr-x. OL6.5 x86_64 Disc 1 20131125
drwxr-xr-x. repodata
```

- Notice that the `repodata` directory has been created.

- d. View the contents of the `repodata` directory.

```
ls -l repodata
-rw-r--r-- ...other.sqlite.bz2
-rw-r--r-- ...other.xml.gz
-rw-r--r-- ...filelists.sqlite.bz2
-rw-r--r-- ...filelists.xml.gz
...
-rw-r--r-- repomd.xml
```

- e. Run the `yum clean all` command to clean up the `yum` cache.

```
yum clean all
Loaded plugins: refresh-packagekit, security
Cleaning repos: InstallMedia
Cleaning up Everything
```

- f. Change to the `yum` repository directory and use the `vi` editor to create the `iso.repo` file:

```
cd /etc/yum.repos.d
vi iso.repo
[Myrepo]
name=Oracle Linux
baseurl=file:///media
enabled=1
gpgkey=file:///media/RPM-GPG-KEY-oracle
gpgcheck=1
```

- g. In this example, the GPG key is located on the Oracle Linux `dvd.iso` image.

```
cd /media/OL*
ls *GPG*
RPM-GPG-KEY RPM-GPG-KEY-oracle
```

- There are two files that contain the GPG key.
- The files are the same, so either can be used.

```
diff RPM-GPG-KEY RPM-GPG-KEY-oracle
```

- h. Use the `cp` command to copy the `RPM-GPG-KEY-oracle` file from `/media/OL*` to `/media`.

```
cp /media/OL*/RPM-GPG-KEY-oracle /media
```

- i. Manually install the public key.

```
rpm --import RPM-GPG-KEY-oracle
```

- j. Confirm the import of the public key.

```
rpm -qa gpg-pubkey
gpg-pubkey-...
```

- k. Use the `yum repolist` command to list the configured repositories.

```
yum repolist
...
repo id repo name status
Myrepo Oracle Linux 26,278
repolist: 26,278
```

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 6-3: Using the yum Utility

---

### Overview

In this practice, you use the `yum` utility to list both installed packages and packages available to be installed, install a package, check for packages with updates available, update a package, and remove a package.

### Assumptions

- This practice assumes that you have a `yum` repository configured.
- This practice is performed from the **host03** VM.
- You are logged on as `root` on **host03**.

### Tasks

1. List packages by using `yum`.
  - a. List all packages installed on your system and all packages available in all configured repositories.

```
yum list
Loaded plugins: refresh-packagekit, security
Installed Packages
ConsoleKit.x86_64 0.4.1-3.el6 @anaconda-
Ora...
ConsoleKit.libs.x86_64 0.4.1-3.el6 @anaconda-
Ora...
...
zlib-devel.i686 1.2.3-27.el6 Myrepo
zlib-devel.x86_64 1.2.3-27.el6 Myrepo
zsh.x86_64 4.3.10-7.el6 Myrepo
```

- Notice that some packages were installed during installation (`@anaconda...`) and some are available in the new repository (`Myrepo`).

- b. List only the installed packages.

```
yum list installed
Loaded plugins: refresh-packagekit, security
Installed Packages
ConsoleKit.x86_64 0.4.1-3.el6 @anaconda-Ora...
...
zlib.x86_64 1.2.3-29.el6 @anaconda-Ora...
```

- c. List only the packages that are available to be installed from enabled repositories.

```
yum list available
Available Packages
389-ds-base.x86_64 1.2.11.15-29.el6 Myrepo
389-ds-base-libs.i686 1.2.11.15-29.el6 Myrepo
...
```

- d. List the name of the package to which the /etc/sysconfig/crond file belongs.

```
yum provides /etc/sysconfig/crond
cronie-1.4.4-12.el6.x86_64
...
```

2. Install packages by using yum.

Install the **389-ds-base** package.

- Answer **y** when prompted.

```
yum install 389-ds-base
Setting up Install Process
Resolving Dependencies
--> Running transaction check
...
--> Finished Dependency Resolution
Dependencies Resolved
=====
Package
=====
Installing:
 389-ds-base
Installing for dependencies:
 389-ds-base-libs
...
Transaction Summary
=====
Install 14 Package(s)
Total download size: 3.5 M
Installed size: 12 M
Is this ok [y/N]: y
Downloading packages
...
Running rpm_check debug
Running Transaction Test
...
 Installing : svrcore-4.0.4-5.1.el6.x86_64
...
Installed:
 389-ds-base.x86_64 0:1.2.11.15-29.el6
Dependency Installed:
 389-ds-base-libs.x86_64 0:1.2.11.15-29.el6
...
Complete!
```

3. Update packages by using `yum`.

- You have no packages with updates available.
  - This exercise is an example and shows only the steps needed to update a package.
- a. Check which installed packages have updates available.

```
yum check-update
DeviceKit-power.x86_64 014-1.el6 Myrepo
...
```

b. Update the **DeviceKit-power** package.

```
yum update DeviceKit-power
Setting up Update Process
Resolving Dependencies
--> Running transaction check
...
--> Finished Dependency Resolution
Dependencies Resolved
=====
Package
=====
Updating:
 DeviceKit-power
Transaction Summary
=====
Upgrade 1 Package(s)
Total download size: 90 k
Is this ok [y/N] : y
Downloading packages
DeviceKit-power-014-3.el6.x86_64.rpm
Running rpm_check debug
Running Transaction Test
...
Updating : DeviceKit-power-014-3.el6.x86_64
Cleanup : DeviceKit-power-014-3.el6.x86_64

Updated:
 DeviceKit-power-014-3.el6.x86_64
Complete!
```

4. Remove packages by using `yum`.

Remove the **389-ds-base** package.

- If the following “`yum remove`” command fails, run the “`yum clean all`” command to clean up the `yum` cache, then run the “`yum remove`” command again.

- Answer **y** when prompted.

```
yum remove 389-ds-base
Setting up Remove Process
Resolving Dependencies
--> Running transaction check
...
--> Finished Dependency Resolution
Dependencies Resolved
=====
Package
=====
Removing:
 389-ds-base
Transaction Summary
=====
Remove 1 Package(s)
Installed size: 4.9 M
Is this ok [y/N]: y
Downloading packages
Running rpm_check debug
Running Transaction Test
...
Erasing : 389-ds-base-1.2.9.14-1.el6.x86_64

Removed:
 389-ds-base.x86_64 0:1.2.9.14-1.el6
Complete!
```

## Practice 6-4: Unbreakable Linux Network (ULN)

### Overview

In this practice, you become familiar with the ULN web interface. You cannot access the ULN without a Customer Support Identifier (CSI). All you can do is read through the tasks in this practice to help understand the capabilities of ULN.

### Assumptions

This practice is not intended to be a hands-on exercise.


### Tasks

1. Access the Unbreakable Linux Network (ULN).
  - a. From a browser, enter the following URL for the Unbreakable Linux Network (ULN):  
<https://linux.oracle.com>.




- b. Click **Sign On** to be prompted for login credentials.

- c. The first time you access the ULN, use your email address and your Customer Support Identifier (CSI) to log in. You are then required to create a password.




- In the future, use your email address and password to log in to ULN.
  - When you register for the first time, you are subscribed to the default channel of ol6\_<arch>\_latest.
  - After registering, you can specify a different channel by using the web interface.
2. Use the ULN web interface.

- a. You are first taken to the **ULN Home** page.



- The ULN web interface allows subscribers to register systems and subscribe to channels.
- b. Click the **Channels** tab to list of all the channels available via ULN.



The screenshot shows the Oracle Unbreakable Linux Network (ULN) web interface. The top navigation bar includes links for Home, Channels, Systems, Errata, CVE, and CSI Administration. The Channels tab is selected. Below the navigation, a breadcrumb trail shows Home > Channels. The main content area displays a table titled "Channels" with the following columns: Name, Label, Description, and Packages. The table lists various channels for Oracle Linux 6 x86\_64, such as "Oracle Linux 6 Latest (x86\_64)", "Latest Unbreakable Enterprise Kernel for Oracle Linux 6 (x86\_64)", and "Dtrace for Oracle Linux 6 (x86\_64) - Latest". Each row provides a detailed description of the channel and the number of packages available for download.

| Name                                                                                         | Label                              | Description                                                                                                              | Packages             |
|----------------------------------------------------------------------------------------------|------------------------------------|--------------------------------------------------------------------------------------------------------------------------|----------------------|
| <a href="#">Oracle Linux 6 Latest (x86_64)</a>                                               | ol6_x86_64_latest                  | All packages released for Oracle Linux 6 (x86_64) including the latest errata packages. (x86_64)                         | <a href="#">6768</a> |
| <a href="#">Latest Unbreakable Enterprise Kernel for Oracle Linux 6 (x86_64)</a>             | ol6_x86_64_UEK_latest              | Latest Unbreakable Enterprise Kernel packages for Oracle Linux 6 (x86_64)                                                | <a href="#">15</a>   |
| <a href="#">Dtrace for Oracle Linux 6 (x86_64) - Latest</a>                                  | ol6_x86_64_Dtrace_latest           | Latest packages required for Dtrace on Oracle Linux 6 (x86_64)                                                           | <a href="#">11</a>   |
| <a href="#">Unbreakable Enterprise Kernel Release 3 for Oracle Linux 6 (x86_64) - Latest</a> | ol6_x86_64_UER3_latest             | Latest packages for Unbreakable Enterprise Kernel Release 3 for Oracle Linux 6 (x86_64)                                  | <a href="#">22</a>   |
| <a href="#">Oracle Linux 6 GA (x86_64)</a>                                                   | ol6_ga_x86_64_base                 | All packages released for Oracle Linux 6 GA (x86_64). No errata included.                                                | <a href="#">6010</a> |
| <a href="#">Oracle Linux 6 GA Patches (x86_64)</a>                                           | ol6_ga_x86_64_patch                | Updated packages published after release of Oracle Linux 6 (x86_64)                                                      | <a href="#">724</a>  |
| <a href="#">Oracle Linux 6 Dtrace Userspace Tools (x86_64) - Latest</a>                      | ol6_x86_64_Dtrace_userspace_latest | The latest Dtrace userspace tools for Oracle Linux 6 (x86_64).                                                           | <a href="#">2</a>    |
| <a href="#">Oracle Linux 6 Update 1 installation media copy (x86_64)</a>                     | ol6_u1_x86_64_base                 | All packages released on the Oracle Linux 6 Update 1 (x86_64) installation media. This channel does not contain updates. | <a href="#">6148</a> |
| <a href="#">Oracle Linux 6 Update 1 Patch (x86_64)</a>                                       | ol6_u1_x86_64_patch                | Updated packages published after release of Oracle Linux 6 Update 1 (x86_64).                                            | <a href="#">631</a>  |
| <a href="#">Oracle Linux 6 Update 2 installation media copy (x86_64)</a>                     | ol6_u2_x86_64_base                 | All packages released on the Oracle Linux 6 Update 2 (x86_64) installation media. This channel does not contain updates. | <a href="#">6279</a> |
| <a href="#">Oracle Linux 6 Update 2 Patch (x86_64)</a>                                       | ol6_u2_x86_64_patch                | Updated packages published after release of Oracle Linux 6 Update 2 (x86_64).                                            | <a href="#">782</a>  |
| <a href="#">Oracle Linux 6 Update 3 installation media copy (x86_64)</a>                     | ol6_u3_x86_64_base                 | All packages released on the Oracle Linux 6 Update 3 (x86_64) installation media. This channel does not contain updates. | <a href="#">6324</a> |
| <a href="#">Oracle Linux 6 Update 3 Patch (x86_64)</a>                                       | ol6_u3_x86_64_patch                | Updated packages published after release of Oracle Linux 6 Update 3 (x86_64)                                             | <a href="#">743</a>  |
| <a href="#">Oracle Linux 6 Update 4 installation media copy (x86_64)</a>                     | ol6_u4_x86_64_base                 | All packages released on the Oracle Linux 6 Update 4 (x86_64) installation media. This channel does not contain updates. | <a href="#">6245</a> |
| <a href="#">Oracle Linux 6 Update 4 Patch (x86_64)</a>                                       | ol6_u4_x86_64_patch                | Updated packages published after release of Oracle Linux 6 Update 4 (x86_64)                                             | <a href="#">846</a>  |
| <a href="#">Oracle Linux 6 Update 5 installation media copy (x86_64)</a>                     | ol6_u5_x86_64_base                 | All packages released on the Oracle Linux 6 Update 5 (x86_64) installation media. This channel does not contain updates. | <a href="#">6421</a> |
| <a href="#">Oracle Linux 6 Update 5 Patch (x86_64)</a>                                       | ol6_u5_x86_64_patch                | Updated packages published after release of Oracle Linux 6 Update 5 (x86_64)                                             | <a href="#">115</a>  |
| <a href="#">HA Utilities for MySQL and Oracle Linux 6 (x86_64)</a>                           | ol6_x86_64_mysql-ha-utils          | Management Utilities for MySQL HA with Oracle Linux 6                                                                    | <a href="#">2</a>    |

- For each channel, you can see its detailed description and how many packages are available for download.
- You can also view only the channels for a specific Release and Architecture. The screenshot displays the channels for Oracle Linux 6 x86\_64 architecture.

- c. Click the **Systems** tab to see registered systems and the number of RPMs available on the subscribed channels to be downloaded and installed.

The screenshot shows the Oracle Unbreakable Linux Network interface. At the top, there's a navigation bar with links for Home, Channels, Systems, and Errata. Below that, a breadcrumb trail indicates the current location: Home > Systems > System Detail. The main content area is titled "System Details" and displays the following system information:

- Name: mouser
- Architecture: x86\_64
- OS Release: 6
- Release Name: oraclelinux-release
- Registered: 2011-09-20

Below this, there's a section titled "Subscribed Channels" with a "Manage Subscriptions" button. A table lists one channel:

| Name                           | Description                                                                                      |
|--------------------------------|--------------------------------------------------------------------------------------------------|
| Oracle Linux 6 Latest (x86_64) | All packages released for Oracle Linux 6 (x86_64) including the latest errata packages. (x86_64) |

Finally, there's a large table titled "Available Errata" with columns for Type, Severity, Advisory, Summary, Release Date, and Num. of Downloadable Packages. The table contains 13 rows of errata entries.

- From this screen, you can also subscribe to additional channels.
- For each channel, you can see a comprehensive listing of each available package. You can also search the list.
- You can get many package details, including a list of the files installed on the system when the RPM is downloaded and a list of the other packages that are necessary for this one to be installed.
- You can also select to download the source RPMs in addition to the binary RPMs.

- d. Click the **Errata** tab to list all errata that have been published.

The screenshot shows the Oracle Unbreakable Linux Network website. The top navigation bar includes links for Home, Channels, Systems, and Errata. The Errata tab is selected. Below the navigation, a breadcrumb trail shows 'Home > Errata'. The main content area is titled 'Errata' and displays a table of published errata. The table has columns for Type, Severity, Advisory, Summary, Systems Affected, and Release Date. The 'Type' column uses icons to represent different advisory types. The 'Severity' column includes categories like 'Important' and 'Moderate'. The 'Advisory' column lists specific update identifiers. The 'Summary' column provides a brief description of each update. The 'Systems Affected' column shows the count of systems impacted, which is consistently 0 for all entries. The 'Release Date' column shows dates ranging from September 10, 2011, to September 14, 2011. A legend on the right side defines the icons: a yellow circle for Enhancement, a red circle for Bug fix, and a green circle for Security. A task menu on the right offers options to subscribe to errata mailing lists.

| Type | Severity  | Advisory                       | Summary                                | Systems Affected | Release Date |
|------|-----------|--------------------------------|----------------------------------------|------------------|--------------|
| !    | Important | <a href="#">ELSA-2011-1317</a> | cyrus-imapd security update            | 0                | 2011-09-10   |
| !    | -         | <a href="#">ELBA-2011-1316</a> | cups bug fix update                    | 0                | 2011-09-10   |
| !    | -         | <a href="#">ELBA-2011-1314</a> | ipa bug fix update                     | 0                | 2011-09-15   |
| !    | Moderate  | <a href="#">ELSA-2011-1293</a> | squid security update                  | 0                | 2011-09-14   |
| !    | -         | <a href="#">ELBA-2011-1290</a> | ipa-client bug fix update              | 0                | 2011-09-14   |
| !    | -         | <a href="#">ELBA-2011-1284</a> | xmlrpc-c bug fix update                | 0                | 2011-09-13   |
| !    | -         | <a href="#">ELBA-2011-1285</a> | portreserve bug fix update             | 0                | 2011-09-13   |
| !    | -         | <a href="#">ELBA-2011-1287</a> | device-mapper-multipath bug fix update | 0                | 2011-09-13   |
| !    | Moderate  | <a href="#">ELSA-2011-1289</a> | librsvg2 security update               | 0                | 2011-09-13   |
| !    | -         | <a href="#">ELBA-2011-1286</a> | ghostscript bug fix update             | 0                | 2011-09-13   |
| !    | Important | <a href="#">ELSA-2011-1282</a> | nss and nspr security update           | 0                | 2011-09-12   |
| !    | -         | <a href="#">ELBA-2011-2026</a> | up2date bug fix update                 | 0                | 2011-09-12   |
| !    | -         | <a href="#">ELBA-2011-2027</a> | enterprise-release bug fix             | 0                | 2011-09-12   |
| !    | -         | <a href="#">ELBA-2011-1280</a> | certmonger bug fix update              | 0                | 2011-09-07   |
| !    | Important | <a href="#">ELSA-2011-1267</a> | thunderbird security update            | 0                | 2011-09-07   |
| !    | Important | <a href="#">ELSA-2011-1268</a> | firefox security update                | 0                | 2011-09-07   |
| !    | Important | <a href="#">ELSA-2011-1212</a> | kernel security and bug fix update     | 0                | 2011-09-07   |
| !    | Important | <a href="#">ELSA-2011-1264</a> | gstreamer-plugins security update      | 0                | 2011-09-07   |
| !    | Important | <a href="#">ELSA-2011-1266</a> | seamonkey security update              | 0                | 2011-09-07   |
| !    | -         | <a href="#">ELBA-2011-1255</a> | binutils bug fix update                | 0                | 2011-09-06   |
| !    | -         | <a href="#">ELBA-2011-1262</a> | openafs bug fix update                 | 0                | 2011-09-06   |

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## **Practices for Lesson 7: Ksplice**

**Chapter 7**

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practices for Lesson 7: Ksplice

---

### Practices Overview

In these practices, you do the following:

- View the Ksplice Offline Client Packages from the Unbreakable Linux Network (ULN).
- Upload the Ksplice packages from **dom0** to the **host03** VM.
- Install the Ksplice Offline Client and Kernel Updates to **host03**.
- View the effective kernel version after updates are applied.
- List the active Oracle Ksplice updates in your running kernel.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 7-1: Viewing Ksplice Offline Client Packages from the Unbreakable Linux Network (ULN)

### Overview

In this practice, you:

- View the Ksplice Offline Client files available from ULN
- View the `167283.sh` file

### Assumptions

- You are logged on as the `root` user on `dom0`.
- To access ULN, you must be a licensed Oracle customer with an active Oracle Linux support subscription.

### Tasks

1. View Ksplice Offline Client files.

a. The following screenshot shows ULN channels for Oracle Linux 6 x86\_64 architecture.

| Channel                                                                                                        | Description                                                                                                                                                                                                             | Package Count |
|----------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------|
| <a href="#">Oracle Linux 6 (x86_64) - Latest</a>                                                               | All packages released for Oracle Linux 6 GA (x86_64). No errata included.                                                                                                                                               | 6010          |
| <a href="#">Oracle Linux 6 GA (x86_64)</a>                                                                     | Updated packages published after release of Oracle Linux 6 (x86_64)                                                                                                                                                     | 724           |
| <a href="#">Oracle Linux 6 Dtrace Userspace Tools (x86_64) - Latest</a>                                        | The latest Dtrace userspace tools for Oracle Linux 6 (x86_64).                                                                                                                                                          | 2             |
| <a href="#">Oracle Linux 6 Update 1 installation media copy (x86_64)</a>                                       | All packages released on the Oracle Linux 6 Update 1 (x86_64) installation media. This channel does not contain updates.                                                                                                | 6148          |
| <a href="#">Oracle Linux 6 Update 1 Patch (x86_64)</a>                                                         | Updated packages published after release of Oracle Linux 6 Update 1 (x86_64).                                                                                                                                           | 631           |
| <a href="#">Oracle Linux 6 Update 2 installation media copy (x86_64)</a>                                       | All packages released on the Oracle Linux 6 Update 2 (x86_64) installation media. This channel does not contain updates.                                                                                                | 6279          |
| <a href="#">Oracle Linux 6 Update 2 Patch (x86_64)</a>                                                         | Updated packages published after release of Oracle Linux 6 Update 2 (x86_64).                                                                                                                                           | 787           |
| <a href="#">Oracle Linux 6 Update 3 installation media copy (x86_64)</a>                                       | All packages released on the Oracle Linux 6 Update 3 (x86_64) installation media. This channel does not contain updates.                                                                                                | 6324          |
| <a href="#">Oracle Linux 6 Update 3 Patch (x86_64)</a>                                                         | Updated packages published after release of Oracle Linux 6 Update 3 (x86_64)                                                                                                                                            | 743           |
| <a href="#">Oracle Linux 6 Update 4 installation media copy (x86_64)</a>                                       | All packages released on the Oracle Linux 6 Update 4 (x86_64) installation media. This channel does not contain updates.                                                                                                | 6245          |
| <a href="#">Oracle Linux 6 Update 4 Patch (x86_64)</a>                                                         | Updated packages published after release of Oracle Linux 6 Update 4 (x86_64)                                                                                                                                            | 846           |
| <a href="#">Oracle Linux 6 Update 5 installation media copy (x86_64)</a>                                       | All packages released on the Oracle Linux 6 Update 5 (x86_64) installation media. This channel does not contain updates.                                                                                                | 6421          |
| <a href="#">Oracle Linux 6 Update 5 Patch (x86_64)</a>                                                         | Updated packages published after release of Oracle Linux 6 Update 5 (x86_64)                                                                                                                                            | 115           |
| <a href="#">HA Utilities for MySQL and Oracle Linux 6 (x86_64)</a>                                             | Management Utilities for MySQL HA with Oracle Linux 6                                                                                                                                                                   | 2             |
| <a href="#">Oracle Linux 6 GDM Multiseat</a>                                                                   | GDM Multiseat packages required for Sun Ray Software on Oracle Linux 6 (x86_64)                                                                                                                                         | 7             |
| <a href="#">Oracle Software for Oracle Linux 6 (x86_64)</a>                                                    | Oracle Software for Oracle Linux 6 (x86_64)                                                                                                                                                                             | 20            |
| <a href="#">Oracle Linux 6 Add ons (x86_64)</a>                                                                | Oracle Linux 6 Add ons (x86_64)                                                                                                                                                                                         | 33            |
| <a href="#">Unbreakable Enterprise Kernel for Oracle Linux 6 (x86_64)</a>                                      | Unbreakable Enterprise Kernel for Oracle Linux 6 (x86_64) released on the install media.                                                                                                                                | 15            |
| <a href="#">Unbreakable Enterprise Kernel Release 3 (3.8 based) for Oracle Linux 6 (x86_64) - Beta release</a> | All packages required for beta release of 3.8 based Unbreakable Enterprise Kernel Release 3 for Oracle Linux 6 (x86_64). Packages from this channel are not supported and are not intended for production environments. | 99            |
| <a href="#">Dtrace for Oracle Linux 6 (x86_64) - Beta release</a>                                              | All packages required for the beta release of Dtrace for Oracle Linux 6 (x86_64). Packages from this channel are not supported and are not intended for production environments.                                        | 22            |
| <a href="#">MySQL for Oracle Linux 6 (x86_64)</a>                                                              | Latest MySQL packages released on the Oracle Linux 6 (x86_64).                                                                                                                                                          | 23            |
| <a href="#">Ksplice for Oracle Linux 6 (x86_64)</a>                                                            | Oracle Ksplice clients, updates, and dependencies for Oracle Linux 6 (x86_64).                                                                                                                                          | 121           |
| <a href="#">OFED supporting tool packages for Unbreakable Enterprise Kernel on Oracle Linux 6 (x86_64)</a>     | Latest OpenFabrics Enterprise Distribution (OFED) supporting tools for the Unbreakable Enterprise Kernel (UEK) on Oracle Linux 6 (x86_64).                                                                              | 38            |
| <a href="#">ol6_x86_64_ofed_UEK</a>                                                                            |                                                                                                                                                                                                                         |               |

- Notice the “Ksplice for Oracle Linux 6 (x86\_64)” entry towards the bottom of the screen.
- Also notice that this channel contains 121 packages.

- b. Clicking the “121” packages link displays the following screen:

| Package                                                                         | Description                                                                                         |
|---------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------|
| <a href="#">python-ksplice-uptrack-0.2.1-1.el6.noarch</a>                       | -                                                                                                   |
| <a href="#">uptrack-1.2.14-0.el6.noarch</a>                                     | Client for the Ksplice Uptrack rebootless kernel update service                                     |
| <a href="#">uptrack-libyaml-0.1.3-1.el6.x86_64</a>                              | -                                                                                                   |
| <a href="#">uptrack-offline-1.2.15.offline-0.el6.noarch</a>                     | Oracle Linux Support tool - offline client for the Ksplice Uptrack rebootless kernel update service |
| <a href="#">uptrack-python-cjson-1.0.5-5.el6.x86_64</a>                         | -                                                                                                   |
| <a href="#">uptrack-PyYAML-3.08-4.el6.x86_64</a>                                | -                                                                                                   |
| <a href="#">uptrack-updates-2.6.32-100.28.11.el6.x86_64-20131014-0.noarch</a>   | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-100.28.15.el6.x86_64-20131014-0.noarch</a>   | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-100.28.17.el6.x86_64-20131014-0.noarch</a>   | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-100.28.9.el6.x86_64-20131014-0.noarch</a>    | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-100.34.1.el6uek.x86_64-20131014-0.noarch</a> | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-100.35.1.el6uek.x86_64-20131014-0.noarch</a> | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-100.36.1.el6uek.x86_64-20131014-0.noarch</a> | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-100.37.1.el6uek.x86_64-20131014-0.noarch</a> | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-131.0.15.el6.x86_64-20131018-0.noarch</a>    | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-131.12.1.el6.x86_64-20131018-0.noarch</a>    | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-131.17.1.el6.x86_64-20131018-0.noarch</a>    | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-131.21.1.el6.x86_64-20131018-0.noarch</a>    | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-131.21.1.el6.x86_64-20131018-0.noarch</a>    | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-131.4.1.el6.x86_64-20131018-0.noarch</a>     | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-131.6.1.el6.x86_64-20131018-0.noarch</a>     | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-200.16.1.el6uek.x86_64-20131014-0.noarch</a> | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-200.19.1.el6uek.x86_64-20131014-0.noarch</a> | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |
| <a href="#">uptrack-updates-2.6.32-200.20.1.el6uek.x86_64-20131014-0.noarch</a> | Rebootless updates for the Ksplice Uptrack rebootless kernel update service                         |

- This is only a partial list of the 121 packages available from this channel.
- Notice there is an “uptrack-updates” RPM for each supported kernel.
- The first six packages are not specific to a kernel version.
- The uptrack-offline package is the Ksplice Offline Client.

**Note:** The uptrack-offline package will not install if the uptrack-1.2.14-0.el6.noarch.rpm package is already installed.

## 2. View downloaded packages.

As the root user on dom0, use the cd command to change to the /ovs/seed\_pool/ksplice directory. Use the ls command to display the contents of the directory.

```
[dom0]# cd /ovs/seed_pool/ksplice
[dom0]# ls
167283.sh
uptrack-libyaml-0.1.3-1.el6.x86_64.rpm
uptrack-offline-1.2.15.offline-0.el6.noarch.rpm
uptrack-python-cjson-1.0.5-5.el6.x86_64.rpm
uptrack-PyYAML-3.08-4.el6.x86_64
uptrack-updates-3.8.13-16.2.1.el6uek.x86_64-20131219-0.noarch.rpm
```

- All files in this directory, with the exception of the 167283.sh file, have been downloaded from ULN.

- The 167283.sh file has been downloaded from <http://www.oracle.com/ocom/groups/public/@otn/documents/webcontent/167283.sh>.
- The uptrack-offline package and dependencies are provided.
- The package containing kernel updates for the 3.8.13-16.2.1.el6uek.x86\_64 kernel is also provided.

3. View the 167283.sh file.

Use the less command to view the 167283.sh file.

- Some of the comments in the script are shown as follows:

```
[dom0]# less 167283.sh
#!/bin/sh
...
yum repository paths
...
Identify what OS the script is running on.
...
create the directory for the generated repo files
...
easiest way to get the list of channel labels ...
...
determine the correct yum repository directory
...
Create the required yum repository directory.
...
figure out which rpms are new
...
figure out which rpms are there currently
...
actually get the rpms
...
run createrepo to generate yum repo metadata
...
```

- This 167283.sh file is not required for this practice. It is only required to create the Yum repositories for the registered channels.
- The purpose of this task is just to familiarize you with the contents of this 167283.sh script.

## Practice 7-2: Using sftp to Upload Ksplice Packages

### Overview

In this practice, you perform the following:

- Use sftp to upload the Ksplice packages from **dom0** to the **host03** VM.
- The sftp command is discussed in Lesson 15 – OpenSSH.

### Assumptions

You are the `root` user on **dom0**.

### Tasks

1. Use the sftp command to transfer the Ksplice packages from **dom0** to **host03**.
  - a. From the `/OVS/seed_pool/ksplice` directory on **dom0**, use the `sftp host03` command to connect to **host03** as `root`.
    - Provide `root` password of `oracle` when prompted.
  - b. Use the `mput *` command to upload all files from the current directory on **dom0** to the `root` user's home directory on **host03**.

```
[dom0] # sftp host03
root@host03's password: oracle
sftp>
```

- b. Use the `mput *` command to upload all files from the current directory on **dom0** to the `root` user's home directory on **host03**.

```
sftp> mput *
Uploading 167283.sh to /root/167283.sh
...
Uploading uptrack-updates-3.8.13-16.2.1.el6uek.x86_64-20131219-
0.noarch.rpm to /root/uptrack-updates-3.8.13-
16.2.1.el6uek.x86_64-20131219-0.noarch.rpm
...
sftp>
```

- c. Use the `quit` command to exit sftp.

```
sftp> quit
```

## Practice 7-3: Installing the Ksplice Offline Client and Kernel Updates

### Overview

In this practice, you perform the following:

- Install the Ksplice Offline Client packages on **host03**.
- Install the kernel updates package on **host03**.
- View the effective kernel version after updates have been installed.
- List the active Oracle Ksplice updates in your running kernel.

### Assumptions

You are the `root` user on **host03**.

### Tasks

1. Ensure you have the correct version of the Ksplice Uptrack kernel update.
  - a. From **host03**, use the `uname -r` command to display the version of the running kernel.

```
uname -r
3.8.13-16.2.1.el6uek.x86_64
```

- b. Use the `ls` command to display the contents of the directory.

```
cd
ls
...
uptrack-updates-3.8.13-16.2.1.el6uek.x86_64-20131219-
0.noarch.rpm
```

- Notice the kernel version on **host03** is the same as the Ksplice Uptrack kernel package, `3.8.13-16.2.1.el6uek.x86_64`.

2. Install the Ksplice Offline Client packages.

- a. Use the `rpm -Uvh` command to install the `uptrack-offline` package.

```
rpm -Uvh uptrack-offline-1.2.15.offline-0.el6.noarch.rpm
error: Failed dependencies:
 uptrack-PyYAML is needed by uptrack-offline-
1.2.15.offline-0.el6.noarch
 uptrack-python-cjson is needed by uptrack-offline-
1.2.15.offline-0.el6.noarch
```

- Notice the install failed because two dependent packages are missing.

- b. Use the `rpm -Uvh` command to install the `uptrack-PyYAML` package.

```
rpm -Uvh uptrack-PyYAML-3.08-4.el6.x86_64.rpm
error: Failed dependencies:
 uptrack-libyaml >= 0.1.3-1 is needed by uptrack-
PyYAML...
```

- Notice the install failed because the `uptrack-libyaml` package is not installed.

- c. Use the `rpm -Uvh` command to install the `uptrack-libyaml` package.

```
rpm -Uvh uptrack-libyaml-0.1.3-1.el6.x86_64.rpm
Preparing... ## [100%]
1:uptrack-libyaml ## [100%]
```

- This package installed successfully with no dependency errors.

- d. Perform Task 2b to install the `uptrack-PyYAML` package.

```
rpm -Uvh uptrack-PyYAML-3.08-4.el6.x86_64.rpm
Preparing... ## [100%]
1:uptrack-PyYAML ## [100%]
```

- This package installed successfully with no dependency errors.

- e. Use the `rpm -Uvh` command to install the `uptrack-python-cjson` package.

```
rpm -Uvh uptrack-python-cjson-1.0.5-5.el6.x86_64.rpm
Preparing... ## [100%]
1:uptrack-python-cjson ## [100%]
```

- This package installed successfully with no dependency errors.

- f. Perform Task 2a to install the `uptrack-offline` package.

```
rpm -Uvh uptrack-offline-1.2.15.offline-0.el6.noarch.rpm
Preparing... ## [100%]
1:uptrack-offline ## [100%]
There are no existing modules on disk that need basename
migration.
```

- All Ksplice Offline Client packages are now installed.

3. Install the kernel updates package.

Use the `rpm -Uvh` command to install the `uptrack-updates` package.

```
rpm -Uvh uptrack-updates-3.8.13-16.2.1.el6uek.x86_64-20131219-
0.noarch.rpm
Preparing... ## [100%]
1:uptrack-updates-3.8.13-##### ## [100%]
The following steps will be taken:
Install [27mo8t1x] CVE-2013-4387: Memory corruption in IPv6 ...
Install [spxsnial] CVE-2013-0342: Denial of service in IPv6 ...
...
Installing [rz5tgzrd] Incorrect NumVFs reporting in PCI I/O ...
Your kernel is fully up to date.
Effective kernel version is 3.8.13-16.3.1.el6uek
```

- All kernel updates are installed.

4. Show effective kernel.

- a. Use the `uname -r` command to display the version of the running kernel.

```
uname -r
3.8.13-16.2.1.el6uek.x86_64
```

- Notice the kernel version remains at version `3.8.13-16.2.1.el6uek`.

- b. Use the `uptrack-uname -r` command to display the effective kernel version.

```
uptrack-uname -r
3.8.13-16.3.1.el6uek.x86_64
```

- Notice the effective kernel version is `3.8.13-16.3.1.el6uek`.

- c. Use the `uptrack-show` command to list the active Oracle Ksplice updates in your running kernel.

```
uptrack-show
Installed updates:
[27mo8t1x] CVE-2013-4387: Memory corruption in IPv6 UDP ...
[spxsnial] CVE-2013-0342: Denial of service in IPv6 privacy ...
...
[rz5tgzrd] Incorrect NumVFs reporting in PCI I/O Single Root ...

Effective kernel version is 3.8.13-16.3.1.el6uek
```

## **Practices for Lesson 8: Automating Tasks**

**Chapter 8**

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practices for Lesson 8: Automating Tasks

---

### Practices Overview

In these practices, you use and configure utilities to automate tasks.

## Practice 8-1: Automating Tasks

### Overview

In this practice, you use the `crontab` utility and the `at` utility to automate tasks. You also prevent the `oracle` user from being able to use the `at` utility.

### Assumptions

You are the `root` user on the **host03** VM.

### Tasks

1. Create a `crontab` for the `root` user.

- a. Use the `crontab -l` command to list the contents of your `crontab` file.

```
crontab -l
no crontab for root
```

- b. Use the `ls` command to view the contents of the `/var/spool/cron` directory.

```
ls /var/spool/cron
```

- Notice that there are no files in this directory.

- c. Use the `crontab -e` command to create a `crontab` job that runs the `ls` command every minute.

```
crontab -e
* * * * * ls
```

- The `crontab -e` command uses the `vi` editor. Save your changes and exit `vi`. The following messages will be displayed:

```
no crontab for root - using an empty one
crontab: installing new crontab
```

- d. Use the `ls` command to view the contents of the `/var/spool/cron` directory.

```
ls /var/spool/cron
root
You have new mail in /var/spool/mail/root
```

- Now there is a `root` file in the `/var/spool/cron` directory.
- You also have mail, because the output from cron jobs is sent to the user's mailbox.

- e. Use the `cat` command to display the contents of the `/var/spool/cron/root` file.

```
cat /var/spool/cron/root
* * * * * ls
```

- f. Use the `mail` command to view the results of your job.

```
mail
...
>N 1 Cron Daemon Tue Dec 13 07:17 23/726 "Cron <root@host03>
ls"
N 2 Cron Daemon Tue Dec 13 07:18 23/726 "Cron <root@host03>
ls"
N 3 Cron Daemon Tue Dec 13 07:19 23/726 "Cron <root@host03>
ls"
&
```

- To view details of mailbox entries, press the associated number and then press **Enter**.
- To redisplay the header, press **h** and then press **Enter**.
- To quit the mail program, press **q** and then press **Enter**.

- g. Press **q** and then **Enter** to quit the mail program.

```
& q
```

- h. Use the `crontab -r` to remove the crontab.

```
crontab -r
crontab -l
no crontab for root
```

- i. Use the `ls` command to view the contents of the `/var/spool/cron` directory.


```
ls /var/spool/cron
```

- Notice that the `root` file has been removed from this directory.

2. Create a crontab for the `oracle` user.

- In this task, you switch back and forth between the `root` user and the `oracle` user.

- a. Open another tab in your terminal window by pressing **Shift + Ctrl + T**, or right-click and select **Open Tab**.
  - Your terminal window now has a tab where you are logged in as `root` and a tab where you are logged in as the `oracle` user:



- b. Click the `oracle@host03` tab and confirm that you are logged in as the `oracle` user.

```
[oracle@host03 ~]$ whoami
oracle
```

- c. Use the `crontab -l` command to list the contents of your `crontab` file.

```
[oracle@host03 ~]$ crontab -l
no crontab for oracle
```

- d. Use the `crontab -e` command to create a `cron` job that runs the `echo "Hello World"` command every other minute.

```
[oracle@host03 ~]$ crontab -e
*/2 * * * * echo "Hello World"
```

- e. Use the `crontab -l` command to list the contents of your `crontab` file.

```
[oracle@host03 ~]$ crontab -l
*/2 * * * * echo "Hello World"
```

- f. Use the `mail` command to view the results of your job.

```
[oracle@host03 ~]$ mail
...
>N 1 Cron Daemon Tue Dec 13 08:40 21/728 "Cron
<oracle@host03>"
 N 2 Cron Daemon Tue Dec 13 08:42 23/726 "Cron
<oracle@host03>"
&
```

- g. To view details of mailbox entries, press the associated number and then press **Enter**.  
h. After viewing the details, press **q** and then **Enter** to quit.

```
& 1
...
Hello World
& q
```

- i. Use the `ls` command to view the contents of the `/var/spool/cron` directory.

```
[oracle@host03 ~]$ ls /var/spool/cron
ls: cannot open directory /var/spool/cron/: Permission denied
```

- Notice that only the `root` user has permission to view this directory.

- j. Click the `root@host03` tab and confirm that you are logged in as the `root` user.

```
[root@host03 ~]# whoami
root
```

- k. Use the `ls` command to view the contents of the `/var/spool/cron` directory.

```
[root@host03 ~]# ls /var/spool/cron
oracle
```

- Notice that there is an `oracle` file in the `/var/spool/cron` directory.

- l. Use the `cat` command to display the contents of the `oracle` user's crontab.

```
[root@host03 ~]# cat /var/spool/cron/oracle
*/2 * * * * echo "Hello World"
```

- m. Use `crontab -r` to remove the crontab for the `oracle` user.

```
[root@host03 ~]# crontab -u oracle -r
[root@host03 ~]# crontab -u oracle -l
no crontab for oracle
```

3. Use the `at` command to schedule a one-time task at a specific time.

- a. Click the **oracle@host03** tab and confirm that you are logged in as the `oracle` user.

```
[oracle@host03 ~]$ whoami
oracle
```

- b. Use `at` to run the `pwd` command in 1 minute.

```
[oracle@host03 ~]$ at now + 1 minute
at> pwd
at> <EOT>
```

- Press **Ctrl + D** for `<EOT>`.

- c. Use the `atq` command to list the pending job.

```
[oracle@host03 ~]$ atq
1 2013-11-15 11:19 a oracle
```

- d. After 1 minute has passed, run the `atq` command again.

```
[oracle@host03 ~]$ atq
```

- Notice that this time there are no pending jobs.

- e. Use the `mail` command to view the results of your job.

```
[oracle@host03 ~]$ mail
...
>N 2 Oracle Student Tue Dec 13 10:16 14/509 "Output from your
job"
&
```

- f. To view details of mailbox entries, press the associated number and then press **Enter**.

- g. After viewing the details, press `q` and then **Enter** to quit.

```
& 2
...
/home/oracle
& q
```

4. Restrict usage of the `at` command.

- In this task, you also switch back and forth between the `root` user and the `oracle` user.

- a. Click the **root@host03** tab and confirm that you are logged in as the `root` user.

```
[root@host03 ~]# whoami
root
```

- b. Use the `vi` editor to add user `oracle` to the `/etc/at.deny` file.

```
[root@host03 ~]# vi /etc/at.deny
oracle
```

- c. Click the **oracle@host03** tab and confirm that you are logged in as the **oracle** user.

```
[oracle@host03 ~]$ whoami
oracle
```

- d. Attempt to use **at** to run a command in 1 minute.

```
[oracle@host03 ~]$ at now + 1 minute
You do not have permission to use at.
```

5. In the terminal window, click the **X** in the **oracle@host03** tab to close this tab.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## **Practices for Lesson 9: Kernel Module Configuration**

**Chapter 9**

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## **Practices for Lesson 9: Kernel Module Configuration**

---

### **Practices Overview**

In these practices, you work with loadable kernel modules.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 9-1: Using Loadable Kernel Modules

---

### Overview

In this practice, you perform the following tasks:

- Use kernel module utilities to list modules.
- Get detailed information about modules.
- Load and unload kernel modules.
- Explore the kernel module configuration directory.

### Assumptions

- You are the `root` user on the **host03** VM.
- Sample output is provided. Kernel modules shown might differ from actual output.

### Tasks

1. List kernel modules.
  - a. Use the `lsmod` command to list kernel modules currently loaded into the kernel.
    - You might see numerous Ksplice modules as a result of completing Practice 7-3.
    - These Ksplice modules are not included in the sample output.

```
lsmod
Module Size Used by
nls_utf8 1421 1
fuse 78015 0
autofs4 33910 3
...
nf_conntrack 84594 3 nf_conntrack_ipv4, ...
...
```

- The “Size” column displays the amount of memory the module uses.
- The “Used by” column gives the total number of processes that are using the module and other modules that it depends on, followed by a list of those dependent modules.
- In the sample output, `nf_conntrack` depends on modules `nf_conntrack_ipv4`, `nf_conntrack_ipv6`, and other listed modules being loaded before `nf_conntrack` loads.

- b. Use the `modinfo` command to list detailed information about a specific kernel module (for example, `nls_utf8`).

```
modinfo nls_utf8
filename: /lib/modules/3.8.13-
16.2.1.el6uek.x86_64/kernel/fs/nls/nls_utf8.ko
license: Dual BSD/GPL
srcversion:
...
```

- Notice that the kernel modules are loaded from the `/lib/modules/<kernel_version>/kernel` directory.
- c. Use the `uname -r` command to display the kernel version.

```
uname -r
3.8.13-16.2.1.el6uek.x86_64
```

- d. List the kernel modules for the kernel version.

```
ls -R /lib/modules/`uname -r`/kernel
arch
crypto
Documentation
drivers
fs
kernel
lib
mm
net
sound
...
```

- The actual kernel modules have a `.ko` (kernel object) extension.

## 2. Load kernel modules.

- a. Determine if the `nfs` kernel module is currently loaded.

```
lsmod | grep nfs
```

- In this example, `nfs` is not loaded.
- b. Use the `modprobe` command to load the `nfs` kernel module.
- Sample output is provided. Yours might be different with Ksplice kernel updates applied.

```
modprobe -v nfs
...
insmod /lib/modules/3.8.13-
16.2.1.el6uek.x86_64/kernel/net/sunrpc/sunrpc.ko
insmod /lib/modules/3.8.13-
16.2.1.el6uek.x86_64/kernel/fs/lockd/lockd.ko
insmod /lib/modules/3.8.13-
16.2.1.el6uek.x86_64/kernel/fs/fscache/fscache.ko
insmod /lib/modules/3.8.13-
16.2.1.el6uek.x86_64/kernel/fs/nfs/nfs.ko
```

- The preceding example includes the `-v` (verbose) option.
- Notice that the dependent modules are loaded prior to loading the `nfs` module.
- Also notice that `modprobe` uses the `insmod` command to load the modules.

c. Confirm that the **nfs** module is loaded.

- Sample output is provided. Yours might be different with Ksplice kernel updates applied.

```
lsmod | grep nfs
nfs 184416 0
fscache 52962 1 nfs
lockd 84481 1 nfs
sunrpc 267672 3 nfs, lockd
```

3. List module dependencies.

- Kernel module dependencies are listed in `/lib/modules/<kernel_version>/modules.dep`.

Use the `grep` command to list the kernel module dependencies for **nfs**.

```
grep nfs.ko /lib/modules/`uname -r`/modules.dep
kernel/fs/nfs/nfs.ko: kernel/fs/fscache/fscache.ko
kernel/fs/lockd/lockd.ko kernel/net/sunrpc/sunrpc.ko
...
```

- Notice that the dependencies listed in the `modules.dep` file for **nfs** correspond with the modules loaded when running the `modprobe nfs` command.

4. Unload kernel modules.

Use the `modprobe -r` command to unload the **nfs** kernel module.

```
modprobe -rv nfs
rmmod /lib/modules/3.8.13-
16.2.1.el6uek.x86_64/kernel/fs/nfs/nfs.ko
rmmod /lib/modules/3.8.13-
16.2.1.el6uek.x86_64/kernel/fs/fscache/fscache.ko
rmmod /lib/modules/3.8.13-
16.2.1.el6uek.x86_64/kernel/fs/lockd/lockd.ko
...
rmmod /lib/modules/3.8.13-
16.2.1.el6uek.x86_64/kernel/net/sunrpc/sunrpc.ko
```

- The preceding example includes the `-v` (verbose) option.
- Notice that `modprobe -r` uses the `rmmod` command to remove the modules.
- Also notice that the dependent modules are removed when possible.
- In some cases, a module cannot be removed because it is a dependent module for other kernel modules.

5. Explore the `/etc/sysconfig/modules` directory.

- This directory contains files that specify kernel modules to be loaded at boot time.

- a. List the contents of the /etc/sysconfig/modules directory.

```
ls -l /etc/sysconfig/modules
...
-rwxr-xr-x. bluez-uinput.modules
```

- Files in this directory must end in .modules.

- b. View the contents of the .modules file.

```
cat /etc/sysconfig/modules/bluez-uinput.modules
#!/bin/sh
if [! -c /dev/input/uinput]; then
 exec /sbin/modprobe uinput
fi
```

- Notice that the .modules files are executable shell scripts.

6. Explore the /etc/modprobe.d directory.

- Files in this directory end with .conf and are used for specifying the options that are to be used with kernel modules.
- a. List the contents of the /etc/modprobe.d directory.
- Sample output is shown, your output might be different.

```
ls -l /etc/modprobe.d
-rw-r--r--. anaconda.conf
-rw-r--r--. blacklist.conf
-rw-r--r--. dist-alsa.conf
-rw-r--r--. dist.conf
-rw-r--r--. dist-oss.conf
-rw-r--r--. openfwf.conf
```

- b. Use the grep command to search for the occurrence of **alias** in the dist.conf file.

```
grep alias /etc/modprobe.d/dist.conf
...
alias nfs4 nfs
alias rpc_pipefs sunrpc
...
```

- The **alias** command is used to create alternative names for kernel modules.
  - The alias name is the first argument; kernel module name is the second argument.
- c. Load the **nfs** module by using the **nfs4** alias name.

```
modprobe nfs4
lsmod | grep nfs
nfs 184416 0
fscache 52962 1 nfs
lockd 84481 1 nfs
sunrpc 267672 3 nfs, lockd
```

- d. Unload the **nfs** module by using the **nfs4** alias name.

```
modprobe -r nfs4
lsmod | grep nfs
```

- e. Use the **grep** command to search for the occurrence of **options** in the **openfwf.conf** file.

```
grep options /etc/modprobe.d/openfwf.conf
options b43 nohwcrypt=1 qos=0
```

- Notice the **nohwcrypt** and **qos** options for the **b43** kernel module.

- f. Load the **b43** kernel module.

```
modprobe b43
```

- g. Use the **modinfo** command to list detailed information about the **b43** kernel module.

```
modinfo b43
...
parm: nohwcrypt:Disable hardware encryption. (int)
...
parm: qos:Enable QOS support (default on) (int)
...
```

- Notice the **nohwcrypt** and **qos** parameters for the **b43** kernel module.

- h. Unload the **b43** kernel module.

```
modprobe -r b43
```

- i. Use the **grep** command to search for the occurrence of **install** in the **dist.conf** file.

```
grep install /etc/modprobe.d/dist.conf
...
install net-pf-3 /bin/true
...
```

- For the **net-pf-3** kernel module, the **/bin/true** command runs instead of loading the kernel module when using **modprobe**.

- j. Determine if the **net-pf-3** kernel module is loaded Attempt to load the **net-pf-3** kernel module. Then determine if the module loaded.

```
lsmod | grep net-pf-3
modprobe net-pf-3
lsmod | grep net-pf-3
```

- Notice that the **net-pf-3** kernel module did not load because of the setting in the **/etc/modprobe.d/dist.conf** file.

## **Practices for Lesson 10: User and Group Administration**

**Chapter 10**

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practices for Lesson 10: User and Group Administration

---

### Practices Overview

In these practices, you:

- Create user and group accounts
- Understand the benefits of user private groups
- Configure password aging
- Use the User Manager GUI
- Restrict the use of the `su` command
- Allow the use of the `sudo` command

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 10-1: User Account Administration

---

### Overview

In this practice, you use command-line utilities to create new user accounts, view files that are updated when adding a new user, modify a user account, set a password for the new user, and log in as the new user.

### Assumptions

You are the `root` user on `host03` VM.

### Tasks

1. Add a user.

- a. Use the command-line utility to add `student1` user.

```
useradd student1
```

- b. Use the `cat` command to view the new `student1` entry in the `/etc/passwd` file.

```
cat /etc/passwd
...
oracle:x:500:500:Oracle Student:/home/oracle:/bin/bash
student1:x:501:501::/home/student1:/bin/bash
```

- Notice that the new user's UID and GID incremented by one.
- Notice that a home directory was created for the new user (`/home/student1`).
- Notice that the default shell for the new user is `bash` (`/bin/bash`).

- c. Use the `ls` command to verify that the new user's home directory was created.

```
ls /home
lost+found oracle student1
```

- A home directory was created because `CREATE_HOME` in `/etc/login.defs` is set to yes.

- d. View the value of `CREATE_HOME` in `/etc/login.defs`.

```
grep CREATE_HOME /etc/login.defs
CREATE_HOME yes
```

- e. View the default settings for a new user, stored in `/etc/default/useradd`.

```
cat /etc/default/useradd
GROUP=100
HOME=/home
INACTIVE=-1
EXPIRE=
SHELL=/bin/bash
SKEL=/etc/skel
CREATE_MAIL_SPOOL=yes
```

- Notice that the `SKEL` directive is set to `/etc/skel`.

- f. View the contents of /etc/skel.

```
ls -la /etc/skel
-rw-r--r--. .bash_logout
-rw-r--r--. .bash_profile
-rw-r--r--. .bashrc
drwxr-xr-x. .gnome2
drwxr-xr-x. .mozilla
```

- g. View the contents of the new user's home directory.

```
ls -la /home/student1
-rw-r--r--. .bash_logout
-rw-r--r--. .bash_profile
-rw-r--r--. .bashrc
drwxr-xr-x. .gnome2
drwxr-xr-x. .mozilla
```

- Notice that the contents of **SKEL** (/etc/skel) are copied to the new user's home directory.

- h. View the new student1 entry in the /etc/group file.

```
cat /etc/group
...
oracle:x:500:
student1:x:501:
```

- Because Oracle Linux uses a user private group (UPG) scheme, a new private group (student1, GID=501) was created when the user student1 was created.

- i. View the new student1 entry in the /etc/shadow file.

```
cat /etc/shadow
...
oracle:6....:16048:0:99999:7:::
student1:!:16049:0:99999:7:::
```

- j. View the new student1 entry in the /etc/gshadow file.

```
cat /etc/gshadow
...
oracle:!::
student1:!::
```

- k. Add a new user with the following characteristics:

- Username=student2
- UID=555
- GECOS information="Oracle Student2"
- Default shell=/bin/sh (Bourne shell)

```
useradd -u 555 -c "Oracle Student2" -s /bin/sh student2
```

I. View the new student2 entry in /etc/passwd.

```
tail -2 /etc/passwd
student1:x:501:501::/home/student1:/bin/bash
student2:x:555:555:Oracle Student2:/home/student2:/bin/sh
```

- Notice that the UID and GID are 555.
- Notice the GECOS information.
- Notice the default shell.

2. Modify a user account.

- a. Use the usermod command to modify GECOS information for student1 user as follows:

```
usermod -c "Oracle Student1" student1
```

- b. View the student1 entry in the /etc/passwd file.

```
grep student1 /etc/passwd
student1:x:501:501:Oracle Student1:/home/student1:/bin/bash
```

3. Assign a password to the new user.

- a. Use the cat command to view the /etc/shadow file.

```
cat /etc/shadow
...
oracle:6...:16048:0:99999:7:::
student1:!:16049:0:99999:7:::
student2:!:16049:0:99999:7:::
```

- Notice the “! !” in the student1 and student2 records, indicating that no password has been assigned (and that the accounts are locked).
- b. Use the passwd command to create a password (of password) for the student1 user.

```
passwd student1
Changing password for user student1.
New password: password
BAD PASSWORD: it is based on a dictionary word
Retype new password: password
passwd: all authentication tokens updated successfully.
```

- Ignore the “BAD PASSWORD” warning, continuing to use password as the password.


- c. View the /etc/shadow file.

```
cat /etc/shadow
...
oracle:6...:16048:0:99999:7:::
student1:6...: 16049:0:99999:7:::
student2:!: 16049:0:99999:7:::
```


- Now the “! !” for student1 has been replaced with a hashed password value.
- Notice that the student2 account is still locked because a password has not been assigned.

4. Log in as new user.


- a. Log out as user oracle by selecting Log out oracle from the System menu, as shown below.



- b. Click **Log Out** in the pop-up, as shown below.




- c. Attempt to log in as student2 by selecting Oracle Student2 from the menu, as shown below.




- You are prompted for a password, but because no password has been assigned to student2, the account is locked.
- Regardless of what you enter for a password, the following screen appears, and you cannot log in as student2.

- d. Log in as student1 by selecting Oracle Student1 from the login menu.
- e. Enter password when prompted for the password, as shown below.



- f. Click **Log In** to successfully log in as student1.

- g. Right-click the desktop to display the pop-up menu as shown below.



- h. Select **Open in Terminal** to open a terminal window.  
i. From the terminal window, use the `whoami` command to verify that you are logged in as student1.

```
$ whoami
student1
```

- j. Use the `su -` command to become the root. The password is oracle.

```
$ su -
Password: oracle
whoami
root
```

## Practice 10-2: Group Account Administration

### Overview

In this practice, you create a new group account and add users to this new group.

### Tasks

1. Add a group.

- Use the `groupadd` command without any options to add the `students` group.

```
groupadd students
```

- View the new `students` entry in the `/etc/group` file.

```
cat /etc/group
...
oracle:x:500:
student1:x:501:
student2:x:555:
students:x:556:
```

- Notice that the GID for the new group is incremented by one.

2. Add users to the new group.

- Use the `usermod` command to add users `student1` and `student2` to the `students` group.

```
usermod -aG 556 student1
tail -1 /etc/group
students:x:556:student1
usermod -aG 556 student2
tail -1 /etc/group
students:x:556:student1,student2
```

- Notice that both `student1` and `student2` have a secondary group membership in the `students` group.

- View the primary group membership for `student1` and `student2`.

```
grep student /etc/passwd
student1:x:501:501:Oracle Student1:/home/student1:/bin/bash
student2:x:555:555:Oracle Student2:/home/student2:/bin/sh
```

- Notice that `student1` primary group is still 501.
- Notice that `student2` primary group is still 555.

## Practice 10-3: User Private Groups

---

### Overview

In this practice, you use the UPG scheme to give different users write access to files in a single directory.

### Tasks

- Set up a shared /students directory.

- Use the `mkdir` command to create the `/students` directory.

```
mkdir /students
```

- Use the `chgrp` command to change the group for this `/students` directory to the `students` group.

```
ls -ld /students
drwxr-xr-x. root root /students
chgrp -R students /students
ls -ld /students
drwxr-xr-x. root students /students
```

- Notice that the group was set to `root` before issuing the `chgrp` command.
- The `-R` option (recursive) sets the new group for files and directories within `/students`.

- Use the `chmod` command to set the `setgid` bit on the directory, and give write permissions to the group.

```
chmod -R 2775 /students
ls -ld /students
drwxrwsr-x. root students /students
```

- Notice the new group permissions on the `/students` directory.

- Create files in the `/students` directory as the `oracle` user.

- Use the `usermod` command to add the `oracle` user to the `students` group.

```
usermod -aG students oracle
```

- Use the `su` command to become the `oracle` user.

```
su - oracle
$ whoami
oracle
```

- Use the `groups` command to display group membership.

```
$ groups
oracle students
```

- Notice that the `oracle` user belongs to two groups, `oracle` and `students`.

- d. Use the `cd` command to change to the `/students` directory.

```
$ cd /students
```

- e. Use the `touch` command to create a file in the `/students` directory.

```
$ touch oracle_file
```

- f. Use the `ls` command to display the permissions and ownership of the new file.

```
$ ls -l oracle_file
rw-rw-r-- oracle students oracle_file
```

- Notice that the permissions are read-write for the `students` group.

3. Create and edit files within the shared directory as a different user.

- a. Use the `su` command to become the `student1` user. The password is `password`.

```
$ su - student1
Password: password
$ whoami
student1
```

- b. Use the `groups` command to display group membership.

```
$ groups
student1 students
```

- Notice that the `student1` user belongs to two groups, `student1` and `students`.

- c. Use the `cd` command to change to the `/students` directory.

```
$ cd /students
```

- d. Use the `touch` command to create a file in the `/students` directory.

```
$ touch student1_file
```

- e. Use the `ls` command to display the permissions and ownership of the files.

```
$ ls -l
rw-rw-r-- oracle students oracle_file
rw-rw-r-- student1 students student1_file
```

- Notice that the permissions are read-write for the `students` group.

- f. Use the `touch` command to update the time stamp on the `oracle_file`.

```
$ touch oracle_file
```

- Updating the time stamp implies write permissions on the file.

- g. Use the `exit` command twice to log out as the `student1` user, and then as the `oracle` user.

```
$ exit
$ whoami
oracle
$ exit
whoami
root
```

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 10-4: Password Aging

### Overview

In this practice, you modify the password aging parameters for a user.

### Tasks

1. View the password aging information.

- a. List password aging information in /etc/shadow for user student1.

```
grep student1 /etc/shadow
student1:6...:16049:0:99999:7:::
```

- This example gives 16049 days since the password changed (counted in days since Jan 1, 1970). This will differ depending on when you created the account.
- 0 is the number of days that need to pass before the password must be changed by the user (0 means **never**).
- 99999 is the maximum number of days since the password changed that the password can be used. After this amount of days, the password needs to be changed by the user.
- 7 is the number of days before expire date that the user is warned about the pending password change policy. If the password is not changed after this number of days, the user account is locked.

- b. Use the chage command to view password aging information for user student1.

```
chage -l student1
Last password change : Dec 10, 2013
Password expires : never
Password inactive : never
Account expires : never
Minimum number of days between password change : 0
Maximum number of days between password change : 99999
Number of days or warning before password expires: 7
```

2. Modify the password aging parameters.

a. Make the following changes to password aging for user student1:

- Change the minimum password age value to 14.
- Change the maximum password age value to 30.

```
chage student1
Changing the aging information for student1
Enter the new value, or press ENTER for the default
 Minimum Password Age [0]: 14
 Maximum Password Age [99999]: 30
 Last Password Change (YYYY-MM-DD) [2013-12-10]: ENTER
 Password Expiration Warning [7]: ENTER
 Password Inactive [-1]: ENTER
 Account Expiration Date (YYYY-MM-DD) [1969-12-31]: ENTER
```

- Making these changes means that in 14 days, the user has 30 days to change his or her password.

b. View the results of the change for student1 in /etc/shadow.

```
grep student1 /etc/shadow
student1:6...:16049:14:30:7:::
```

c. Use the chage command to view password aging information for user student1.

```
chage -l student1
Last password change : Dec 10, 2013
Password expires : Jan 09, 2014
Password inactive : never
Account expires : never
Minimum number of days between password change : 14
Maximum number of days between password change : 30
Number of days or warning before password expires: 7
```

## Practice 10-5: Using the User Manager GUI

### Overview


In this practice, you use the User Manager GUI to add a user, and add and delete a group account.

### Tasks


1. Add a new user by using the User Manager GUI.

- a. Use the `system-config-users` command to display the User Manager Tool GUI.

```
system-config-users
```




- b. Click **Add User** and provide the new user information, as shown below.
- Use password for the **Password** and **Confirm Password** fields.



- c. Click **OK**.
- d. Click **Yes** to use the weak password.
- The updated list of user accounts is displayed.

2. Add a new group.


- a. Click the **Groups** tab to display the list of group accounts, as shown below.



- b. Click **Add Group** and provide a **Group Name** of your choice when prompted.  
c. Click **OK** and notice the updated group list.

3. Delete a group


- Select the newly added group and click the **Delete** button.
- Click **Yes** to confirm the delete, as shown below.




- Select the **student1** group and click the **Delete** button.
  - Notice that you cannot delete this group, because the user **student1** still exists.
- Exit the User Manager tool.

Select **Quit** from the **File** menu.


5. Log out as user student1.
  - a. Log out as user student1 by selecting **Log out student1...** from the **System** menu as shown below.



- b. Click **Log Out** in the pop-up as shown below.



6. Log in as user oracle.
  - a. Log in as the oracle user by selecting Oracle Student from the menu as shown below.



- b. Enter the password oracle when prompted for the password.
- c. Right-click on the desktop and select **Open in Terminal** from the pop-up menu.
- d. From the command prompt in the terminal window, use the su - command to become the root user. Password is oracle.

```
$ su -
Password: oracle
whoami
root
```

## Practice 10-6: Restricting the Use of the su Command

---

### Overview

In this practice, you limit the use of the `su` command to members of the `wheel` group.

### Tasks

- Add the `oracle` user to the `wheel` group.

- Use the `grep` command to display the `wheel` entry in the `/etc/group` file.

```
grep wheel /etc/group
wheel:x:10:
```

- Use the `usermod` command to add the `oracle` user to the `wheel` group. Re-run the `grep` command in Task 1a to verify the modification occurred.

```
usermod -aG wheel oracle
grep wheel /etc/group
wheel:x:10:oracle
```

- Notice the `oracle` user is now a member of the `wheel` group.

- Restrict the use of the `su` command to members of the `wheel` group.

- The procedure shown below involves the modification of a Pluggable Authentication Module (PAM) configuration file. PAM is discussed in Lesson 16.

- Use the `vi` editor to uncomment (remove the `#` sign) the following entry in the `/etc/pam.d/su` file:

```
vi /etc/pam.d/su
...
#auth required pam_wheel.so use_uid (old entry)
auth required pam_wheel.so use_uid (new entry)
...
```

- With this entry uncommented, only members of the `wheel` group can use the `su` command (with the exception of the `root` user).

- Verify that only members of the `wheel` group can use the `su` command.

- As the `root` user, use the `su - student1` command to become the `student1` user.

```
su - student1
```

- The `su` command is successful because you are the `root` user.

- b. Use the `whoami` command to confirm you are the `student1` user. Use the `su` command to become the `root` user. The `root` user password is `oracle`.

```
$ whoami
student1
$ su -
Password: oracle
su: incorrect password
```

- The `su` command fails with an “incorrect password” message, but the real reason the `su` command fails is because your system is configured to allow only members of the `wheel` group to use the `su` command.

- c. Use the `exit` command to log out as the `student1` user. Use the `whoami` command to confirm you are the `root` user. Use the `exit` command again, and then `whoami` to verify that you are the `oracle` user.

```
$ exit
logout
whoami
root
exit
logout
$ whoami
oracle
```

- d. Use the `su` command to become the `root` user. The `root` user password is `oracle`.

```
$ su -
Password: oracle
whoami
root
```

- The `su` command succeeded this time because the `oracle` user is a member of the `wheel` group.

4. Remove restriction of the `su` command to members of the `wheel` group.

Use the `vi` editor to comment out (insert the `#` sign) the following entry in the `/etc/pam.d/su` file:

```
vi /etc/pam.d/su
...
auth required pam_wheel.so use_uid (old entry)
#auth required pam_wheel.so use_uid (new entry)
...
```

- With this entry commented out, all users can use the `su` command.

## Practice 10-7: Allowing the Use of the sudo Command

---

### Overview

In this practice you authorize the `oracle` user to use `sudo` to run administrative commands.

### Tasks

1. Attempt to add a new user without `root` privileges.
  - a. Use the `exit` command to log out as the `root` user. Then run the `whoami` command to verify you are the `oracle` user.

```
exit
logout
$ whoami
oracle
```

- b. Use the `useradd` command to add `new_user` user.

```
$ useradd new_user
bash: /usr/sbin/useradd: Permission denied
```

- Notice you do not have permission to add a new user.
- c. Insert the `sudo` command before the previous `useradd` command (in Task 1b) to add `new_user` user. Enter `oracle` for password when prompted.

```
$ sudo useradd new_user
```

We trust you have received the usual lecture from the local System Administrator. It usually boils down to these three things:

- #1) Respect the privacy of others.
- #2) Think before you type.
- #3) With great power comes great responsibility.

```
[sudo] password for oracle: oracle
oracle is not in the sudoers file. This incident will be reported.
```

- Notice that the command fails because the `oracle` user is not in the `sudoers` file.
- Also notice the incident is reported (in the `/var/log/secure` file).
- d. Use the `su -` command to become the `root` user. The password is `oracle`.

```
$ su -
Password: oracle
whoami
root
```

- e. Use the `tail` command to view the end of the `/var/log/secure` file.

```
tail /var/log/secure
...
<date_time> host03 sudo: oracle : user NOT in sudoers ...
...
```

- Notice the “`sudo`” entry in the `/var/log/secure` file.

2. Add the `oracle` user to the `/etc/sudoers` file.

- a. Use the `visudo` command to edit the `/etc/sudoers` file.

- This command opens the `/etc/sudoers` file with the `vi` editor but also locks the file.

```
visudo
```

- b. Add the following line to `/etc/sudoers`, which allows the `oracle` user to use `sudo` to run administrative commands. After adding the line, save the file and exit `visudo`.

```
oracle ALL=(ALL) ALL
```

3. Add a new user by using `sudo`.

- a. Use the `exit` command to log out as the `root` user, then run the `whoami` command to verify you are the `oracle` user.

```
exit
logout
$ whoami
oracle
```

- b. Use the `sudo useradd` command to add `new_user` user. Enter `oracle` for password when prompted.

```
$ sudo useradd new_user
[sudo] password for oracle:
```

- No error message implies the command was successful.

- c. Use the `grep` command to search for `new_user` in the `/etc/passwd` file.

```
$ grep new_user /etc/passwd
new_user:x:558:558::/home/new_user:/bin/bash
```

4. Run another command, which requires administrative privileges.

- a. Use the `service` command to restart the `atd` service.

```
$ service atd restart
rm: cannot remove 'var/run/atd.pid': Permission denied [FAILED]
...
```

- Notice the command failed with “Permission denied”

- b. Use sudo and run the service command again, to restart the atd service.

```
$ sudo service atd restart
Stopping atd: [OK]
Starting atd: [OK]
```

- Notice the command succeeded this time.
- Also notice you were not prompted to enter the oracle user's password.
- When a user has been authenticated, the user may then use sudo without a password for a short period of time (the default timeout is 5 minutes).
- Do not perform this, but to override the default 5 minute timeout, set the "timestamp\_timeout" value in /etc/sudoers to change the default.

- c. In preparation for the next practice, use the su - command to become the root user. The password is oracle.

```
$ su -
Password: oracle
whoami
Root
```

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## **Practices for Lesson 11: File Systems**

**Chapter 11**

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practices for Lesson 11: File Systems

---

### Practices Overview

In these practices, you:

- Display the partition table, list the mounted file systems, and display the swap space configured during installation
- Partition disk devices, and create and mount file systems on the partitions
- Increase the amount of swap space by creating a swap file

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 11-1: Listing the Current Disk Partitions

---

### Overview


In this practice, you verify the selections made during installation regarding disk partitioning, mount points for file systems, and swap space.

### Assumptions

You are the root user on **host03** VM.

### Tasks

- Relate the partition table to selections made during installation.
  - Recall that three virtual disk images were created before initiating the installation:
    - A 12-GB disk image (**system.img**) for the operating system
    - A 5-GB disk image (**u01.img**) for the various storage administration practices
    - A 5-GB disk image (**u02.img**) for the various storage administration practices
  - These three Storage Devices were presented during the installation, as shown in the following:




Use the `fdisk` command to display the partition table.

```
fdisk -l | grep /dev
Disk /dev/xvda: 12.9 GB, 12884901888 bytes
/dev/xvda1 * 1 64 512000 83 Linux
/dev/xvda2 64 829 6144000 83 Linux
/dev/xvda3 829 1084 2048000 83 Linux
/dev/xvda4 1084 1567 3877888 5 Extended
/dev/xvda5 1084 1567 3876864 82 Linux swap / Solaris
Disk /dev/xvdb: 5368 MB, 5368709120 bytes
Disk /dev/xvdd: 5368 MB, 5368709120 bytes
```

- This lists the three storage devices:
  - `/dev/xvda`, approximately 12 GB in size
  - `/dev/xvdb`, approximately 5 GB in size
  - `/dev/xvdd`, approximately 5 GB in size
- This indicates that the first partition on `xvda` contains **Boot** files (marked with \*).
- This also indicates that the fourth partition on `xvda` is an **Extended** partition.
  - A maximum of four primary partitions can be placed on any hard disk.
  - One of the four partitions may be designated as an Extended partition.
  - This Extended partition may then be subdivided into multiple logical partitions.
  - The `/dev/xvda5` is a logical partition as shown as follows:

2. Relate the mounted partitions to selections made during installation.
- Recall the final partition table created during the installation, shown in the following:



Use the `df -h` command to list the mounted partitions.

```
df -h
Filesystem Size Used Avail Use% Mounted on
/dev/xvda2 5.7G 3.2G 2.2G 60% /
tmpfs 1002M 384K 1001M 1% /dev/shm
/dev/xvda1 477M 55M 397M 13% /boot
/dev/xvda3 1.9G 5.1M 1.8G 1% /home
/dev/sr0 3.7G 3.7G 0 100% /media/OL6.5_x86_64 ...
```

- The first partition on hard drive xvda (**xvda1**) contains **boot** files, mounted on `/boot`.
- The second partition on hard drive xvda (**xvda2**) is the **root** file system, mounted on `/`.
- The third partition on hard drive xvda (**xvda3**) is for user **home** directories, mounted on `/home`.

3. Relate the swap space to selections made during installation.

Use the **swapon** command to display the swap space.

| # <b>swapon -s</b> |           |         |      |          |
|--------------------|-----------|---------|------|----------|
| Filename           | Type      | Size    | Used | Priority |
| /dev/xvda5         | partition | 3876860 | 0    | -1       |

- The fifth partition (the first logical partition) on hard drive xvda (**xvda5**) is **swap** space.
- The **/proc/swaps** file provides the same information.

| # <b>cat /proc/swaps</b> |           |         |      |          |
|--------------------------|-----------|---------|------|----------|
| Filename                 | Type      | Size    | Used | Priority |
| /dev/xvda5               | partition | 3876860 | 0    | -1       |

## Practice 11-2: Partitioning a Storage Device

---

### Overview

In this practice, you create a partition by using the `fdisk` utility, and create a second partition by using the `parted` utility. Recall from the lecture that if the file system is greater than 2 TB, you cannot use `fdisk` and you must use `parted`.

### Assumptions

You are the root user on **host03** VM.

### Tasks

1. Partition a storage device by using `fdisk`.

- a. Use the `fdisk` command to partition `/dev/xvdb`.

```
fdisk /dev/xvdb
```

WARNING: DOS-compatible mode is deprecated. It's strongly recommended to switch off the mode (command 'c') and change display units to sectors (command 'u').

Command (m for help) :

- DOS does not allow a partition to start (or end) mid-cylinder. Toggling those two items ('c' and 'u') allows partition start and end points to be entered in a sector-exact way, rather than by cylinder.

- b. Perform the recommended actions (switch off DOS-compatible mode and change display units to sectors).

Command (m for help) : c

DOS Compatibility flag is not set

Command (m for help) : u

Changing display/entry units to sectors

- c. Display the `fdisk` menu.

Command (m for help) : m

- a toggle a bootable flag
- b edit bsd disklabel
- c toggle the dos compatibility flag
- d delete a partition
- l list known partition types
- m print this menu
- n add a new partition
- o create a new empty DOS partition table
- p print the partition table
- q quit without saving changes
- s create a new empty sun disklabel
- t change a partition's system id
- u change display/entry units

```
v verify the partition table
w write table to disk and exit
x extra functionality (experts only)
```

- d. Add a new **primary** partition, giving the partition number **1**.

```
Command (m for help): n
Command action
 e extended
 p primary partition (1-4)
p
Partition number (1-4): 1
```

- As was previously discussed, a maximum of four primary partitions can be placed on any hard disk.
- One of the four partitions may be designated as an extended partition. This partition may then be subdivided into multiple logical partitions.

- e. Continue adding the new partition, using the following parameters:

```
First sector (2048-10485759, default 2048): ENTER
Using default value 2048
Last sector, +sectors or +size{K,M,G} (2048-10485759, default
10485759): 2100000
```

- f. Display the new partition table.

```
Command (m for help): p
Disk /dev/xvdb: 5368 MB, 5368709120 bytes
...
Device Boot Start End Blocks Id System
/dev/xvdb1 2048 2100000 1048976+ 83 Linux
```

- g. Save the new partition table.

```
Command (m for help): w
The partition table has been altered!
Calling ioctl() to re-read partition table.
Syncing disks.
```

- h. Use the **fdisk** command to list the partition table on **/dev/xvdb**.

```
fdisk -l /dev/xvdb
Disk /dev/xvdb: 5368 MB, 5368709120 bytes
...
Device Boot Start End Blocks Id System
/dev/xvdb1 1 519 1048976+ 83 Linux
```

2. Partition a storage device by using `parted`.

- a. Use the `parted` command to partition `/dev/xvdd`.

```
parted /dev/xvdd
GNU Parted 2.1
Using /dev/xvdd
Welcome to GNU Parted! Type 'help' to view a list of commands.
(parted)
```

- b. Enter `help` to view a list of commands.

```
(parted) help
 align-check TYPE N check partition N for TYPE(min|opt)
 alignment
 check NUMBER do a simple check on the file system
 ...
```

- c. Enter `print` to print the partition table.

```
(parted) print
Model: Xen Virtual Block Device (xvd)
Disk /dev/xvdd: 5369MB
Sector size (logical/physical): 512B/512B
Partition Table: msdos
Number Start End Size Type Filesystem Flags

```

- Notice that the disk has not been partitioned.

- d. Get `help` for the `mkpart` command.

```
(parted) help mkpart
 mkpart PART-TYPE [FS-TYPE] START END make a partition
 PART-TYPE is one of: primary, logical, extended
 FS-TYPE is one of: ext4, ext3, ext2, fat32, fat16...
 START and END are disk locations, such as 4GB or 10%...
 'mkpart' makes a partition without creating a new file
 system on the partition. FS-TYPE may be specified to set
 an appropriate partition ID.
```

- e. Use the `mkpart` command to create a partition, using the following parameters:

```
(parted) mkpart
Partition type? primary/extended? primary
Filesystem type? [ext2] ENTER
Start? 0
End? 20%
Warning: The resulting partition is not properly aligned for
best performance. Ignore/Cancel? i
```

- f. Print the partition table.

```
(parted) print
Model: Xen Virtual Block Device (xvd)
Disk /dev/xvdd: 5369MB
Sector size (logical/physical): 512B/512B
Partition Table: msdos
Number Start End Size Type Filesystem Flags
 1 512B 1074MB 1074MB primary
```

- g. Save the partition table and exit.

```
(parted) quit
Information: You may need to update /etc/fstab
```

- h. Use the fdisk command to list the partition table on /dev/xvdd.

- Remember that Linux is case-sensitive. Use lowercase /dev in the command as shown.

```
fdisk -l /dev/xvdd
Disk /dev/xvdd: 5368 MB, 5368709120 bytes
...
Device Boot Start End Blocks Id System
/dev/xvdd1 1 1886 1048575+ 83 Linux
```

3. Display the major and minor numbers.

- a. Display the partitions in /proc/partitions.

```
cat /proc/partitions
major minor #blocks name
 202 0 12582912 xvda
 202 1 512000 xvda1
 202 2 6144000 xvda2
 202 3 2048000 xvda3
 202 4 1 xvda4
 202 5 3876864 xvda5
 202 16 5242880 xvdb
 202 17 1048976 xvdb1
 202 48 5242880 xvdd
 202 49 1048575 xvdd1
...
```

- The kernel uses the major and minor numbers to access a device.

- b. Use the `ls` command to list the major and minor numbers for devices in `/dev`.

```
ls -l /dev/xvd*
brw-rw----. 1 root disk 202, 0 Dec 22 05:19 /dev/xvda
brw-rw----. 1 root disk 202, 1 Dec 22 03:18 /dev/xvda1
brw-rw----. 1 root disk 202, 2 Dec 22 03:18 /dev/xvda2
brw-rw----. 1 root disk 202, 3 Dec 22 03:18 /dev/xvda3
brw-rw----. 1 root disk 202, 4 Dec 22 03:18 /dev/xvda4
brw-rw----. 1 root disk 202, 5 Dec 22 03:18 /dev/xvda5
brw-rw----. 1 root disk 202, 16 Dec 22 05:19 /dev/xvdb
brw-rw----. 1 root disk 202, 17 Dec 22 03:18 /dev/xvdb1
brw-rw----. 1 root disk 202, 48 Dec 22 06:22 /dev/xvdd
brw-rw----. 1 root disk 202, 49 Dec 22 10:25 /dev/xvdd1
```

## Practice 11-3: Creating ext File Systems

---

### Overview

In this practice, you create ext3 and ext4 file systems on the new partitions, mount the file systems, and update the file system mount table.

### Assumptions

You are the `root` user on **host03** VM.

### Tasks

1. Make an **ext3** file system on `/dev/xvdb1`.

- a. Use the `mkfs` command to make an **ext3** file system on `/dev/xvdb1`.

```
mkfs -t ext3 /dev/xvdb1
mke2fs 1.43-WIP (20-Jun-2013)
warning: 100 blocks unused.
Filesystem label=
OS type: Linux
...
Writing inode tables: done
Creating journal (8192 blocks): done
...
```

- b. Use the `blkid` command to display the attributes of the `/dev/xvdb1` block device.

  - UUID values are not shown throughout this practice.

```
blkid /dev/xvdb1
/dev/xvdb1: UUID="..." SEC_TYPE="ext2" TYPE="ext3"
```

2. Make an **ext4** file system on `/dev/xvdd1`.

- a. Use the `mkfs` command to make an **ext4** file system on `/dev/xvdd1` and assign a label of **Test**.

```
mkfs -t ext4 -L Test /dev/xvdd1
mke2fs 1.43-WIP (20-Jun-2013)
Filesystem label=Test
OS type: Linux
...
Writing inode tables: done
Creating journal (4096 blocks): done
...
```

- b. Display the attributes of the `/dev/xvdd1` block device.

```
blkid /dev/xvdd1
/dev/xvdd1: LABEL="Test" UUID="..." TYPE="ext4"
```

3. Set the file system label on /dev/xvdb1.

- a. Use the `e2label` command to set the file system label on /dev/xvdb1 to **Dev**.
- Remember that Linux is case-sensitive. Use uppercase **Dev** in the command as shown.

```
e2label /dev/xvdb1 Dev
```

- b. Display the attributes of all the block devices.

```
blkid
/dev/xvda1: UUID="..." TYPE="ext4"
/dev/xvda2: UUID="..." TYPE="ext4"
/dev/xvda3: UUID="..." TYPE="ext4"
/dev/xvda5: UUID="..." TYPE="swap"
...
/dev/xvdb1: UUID="..." SEC_TYPE="ext2" TYPE="ext3" LABEL="Dev"
/dev/xvdd1: LABEL="Test" UUID="..." TYPE="ext4"
```

- Notice that both /dev/xvdb1 and /dev/xvdd1 now have labels.

4. Mount the file systems.

- a. Use the `mkdir` command to create mount points.

```
mkdir /Test /Dev
```

- b. Use the `mount` command to mount /dev/xvdb1 on /Dev.

```
mount /dev/xvdb1 /Dev
```

- c. Mount /dev/xvdd1 on /Test.

```
mount /dev/xvdd1 /Test
```

- d. Use the `df` command to display the mounted file systems.

```
df -h
Filesystem Size Used Avail Use% Mounted on
/dev/xvda2 5.7G 3.2G 2.2G 60% /
tmpfs 1002M 384K 1001M 1% /dev/shm
/dev/xvda1 477M 55M 397M 13% /boot
/dev/xvda3 1.9G 5.1M 1.8G 1% /home
/dev/sr0 3.7G 3.7G 0 100% /media/OL6.5_x86_64 ...
/dev/xvdb1 1008M 34M 924M 4% /Dev
/dev/xvdd1 992M 1.3M 940M 1% /Test
```

- e. Use the `mount` command to display the mounted file systems.

```
mount
/dev/xvda2 on / type ext4 (rw)
...
/dev/xvda1 on /boot type ext4 (rw)
/dev/xvda3 on /home type ext4 (rw)
...
/dev/xvdb1 on /Dev type ext3 (rw)
/dev/xvdd1 on /Test type ext4 (rw)
```

- Notice that the file systems are mounted read-write (**rw**) by default.

- f. Display the mounts in `/proc/mounts`.

```
cat /proc/mounts
rootfs / rootfs rw 0 0
proc /proc proc rw,relatime 0 0
...
/dev/xvdb1 /Dev ext3 rw,seclabel,relatime,errors=continue...
/dev/xvdd1 /Test ext4 rw,seclabel,relatime,data=ordered 0 0
```

5. Update the file systems mount table.

Use the `vi` editor to add the following entries to `/etc/fstab`.

```
vi /etc/fstab
LABEL=Dev /Dev ext3 defaults 0 0
LABEL=Test /Test ext4 defaults 0 0
```

- After updating `/etc/fstab`, the new file systems will mount after a reboot.

## Practice 11-4: Increasing Swap Space

---

### Overview

In this practice, you increase the amount of swap space by creating, initializing, and enabling a swap file.

### Assumptions

You are the root user on **host03** VM.

### Tasks

1. Display the current amount of swap space.
  - a. Use the `swapon` command to display the current amount of swap space.

```
swapon -s
Filename Type Size Used Priority
/dev/xvda5 partition 3876860 0 -1
```

- b. Display the amount of swap space used by viewing the `/proc/meminfo` file.

```
grep -i swap /proc/meminfo
SwapCached: 0 kB
SwapTotal: 3876860 kB
SwapFree: 3876860 kB
```

- c. Use the `free` command to display the amount of swap space.

```
free
 total used free
...
Swap 3876860 0 3876860
```

2. Create and initialize a swap file.
  - a. Use the `dd` command to create a 1 GB swap file, `/swapfile`.

```
dd if=/dev/zero of=/swapfile bs=1024 count=1000000
1000000+1 records in
1000000+1 records in
1024000000 bytes (1.0 GB) copied, 7.28291 s, 141 MB/s
```

- b. Use the `mkswap` command to initialize the swap file.

```
mkswap /swapfile
mkswap: /swapfile: warning: don't erase bootbits sectors
on whole disk. Use -f to force.
Setting up swapspace version 1, size = 999996 KiB
no label, UUID=...
```

3. Enable swapping on the swap file.

- a. Use the `swapon` command to enable swapping on the swap file.

```
swapon /swapfile
```

- b. Display the updated amount of swap space (repeat step 1).

```
swapon -s
Filename Type Size Used Priority
/dev/xvda5 partition 3876860 0 -1
/swapfile file 999996 0 -2
grep -i swap /proc/meminfo
SwapCached: 0 kB
SwapTotal: 4876856 kB
SwapFree: 4876856 kB
free
 total used free
...
Swap 4876856 0 4876856
```

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## **Practices for Lesson 12: Storage Administration**

**Chapter 12**

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## **Practices for Lesson 12: Storage Administration**

---

### **Practices Overview**

In these practices, you create and work with logical volumes and RAID devices.

## Practice 12-1: Creating Linux LVM Partitions

---

### Overview

In this practice, you create new partitions and change the system ID to Linux LVM.

### Assumptions

You are the `root` user on **host03** VM.

### Tasks

1. Un-mount all file systems currently mounted on `/dev/xvdb`.

- a. Determine what is currently mounted on `/dev/xvdb`.

```
df -h | grep xvdb
/dev/xvdb1 1008M 34M 924M 4% /Dev
```

- b. Un-mount `/dev/xvdb1`.

```
umount /dev/xvdb1
```

2. Create a new partition on `/dev/xvdb`.

- a. Use the `fdisk` command to create a new **primary** partition on `/dev/xvdb` by using the following parameters:

```
fdisk /dev/xvdb
WARNING: DOS-compatible mode is deprecated. It's strongly
recommended to switch off the mode (command 'c') and change
display units to sectors (command 'u').
Command (m for help): n
Command action
 e extended
 p primary partition (1-4)
p
Partition number (1-4): 2
First cylinder (519-2590, default 519): ENTER
Using default value 519
Last cylinder, +cylinders or +size{K,M,G} (519-2590, default
2590): +1G
```

- b. Use the “t” command to change the system ID on the new partition.

```
Command (m for help): t
Partition number (1-4): 2
Hex code (type L to list codes): L
 0 Empty 24 NEC DOS 81 Minix / old Lin ...
 1 FAT12 39 Plan 9 82 Linux swap ...
 2 XENIX root 3c PartitionMagic 83 Linux ...
...
 8 AIX 4f QNX4.x 3rd part 8e Linux LVM ...
...
```

```
...
Hex code (type L to list codes) : 8e
Changed system type of partition 2 to 8e (Linux LVM)
```

- c. Print the new partition table.

```
Command (m for help) : p
Disk /dev/xvdb: 5368 MB, 5368709120 bytes
...
 Device Boot Start End Blocks Id System
/dev/xvdb1 1 519 1048976+ 83 Linux
/dev/xvdb2 519 1037 1048887+ 8e Linux LVM
```

- d. Save the new partition table.

```
Command (m for help) : w
The partition table has been altered!
Calling ioctl() to re-read partition table.
Syncing disks.
```

3. Unmount all file systems currently mounted on /dev/xvdd.

- a. Determine what is currently mounted on /dev/xvdd.

```
df -h | grep xvdd
/dev/xvdd1 1008M 18M 940M 2% /Test
```

- b. Unmount /dev/xvdd1.

```
umount /dev/xvdd1
```

4. Create a new partition on /dev/xvdd.

- a. Use the fdisk command to create a new **primary** partition on /dev/xvdd by using the following parameters:

```
fdisk /dev/xvdd
WARNING: DOS-compatible mode is deprecated. It's strongly
recommended to switch off the mode (command 'c') and change
display units to sectors (command 'u').
Command (m for help) : n
Command action
 e extended
 p primary partition (1-4)
p
Partition number (1-4) : 2
First cylinder (1886-9429, default 1886) : ENTER
Using default value 1886
Last cylinder, +cylinders or +size{K,M,G} (1886-9429, default
9429) : +1G
```

- b. Use the “t” command to change the system ID on the new partition.

```
Command (m for help): t
Partition number (1-4): 2
Hex code (type L to list codes): 8e
Changed system type of partition 2 to 8e (Linux LVM)
```

- c. Save the new partition table.

```
Command (m for help): w
The partition table has been altered!
Calling ioctl() to re-read partition table.
Syncing disks.
```

- d. List the partition table on /dev/xvdd.

```
fdisk -l /dev/xvdd
Disk /dev/xvdd: 5368 MB, 5368709120 bytes
...
Device Boot Start End Blocks Id System
/dev/xvdd1 1 1886 1048575+ 83 Linux
/dev/xvdd2 1886 3772 1048656+ 8e Linux LVM
```

5. Re-mount the file systems.

Use the `mount -a` command to mount all file systems mentioned in `/etc/fstab`.

```
mount -a
```

6. Display the mounted file systems.

```
df -h
...
/dev/xvdb1 1008M 34M 924M 4% /Dev
/dev/xvdd1 992M 1.3M 940M 1% /Test
```

## Practice 12-2: Creating a Logical Volume

---

### Overview

In this practice, you create physical volumes, a volume group, and a logical volume. You also use LVM utilities to display information about the LVM entities.

### Tasks

1. Initialize the new partitions for use by LVM (create physical volumes).

- a. List the partitions with the Linux LVM (**8e**) system ID.

```
fdisk -l | grep 8e
/dev/xvdb2 519 1037 1048887+ 8e Linux LVM
/dev/xvdd2 1886 3772 1048656 8e Linux LVM
```

- b. Use the `pvcreate` command to create physical volumes on both partitions.

```
pvcreate -v /dev/xvdb2 /dev/xvdd2
 Set up physical volume for "/dev/xvdb2" with ...
 Zeroing start of device /dev/xvdb2
 Writing physical volume data to disk "/dev/xvdb2"
 Physical volume "/dev/xvdb2" successfully created
 Set up physical volume for "/dev/xvdd2" with ...
 Zeroing start of device /dev/xvdd2
 Writing physical volume data to disk "/dev/xvdd2"
 Physical volume "/dev/xvdd2" successfully created
```

2. Display information about the physical volumes.

- a. Use the `pvdisplay` command to display attributes of the physical volumes.

```
pvdisplay
"/dev/xvdb2" is a new physical volume of "1.00 GiB"
--- NEW Physical volume ---
PV Name /dev/xvdb2
VG Name
PV Size 1.00 GiB
...
"/dev/xvdd2" is a new physical volume of "1.00 GiB"
--- NEW Physical volume ---
PV Name /dev/xvdd2
VG Name
PV Size 1.00 GiB
...
```

- b. Use the `pvs` command to report information in a more condensed form.

```
pvs
PV VG Fmt Attr PSize PFree
/dev/xvdb2 lvm2 a- 1.00g 1.00g
/dev/xvdd2 lvm2 a- 1.00g 1.00g
```

- c. Use the `pvscan` command to scan all disks for physical volumes.

```
pvscan
PV /dev/xvdb2 lvm2 [1.00 GiB]
PV /dev/xvdd2 lvm2 [1.00 GiB]
Total: 2 [2.00 GiB] / in use: 0 [0] / in no VG: 2 [2.00...]
```

3. Create a volume group.

Use the `vgcreate` command to create a volume group named **myvolg** from the `/dev/xvdb2` physical volume.

```
vgcreate -v myvolg /dev/xvdb2
Wiping cache of LVM-capable devices
Adding physical volume '/dev/xvdb2' to volume group...
Archiving volume group "myvolg" metadata (seqno 0).
Creating volume group backup "/etc/lvm/backup/myvolg"...
Volume group "myvolg" successfully created
```

4. Display information about the volume group.

- a. Use the `vgdisplay` command to display attributes of the volume group.

```
vgdisplay
--- Volume group ---
VG Name myvolg
System ID
Format lvm2
...
VG Size 1020.00 MiB
...
```

- b. Use the `vgs` command to report information in a more condensed form.

```
vgs
VG #PV #LV #SN Attr VSize VFree
myvolg 1 0 0 wz--n- 1020.00m 1020.00m
```

- c. Use the `vgscan` command to scan all disks for volume groups and rebuild caches.

```
vgscan
Reading all physical volumes. This may take a while...
Found volume group "myvolg" using metadata type lvm2
```

- d. Display information about the physical volumes.

```
pvs
PV VG Fmt Attr PSize PFree
/dev/xvdb2 myvolg lvm2 a- 1020.00m 1020.00m
/dev/xvdd2 lvm2 a- 1.00g 1.00g
```

- Notice that the /dev/xvdb2 physical volume is assigned to the **myvolg** volume group.

5. Create a logical volume.

Use the **lvcreate** command to create a 500 MB logical volume named **myvol** from the **myvolg** volume group.

```
lvcreate -v -L 500m -n myvol myvolg
Setting logging type to disk
Finding volume group "myvolg"
Archiving volume group "myvolg" metadata (seqno 1).
Creating logical volume myvol
...
Logical volume "myvol" created
```

6. Display information about the logical volume.

- a. Use the **lvdisplay** command to display attributes of the logical volume.

```
lvdisplay
--- Logical volume ---
LV Name /dev/myvolg/myvol
VG Name myvolg
...
LV Size 500.00 MiB
...
```

- b. Use the **lvs** command to report information in a more condensed form.

```
lvs
LV VG Attr LSize Pool Origin Data% Move Log Cpy%...
myvol myvolg -wi-a- 500.00m
```

- c. Use the **lvscan** command to scan all disks for logical volumes.

```
lvscan
ACTIVE '/dev/myvolg/myvol' [500.00 MiB] inherit
```

- d. Display information about the physical volumes.

```
pvs
PV VG Fmt Attr PSize PFree
/dev/xvdb2 myvolg lvm2 a- 1020.00m 520.00m
/dev/xvdd2 lvm2 a- 1.00g 1.00g
```

- Notice that the free space in the /dev/xvdb2 physical volume has been reduced.

- e. Display information about the volume group.

```
vgs
VG #PV #LV #SN Attr VSize VFree
myvolg 1 1 0 wz--n- 1020.00m 520.00m
```

- Notice that the free space in the **myvolg** volume group has also been reduced.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 12-3: Creating a File System and Mounting a Logical Volume

### Overview

In this practice, you create a file system on the logical volume and mount the logical volume.

### Tasks

1. Display the block device name that was automatically created.

- a. List the /dev entries for the **myvol** logical volume.

```
ls -l /dev/myvolg/myvol
lrwxrwxrwx. /dev/myvolg/myvol -> ../dm-0
ls -l /dev/mapper/myvolg-myvol
lrwxrwxrwx. /dev/mapper/myvolg-myvol -> ../dm-0
```

- Notice that two entries were automatically created.
- Notice that both entries are symbolic links to /dev/dm-0.

- b. List the /dev/dm-0 entry.

```
ls -l /dev/dm-0
brw-rw----. /dev/dm-0
```

- Notice that /dev/dm-0 is a block device.

2. Create a file system on the logical volume.

- a. Create an **ext4** file system on the **myvol** logical volume.

```
mkfs.ext4 /dev/mapper/myvolg-myvol
mke2fs 1.43-WIP (20-Jun-2013)
Filesystem label=
OS type: Linux
...
Writing inode tables: done
Creating journal (8192 blocks): done
...
```

- b. Create a /myvol mount point.

```
mkdir /myvol
```

- c. Mount the file system.

```
mount /dev/mapper/myvolg-myvol /myvol
```

- d. Display the mounted file systems.

```
df -h
Filesystem Size Used Avail Use% Mounted on
...
/dev/mapper/myvolg-myvol
 477M 2.3M 449M 1% /myvol
```

3. Update the file systems mount table.

Use the `vi` editor to add the following line to `/etc/fstab`.

```
/dev/mapper/myvolg-myvol /myvol ext4 defaults 0 0
```

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 12-4: Backing Up Volume Group Metadata

---

### Overview

In this practice, you view the LVM configuration file, view the automatically created metadata backups and archives, and manually create a volume group metadata backup.

### Tasks

- View the configuration of metadata backups and archiving.


- Use the `less` command to view the logical volume configuration, `/etc/lvm/lvm.conf`.

```
less /etc/lvm/lvm.conf
This is an example configuration file for the LVM2 system.
...
```

- Use the search character, `/`, and search for the word `backup`.

```
/backup
```

- The search results in displaying the following screenshot:



```
root@host03:/etc/lvm
File Edit View Search Terminal Help
Configuration of metadata backups and archiving. In LVM2 when we
talk about a 'backup' we mean making a copy of the metadata for the
current system. The 'archive' contains old metadata configurations.
Backups are stored in a human readable text format.
backup {
 # Should we maintain a backup of the current metadata configuration ?
 # Use 1 for Yes; 0 for No.
 # Think very hard before turning this off!
 backup = 1

 # Where shall we keep it ?
 # Remember to back up this directory regularly!
 backup_dir = "/etc/lvm/backup"

 # Should we maintain an archive of old metadata configurations.
 # Use 1 for Yes; 0 for No.
 # On by default. Think very hard before turning this off.
 archive = 1

 # Where should archived files go ?
 # Remember to back up this directory regularly!
 archive_dir = "/etc/lvm/archive"
:
```

- Notice that metadata backups are enabled (`backup = 1`) and backups are stored in the `/etc/lvm/backup` directory.
- Also notice that archives are enabled (`archive = 1`) and the archives are stored in the `/etc/lvm/archive` directory.

- c. Press q to quit the less command.

```
:q
#
```

- d. You can also use the lvm dumpconfig command to view the configuration settings.

```
lvm dumpconfig
config {
 checks=1
 abort_on_errors=0
 profile_dir="/etc/lvm/profile"
}
...
backup {
 backup=1
 backup_dir="/etc/lvm/backup"
 archive=1
 archive_dir="/etc/lvm/archive"
}
```

2. View the metadata backup and archive files.

- a. Use the cd command to change to the /etc/lvm/backup directory. Use the ls command to display the contents of the directory.

```
cd /etc/lvm/backup
ls
myvolg
```

- Notice the backup of the myvolg volume group was created automatically.

- b. Use the less command to view the contents of the myvolg metadata backup file.

```
less myvolg
...
description = "Created *after* executing 'lvcreate -v -L 500...
...
myvolg {
...
 physical_volumes {
 pv0 {
...
 device = "/dev/xvdb2"
...
 logical_volumes {
 myvol { ...

```

- Notice the description states the backup file was created “after” executing the `lvcreate` command string.
  - Also notice the metadata backup file includes information on the physical and logical volumes.
- c. Use the `cd` command to change to the `/etc/lvm/archive` directory. Use the `ls` command to display the contents of the directory.

```
cd /etc/lvm/archive
ls
myvolg_00000-... myvolg_00001-...
```

- Notice the archive files for the `myvolg` volume group were created automatically.
- d. Use the `less` command to view the contents of the `myvolg-00000*` archive file.

```
less myvolg_00000*
...
description = "Created *before* executing 'vgcreate -v myvolg...
...
myvolg {
...
 physical_volumes {
 pv0 {
...
 device = "/dev/xvdb2"
...
...
```

- Notice the description states the archive file was created “before” executing the `vgcreate` command string.
- e. Use the `less` command to view the contents of the `myvolg-00001*` archive file.

```
less myvolg_00001*
...
description = "Created *before* executing 'lvcreate -v -L 500...
...
myvolg {
...
 physical_volumes {
 pv0 {
...
 device = "/dev/xvdb2"
...
...
```

- Notice the description states the archive file was created “before” executing the `lvcreate` command string.

3. Create a metadata backup of the myvolg volume group.
  - a. Use the vgcfgbackup command to back up the metadata for the myvolg volume group.
    - Include the -f file\_today argument to the vgcfgbackup command.

```
vgcfgbackup -f file_today myvolg
Volume group "myvolg" successfully backed up.
```

- b. Use the diff command to display the differences in the newly created backup file and the existing myvolg metadata backup file.
    - Note that the newly created backup was created in the current directory.

```
diff file_today /etc/lvm/backup/myvolg
...
< description = "vgcfgbackup -f file_today myvolg"

> description = "Created *after* executing 'lvcreate -v -L ...
..."
```

- Notice that only the description and the creation information are different.

## Practice 12-5: Creating a Logical Volume Snapshot

---

### Overview

In this practice, you create a snapshot volume, mount the snapshot, and remove the snapshot volume.

### Tasks

1. Create a snapshot volume.

- a. Copy `/boot/init*` to the mounted logical volume, `/myvol`.

```
cp /boot/init* /myvol
ls /myvol
initramfs-2.6.32-431.el6.x86_64.img
initramfs-3.8.13-16.2.1.el6uek.x86_64.img
lost+found
```

- b. Create a 500 MB snapshot named `myvol-snapshot` of the `myvol` logical volume.

```
lvcreate -L 500m -s -n myvol-snapshot myvolg/myvol
Logical volume "myvol-snapshot" created
```

- c. Use the `ls -l` command to list the contents of the `/etc/lvm/backup` directory and the contents of the `/etc/lvm/archive` directory.

```
ls -l /etc/lvm/backup
-rw-----. <date_time> myvolg
ls -l /etc/lvm/archive
-rw-----. <date_time> file_today
-rw-----. <date_time> myvolg_00000...
-rw-----. <date_time> myvolg_00001...
-rw-----. <date_time> myvolg_00002...
```

- Notice a new `myvolg` backup file was automatically created when the snapshot was created (note the time stamp).
- Also notice a new archive file was automatically created (note the time stamp on the `myvolg_00002*` file).

- d. List the logical volumes.

```
lvs
 LV VG Attr LSize Pool Origin Data% Move...
 myvol myvolg owi-aos 500.00m
 myvol-snapshot myvolg swi-a-s 500.00m myvol 0.00
```

- e. List the contents of the /dev/myvolg and /dev/mapper directories.

```
ls -l /dev/myvolg
lrwxrwxrwx. myvol -> ../dm-0
lrwxrwxrwx. myvol-snapshot -> ../dm-1
ls -l /dev/mapper
crw-rw----. control
lrwxrwxrwx. myvolg-myvol -> ../dm-0
lrwxrwxrwx. myvolg-myvol-real -> ../dm-2
lrwxrwxrwx. myvolg-myvol--snapshot -> ../dm-1
lrwxrwxrwx. myvolg-myvol--snapshot-cow -> ../dm-3
```

2. Mount the snapshot.

- a. Mount the snapshot onto /mnt.

```
mount -t ext4 /dev/myvolg/myvol-snapshot /mnt
```

- b. Display the mounted file systems.

```
df -h
Filesystem Size Used Avail Use% Mounted on
...
/dev/mapper/myvolg-myvol
 477M 41M 411M 9% /myvol
/dev/mapper/myvolg-myvol--snapshot
 477M 41M 411M 9% /mnt
```

- c. List the files on /mnt.

```
ls /mnt
initramfs-2.6.32-431.el6.x86_64.img
initramfs-3.8.13-16.2.1.el6uek.x86_64.img
lost+found
```

- Notice that these are the same files that were copied onto /myvol.

- d. Remove the files on /myvol.

```
rm /myvol/init*
rm: remove regular file '/myvol/init...'? y
...
```

- e. List the files on /mnt.

```
ls /mnt
initramfs-2.6.32-431.el6.x86_64.img
initramfs-3.8.13-16.2.1.el6uek.x86_64.img
lost+found
```

- Notice that these files are still present.
- Removing files from the original volume does not change the snapshot's content.

3. Remove the snapshot.

- a. Unmount the snapshot from /mnt.

```
umount /mnt
```

- b. Use the lvremove command to remove the snapshot.

- Answer **y** when asked, “Do you really want to ...”

```
lvremove -v myvolg/myvol-snapshot
Using logical volume(s) on command line
Do you really want to remove active logical volume myvol-
snapshot? [y/n]: y
...
Logical volume "myvol-snapshot" successfully removed
```

- c. List the logical volumes.

```
lvs
 LV VG Attr LSize Pool Origin Data% Move...
 myvol myvolg -wi-ao 500.00m
```

- Notice that the **myvol-snapshot** logical volume has been deleted.

- d. List the contents of the /dev/myvolg and /dev/mapper directories.

```
ls -l /dev/myvolg
lrwxrwxrwx. myvol -> ../dm-0
ls -l /dev/mapper
crw-rw----. control
lrwxrwxrwx. myvolg-myvol -> ../dm-0
```

- Notice that the “snapshot” entries have been deleted.

- e. Use the ls -l command to list the contents of the /etc/lvm/backup directory and the contents of the /etc/lvm/archive directory.

```
ls -l /etc/lvm/backup
-rw-----. <date_time> myvolg
ls -l /etc/lvm/archive
-rw-----. <date_time> file_today
-rw-----. <date_time> myvolg_00000...
-rw-----. <date_time> myvolg_00001...
-rw-----. <date_time> myvolg_00002...
-rw-----. <date_time> myvolg_00003...
```

- Notice a new **myvolg** backup file was automatically created when the snapshot was removed (note the time stamp).
- Also notice a new archive file was automatically created (note the time stamp on the **myvolg\_00003\*** file).

## Practice 12-6: Increasing the Capacity of a Logical Volume

---

### Overview

In this practice, you increase the size of a logical volume and add a physical volume to a volume group.

### Tasks

1. Increase the size of a logical volume.

- a. List the volume group.

```
vgs
VG #PV #LV #SN Attr VSize VFree
myvolg 1 1 0 wz--n- 1020.00m 520.00m
```

- Notice that the **myvolg** volume group has 520 MB of free space.

- b. List the logical volume.

```
lvs
LV VG Attr LSize Pool Origin Data% Move...
myvol myvolg -wi-ao 500.00m
```

- Notice that the **myvol** logical volume is 500 MB in size.

- c. Display the mounted file systems.

```
df -h
Filesystem Size Used Avail Use% Mounted on
...
/dev/mapper/myvolg-myvol
 477M 2.3M 449M 1% /myvol
```

- Notice that the size of the file system is 477 MB.

- d. Use the **lvextend** command to increase the size of the **myvolg/myvol** logical volume and the file system by 500 MB.

```
lvextend -L +500M -r myvolg/myvol
Extending logical volume myvol to 1000.00 MiB
Logical volume myvol successfully resized
resize2fs 1.43-WIP (20-Jun-2013)
Filesystem at /dev/mapper/myvolg-myvol is mounted on /myvol; online resize required
...
The filesystem on /dev/mapper/myvolg-myvol is now 1024000 blocks long.
```

- The **-r** option causes the file system to be resized.

- e. List the volume group.

```
vgs
VG #PV #LV #SN Attr VSize VFree
myvolg 1 1 0 wz--n- 1020.00m 20.00m
```

- Notice that the **myvolg** volume group now has only 20 MB of free space.

- f. List the logical volume.

```
lvs
LV VG Attr LSize Pool Origin Data% Move...
myvol myvolg -wi-ao 1000.00m
```

- Notice that the **myvol** logical volume is now 1,000 MB in size.

- g. Display the mounted file systems.

```
df -h
Filesystem Size Used Avail Use% Mounted on
...
/dev/mapper/myvolg-myvol
 961M 2.5M 914M 1% /myvol
```

- Notice that the size of the file system is now 961 MB.

- h. Use the `ls -l` command to list the contents of the `/etc/lvm/backup` directory and the contents of the `/etc/lvm/archive` directory.

```
ls -l /etc/lvm/backup
-rw-----. <date_time> myvolg
ls -l /etc/lvm/archive
-rw-----. <date_time> file_today
-rw-----. <date_time> myvolg_00000...
-rw-----. <date_time> myvolg_00001...
-rw-----. <date_time> myvolg_00002...
-rw-----. <date_time> myvolg_00003...
-rw-----. <date_time> myvolg_00004...
```

- Notice a new `myvolg` backup file was automatically created when the logical volume was extended (note the time stamp).
- Also notice a new archive file was automatically created (note the time stamp on the `myvolg_00004*` file).

## 2. Add a physical volume to a volume group.

- a. List the physical volumes.

```
pvs
PV VG Fmt Attr PSize PFree
/dev/xvdb2 myvolg lvm2 a-- 1020.00m 20.00m
/dev/xvdd2 lvm2 a-- 1.00g 1.00g
```

- Notice that the `/dev/xvdd2` physical volume is not assigned to a volume group.

- b. List the volume group.

```
vgs
VG #PV #LV #SN Attr VSize VFree
myvolg 1 1 0 wz--n- 1020.00m 20.00m
```

- Notice that the **myvolg** volume group is 1020 MB in size.

- c. Use the `vgextend` command to add the `/dev/xvdd2` physical volume to the **myvolg** volume group.

```
vgextend -v myvolg /dev/xvdd2
 Checking for volume group "myvolg"
...
 Adding physical volume '/dev/xvdd2' to volume group 'myvolg'
...
 Volume group "myvolg" successfully extended
```

- d. List the physical volumes.

```
pvs
PV VG Fmt Attr PSize PFree
/dev/xvdb2 myvolg lvm2 a-- 1020.00m 20.00m
/dev/xvdd2 myvolg lvm2 a-- 1020.00m 1020.00m
```

- Notice that the `/dev/xvdd2` physical volume is now assigned to the **myvolg** volume group.

- e. List the volume group.

```
vgs
VG #PV #LV #SN Attr VSize VFree
myvolg 2 1 0 wz--n- 1.99g 1.02g
```

- Notice that the **myvolg** volume group now has two PVs (physical volumes).
- Notice that VSize and VFree have increased.

- f. List the logical volume.

```
lvs
LV VG Attr LSize Pool Origin Data% Move...
myvol myvolg -wi-ao 1000.00m
```

- g. Use the `ls -l` command to list the contents of the `/etc/lvm/backup` directory and the contents of the `/etc/lvm/archive` directory.

```
ls -l /etc/lvm/backup
-rw-----. <date_time> myvolg
ls -l /etc/lvm/archive
-rw-----. <date_time> file_today
-rw-----. <date_time> myvolg_00000...
-rw-----. <date_time> myvolg_00001...
-rw-----. <date_time> myvolg_00002...
-rw-----. <date_time> myvolg_00003...
-rw-----. <date_time> myvolg_00004...
-rw-----. <date_time> myvolg_00005...
```

- Notice a new `myvolg` backup file was automatically created when the physical volume was added to the volume group (note the time stamp).
- Also notice a new archive file was automatically created (note the time stamp on the `myvolg_00005*` file).

3. Increase the size of the **myvolg/myvol** logical volume and the file system.

- a. Increase the size of the **myvolg/myvol** logical volume and the file system by 500 MB.

```
lvextend -L +500M -r myvolg/myvol
Extending logical volume myvol to 1.46 GiB
Logical volume myvol successfully resized
resize2fs 1.43-WIP (20-May-2013)
...
The file system on /dev/mapper/myvolg-myvol is now 1536000
blocks long.
```

- b. List the logical volume.

```
lvs
LV VG Attr LSize Pool Origin Data% Move...
myvol myvolg -wi-ao 1.46g
```

- c. Display the mounted file systems.

```
df -h
Filesystem Size Used Avail Use% Mounted on
...
/dev/mapper/myvolg-myvol
 1.5G 2.8M 1.4G 1% /myvol
```

- Notice that the size of the file system is now 1.5 GB.

- d. Use the `ls -l` command to list the contents of the `/etc/lvm/backup` directory and the contents of the `/etc/lvm/archive` directory.

```
ls -l /etc/lvm/backup
-rw----- . <date_time> myvolg
ls -l /etc/lvm/archive
-rw----- . <date_time> file_today
-rw----- . <date_time> myvolg_00000...
-rw----- . <date_time> myvolg_00001...
-rw----- . <date_time> myvolg_00002...
-rw----- . <date_time> myvolg_00003...
-rw----- . <date_time> myvolg_00004...
-rw----- . <date_time> myvolg_00005...
-rw----- . <date_time> myvolg_00006...
```

- Notice a new `myvolg` backup file was automatically created when the logical volume was extended (note the time stamp).
- Also notice a new archive file was automatically created (note the time stamp on the `myvolg_00006*` file).

## Practice 12-7: Restoring Volume Group Metadata

---

### Overview

In this practice, you restore volume group metadata from a backup.

### Tasks

1. View the current physical volumes, volume group, and logical volumes configuration.

- a. Use the `pvs` command to display information about the physical volumes.

```
pvs
PV VG Fmt Attr PSize PFree
/dev/xvdb2 myvolg lvm2 a- 1020.00m 0
/dev/xvdd2 myvolg lvm2 a- 1020.00m 540.00m
```

- b. Use the `vgs` command to display information about the volume groups.

```
vgs
VG #PV #LV #SN Attr VSize VFree
myvolg 2 1 0 wz--n- 1.99g 540.00m
```

- c. Use the `lvs` command to display information about the logical volumes.

```
lvs
LV VG Attr LSize Pool Origin Data% Move Log Cpy%...
myvol myvolg -wi-ao- 1.46g
```

2. Restore the volume group metadata from a backup.

- a. Use the `vgcfgrestore` command to restore the `myvolg` volume group metadata.

```
vgcfgrestore myvolg
Restored volume group myvolg
```

- b. Repeat Task 1a, 1b, 1c (execute the `pvs`, `vgs`, `lvs` commands).

```
pvs
PV VG Fmt Attr PSize PFree
/dev/xvdb2 myvolg lvm2 a- 1020.00m 0
/dev/xvdd2 myvolg lvm2 a- 1020.00m 540.00m
vgs
VG #PV #LV #SN Attr VSize VFree
myvolg 2 1 0 wz--n- 1.99g 540.00m
lvs
LV VG Attr LSize Pool Origin Data% Move Log Cpy%...
myvol myvolg -wi-ao- 1.46g
```

- Notice that the LVM information is the same as the output in Task 1.

## Practice 12-8: Creating a RAID Device

### Overview

In this practice, you remove the logical volume and LVM entities, create a RAID array device, and create a file system and mount the RAID device. You remove the RAID device at the end of the practice.

### Tasks

1. Remove the logical volume, volume group, and physical volume.

- a. Unmount the **myvol** logical volume.

```
umount /myvol
```

- b. Use the **lvremove** command to remove the **myvol** logical volume.

- Answer **y** when asked, “Do you really want to ...”

```
lvremove myvolg/myvol
```

```
Do you really want to remove active logical volume myvol? [y/n] :
```

```
y
```

```
Logical volume "myvol" successfully removed
```

- c. Use the **vgremove** command to remove the **myvolg** volume group.

```
vgremove myvolg
```

```
Volume group "myvolg" successfully removed
```

- d. Use the **pvremove** command to remove the physical volumes.

```
pvremove /dev/xvdb2 /dev/xvdd2
```

```
Labels on physical volume "/dev/xvdb2" successfully wiped
```

```
Labels on physical volume "/dev/xvdd2" successfully wiped
```

- e. Remove the **/myvol** mount point.

```
rmdir /myvol
```

- f. Using the **vi** editor, remove the following line from **/etc/fstab**.

```
vi /etc/fstab
```

```
/dev/mapper/myvolg-myvol /myvol ext4 defaults 0 0
```

2. Create a new RAID array.

- Disk partitions used in a RAID device should have the partition type changed to **fd** (Linux raid autodetect). This step is omitted however in this practice.

- a. Use the `mdadm` command to create a RAID-1 device, `/dev/md0`, from the `/dev/xvdb2` and `/dev/xvdd2` partitions.

- Answer `y` when asked, “Continue creating array?”

```
mdadm --create /dev/md0 --level=1 --raid-devices=2 /dev/xvdb2
/dev/xvdd2
...
Continue creating array? y
mdadm: Defaulting to version 1.2 metadata
mdadm: array /dev/md0 started.
```

- b. View the `/proc/mdstat` file to check the status of the MD RAID devices.

```
cat /proc/mdstat
Personalities : [raid1]
md0 : active raid1 xvdd2[1] xvdb2[0]
 1047616 blocks super 1.2 [2/2] [UU]
 [= =====>.....] resync = ...
unused devices: <none>
```

- c. Run the previous command again to confirm the resync has completed.

```
cat /proc/mdstat
Personalities : [raid1]
md0 : active raid1 xvdd2[1] xvdb2[0]
 1047616 blocks super 1.2 [2/2] [UU]
unused devices: <none>
```

- Ensure the resync is complete before continuing.

- d. Use the `mdadm` command to view information about the RAID device.

```
mdadm --query /dev/md0
/dev/md0: 1023.06MiB raid1 2 devices, 0 spares. Use mdadm --
detail for more detail.
mdadm --detail /dev/md0
/dev/md0:
 Version : 1.2
 Creation Time : Tue Dec 10...
 Raid Level : raid1
 Array Size : 1047616 (1023.23 MiB 1072.76 MB)
 Used Dev Size : 1047616 (1023.243MiB 1072.76 MB)
 Raid Devices : 2
 Total Devices : 2
 ...
 Number Major Minor RaidDevice State
 0 202 18 0 active sync /dev/xvdb2
 1 202 50 1 active sync /dev/xvdd2
```

### 3. Create a file system on the RAID device.

- a. Create an **ext4** file system on `/dev/md0`.

```
mkfs.ext4 /dev/md0
mke2fs 1.43-WIP (20-Jun-2013)
Filesystem label=
OS type: Linux
...
Writing inode tables: done
Creating journal (4096 blocks): done
...
```

- b. Create a mount point named `/raid`.

```
mkdir /raid
```

- c. Mount the file system.

```
mount /dev/md0 /raid
```

- d. Display the mounted file systems.

```
df -h
Filesystem Size Used Avail Use% Mounted on
...
/dev/md0 991M 1.3M 939M 1% /raid
```

4. Create the **mdadm** configuration file, `/etc/mdadm.conf`.

Use the `vi` editor to create `/etc/mdadm.conf` and add the following entry.

- This step does not need to be performed in this practice. You are directed to remove this file in the next task.

```
vi /etc/mdadm.conf
ARRAY /dev/md0 devices=/dev/xvdb2,/dev/xvdd2
```

- With this file, the RAID array is properly detected and initialized after a reboot.

5. Remove the RAID array.

- a. Remove the `/etc/mdadm.conf` file.

```
rm /etc/mdadm.conf
rm: remove regular file '/etc/mdadm.conf'? y
```

- b. Unmount the `raid` volume.

```
umount /raid
```

- c. Use the `mdadm` command to deactivate the array and release all resources.

```
mdadm --stop /dev/md0
mdadm: stopped /dev/md0
```

- d. Display the attributes of the /dev/xvdb3 block device.

```
blkid | grep raid
/dev/xvdb2: UUID="..." UUID_SUB="..."
LABEL="host03.example.com:0" TYPE="linux_raid_member"
/dev/xvdd2: UUID="..." UUID_SUB="..."
LABEL="host03.example.com:0" TYPE="linux_raid_member"
```

- Notice /dev/xvdb2 and /dev/xvdd2 have a type of "linux\_raid\_member."

- e. Use the mdadm command to overwrite the md superblock on /dev/xvdb2 and /dev/xvdd2.

```
mdadm --zero-superblock /dev/xvdb2
mdadm --zero-superblock /dev/xvdd2
```

- f. Display the attributes of /dev/xvdb2 and /dev/xvdd2 to ensure that the md superblock was overwritten.

```
blkid /dev/xvdb2
blkid /dev/xvdd2
```

- No output indicates the superblock was overwritten.

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## **Practices for Lesson 13: Network Configuration**

**Chapter 13**

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practices for Lesson 13: Network Configuration

---

### Practices Overview

In these practices, you:

- Configure the `eth1` network interface by editing network interface configuration files
- Use NetworkManager to configure network interfaces
- Use the `system-config-network` text-based utility
- Access the Public Yum Repository from **host03** VM
- Simulate upgrading your system (you *do not* actually perform the upgrade)

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 13-1: Configuring the `eth1` Network Interface

### Overview

In this practice, you:

- Display the available network interfaces on your system
- View the network interface configuration files
- Configure a static IP address for the `eth1` network interface
- Update your `/etc/hosts` file
- Create a default gateway and display your route table
- Ensure that connectivity to `dom0` and the other VM guests

### Assumptions

You are the `root` user on **host03** VM.

If you were unable to complete the OS install and FirstBoot configuration on **host03** in Practice 3, substitute **host02** for **host03** in this practice and all future practices in this course. Use the IP address of 192.168.1.102 for the `eth1` interface if configuring **host02**.

### Tasks

1. Use the `ifconfig` command to display your available network interfaces.

```
ifconfig
eth0 Link encap:Ethernet HWaddr 00:16:3E:00:01:03
 inet addr:192.0.2.103 ...
 inet6 addr: ...
...
eth1 Link encap:Ethernet HWaddr 00:16:3E:00:02:03
 inet6 addr: ...
...
lo Link encap:Local Loopback
 inet addr:127.0.0.1
...
```

- Notice that you have two Ethernet interfaces (`eth0` and `eth1`) and the loopback interface (`lo`).
  - The `eth0` Ethernet interface has an IPv4 address, but `eth1` does not.
2. View the network interface configuration files.
    - a. Use the `cd` command to change to the `/etc/sysconfig/network-scripts` directory.

```
cd /etc/sysconfig/network-scripts
```

- b. Use the `ls` command to view the contents of this directory.

```
ls
ifcfg-eth0 ifdown-ipv6 ifup ifup-plip...
ifcfg-eth1 ifdown-isdn ifup-aliases ifup-plush...
ifcfg-lo ifdown-post...
```

- Notice that you have a configuration file for eth0, `ifcfg-eth0`.
- Notice that you have a configuration file for eth1, `ifcfg-eth1`.
- Notice that you have a configuration file for the loopback interface, `ifcfg-lo`.
- Several interface control scripts exist in this directory to activate and deactivate network interfaces.

3. Configure `eth1` with a static IP address of 192.168.1.103.

Use the `vi` editor to edit the `ifcfg-eth1` file as follows (the use of quotes around the assigned values is optional):

```
DEVICE=eth1
BOOTPROTO=none
ONBOOT=yes
USERCTL=no
IPV6INIT=no
PEERDNS=yes
TYPE=Ethernet
HWADDR=00:16:3e:00:02:03
IPADDR=192.168.1.103
NETMASK=255.255.255.0
BROADCAST=192.168.1.255
UUID=... (leave this as is)
NM_CONTROLLED=yes
```

4. Specify routing and host information for all network interfaces.

- a. Use the `cd` command to change to the `/etc/sysconfig` directory.

```
cd /etc/sysconfig
```

- b. If necessary, use the `vi` editor to edit the `network` file as follows:

- This file might not require editing.

```
NETWORKING=yes
HOSTNAME=host03.example.com
GATEWAY=192.0.2.1
```

5. Use the `vi` editor to edit the `/etc/hosts` file as follows:

```
127.0.0.1 localhost.localdomain localhost
192.0.2.1 example.com dom0
192.0.2.101 host01.example.com host01
192.0.2.102 host02.example.com host02
192.0.2.103 host03.example.com host03
```

6. Use the `service` command to restart the network.

```
service network restart
Shutting down interface eth0...
Shutting down interface eth1...
Shutting down loopback interface...
Bringing up loopback interface...
Bringing up interface eth0...
Bringing up interface eth1...
...
```

7. Use the `ifconfig` command to display the status of the interfaces.

```
ifconfig
eth0 Link encap:Ethernet HWaddr 00:16:3E:00:01:03
 inet addr:192.0.2.103 Bcast:192.0.2.255 Mask:...
 inet6 addr: ...
 UP BROADCAST RUNNING MULTICAST MTU:1500...
...
eth1 Link encap:Ethernet HWaddr 00:16:3E:00:02:03
 inet addr:192.168.1.103 Bcast:192.168.1.255 Mask:...
 inet6 addr: ...
 UP BROADCAST RUNNING MULTICAST MTU:1500...
...
```

- Notice that both `eth0` and `eth1` now have IP addresses.
8. Use the `netstat -r` command (or `route`) to display the route table.

```
netstat -r
Destination Gateway Genmask Flags ... Iface
default example.com 0.0.0.0 UG eth0
192.0.2.0 * 255.255.255.0 U eth0
192.168.1.0 * 255.255.255.0 U eth1
```

- Notice that all packets destined for the `192.168.1` subnet use the `eth1` interface.
- Notice that all packets destined for the `192.0.2` subnet use the `eth0` interface.
- Notice that all other packets are routed through the `192.0.2.1` default gateway (`example.com`), via `eth0`.

9. Use the `ping` command to verify that you can communicate to **dom0** and the other VM guests.
- Press **Ctrl + C** to kill the `ping` command.

```
ping dom0
64 bytes from example.com (192.0.2.1)...
CTRL-C
ping host01
64 bytes from host01.example.com (192.0.2.101)...
CTRL-C
ping host02
64 bytes from host02.example.com (192.0.2.102)...
CTRL-C
```

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 13-2: Using NetworkManager

### Overview

In this practice, you:

- Use the NetworkManager utility to disable and enable network connections
- View the configuration capabilities of NetworkManager
- Run the `nm-tool` utility that is included with NetworkManager


### Assumptions

You are the `root` user on **host03** VM.

### Tasks

1. Install NetworkManager if necessary.

- NetworkManager is installed and running if you see this computer screen icon on the GNOME panel as follows:



- a. Use the `rpm` command to verify that the NetworkManager package is installed.

```
rpm -qa | grep -i networkmanager
NetworkManager-...
NetworkManager-glib-...
NetworkManager-gnome-...
```

- b. If NetworkManager is not installed, use the `yum` command to install the package.

```
yum install NetworkManager
...
```


- c. Use the `service` command to verify that NetworkManager is running.

```
service NetworkManager status
NetworkManager (pid ...) is running...
```

- d. If NetworkManager is not running, use the `service` command to start it.

```
service NetworkManager start
...
```

2. Use NetworkManager to enable/disable networking.
  - a. Right-click the computer screens icon to display the following drop-down menu:



- b. Deselect the **Enable Networking** check box to disable networking.
    - Notice that an X in red appears on the icon, as shown:



- c. From a command window, use the ifconfig command to display the network status.

```
ifconfig
lo Link encap: Local Loopback
...


```

- Notice that eth0 and eth1 interfaces do not appear in the output.
- d. Right-click the computer screens icon to display the drop-down menu.
- e. Click **Enable Networking**. Ensure that the check box is selected and the X in red disappears.
- f. Run the ifconfig command again to ensure that the eth0 and eth1 interfaces appear.

```
ifconfig
eth0 Link encap:Ethernet HWaddr 00:16:3E:00:01:03
...
eth1 Link encap:Ethernet HWaddr 00:16:3E:00:02:03
...
lo Link encap:Local Loopback
...
```

- Notice that the eth0 and eth1 interfaces now appear.

- g. Click the computer screens icon to display the following drop-down menu:



- Notice that **System eth0** is in bold font, meaning it is connected.
- Notice that **System eth1** is in bold font, meaning it is also connected.

- h. Click **Disconnect** under **System eth0** to disconnect eth0. The menu now appears as follows:



- Notice that **Wired Network (eth0)** is disconnected, but **System eth0** (not in bold) is listed under the **Available** heading.

- i. From a command window, use the ifconfig command to display the network status.


```
ifconfig
...
```

- Notice that both eth0 and eth1 are listed, but eth0 does not include an IPv4 address.

- j. Click the computer screens icon and click **System eth0** to re-connect the eth0 interface.

3. Use NetworkManager to view network connections.

- a. Right-click the computer screens icon to display the following drop-down menu:



- b. Click **Edit Connections** to display the following screen:



- Do not run the following command but note that you can also access this Network Connections window by entering `nm-connection-editor&` from a command line.

```
nm-connection-editor&
```

- c. Select **System eth0**, then click **Edit**, and then select the **IPv4 Settings** tab to display the following screen:



- Notice that this information displayed is included in the configuration file, /etc/sysconfig/network-scripts/ifcfg-eth0.
  - Any updates made from this screen are written to the configuration file.
- d. Do not make any changes; click **Cancel** after viewing the different tabs and settings.
- e. Click **Close** to close the Network Connections window.
4. Run the `nm-tool` utility that is included with NetworkManager.

```
nm-tool
NetworkManager Tool
State: connected
- Device: eth0 [System eth0] --
...
```

## Practice 13-3: Using the system-config-network Utility

### Overview

In this practice, you use the `system-config-network` text-based utility to configure the network interfaces.

### Assumptions

You are the `root` user on `host03` VM.

### Tasks

1. Install `system-config-network` if necessary.
  - a. Use the `rpm` command to verify that the `system-config-network` package is installed.


```
rpm -qa | grep system-config-network
system-config-network-tui-...
```
  - b. If the utility is not installed, use the `yum` command to install the package.

```
yum install system-config-network-tui
...
```

2. Use system-config-network to configure network interfaces.
  - a. Enter system-config-network from the command line.

```
system-config-network
```


- The screen appears as follows:



- b. Use the up/down arrows to select **Device configuration** and press Enter. The screen appears as follows:



- c. Use the up/down arrows to select **eth1 (eth1) - Ethernet** and press Enter. The screen appears as follows:



- Notice that this information displayed is included in the configuration file, /etc/sysconfig/network-scripts/ifcfg-eth1.
  - Any updates made from this screen are written to the configuration file.
- d. Do not make any changes; use the arrow keys to highlight **Cancel** and then press Enter.
- e. Highlight **Cancel** on the previous screen and press Enter.
- f. Highlight **Quit** on the initial screen and press Enter to exit the utility.

## Practice 13-4: Accessing the Public Yum Repository

---

### Overview

In this practice, you configure your system to access the Public Yum Repository from **host03** VM. You then simulate upgrading your system (you *do not* actually perform the upgrade).

### Assumptions

- You are the root user on **host03** VM.
- You are the root user on **dom0**.

### Tasks

1. Configure the Domain Name Service (DNS) resolver configuration file.
  - This file, `/etc/resolv.conf`, provides access to DNS.
  - a. From **dom0**, use the `cat` command to view the `/etc/resolv.conf` file.

```
[dom0]# cat /etc/resolv.conf
search us.oracle.com example.com
nameserver 192.0.2.1
nameserver 152.68.154.3
nameserver 10.216.106.3
nameserver 193.32.3.252
```

- This is sample output only; your system may be different.
- b. On **host03**, use the `cat` command to view the `/etc/resolv.conf` file.

```
[host03]# cat /etc/resolv.conf
Generated by NetworkManager
...
```

- The first line in `/etc/resolv.conf` states, “Generated by NetworkManager”.
- If NetworkManager is running, which it is, any time you restart the network service, the contents of `/etc/resolv.conf` is updated from information in the `ifcfg-eth*` file.
- Rather than update `/etc/resolv.conf`, update the `ifcfg-eth*` file to include DNS information.
- c. On **host03**, add the `DOMAIN` directive and add the `DNS1`, `DNS2`, `DNS3`, and `DNS4` entries to the `/etc/sysconfig/network-scripts/ifcfg-eth0` file as shown.
- These entries were obtained from the `/etc/resolv.conf` file on **dom0**.

```
[host03]# vi /etc/sysconfig/network-scripts/ifcfg-eth0
...
DOMAIN="example.com us.oracle.com"
DNS1=192.0.2.1
DNS2=152.68.154.3
DNS3=10.216.106.3
DNS4=193.32.3.252
```

- d. On **host03**, use the **service** command to restart the network service.

```
[host03]# service network restart
...
```

- e. On **host03**, use the **cat** command to display the contents of the **/etc/resolv.conf** file.


```
[host03]# cat /etc/resolv.conf
Generated by NetworkManager
search example.com us.oracle.com
nameserver 192.0.2.1
nameserver 152.68.154.3
nameserver 10.216.106.3
...
nameserver 193.32.3.252
```

- Notice that NetworkManager generated the **/etc/resolv.conf** file.
2. Determine the **HTTP Proxy** server on **dom0**.
- a. On **dom0**, open Firefox by double-clicking the Firefox icon on the desktop.



- b. From the browser menu bar, select **Edit > Preferences** to display the Firefox Preferences window.
- c. Click the **Advanced** menu option on the Firefox Preferences window, and then select the **Network** tab.

- d. Click the **Settings** button. The Connection Settings window appears.




- e. Note the **HTTP Proxy** setting.
- In this example, the setting is ges-proxy.us.oracle.com on Port 80.
- f. Click **Cancel** to close the **Connection Settings** window.
- g. Click **Close** to close the **Firefox Preferences** window.
- h. Click **File > Quit** to close the browser on **dom0**.


All remaining commands are issued from **host03** VM.

3. Set the **HTTP Proxy** in Firefox to the setting configured on **dom0**.
- In this example, the **HTTP Proxy** is ges-proxy.us.oracle.com; yours may be different.

- On the GNOME menu bar, click the Firefox Web Browser icon to the right of the System menu option to start the Firefox web browser.



- The browser appears.
- From the browser menu bar, select **Edit > Preferences** to display the Firefox Preferences window.
  - Click the **Advanced** menu option and then select the **Network** tab.
  - Click the **Settings** tab and enter the following settings:
    - 1) Click **Manual proxy configuration**.
    - 2) Set **HTTP Proxy** to ges-proxy.us.oracle.com on Port 80 (or whatever was set on dom0).
    - 3) Select **Use this proxy server for all protocols**.
      - The Connection Settings window should look like the following (substituting the **HTTP Proxy** settings from **dom0**):



- e. Click **OK**.
  - f. Click **Close** in the Firefox Preferences window.
4. Access the Oracle Public Yum Server URL.
    - a. From a browser, enter the URL for the Public Yum Server, <http://public-yum.oracle.com>.
      - The page looks like the following:

The screenshot shows a Mozilla Firefox browser window titled "Oracle Public Yum Server - Mozilla Firefox". The address bar contains "public-yum.oracle.com". The page content is the Oracle Public Yum Server landing page. It features a penguin logo, the title "Public Yum Server", and a note that it was last updated on September 25, 2013. Below this, there's an "Introduction" section with information about Oracle Linux packages and Oracle VM installation media. It also mentions the Oracle Linux errata mailing list and support options. Under the "Getting Started" section, there are two bullet points: "Download and Install Oracle Linux" and "Download and copy the appropriate yum configuration file in place, by running the following commands as root". For Oracle Linux 4, 5, and 6, specific command examples are provided. The Oracle Linux 6 section includes the following commands:

```
cd /etc/yum.repos.d
mv Oracle-Base.repo Oracle-Base.repo.disabled
wget http://public-yum.oracle.com/public-yum-el6.repo
```

For Oracle Linux 5:

```
cd /etc/yum.repos.d
wget http://public-yum.oracle.com/public-yum-el5.repo
```

For Oracle Linux 6:

```
Oracle Linux 6
```

- b. Scroll down to see the instructions for **Oracle Linux 6**, which are:

```
cd /etc/yum.repos.d
wget http://public-yum.oracle.com/public-yum-el6.repo
```

    - This is information only; you do not need to run these commands.
    - The `public-yum-el6.repo` file is already present on **host03**.
- c. Close the browser by selecting **File > Quit**.
  5. Configure **host03** to access the Public Yum Server repositories.
    - You are configuring your system to access the latest Oracle Linux 6 repository on the Public Yum Server, [public\_el6\_latest] and the latest Oracle Linux 6 UEK repositories, [public\_el6\_UEK\_latest] and [public\_el6\_UEKR3\_latest].

- a. Use the `ls` command to display the `/etc/yum.repos.d` directory.

```
ls /etc/yum.repos.d
iso.repo packagekit-mdeia.repo public-yum-ol6.repo
```

- Notice the `public-yum-ol6.repo` file already exists in this directory.

- b. Use the `vi` editor to edit the `public-yum-ol6.repo` file.

- Change `enabled=0` to `enabled=1` for `[public_ol6_latest]`.
- Change `enabled=0` to `enabled=$uek` for `[public_ol6_UEK_latest]`.
- Change `enabled=0` to `enabled=$uekr3` for `[public_ol6_UEKR3_latest]`.

```
cd /etc/yum.repos.d
vi public-yum-ol6.repo
[public_ol6_latest]
...
enabled=0 (old value)
enabled=1 (new value)
...
[public_ol6_UEK_latest]
...
enabled=0 (old value)
enabled=$uek (new value)
...
[public_ol6_UEKR3_latest]
...
enabled=0 (old value)
enabled=$uekr3 (new value)
...
```

- Leave all other repositories in this file disabled (`enabled=0`).

- c. Run the `yum repolist` command to list the configured repositories.

```
yum repolist
Loaded plugins: ...
... [Errno 14] PYCURL ERROR 6 - "Couldn't resolve host 'public-yum.oracle.com'"
...
repo id Repo name status
Myrepo ...
ol6_UEK_latest ...
ol6_latest ...
```

- Notice two things:

- You received an error because you are unable to resolve <http://public-yum.oracle.com>.
- You still have your local repository – `Myrepo` – enabled.

- d. To resolve the Public Yum Server, use the `export` command to set the `http_proxy` from the command line.

- Recall the proxy setting was obtained from `dom0`.

```
export http_proxy=http://ges-proxy.us.oracle.com:80
```

- e. To disable the local yum repository, use the `vi` editor to edit the `iso.repo` file and change `enabled=1` to `enabled=0` for [Myrepo].

```
vi /etc/yum.repos.d/iso.repo
[Myrepo]
...
enabled=0
...
```

- f. Run the `yum clean all` command to clean up various things that accumulate in the yum cache directory.

```
yum clean all
...
Cleaning up Everything
```

- g. Run the `yum repolist` command to list the configured repositories.

- This command takes a few minutes to complete.

```
yum repolist
...
repo id Repo name status
public_ol6_UEKR3_latest ...
public_ol6_latest ...
...
```

- Notice that you now have `public_ol6_UEKR3_latest` and `public_ol6_latest` included in the output.

6. Run the `yum update` command to display available updates. *Do not apply updates!*

- Answer `n` when asked, "Is this ok".

```
yum update
...
Transaction Summary
=====
Install ...
Upgrade ...
Total download size: ...
Is this ok [y/N] : n
Exiting on user Command
```

- You should see that a number of new packages are available to be installed.
- You should also see a number of packages with updates available.
- *Do not* update at this time. It takes a couple of hours to update.

7. Enable your local yum repository and disable all repositories in the `public-yum-ol6.repo` file.
  - a. Use the `vi` editor to edit the `public-yum-ol6.repo` file and change `enabled=0` for ALL repositories.

```
vi /etc/yum.repos.d/public-yum-ol6.repo
[public_ol6_latest]
...
enabled=0 (new value)

...
[public_ol6_UEK_latest]
...
enabled=0 (new value)

...
[public_ol6_UEKR3_latest]
...
enabled=0 (new value)
...
```

- b. To ensure all repositories are disabled, run the following `grep` command.

```
grep enabled /etc/yum.repos.d/public-yum-ol6.repo
enabled=0
enabled=0
enabled=0
...
...
```

- Ensure all “enabled” settings are 0.
- c. To enable the local yum repository, use the `vi` editor to edit the `iso.repo` file and change `enabled=0` to `enabled=1` for [Myrepo].

```
vi /etc/yum.repos.d/iso.repo
[Myrepo]
...
enabled=1
...
```

- d. Run the command `yum clean all`, to clean up various things which accumulate in the yum cache directory.

```
yum clean all
...
Cleaning up Everything
```

- e. Run the `yum repolist` command to list the configured repositories.

```
yum repolist
...
repo id Repo name status
Myrepo
...
```

- Notice that you now have only the Myrepo enabled.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## **Practices for Lesson 14: File Sharing**

**Chapter 14**

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practices for Lesson 14: File Sharing

---

### Practices Overview

In these practices, you:

- Configure an NFS server and client and mount an exported file system
- Use automounter to mount the virtual CD drive
- Configure and use an FTP server

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 14-1: Configuring an NFS Server and an NFS Client

---

### Overview

In this practice, you:

- Export a file system from **host03** VM and mount it on **host01** VM
- Ensure that the required package is installed and that services are running
- Use various NFS-related commands and files to share file systems using NFS

### Assumptions

- This practice is performed on **host01** and on **host03** VMs.
- Open a terminal window on each system.
- Log in as the `root` user on each system.
- The prompts in the solution section include either **host01** or **host03** to indicate which system to enter the command from.

### Tasks

1. Use the `rpm` command to verify that the `nfs-utils` package is installed on **host03**.

```
[host03]# rpm -qa | grep nfs-utils
nfs-utils-lib-...
nfs-utils-...
```

- In this example, the package is already installed along with the `nfs-utils-lib` package.
2. Use the `vi` editor and edit `/etc/exports` on **host03** to add an entry to export `/Dev` to all client systems.

```
[host03]# vi /etc/exports
/Dev *
```

3. Use the `service` command to verify that the required services are started on **host03**.

- a. Verify that the `rpcbind` service is started.

```
[host03]# service rpcbind status
rpcbind (pid ...) is running...
```

- In this example, the `rpcbind` service is running.

- b. Verify that the `nfs` service is started.

```
[host03]# service nfs status
rpc.svcgssd is stopped
rpc.mountd is stopped
nfsd is stopped
rpc.rquotad is stopped
```

- In this example, the services are not running and must be started.

- c. Start the `nfs` service and associated services.

```
[host03]# service nfs start
Starting NFS services: exportfs: No options for /Dev *:...
 [OK]
Starting NFS quotas: [OK]
Starting NFS mountd: [OK]
Starting NFS daemon: [OK]
Starting RPC idmapd: [OK]
```

- d. Verify that the `nfslock` service is started.

```
[host03]# service nfslock status
rpc.statd (pid ...) is running...
```

- In this example, the service is running.

4. Use the `chkconfig` command to configure the services to start at boot time.

```
[host03]# chkconfig rpcbind on
[host03]# chkconfig nfs on
[host03]# chkconfig nfslock on
```

5. Use the `showmount` command on `host03` to display exported file systems.

```
[host03]# showmount -e
Export list for host03.example.com:
/Dev *
```

- If the exported file system is not listed, restart the `nfs` service.
- Whenever a new entry is made to `/etc/exports`, restart the `nfs` service.

6. Use the `exportfs` command on `host03` to export `/Test` to all clients and allow read-write permission. Include the option to ignore `/etc/exports` entries.

```
[host03]# exportfs -i -o rw *:/Test
```

7. Use the `showmount` command on `host03` to display exported file systems.

```
[host03]# showmount -e
Export list for host03.example.com:
/Dev *
/Test *
```

- Notice that both exported file systems are listed.
- You do not need to restart the `nfs` service when using `exportfs`.

8. Use the `cat` command to view the contents of `/var/lib/nfs/etab` on `host03`.

```
[host03]# cat /var/lib/nfs/etab
/Dev * (ro, sync, wdelay, hide, nocrossmnt, secure, root_squash, ...
/Test * (rw, sync, wdelay, hide, nocrossmnt, secure, root_squash, ...
```

- Notice that both exported file systems are listed in this master export table.

- The `rpc.mountd` process reads this file when a client attempts to mount an NFS file system.
9. Use the `vi` editor to edit `/etc/exports` on **host03** and change the entry exporting `/Dev` to the following:

```
[host03]# vi /etc/exports
/Dev * (rw,no_root_squash)
```

- The `rw` option allows client systems to make changes to the file system.
- The `no_root_squash` option allows `root` users on client systems to retain `root` privileges on the file system.

10. Run the `exportfs -r` command on **host03**.

```
[host03]# exportfs -r
```

- This command re-exports the entries in `/etc/exports` and synchronizes `/var/lib/nfs/etab` with `/etc/exports`.

11. Use the `cat` command to view the contents of `/var/lib/nfs/etab` on **host03**.

```
[host03]# cat /var/lib/nfs/etab
/Dev * (rw,sync,wdelay,hide,nocrossmnt,secure,no_root_squash,
...)
```

- Notice the new options on the `/Dev` NFS share.
- Also notice that `/Test` is no longer listed.

12. Use the `showmount` command on **host03** to display exported file systems.

```
[host03]# showmount -e
Export list for host03.example.com:
/Dev *
```

- This confirms that the `/Test` file system is no longer exported.

13. Mount the exported `/Dev` file system on **host01**.

- a. Use the `mkdir` command to create a mountpoint named `/remote_dev` on **host01**.

```
[host01]# mkdir /remote_dev
```

- b. Use the `mount` command to mount the exported file system from **host03**, `/Dev`, with `rw` and `nosuid` options on the local mountpoint, `/remote_dev`.
- The `rw` option mounts the file system with read-write permissions.
  - The `nosuid` option does not allow `setuid` or `setgid` bits to take effect.

```
[host01]# mount -t nfs -o rw,nosuid host03:/Dev /remote_dev
mount.nfs: mount system call failed
```

- In this example, the `mount` command failed.
- It takes several seconds for the `mount` command to fail. Rather than wait for the command to fail, you can press `Ctrl + C` to abort the `mount` command.

- c. If the `mount` command failed, disable the firewall on **host03** by running the following command.

```
[host03]# service iptables stop
iptables: Flushing firewall rules: [OK]
iptables: Setting chains to policy ACCEPT: filter [OK]
iptables: Unloading modules: [OK]
```

- Firewalls, the `iptables` command, and other system security-related issues are covered in the Lesson 17 “Security Administration.”
- You can allow NFS connectivity without disabling the entire firewall, which is covered in Lesson 17.
- For the purposes of this practice, stop the `iptables` service on **host03** to allow **host01** to mount the exported file system from **host03**.

- d. Re-issue the `mount` command from **host01**, if it is necessary to mount the NFS share.

```
[host01]# mount -t nfs -o rw,nosuid host03:/Dev /remote_dev
```

- In this example, the `mount` command was successful.

- e. Run the `mount` command to view the mount information for the NFS share.

```
[host01]# mount | grep nfs
...
host03:/Dev on /remote_dev type nfs
(rw,nosuid,vers=4,addr=192.0.2.103,clientaddr=192.0.2.101)
```

- Notice the mount options (`rw,nosuid`), the NFS version (4), and the server and client IP addresses.

14. Verify that NFS file system is mounted with `rw` (read-write) permissions.

- a. On **host01**, use the `df` command to display the mounted file systems.

```
[host01]# df -h
Filesystem Size Used Avail Use% Mounted on
/dev/xvda2 5.8G 1.9G 3.7G 34% /
...
host03:/Dev 1008M 34M 924M 4% /remote_dev
```

- Notice that the `host03:/Dev` file system is mounted on the local file system `/remote_dev`.

- b. Use the `ls` command to list the contents of `/remote_dev` on **host01**.

```
[host01]# ls /remote_dev
lost+found
```

- c. Use the `ls` command to list the contents of `/Dev` on **host03**.

```
[host03]# ls /Dev
lost+found
```

- Notice the contents of `/Dev` on **host03** are the same as `/remote_dev` on **host01** because they are the same directories.

- d. From **host03**, use the `vi` command to create a file on `/Dev`. Enter some content in the file.

```
[host03]# vi /Dev/test
<enter some content>
```

- e. From **host01**, use the `vi` command to edit the file created from **host03**. Make some changes to the content.

```
[host01]# vi /remote_dev/test
<change the content>
```

- f. From **host03**, use the `cat` command to view the contents of the file in `/Dev`.

```
[host03]# cat /Dev/test
<content has changed>
```

- This again confirms `/Dev` on **host03** is the same as `/remote_dev` on **host01** and that the file system has read-write permissions.

## Practice 14-2: Using Automounter

---

### Overview

In this practice, you:

- Use automounter to mount the virtual CD drive
- Use the `-hosts` map to automount all exports from **host03**

### Assumptions

- This practice is performed on **host01** VM, but you are asked to run a single command from **dom0**.
- Open a terminal window on each system.
- Log in as the `root` user on each system.
- The prompts in the solution section include either **dom0** or **host01** to indicate which system to enter the command from.

### Tasks

1. From **dom0**, use the `grep` command to search for “cd” in the `vm.cfg` file for **host01**.

```
[dom0]# grep cd /ovs/running_pool/host01/vm.cfg
'file:/OVS/seed_pool/OracleLinux-R6-U5-Server-x86_64-
dvd.iso,hdc:cdrom,r'
```

- Notice that the `dvd.iso` image of Oracle Linux 6.5 is on the virtual `cdrom` drive.

All remaining commands are run from the **host01** VM.

2. Use the `rpm` command to verify that the `autofs` package is installed.

```
rpm -q autofs
autofs-5.0.5-88.0.1.el6.x86_64
```

- In this example, the package is already installed.

3. Use the `service` command to verify that the `autofs` service is running.

```
service autofs status
automount (pid ...) is running...
```

- In this example, the service is already running.

4. Use the `grep` command to search for “misc” in the `/etc/auto.master` master map-file.

```
grep misc /etc/auto.master
/misc /etc/auto.misc
...
```

- Notice that the `/misc` mount point is associated with the `/etc/auto.misc` map-file.

5. Use the `grep` command to search for “cd” in the `/etc/auto.misc` file.

```
grep cd /etc/auto.misc
cd -fstype=iso9600,ro,nosuid,nodev :/dev/cdrom
```

- Notice that the `cd` mount point is associated with the `/dev/cdrom` device.

6. Use the `df` command to display the mounted file systems.

```
df -h
Filesystem Size Used Avail Use% Mounted on
/dev/xvda2 5.7G 2.4G 3.1G 44% /
...
...
```

- Notice that the `cdrom` is not mounted.
7. Use the `ls /misc/cd` command, combining the `/misc` mountpoint from `/etc/auto.master` with the `cd` mountpoint from `/etc/auto.misc`.

```
ls /misc/cd
ls
EFI Packages ResilientStorage
EULA README-en RPM-GPG-KEY
eula.en_US README-en.html RPM-GPG-KEY-oracle
...
...
```

- Notice that this is the contents of the Oracle Linux 6.5 dvd-iso image.
8. Use the `df` command to display the mounted file systems.

```
df -h
Filesystem Size Used Avail Use% Mounted on
/dev/xvda2 5.7G 2.4G 3.1G 44% /
...
/dev/sr0 3.7G 3.7G 0 100% /misc/cd
```

- Notice that the `cdrom` was “automounted” simply by accessing it with the `ls` command.
9. Use the `grep` command to search for “net” in the `/etc/auto.master` master map file.

```
grep net /etc/auto.master
/net -hosts
...
...
```

- Notice that the `/net` mount point is associated with the `-hosts` map.
10. Use the `cd` command to change to the `/net/host03` directory to automount all exports from **host03**. Use the `ls` command to list the contents of the directory.

```
cd /net/host03
ls
Dev
```

- Notice that the `/net/host03` directory contains the exported file system from **host03**, `/Dev`.
11. Use the `ls` command to list the contents of `/net/host03/Dev` on **host01**.

```
ls /net/host03/Dev
lost+found test
```

## Practice 14-3: Configuring an FTP Server

### Overview

In this practice, you:

- Install the vsftpd server package on **host03** and start the service
- Install the ftp (client) package and test the setup

### Assumptions

You are the `root` user on **host03** VM.

### Tasks

1. Use the `yum` command to install the `vsftpd` package on **host03**.

```
yum install vsftpd
...
Transaction Summary
=====
Install 1 Package(s)
Total download size: 150 k
Installed size: 331 k
Is this ok [y/N]: y
...
Complete!
```

2. Use the `service` command to start the `vsftpd` service.

```
service vsftpd start
Starting vsftpd for vsftpd: [OK]
```

3. Use the `chkconfig` command to configure `vsftpd` to start at boot time.

```
chkconfig vsftpd on
```

4. Use the `yum` command to install the `ftp` (client) package on **host03**.

```
yum install ftp
...
Transaction Summary
=====
Install 1 Package(s)
Total download size: 57 k
Installed size: 95 k
Is this ok [y/N]: y
...
Complete!
```

5. Test the setup.

- Use the `ftp` command to connect to `localhost` and log in as anonymous.
  - Use any password.
- After connecting, run the `ls` command to display the contents of `/var/ftp`.
- Conclude the test by running the `quit` command to exit.

```
ftp localhost
Connected to localhost (127.0.0.1).
220 (vsFTPd 2.2.2)
Name (localhost:root): anonymous
331 Please specify the password.
Password: your_email_address
230 Login successful.
Remote system type is UNIX.
Using binary mode to transfer files.
ftp> ls
227 Entering Passive Mode (127,0,0,1,54,166).
150 Here comes the directory listing.
drwxr-xr-x 2 0 0 4096 Apr 29 2012 pub
226 Directory send OK.
ftp> quit
221 Goodbye.
```

- Notice that the contents of `/var/ftp` is a `pub` directory for anonymous users.

6. Use the `grep` command to display the `ftp` user information in `/etc/passwd`.

```
grep ftp /etc/passwd
ftp:x:14:50:FTP User:/var/ftp:/sbin/nologin
```

- In this example, the home directory of the `ftp` user is `/var/ftp`.

7. Ensure that the permissions on the home directory are set to 755. Change the settings if necessary.

```
ls -ld /var/ftp
drwxr-xr-x 3 root root 4096 Mar 10 11:01 /var/ftp
```

- In this example, the permissions are correct.

8. Copy the `/root/install.log` file to `/var/ftp/pub` and rename it as `test_file`.

```
cp /root/install.log /var/ftp/pub/test_file
```

- This file is used in the next practice.

## Practice 14-4: Downloading a File from an FTP Server

### Overview

In this practice, you install the `ftp` package on **host01** VM and download a file from the FTP server on **host03**.

### Assumptions

You are the `root` user on **host01** VM.

### Tasks

- From **host01**, use the `cd /misc/cd` command to automount the virtual `cdrom`, which contains Oracle Linux `dvd.iso` image.

```
cd /misc/cd
```

- Use the `cd` command to change to the `Packages` directory.

```
cd Packages
```

- Use the `ls` command to display the `ftp` package name.

```
ls ftp*
ftp-0.17-54.el6.x86_64.rpm
```

- This is an example only; your version may be different.

- Use the `rpm` command to install the `ftp` (client) package name displayed in step 3.

```
rpm -Uvh ftp-0.17-54.el6.x86_64.rpm
...
```

- The `ftp` package may already be installed on your system.

- Use the `cd` command to change back to your home directory.

```
cd
pwd
/root
```

- From **host01**, use the `ftp` utility to connect to the FTP server, **host03**, as anonymous user.

```
[host01]# ftp host03
Connected to localhost (192.0.2.103).
220 (vsFTPd 2.2.2)
Name (host03:root): anonymous
331 Please specify the password.
Password: your_email_address
230 Login successful.
Remote system type is UNIX.
Using binary mode to transfer files.
ftp>
```

7. Enter **help** or **?** to display a list of available commands.

```
ftp> help
```

```
...
```

8. Get help on some of the available commands, for example:

```
ftp> help get
get receive file
ftp> help put
put send one file
ftp> help mget
mget get multiple files
ftp> help cd
cd change remote working directory
ftp> help ls
ls list contents of remote directory
ftp> help !
! escape to the shell
```

9. Use the **cd** command to change to the **pub** directory.

```
ftp> cd pub
```

```
Directory successfully changed.
```

10. Use the **ls** command to list the contents of the **/var/ftp/pub** directory on the FTP server.

```
ftp> ls
227 Entering Passive Mode (192,0,2,103,26,76).
150 Here comes the directory list.
-rw-r--r-- 1 0 0 50204 May 22 22:51 test_file
226 Directory send OK.
```

- Notice that the **test\_file** exists in the **pub** directory on the FTP server.

11. Use the **!ls** command to list the contents of the local directory on **host01**.

```
ftp> !ls
anaconda-ks.cfg install.log install.log.syslog
...
```

12. Use the **get** command to download the **test\_file** file from the FTP server.

```
ftp> get test_file
local: test_file remote: test_file
227 Entering Passive Mode (192,0,2,103,75,78).
150 Opening BINARY mode data connection for test_file ...
226 Transfer complete.
51355 bytes received in ...
```

13. Use the `!ls` command to list the contents of the local directory on **host01**.

```
ftp> !ls
anaconda-ks.cfg install.log install.log.syslog test_file
...
```

- Notice that the `test_file` exists in the local directory on **host01**.

14. Use the `quit` command to exit `ftp`.

```
ftp> quit
221 Goodbye.
```

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 14-5: Configuring xinetd

---

### Overview

In this practice, you configure vsftpd to be controlled by xinetd. Stated another way, you configure the xinetd “superserver” to start vsftpd on demand.

### Assumptions

You are the `root` user on **host03** VM.

### Tasks

- On **host03**, configure vsftpd to run in normal mode.

- The vsftpd server can run in normal mode or it can run in stand-alone mode.
    - In stand-alone mode, vsftpd runs as a daemon and handles connections directly.
    - In normal mode, the xinetd daemon calls vsftpd each time a client tries to make a connection.
- Use the `grep` command to display the “listen” directive in the `/etc/vsftpd/vsftpd.conf` file.

```
grep listen /etc/vsftpd/vsftpd.conf
When "listen" directive is enabled, vsftpd runs in standalone
mode ...
listen=YES
...
```

- Notice that “listen” is enabled (set to YES), which causes vsftpd to run in stand-alone mode.
- Use the `vi` editor to edit `/etc/vsftpd/vsftpd.conf` and set the “listen” directive to NO.

```
vi /etc/vsftpd/vsftpd.conf
listen=NO
```

- Install the xinetd package (if necessary).

- Use the `rpm -qa` command to determine if the xinetd package is installed.

```
rpm -qa | grep xinetd
```

- No output indicates the package is not installed.

- b. Use the `yum` command to install the `xinetd` package on `host03`.

```
yum install xinetd
...
Transaction Summary
=====
Install 1 Package(s)
Total download size: 121 k
Installed size: 259 k
Is this ok [y/N]: y
...
Complete!
```

3. Install and configure a `xinetd` control file for `vsftpd`.

- You must install a `xinetd` control file to run `vsftpd` in normal mode.
- a. Use the `cp` command to copy the sample control file (in `/usr/share/doc/vsftpd*/vsftpd.xinetd`) to the `/etc/xinetd.d/` directory and rename the sample control file `vsftpd`.

```
cp /usr/share/doc/vsftpd*/vsftpd.xinetd /etc/xinetd.d/vsftpd
```

- b. Use the `vi` editor to edit `/etc/xinetd.d/vsftpd` and set the “`disable`” directive to `no`.

```
vi /etc/xinetd.d/vsftpd
disable = no
```

4. Stop the `vsftpd` service and start the `xinetd` service.

- a. Use the `service` command to stop the `vsftpd` service.

```
service vsftpd stop
Shutting down vsftpd: [OK]
```

- b. Use the `service` command to start the `xinetd` service.

```
service xinetd start
Starting xinetd: [OK]
```

- c. Use the `chkconfig` command to display the `xinetd`-based services.

```
chkconfig --list
NetworkManager 0:off 1:off 2:on 3:on 4:on 5:on 6:off
abrt-ccpp 0:off 1:off 2:off 3:on 4:off 5:on 6:off
...
xinetd based services:
 chargen-dgram: off
 chargen-stream: off
...
 time-stream: off
 vsftpd: on
```

- Notice that vsftpd is the only enabled xinetd-based service.
  - Also notice that each xinetd-based service has a configuration file in the /etc/xinetd.d directory
- d. Use the ls command to list the contents of the /etc/xinetd.d directory

```
ls -l /etc/xinetd.d
-rw----- ... chargen-dgram
-rw----- ... chargen-stream
...
-rw-r--r-- ... vsftpd
```

5. Test the xinetd-based service (vsftpd).

- a. From host01, use the ftp utility to connect to the FTP server, host03, as the anonymous user.

```
[host01]# ftp host03
Connected to localhost (192.0.2.103).
220 (vsFTPd 2.2.2)
Name (host03:root): anonymous
331 Please specify the password.
Password: your_email_address
230 Login successful.
Remote system type is UNIX.
Using binary mode to transfer files.
ftp>
```

- b. Use the cd command to change to the pub directory.

```
ftp> cd pub
Directory successfully changed.
```

- c. Use the ls command to list the contents of the /var/ftp/pub directory on the FTP server.

```
ftp> ls
227 Entering Passive Mode (192,0,2,103,26,76).
150 Here comes the directory list.
-rw-r--r-- 1 0 0 50204 May 22 22:51 test_file
226 Directory send OK.
```

- Notice that the test\_file exists in the pub directory on the FTP server.

- d. Use the !ls command to list the contents of the local directory on host01.

```
ftp> !ls
anaconda-ks.cfg install.log install.log.syslog test_file
...
```

- Notice that the test\_file exists in the local directory on host01.

- e. Use the `get` command to download the `test_file` file from the FTP server. Because one copy of `test_file` already exists, include a second argument—`tf2`—to rename the file.

```
ftp> get test_file tf2
local: tf2 remote: test_file
227 Entering Passive Mode (192,0,2,103,149,8).
150 Opening BINARY mode data connection for test_file ...
226 Transfer complete.
50080 bytes received in ...
```

- f. Use the `!ls` command to list the contents of the local directory on **host01**.

```
ftp> !ls
anaconda-ks.cfg install.log install.log.syslog test_file tf2
...
```

- Notice that the `tf2` file exists in the local directory on **host01**.

- g. Use the `quit` command to exit `ftp`.

```
ftp> quit
221 Goodbye.
```

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## **Practices for Lesson 15: OpenSSH**

**Chapter 15**

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practices for Lesson 15: OpenSSH

---

### Practices Overview

In these practices, you do the following:

- You verify that the OpenSSH packages are installed and that the `sshd` service is running.
- You use the `ssh` and `scp` utilities.
- You use the `ssh-keygen` utility to generate keys enabling connectivity without supplying a password.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a  
non-transferable license to use this Student Guide.

## Practice 15-1: Connecting to a Remote Server by Using ssh

---

### Overview

In this practice, you verify that the OpenSSH packages are installed, verify that the `sshd` service is started on the server, and use the `ssh` utility to establish a connection and execute a command on a remote system.

### Assumptions

- This practice is performed on **host01** and **host03** VMs.
- Open a terminal window on each system.
- Log in as the `root` user on each system.
- The prompts in the solution section include either **host01** or **host03** to indicate which system to enter the command from.

### Tasks

1. On **host03**, use the `rpm` command to verify that the `openssh` packages are installed.

```
[host03]# rpm -qa | grep openssh
openssh-clients-...
openssh-server-...
openssh-...
```

- In this example, the packages are already installed.
2. Use the `service` command to verify that the `sshd` service is started on **host03**.
- ```
[host03]# service sshd status
openssh-daemon (pid ...) is running...
```
- In this example, the service is running.
3. Use the `chkconfig` command to verify that the `sshd` service is configured to start at boot time on **host03**.
- ```
[host03]# chkconfig sshd --list
sshd 0:off 1:off 2:on 3:on 4:on 5:on 6:off
```
- In this example, the service is configured to start at boot time.

4. On **host01**, use the `rpm` command to verify that the `openssh` packages are installed.

```
[host01]# rpm -qa | grep openssh
openssh-clients-...
openssh-server-...
openssh-...
```

- In this example, the packages are already installed.
5. On **host01**, use the `useradd` and `passwd` commands to create the `oracle` user. Assign a password of `oracle`. Ignore the “BAD PASSWORD” warning.

```
[host01]# useradd oracle
[host01]# passwd oracle
```

```
Changing password for user oracle.
New password: oracle
...
Retype new password: oracle
passwd: all authentication tokens updated successfully.
```

6. Use the `exit` command to log off as root on **host01**. Log back on as user oracle.

```
[host01]# exit
host01 login: oracle
Password: oracle
```

7. Use the `ls` command to display a long listing of all files in the home directory of user oracle.

```
[oracle@host01 ~]$ ls -la
...
```

- Notice that there is no `~/.ssh` directory.
8. Perform a remote login to **host03** by using the `ssh` command.
- Answer “yes” when asked, “Are you sure.”
  - The oracle user password on **host03** is oracle.

```
[oracle@host01 ~]$ ssh host03
The authenticity of host 'host03 (192.0.2.103)' can't be
established. RSA key fingerprint is ...
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'host03,192.0.2.103' (RSA) to the
list of known hosts.
oracle@host03's password: oracle
[oracle@host03 ~]$
```

9. Use the `hostname` command to display the host name to confirm that you successfully logged on to **host03**.

```
[oracle@host03 ~]$ hostname
host03.example.com
```

10. Use the `logout` command to close the `ssh` connection to **host03**. Use the `hostname` command to confirm that you are back to **host01**.

```
[oracle@host03 ~]$ logout
Connection to host03 closed.
[oracle@host01 ~]$ hostname
host01.example.com
```

11. Use the `ls` command to display a long listing of all files in the home directory of user oracle.

```
[oracle@host01 ~]$ ls -la
...
drwx----- 2 oracle oracle 4096 Mar 8 08:23 .ssh
```

- Notice that there is now a `~/.ssh` directory.
12. Use the `cd` command to change to the `~/.ssh` directory, and then use `ls` to view the contents of the directory.

```
[oracle@host01 ~]$ cd .ssh
[oracle@host01 .ssh]$ ls
known_hosts
```

- Notice that the `known_hosts` file was created.
13. Use the `cat` command to view the contents of the `known_hosts` file.

```
[oracle@host01 .ssh]$ cat known_hosts
host03,192.0.2.103 ssh-rsa ...
```

- Notice that `host03` is now a “known host”.
14. Perform a remote login to `host03` using the `ssh` command.

```
[oracle@host01 .ssh]$ ssh host03
oracle@host03's password: oracle
Last login...
[oracle@host03 ~]$
```

- Notice that you are not asked to confirm this time, because of the existence of the `known_hosts` file.
15. Use the `logout` command to close the `ssh` connection to `host03`. Use the `hostname` command to confirm that you are back to `host01`. Use the `cd` command to change back to the user’s home directory.

```
[oracle@host03 ~]$ logout
Connection to host03 closed.
[oracle@host01 .ssh]$ hostname
host01.example.com
[oracle@host01 .ssh]$ cd
[oracle@host01 ~]$
```

16. Log on to `host03` as user `root` and run the `ls` command with a single `ssh` command.

- The `root` user password on `host03` is `oracle`.
- ```
[oracle@host01 ~]$ ssh root@host03 ls  
root@host03's password: oracle  
...  
[oracle@host01 ~]$
```

- Notice that the `ls` command ran on the remote system displaying the contents of the remote directory, and then the remote connection closed.

17. Use the `hostname` command to confirm that you are back to `host01`.

```
[oracle@host01 ~]$ hostname  
host01.example.com
```

Practice 15-2: Configuring OpenSSH to Connect Without a Password

Overview

In this practice, you use the `ssh-keygen` command to generate an RSA key pair and configure OpenSSH to connect to a remote system without supplying a password. You also use the `scp` command in this practice.

Assumptions

- This practice is performed on **host01** and on **host03** VMs.
- Open a terminal window on each system.
- Log in as the `root` user on **host03**.
- Log in as the `oracle` user on **host01**.
- The prompts in the solution section include either **host01** or **host03** to indicate which system to enter the command from.

Tasks

1. From **host03**, use the `su` command to become the `oracle` user.

```
[root@host03 ~]# su - oracle
[oracle@host03 ~]$ whoami
oracle
[oracle@host03 ~]$ pwd
/home/oracle
```

2. Use the `ls` command to view the contents of the `~/.ssh` directory.

```
[oracle@host03 ~]$ ls ~/.ssh
```

- Notice that the directory is empty.

3. Use the `ssh-keygen` command to create the public and private parts of an RSA key. Accept all the defaults.

```
[oracle@host03 ~]$ ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key (/home/oracle/.ssh/id_rsa):
ENTER
Enter passphrase (empty for no passphrase): ENTER
Enter same passphrase again: ENTER
Your identification has been saved in /home/oracle/.ssh/id_rsa.
Your public key has been saved in /home/oracle/.ssh/id_rsa.pub.
The key fingerprint is:...
The key's randomart image is:...
```

4. Use the `ls` command to view the contents of the `~/.ssh` directory.

```
[oracle@host03 ~]$ ls ~/.ssh
id_rsa      id_rsa.pub
```

- Notice that the `ssh-keygen` command generated two keys.

5. Use the `scp` command to copy `~/.ssh/id_rsa.pub` on the local system (**host03**) to `~/.ssh/authorized_keys` on the remote system (**host01**). Password is `oracle`.

```
[oracle@host03 ~]$ scp .ssh/id_rsa.pub  
host01:~/.ssh/authorized_keys  
The authenticity of host 'host01 (192.0.2.101)' can't be  
established. RSA key fingerprint is ...  
Are you sure you want to continue connecting (yes/no)? yes  
Warning: Permanently added 'host01,192.0.2.101' (RSA) to the  
list of known hosts.  
oracle@host01's password: oracle  
id_rsa.pub 100% 407 0.4KB/s 00:00  
[oracle@host03 ~]$
```

- Because you are connecting to this OpenSSH server for the first time, you are asked to confirm the connection.
 - Notice that a password is required to make the connection.
 - Notice that the file is copied but you are still connected to the local system (**host03**).
6. Perform a remote login to **host01** using the `ssh` command.

```
[oracle@host03 ~]$ ssh host01  
Last login:...
```

- Notice that you no longer need to enter a password.
7. Use the `hostname` command to confirm that you successfully logged on to **host01**.

```
[oracle@host01 ~]$ hostname  
host01.example.com
```

8. Use the `ls` command to view the contents of the `~/.ssh` directory.

```
[oracle@host01 ~]$ ls ~/.ssh  
authorized_keys known_hosts
```

- Notice the existence of the `authorized_keys` file, which allowed you to connect without supplying a password.
9. Use the `logout` command to close the connection to **host01**. Use the `hostname` command to confirm that you are back to **host03**.

```
[oracle@host01 ~]$ logout  
Connection to host01 closed.  
[oracle@host03 ~]$ hostname  
host03.example.com
```

10. Use the `exit` command to log out as `oracle` user and return to the `root` logon. Use the `whoami` command to confirm that you are logged on as `root`.

```
[oracle@host03 ~]$ exit  
logout  
[root@host03 ~]$ whoami  
root
```

Practices for Lesson 16: Pluggable Authentication Modules (PAM)

Chapter 16

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 16: Pluggable Authentication Modules (PAM)

Practices Overview

In these practices, you configure PAM authentication modules first to allow a single login only, and then to disable all non-root logins.

Practice 16-1: Configuring PAM for a Single Login Session

Overview

In this practice, you configure a PAM authentication module on **host03** to allow only a single login session for a user.

Assumptions

- This practice is performed on **host01** and **host03** VMs.
- You open a terminal window on each system.
- You log in as the `root` user on **host03**.
- The prompts in the solution section include either **host01** or **host03** to indicate which system to enter the command from.

Tasks

1. On **host03**, view PAM configuration files and directories.

- a. Use the `ls` command to view the PAM configuration directory, `/etc/pam.d`.

```
[host03]# ls /etc/pam.d
atd          login      sshd
authconfig   newrole   ssh-keycat
...
...
```

- This directory contains files that describe the authentication procedure for an application.

- b. Use the `cat` command to view the `sshd` configuration file in `/etc/pam.d`.

```
[host03]# cat /etc/pam.d/sshd
#%PAM-1.0
auth      required    pam_sepermit.so
auth      include     password-auth
...
...
```

- This file contains a group of directives that define the authentication modules as well as any controls or arguments.
- The authentication modules are listed in the third column.

- c. Use the `find` command to locate the `pam_sepermit.so` authentication module.

```
[host03]# find / -name pam_sepermit.so
/lib64/security/pam_sepermit.so
```

- In this example, the authentication module is located in `/lib64/security`.

- d. Use the `ls` command to view the authentication modules directory.

```
[host03]# ls /lib64/security
pam_access.so      pam_krb5afs.so      pam_sepermit.so
pam_cap.so        pam_krb5.so        pam_shells.so
...
...
```

- Notice that all authentication modules are located in this directory.
2. On **host03**, view the man pages for the `pam_sepermit` authentication module and the associated configuration file.
- Most of the authentication modules have a man page describing their purpose and usage. Use the `man pam_sepermit` command to view the man page for the `pam_sepermit` authentication module.

```
[host03]# man pam_sepermit
...
pam_sepermit - PAM module to allow/deny login depending
On SELinux enforcement state
...
When the user which is logging in matches an entry in the
config file he is allowed access only when the SELinux
is in enforcing mode. Otherwise he is denied access...
...
See sepermit.conf(5) for details.
...
```

- Notice that this module uses a config file, `sepermit.conf`, which controls access when SELinux is in enforcing mode.
 - SELinux stands for “Security-Enhanced Linux” and is covered in a later course.
- b. Use the `man sepermit.conf` command to view the man page for the `sepermit.conf` file.

```
[host03]# man sepermit.conf
...
sepermit.conf - configuration file for the pam_sepermit
module
...
The lines of the configuration file have the following
syntax:
...
```

3. SELinux is covered in a later course but for the purposes of this practice, use the `sestatus` command to display information about SELinux.

```
[host03]# sestatus
SELinux status:      enabled
...
Current mode:        enforcing
...
```

- Output shown is a sample showing that SELinux is enabled and is in enforcing mode.
- With SELinux in enforcing mode, the `pam_sepermit` authentication module allows or denies login.

4. From **host01**, confirm you can remotely log in to host03.
 - a. From the login prompt, log in as **oracle** user. Password is **oracle**.

```
host01 login: oracle
Password: oracle
Last login...
[oracle@host01 ~]$
```

- b. Use the **ssh** command to connect to **host03**. Password is **oracle**.

```
[oracle@host01 ~]$ ssh host03
oracle@host03's password: oracle
Last login:...
[oracle@host03 ~]$
```

- c. Use the **hostname** command to confirm that you successfully logged in to **host03**.

```
[oracle@host03 ~]$ hostname
host03.example.com
```

- Notice that you are successfully able to log in to **host03**.

- d. Use the **logout** command to close the connection to **host03**.

```
[oracle@host03 ~]$ logout
Connect to host03 closed.
[oracle@host01 ~]$ hostname
host01.example.com
```

- Notice that you are now logged off of **host03** and back to **host01**.

5. On **host03**, configure the **pam_sepermit** authentication module to deny login.

- a. Use the **find** command to locate the **sepermit.conf** file.

```
[host03]# find / -name sepermit.conf
/etc/security/sepermit.conf
```

- Notice that the **sepermit.conf** file is located in the **/etc/security** directory.

- b. Use the **vi** editor to add the following entry to **/etc/security/sepermit.conf**.

```
[host03]# vi /etc/security/sepermit.conf
oracle:exclusive
```

- This entry, when read by the PAM module **pam_sepermit.so**, allows only a single login session for the **oracle** user.

6. From **host01**, attempt to log in to **host03**.

Use the **ssh** command to connect to **host03**. Password is **oracle**.

```
[oracle@host01 ~]$ ssh host03
oracle@host03's password: oracle
Permission denied, please try again.
oracle@host03's password: CTRL-C
[oracle@host01 ~]$
```

- Notice that the connection is denied by the PAM authentication module.

To permit the `oracle` user login from **host01**, you can either:

- Remove the entry in the `/etc/pam.d/sshd` file to use the `pam_sepermit.so` module.
- Remove the entry in the `/etc/security/sepermit.conf` file to allow only a single login session.

7. From **host03**, permit user `oracle` to log in from **host01** by using the `vi` editor to comment out the entry to use the `pam_sepermit.so` module from the `/etc/pam.d/sshd` file.
- Comment out this line by inserting a # sign at the beginning of the line as follows:

```
[host03]# vi /etc/pam.d/sshd
auth    required    pam_sepermit.so          (current entry)
#auth   required    pam_sepermit.so          (insert # sign)
```

8. From **host01**, attempt to log in to **host03**.

- a. Use the `ssh` command to connect to **host03**. Password is `oracle`.

```
[oracle@host01 ~]$ ssh host03
oracle@host03's password: oracle
Last login...
[oracle@host03 ~]$ hostname
host03.example.com
```

- Notice that the connection is allowed, and no longer denied by the PAM authentication module.

- b. Use the `logout` command to close the connection to **host03**.

```
[oracle@host03 ~]$ logout
Connect to host03 closed.
[oracle@host01 ~]$ hostname
host01.example.com
```

- Notice that you are now logged off of **host03** and back to **host01**.

- c. From **host01**, log out as `oracle` user.

```
[oracle@host01 ~]$ logout
...
host01 login:
```

9. Return **host03** back to the original state.

- a. From **host03**, use the `vi` editor to edit `/etc/pam.d/sshd` and uncomment the entry to use the `pam_sepermit.so` module (remove the # sign).
- b. From **host03**, use the `vi` editor to edit `/etc/security/sepermit.conf` and remove the entry to allow only a single login for user `oracle`.

Practice 16-2: Configuring PAM to Prevent Non-root Login

Overview

In this practice, you configure a PAM authentication module on **host01** to prevent all non-root user logins.

Assumptions

- This practice is performed on **host01** and **host03** VMs.
- Open a terminal window on each system.
- Log in as the **root** user on **host01**.
- The prompts in the solution section include either **host01** or **host03** to indicate which system to enter the command from.

Tasks

1. On **host01**, configure a PAM authentication module on **host01** to prevent all non-root user logins.

- a. Log in to **host01** as **root**. Password is **oracle**.

```
host01 login: root
Password: oracle
Last login: ...
[root@host01] #
```

- b. Use the `cat` command to view the `login` configuration file in `/etc/pam.d`.

```
# cat /etc/pam.d/login
#%PAM-1.0
auth [user_unknown=ignore success=ok ignore=ignore ...
auth      include      system-auth
account   required    pam_nologin.so
...
```

- The `login` utility uses the `pam_nologin.so` authentication module as well as several other modules.
- c. Use the `man pam_nologin` command to view the man page for the `pam_nologin` authentication module.

```
# man pam_nologin
...
pam_nologin - Prevent non-root users from login
...
pam_nologin is a PAM module that prevents users from
logging into the system when /etc/nologin exists. The
contents of the /etc/nologin file are displayed to the
user...no effect on the root user's ability to log in.
...
```

- Notice that this module uses a configuration file – /etc/nologin – which, if it exists, disables non-root logins.
- d. Use the vi editor and create the /etc/nologin file with the following contents:

```
# vi /etc/nologin
No logins allowed at this time.
```

2. From **host03**, attempt to log in to **host01**.

- a. Use the ssh command to connect to **host01** as user oracle.

```
[host03]# ssh oracle@host01
oracle@host01's password: oracle
No logins allowed at this time.
Connection closed by 192.0.2.101
```

- Notice that the connection is denied by the PAM authentication module.

To permit the non-root user logins, you can do either of the following:

- Delete the /etc/nologin file from **host01**.
- Remove the entry in the /etc/pam.d/login file to use the pam_nologin.so module.

3. From **host01**, permit non-root user logins from **host03** by using the vi editor to comment out the entry to use the pam_nologin.so module from the /etc/pam.d/login file.

- Comment out this line by inserting a # sign at the beginning of the line as follows:

```
# vi /etc/pam.d/login
...
account required pam_nologin.so          (current entry)
#account required pam_nologin.so         (insert # sign)
```

4. From **host03**, attempt to log in to **host01**.

Use the ssh command to connect to **host01** as user oracle.

```
[host03]# ssh oracle@host01
oracle@host01's password: oracle
No logins allowed at this time.
Connection closed by 192.0.2.101
```

- Notice that the connection is still denied by the PAM authentication module.

5. From **host01**, use the grep command to search for the string “pam_nologin” in all the files in the /etc/pam.d directory.

```
[host01]# grep pam_nologin /etc/pam.d/*
/etc/pam.d/login:#account required pam_nologin.so
/etc/pam.d/remote:account required pam_nologin.so
/etc/pam.d/sshd:account required pam_nologin.so
```

- Notice that this module also is called from the remote file and from the sshd file.
- You would need to comment out these lines in these two files as well.
- Alternatively, simply remove the /etc/nologin file to allow non-root logins.

6. Return **host01** back to the original state.

- a. Use the `rm` command to remove the `/etc/nologin` file.

```
# rm /etc/nologin
rm: remove regular file '/etc/nologin'? y
```

- b. Use the `vi` editor to edit `/etc/pam.d/login` and uncomment the entry to use the `pam_nologin.so` module (remove the `#` sign).

```
# vi /etc/pam.d/login
...
#account    required    pam_nologin.so          (current entry)
account    required    pam_nologin.so          (remove # sign)
```

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 17: Security Administration

Chapter 17

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 17: Security Administration

Practices Overview

In these practices, you:

- Create chroot jails
- Configure iptables from the GUI and the command line
- Configure a TCP wrapper

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practice 17-1: Configuring a chroot Jail

Overview

In this practice, you configure a chroot jail and copy all files required to run the /bin/bash shell in the chroot jail.

Assumptions

You are the oracle user on **host03** VM.

Tasks

1. Use the whoami command to confirm that you are logged in as the oracle user.

```
$ whoami  
oracle
```

- If you are the root user, enter exit or logout to become the oracle user.

2. Use the cd command to change to your home directory.

```
$ cd  
$ pwd  
/home/oracle
```

3. Use the mkdir command to make a jail directory in your home directory.

```
$ mkdir jail
```

4. Use the su command to become the root user. Password is oracle.

```
$ su -  
Password: oracle
```

5. Use the echo command to display the value of the SHELL variable.

```
# echo $SHELL  
/bin/bash
```

- In this example, SHELL=/bin/bash.

6. As the root user, use the chroot command to create a chroot jail in the /home/oracle/jail directory.

```
# chroot /home/oracle/jail  
chroot: failed to run command '/bin/bash': No such file or  
directory
```

- Notice that if you do not specify a command as an argument, chroot attempts to run the value of the SHELL variable, /bin/bash, in the chroot jail directory, /home/oracle/jail.
- The command failed because /bin/bash was not found in /home/oracle/jail.

7. Use the `exit` command to log off as `root` and return to `oracle` user login.

```
# exit  
logout  
$ whoami  
oracle
```

8. As the `oracle` user, create a new root directory structure and copy all necessary files into the new root directory.

- a. Within the `jail` directory in your home directory, use the `mkdir` command to make a `bin` directory.

```
$ cd ~/jail  
$ mkdir bin
```

- b. Use the `cp` command to copy `/bin/bash` into `~/jail/bin`.

```
$ cp /bin/bash ~/jail/bin
```

- c. Use the `ldd` command to determine which shared libraries are required by `/bin/bash`.

```
$ ldd /bin/bash  
linux-vdso.so.1 => (0x0000...)  
libtinfo.so.5 => /lib64/libtinfo.so.5 (0x0000...)  
libdl.so.2 => /lib64/libdl.so.2 (0x0000...)  
libc.so.6 => /lib64/libc.so.6 (0x0000...)  
/lib64/ld-linux-x86-64.so.2 (0x0000...)
```

- In this example, there are four shared library files in `/lib64` used by `/bin/bash`.

- d. Within the `jail` directory in your home directory, use the `mkdir` command to make a `lib64` directory.

```
$ cd ~/jail  
$ mkdir lib64
```

- e. Use the `cp` command to copy the four shared library files required for `/bin/bash` from `/lib64` to `~/jail/lib64`.

```
$ cp /lib64/libtinfo.so.5 ~/jail/lib64  
$ cp /lib64/libdl.so.2 ~/jail/lib64  
$ cp /lib64/libc.so.6 ~/jail/lib64  
$ cp /lib64/ld-linux-x86-64.so.2 ~/jail/lib64
```

9. Use the `su` command to become the `root` user. Password is `oracle`.

```
$ su -  
Password: oracle
```

10. As the `root` user, use the `chroot` command to create a `chroot` jail in the `/home/oracle/jail` directory.

```
# chroot /home/oracle/jail
```

- Notice that the `chroot` command was successful—no errors occurred and the `/bin/bash` program executed.

11. Use the `pwd` command to display the current directory.

```
# pwd  
/
```

- Notice that the output indicates that the current directory is the root-level directory even though the actual directory is `/home/oracle/jail`.

12. Use the `exit` command to exit the `chroot` jail.

```
# exit
```

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practice 17-2: Configuring a chroot Jail for ftp Users

Overview

In this practice, you:

- Confirm that anonymous ftp users are placed in a chroot jail on a vsftpd server by default
- Configure an SELinux Boolean to allow local ftp users
- Configure local users to be placed in a chroot jail

Assumptions

- This practice is performed on **host01** and **host03** VMs.
- Open a terminal window on each system.
- Log in as the `root` user on each system.
- The prompts in the solution section include either **host01** or **host03** to indicate which system to enter the command from.
- SELinux is covered in a later course but some functionality is used in this practice.

Tasks

1. Confirm that anonymous ftp users are placed in a chroot jail by default.
 - a. From **host03**, use the `ls` command to list the contents of the `/var/ftp` directory.


```
[host03]# ls -l /var/ftp
drwxr-xr-x... pub
```

 - Notice that `/var/ftp` contains a single directory, `pub`.
 - b. From **host03**, use the `ls` command to list the contents of `/var/ftp/pub`.


```
[host03]# ls /var/ftp/pub
test_file
```

 - Notice that the `/var/ftp/pub` directory contains a single file, `test_file`.
 - c. From **host01**, `ftp` to **host03** as anonymous user. Press Enter when prompted for a password.

```
[host01]# ftp host03
Connected to host03: (192.0.2.103).
220 (vsFTPD 2.2.2)
Name....: anonymous
331 Please specify the password.
Password: ENTER
230 Login successful.
Remote system type is UNIX.
Using binary mode to transfer files.
ftp>
```

- d. Use the `ls` command to list the contents of the current directory.

```
ftp> ls
...
drwxr-xr-x ... pub
226 Directory send OK.
ftp>
```

- Notice that the current directory contains a single directory, `pub`.

- e. Use the `ls` command to list the contents of the `pub` directory.

```
ftp> ls pub
...
-rw-r--r-- ... test_file
226 Directory send OK.
ftp>
```

- Notice that the `pub` directory contains a single file, `test_file`.
- This confirms that the current location of the anonymous FTP user is `/var/ftp` on **host03**.

- f. Use the `pwd` command to display the current directory.

```
ftp> pwd
257 "/"
ftp>
```

- Notice that the output indicates that the current directory is the root-level directory even though the actual directory is `/var/ftp`.
- This confirms that anonymous users are placed in a chroot jail by default.

- g. Use the `quit` command to exit `ftp`.

```
ftp> quit
221 Goodbye.
```

2. From **host01**, attempt to `ftp` to **host03** as the `oracle` user. Password is `oracle`.

- a. Use the `ftp host03` command to connect as the `oracle` user. Password is `oracle`.

```
[host01]# ftp host03
Connected to host03: (192.0.2.103).
220 (vsFTPd 2.2.2)
Name...: oracle
331 Please specify the password.
Password: oracle
500 OOPS: cannot change directory:/home/oracle
Login failed.
ftp>
```

- Notice that login failed for the `oracle` user.

- b. Use the `quit` command to exit `ftp`.

```
ftp> quit  
221 Goodbye.
```

3. On **host03**, use the `setenforce` command to change SELinux to permissive mode. Use the `getenforce` command before and after to confirm the change.

```
[host03]# getenforce  
Enforcing  
[host03]# setenforce 0  
[host03]# getenforce  
Permissive
```

4. From **host01**, attempt to `ftp` to **host03** as the `oracle` user. Password is `oracle`.

- a. From **host01**, `ftp` to **host03** as the `oracle` user. Password is `oracle`.

```
[host01]# ftp host03  
Connected to host03: (192.0.2.103).  
220 (vsFTPd 2.2.2)  
Name...: oracle  
331 Please specify the password.  
Password: oracle  
230 Login successful.  
...  
ftp>
```

- Notice that login was successful this time.
- Setting SELinux to permissive mode allows the `oracle` user to connect to the FTP server.

- b. Use the `quit` command to exit `ftp`.

```
ftp> quit  
221 Goodbye.
```

5. From **host03**, determine which SELinux Boolean(s) denied the `oracle` user connection to the FTP server.

- a. Run the `yum install` command to install the `policycoreutils-python` package.
- This package provides the `semanage` command.

```
# yum install policycoreutils-python  
...  
Transaction Summary  
=====  
Install       6 Package(s)  
Total download size: 1.2 M  
Installed size: 4.4 M  
Is this ok [y/N]: y
```

```
...  
Complete!
```

- b. Use the `semanage` command to list FTP-related Booleans.

```
# semanage boolean -l | grep ftp  
ftp_home_dir          -> off      Allow ftp to read and write  
files in the user home directories  
...
```

- Notice that the `ftp_home_dir` Boolean is set to `off`.
- This Boolean seems to be the cause of the `oracle` user connection denial.

- c. Use the `setenforce` command to change SELinux to enforcing mode. Use the `getenforce` command to confirm the change.

```
# setenforce 1  
# getenforce  
Enforcing
```

- d. Use the `setsebool` command to set the `ftp_home_dir` Boolean to `on`.

```
# setsebool ftp_home_dir on
```

- e. Use the `getsebool` command to confirm that the current status of the `ftp_home_dir` Boolean is `on`.

```
# getsebool ftp_home_dir  
ftp_home_dir --> on
```

6. From `host01`, attempt to `ftp` to `host03` as the `oracle` user. Password is `oracle`.

- a. From `host01`, `ftp` to `host03` as the `oracle` user. Password is `oracle`.

```
[host01]# ftp host03  
Connected to host03: (192.0.2.103).  
220 (vsFTPd 2.2.2)  
Name....: oracle  
331 Please specify the password.  
Password: oracle  
230 Login successful.  
...  
ftp>
```

- Notice that login was successful this time.
- Setting the `ftp_home_dir` SELinux Boolean to `on` allows the `oracle` user to connect to the FTP server even when SELinux policy is enforced.

- b. Use the `pwd` command to display the current directory. Use the `cd` command to move up one level in the directory structure.

```
ftp> pwd  
257 "/home/oracle"  
ftp> cd ..  
250 Directory successfully changed.  
ftp> pwd  
257 "/home"  
ftp>
```

- The output indicates that the `oracle` user was placed in its home directory.
- This is the case for all local users that access a `vsftpd` server; they are placed in their home directory by default, and not in a `chroot` jail.

- c. Use the `quit` command to exit `ftp`.

```
ftp> quit  
221 Goodbye.
```

7. On **host03**, enable options in the `/etc/vsftpd/vsftpd.conf` file to put local users in a `chroot` jail.

- a. Use the `vi` editor to make the following changes:

- Set the `chroot_local_user` directive to `YES`.
- Ensure the following `chroot` directives are commented out (preceded with a `#` sign, as shown):
 - `#chroot_list_enable=YES`
 - `#chroot_list_file=/etc/vsftpd/chroot_list`

```
[host03]# vi /etc/vsftpd/vsftpd.conf  
chroot_local_user=YES  
#chroot_list_enable=YES  
#chroot_list_file=/etc/vsftpd/chroot_list
```

- Normally, after making changes to the `vsftpd.conf` file, you would need to restart the `vsftpd` service.
- However, because you configured `vsftpd` to run under `xinetd` in Practice 14-5, you do not need to restart the `vsftpd` service.

8. From **host01**, attempt to **ftp** to **host03** as the **oracle** user. Password is **oracle**.

- a. From **host01**, **ftp** to **host03** as the **oracle** user. Password is **oracle**.

```
[host01]# ftp host03
Connected to host03: (192.0.2.103).
220 (vsFTPd 2.2.2)
Name...: oracle
331 Please specify the password.
Password: oracle
230 Login successful.
...
ftp>
```

- b. Use the **pwd** command to display the current directory. Use the **ls** command to list the contents of the current directory.

```
ftp> pwd
257 "/"
ftp> ls
...
drwxr-xr-x ... Desktop
drwxr-xr-x ... Documents
drwxr-xr-x ... Downloads
drwxr-xr-x ... Music
...
```

- The output indicates that the current directory is the root-level directory even though the actual directory is `/home/oracle` on the `vsftpd` server, **host03**.

- c. Use the **quit** command to exit **ftp**.

```
ftp> quit
221 Goodbye.
```

Practice 17-3: Configuring iptables

Overview

In this practice, you use the Firewall Configuration GUI as well as the `iptables` command to allow a client system to mount an NFS file system.

Assumptions

- This practice is performed on **host01** and **host03** VMs.
- Open a terminal window on each system.
- Log in as the `root` user on each system.
- The prompts in the solution section include either **host01** or **host03** to indicate which system to enter the command from.

Tasks

1. From **host01**, unmount NFS file systems.

- a. Use the `umount` command to unmount `/remote_dev`.

```
[host01]# umount /remote_dev
```

- b. Use the `df` command to display the mounted file systems.

```
[host01]# df -h
Filesystem      Size  Used  Avail   Use%  Mounted on
/dev/xvda2     5.7G  2.4G  3.1G   44%   /
...

```


- Notice that no NFS mounted file systems are currently mounted.

2. From **host03**, use the `system-config-firewall` command to display the Firewall Configuration GUI.

```
[host03]# system-config-firewall
```

- Click **Close** if the Firewall Configuration Startup message appears.

- The GUI appears.



- In this example, the GUI shows the firewall is enabled. The firewall on your system may be disabled.
- Enable the firewall.
 - If the **Disable** button is active, click **Disable**.
 - Click the **Enable** button.
 - Click the **Apply** button.
 - Click **Yes** to override any existing firewall configuration.
 - The firewall is now enabled.
 - Select **File > Quit** to close the GUI.
 - From **host03**, use the **service** command to view the status of the **iptables** service.

```
[host03]# service iptables status
Table: filter
Chain INPUT (policy ACCEPT)
num  target     prot opt    source        destination
 1    ACCEPT     all  --    0.0.0.0/0  0.0.0.0/0      state ...
...
Chain FORWARD (policy ACCEPT)
num  target     prot opt    source        destination
```

```

1      REJECT    all    --  0.0.0.0/0  0.0.0.0/0      reject...
...
Chain OUTPUT (policy ACCEPT)
num  target     prot opt   source       destination

```

- This shows that the iptables service is running.
5. From **host03**, use the `showmount -e` command to list the NFS exported file systems.

```
[host03]# showmount -e
/Dev          *
```

- In this example, one NFS file system is exported.
6. From **host01**, attempt to mount an NFS file system.

- Use the `mount` command to mount the exported /Dev NFS file system from **host03** with `rw` and `nosuid` options on the local mountpoint, `/remote_dev`.

```
[host01]# mount -t nfs -o rw,nosuid host03:/Dev /remote_dev
```

- With the firewall enabled, the `mount` command hangs.
- Press **Ctrl + C** to abort the `mount` command.

7. From **host03**, run the `iptables -L` command.

```
[host03]# iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt   source       destination
ACCEPT    all    --  0.0.0.0/0  0.0.0.0/0      state RELATED...
```

- Notice that the output is similar to output from the `service iptables status` command.
8. Run the `iptables -L` command again, but this time pipe the output to `grep` and search for “nfs”.

```
[host03]# iptables -L | grep nfs
```


- There are currently no rules containing the string “nfs”.
9. From **host03**, use the Firewall Configuration GUI to configure NFS4 as a trusted service.

- Use the `system-config-firewall` command to start the Firewall Configuration GUI.

```
[host03]# system-config-firewall
```

- The GUI appears.

- b. From the GUI, select the check box next to the NFS4 service, as shown.
- Scroll down if necessary to display the NFS4 check box.



- This allows the NFS services to be trusted, and not blocked by the firewall.
- c. Save the changes made from the GUI.
- 1) Click the **Apply** button.
 - 2) Click **Yes** to override any existing firewall configuration.
 - The change is now active.
- d. Select **File > Quit** to close the GUI.
10. From **host01**, attempt to mount **host03:/Dev** on **/remote_dev**.
- a. Use the `mount` command to mount **host03:/Dev** on **/remote_dev** using the options as shown:

```
[host01]# mount -t nfs -o rw,nosuid host03:/Dev /remote_dev
```

- The `mount` command is successful this time.

- b. Use the `df` command to display the mounted file systems.

```
[host01]# df -h
Filesystem      Size  Used  Avail   Use%  Mounted on
/dev/xvda2      5.7G  2.4G  3.1G   44%   /
...
host03:/Dev    1008M   34M  924M    4%   /remote_dev
```

- Notice that the host03:/Dev file system is mounted on local file system /remote_dev.

- c. Use the `umount` command to unmount /remote_dev.

```
[host01]# umount /remote_dev
```

11. From host03, view and modify the iptable rules.

- a. Run the `iptables -L` command, pipe the output to `grep`, and search for “nfs”.

```
[host03]# iptables -L | grep nfs
ACCEPT  tcp  --  anywhere  anywhere  state NEW tcp dpt:nfs
```

- Now there is an “nfs” rule, which states to accept all incoming TCP traffic from anywhere when the destination service is nfs.
- This rule was created by checking **NFS4** in the Firewall Configuration GUI.

- b. Run `iptables -h` to display all options.

```
[host03]# iptables -h
iptables v1.4.7
Usage: iptables -[AD] chain rule-specification [options]
       iptables -I chain [rulenumber] rule-specification ...
...

```

- From the help, notice that the --line-numbers option displays line numbers.
- Also notice that the -D option deletes a rule.

- c. Run the `iptables` command to list only those rules in the INPUT chain and include line numbers.

```
[host03]# iptables -L INPUT --line-numbers
Chain INPUT (policy ACCEPT)
num  target     prot opt    source        destination
 1    ACCEPT     all  --    0.0.0.0/0  0.0.0.0/0      state ...
...
 4    ACCEPT     tcp  --    0.0.0.0/0  0.0.0.0/0      state ...nfs
...
```

- Notice that the rule to accept nfs traffic is line number 4.

- d. Use the `iptables` command to delete line number 4 from the INPUT chain. Run the previous command to confirm that the rule is deleted.

```
[host03]# iptables -D INPUT 4
[host03]# iptables -L INPUT --line-numbers
Chain INPUT (policy ACCEPT)
num  target     prot opt    source      destination
1    ACCEPT     all  --    0.0.0.0/0  0.0.0.0/0    state ...
```

- Notice that the `nfs` rule in line number 4 has been deleted.
- Any rules below line number 4 have been renumbered.

- e. Use the `service` command to save the `iptables` rules.
- Always save after making any rule changes.

```
[host03]# service iptables save
iptables: Saving firewall rules to /etc/sysconfig/iptables...
```

- f. Use the `service` command to restart the `iptables` service.
- After making any rule changes, restart the `iptables` service.

```
[host03]# service iptables restart
iptables: Setting chains to policy ACCEPT: filter      [  OK  ]
iptables: Flushing firewall rules:                      [  OK  ]
iptables: Unloading modules:                            [  OK  ]
iptables: Applying firewall rules:                     [  OK  ]
```

12. From **host01**, use the `mount` command to mount `host03:/Dev` on `/remote_dev` using the options as shown:

```
[host01]# mount -t nfs -o rw,nosuid host03:/Dev /remote_dev
```

- As expected, the `mount` command hangs on **host01**, when NFS4 is not trusted.
- Press **Ctrl + C** to abort.

13. From **host03**, create an “`nfs`” rule to trust NFS4.

- a. Use the `iptables` command to insert the “`nfs`” rule after line 4 with the following characteristics from the command line:
- Chain = INPUT
 - Protocol = `tcp`
 - State = NEW
 - Destination port = `nfs`
 - Target = ACCEPT

```
[host03]# iptables -I INPUT 4 -p tcp -m state --state NEW --
dport nfs -j ACCEPT
```

- This rule accepts incoming `tcp` traffic for `nfs`.

- b. Use the `service` command to save the `iptables` rules.

```
[host03]# service iptables save
iptables: Saving firewall rules to /etc/sysconfig/iptables...
```

- c. Use the `cat` command to view the `/etc/sysconfig/iptables` file.

```
[host03]# cat /etc/sysconfig/iptables
...
-A INPUT -p tcp -m state --state NEW -m tcp --dport 2049 -j
ACCEPT
...
```

- Notice that the new rule has the actual port number, 2049, for `nfs`.

- d. Use the `grep` command to search for 2049 in the `/etc/services` file.

```
[host03]# grep 2049 /etc/services
nfs 2049/tcp ...
```

- e. Use the `service` command to restart the `iptables` service.

```
[host03]# service iptables restart
iptables: Flushing firewall rules: [ OK ]
iptables: Setting chains to policy ACCEPT: filter [ OK ]
iptables: Unloading modules: [ OK ]
iptables: Applying firewall rules: [ OK ]
```

14. From **host01**, attempt to mount the NFS file system.

- a. Use the `mount` command to mount `host03:/Dev` on `/remote_dev`.

```
[host01]# mount -t nfs -o rw,nosuid host03:/Dev /remote_dev
```

- The `mount` command is successful this time.

- b. From **host01**, use the `df` command to display the mounted file systems.

```
[host01]# df -h
Filesystem      Size  Used  Avail  Use%  Mounted on
/dev/xvda2      5.7G  2.4G   3.1G   44%   /
...
host03:/Dev    1008M   34M   924M    4%   /remote_dev
```

- Notice that the `host03:/Dev` file system is mounted on local file system `/remote_dev`.

Practice 17-4: Configuring a TCP Wrapper

Overview

In this practice, you configure a TCP wrapper to deny one system from using OpenSSH utilities to connect to another system. You also create a custom log file to capture connection attempts that are denied.

Assumptions

- This practice is performed on **host01** and **host03** VMs.
- Open a terminal window on each system.
- Log in as the `root` user on each system.
- The prompts in the solution section include either **host01** or **host03** to indicate which system to enter the command from.

Tasks

1. From **host01**, confirm that you can use the `ssh` command to connect to **host03**. Password is `oracle`. Use the `logout` command to log off after confirming that you can connect.

```
[host01]# ssh host03
root@host03's password: oracle
Last login: ...
[root@host03 ~]# logout
Connection to host03 closed.

[host01]#
```

2. On **host03**, use the `vi` editor to edit `/etc/hosts.deny` and add the following entry.

```
[host03]# vi /etc/hosts.deny
sshd : 192.0.2.101
```

- This entry denies **host01** (192.0.2.101) from using the OpenSSH utilities to connect to **host03**.
- 3. From **host01**, attempt to use the `ssh` command to connect to **host03**.

```
[host01]# ssh host03
ssh_exchange_identification: Connection closed by remote host
```

- This time you are denied connecting using `ssh`.
- 4. On **host03**, use the `vi` editor to edit `/etc/hosts.deny` and modify the entry as follows.

```
[host03]# vi /etc/hosts.deny
sshd : 192.0.2.101 : spawn /bin/echo "%c tried to connect to %d
and was blocked." >> /var/log/tcpwrappers.log
```

5. From **host01**, attempt to use the `ssh` command to connect to **host03**.

```
[host01]# ssh host03
ssh_exchange_identification: Connection closed by remote host
```

- You are still denied connecting and a message was written to a log file.

6. On **host03**, use the `cat` command to view the `/var/log/tcpwrappers.log` file.

```
[host03]# cat /var/log/tcpwrappers.log  
192.0.2.101 tried to connect to sshd and was blocked.
```

7. Restore the system.

- a. On **host03**, use the `vi` editor to edit `/etc/hosts.deny` and delete the entry created earlier in this practice.

- Delete the strike-through line as follows:

```
[host03]# vi /etc/hosts.deny  
sshd : 192.0.2.101 : spawn /bin/echo "%c tried to connect to %d  
and was blocked." >> /var/log/tcpwrappers.log
```

- b. From **host01**, confirm that you can use the `ssh` command to connect to **host03**.

Password is `oracle`. Use the `logout` command to log off after confirming that you can connect.

```
[host01]# ssh host03  
root@host03's password: oracle  
Last login: ...  
[root@host03 ~]# hostname  
host03.example.com  
[root@host03 ~]# logout  
Connection to host03 closed.  
[host01]#
```

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 18: Oracle on Oracle

Chapter 18

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 18: Oracle on Oracle

Practices Overview

In these practices, you install and run Oracle RDBMS Pre-Install RPM for Oracle Linux 6. You also configure ASMLib.

Practice 18-1: Using sftp to Upload oracle-rdbms-server Package

Overview

In this practice, you use `sftp` to upload the `oracle-rdbms-server-12cR1-preinstall` package from the `dom0` to the `host03` VM. Normally, you obtain these packages from ULN or from the Public Yum Server.

Assumptions

- You are the root user on the `host03` VM.
- You are the root user on `dom0`.

Tasks

1. Use the `sftp` to transfer the `oracle-rdbms-server-12cR1-preinstall` package from `dom0` to `host03`.

- a. From `dom0`, use the `cd` command to change to the `/OVS/seed_pool/sfws` directory.

```
[dom0]# cd /OVS/seed_pool/sfws
```

- b. Use the `ls` command to view the directory for the `oracle*` package.

```
[dom0]# ls oracle*
oracle-rdbms-server-12cR1-preinstall-1.0-8.el6.x86_64.rpm
```

- c. Use the `sftp` command to connect to `host03` as root. Password is `oracle`.

```
[dom0]# sftp host03
root@host03's password: oracle
sftp>
```

- d. Use the `mput` command to copy the `oracle*` packages to `host03`.

```
sftp> put oracle*
Uploading oracle-rdbms-server-12cR1-preinstall-1.0-
8.el6.x86_64.rpm to /root/...
...
```

- e. Use the `quit` command to exit `sftp`.

```
sftp> quit
```

Practice 18-2: Installing and Running Oracle RDBMS Pre-Install

Overview

In this practice, you install the `oracle-rdbms-server-12cR1-preinstall` package, run the “verify” program, and view the results.

Assumptions

- You are the root user on **host03** VM.
- All commands are issued from **host03** VM.

Tasks

1. Install the `oracle-rdbms-server-12cR1-preinstall` package and view the results of the installation.
 - a. Use the `cd` command to change to the root user's home directory, and then use the `yum` command to install the `oracle-rdbms-server-12cR1-preinstall` package.

```
# cd
# yum install oracle-rdbms-server-12cR1-preinstall-1.0-
8.el6.x86_64.rpm
...
Setting up Install Process
...
Resolving Dependencies
...
Transaction Summary
=====
Install    7 Package(s)
Total size: 7.3 M
Total download size: 7.3 M
Installed size: 23 M
Is this ok [y/N] y
Downloading Packages:
...
Installed:
  oracle-rdbms-server-12cR1-preinstall.x86_64 0:1.0-8.el6
Dependency Installed:
...
Complete!
```

- b. Use the `find` command to locate all `oracle-rdbms*` files.

```
# find / -name "*oracle-rdbms*"
/usr/bin/oracle-rdbms-server-12cR1-preinstall-verify
/etc/sysconfig/oracle-rdbms-server-12cR1-preinstall
/etc/sysconfig/oracle-rdbms-server-12cR1-preinstall/oracle-
rdbms-server-12cR1-preinstall.conf
...
/etc/rc.d/init.d/oracle-rdbms-server-12cR1-preinstall-firstboot
/etc/security/limits.d/oracle-rdbms-server-12cR1-preinstall.conf
/var/log/oracle-rdbms-server-12cR1-preinstall
...
```

- c. Use the `ls -l` command to display the file type of the `/usr/bin/oracle-rdbms*` file.

```
# ls -l /usr/bin/oracle-rdbms*
lrwxrwxrwx ... /usr/bin/oracle... -> /etc/sysconfig/oracle...
```

- Notice that this file is a symbolic link to a file in the `/etc/sysconfig/oracle-rdbms*` directory.

- d. Change to the `/etc/sysconfig/oracle-rdbms*` directory and view the contents of the directory.

```
# cd /etc/sysconfig/oracle-rdbms*
# ls
oracle-rdbms-server-12cR1-preinstall.conf
oracle-rdbms-server-12cR1-preinstall.param
oracle-rdbms-server-12cR1-preinstall-verify
```

- e. Use the `less` command to view each of the files in the `/etc/sysconfig/oracle-rdbms*` directory.

```
# less oracle-rdbms-server-12cR1-preinstall.conf
...
# less oracle-rdbms-server-12cR1-preinstall.param
...
# less oracle-rdbms-server-12cR1-preinstall-verify
...
```

- Notice that the `*.param` file is the main configuration file.
- Notice that the `*.verify` file is the Bash script that modifies settings.

2. View the modifications made by the `oracle-rdbms-server-12cR1-preinstall-verify` script.
- Note that previous versions of the `oracle-rdbms-server-...-verify` script needed to be run manually. With this version, however, the script is executed automatically when the RPM package is installed.

- a. Use the less command to view the Oracle RDBMS Pre-install log file, /var/log/oracle-rdbms-server-12cR1-preinstall/results/orakernel.log (sample output shown).

```
# cd /var/log/oracle-rdbms-server-12cR1-preinstall/results
# less orakernel.log
Adding group oinstall with gid 54321
Adding group dba
User oracle is already present
uid=500(oracle) gid=500(oracle) ...
Creating oracle user passed

Verifying kernel parameters as per Oracle recommendations...
fs.file-max          6815744
kernel.sem            250 32000 100 128
kernel.shmmni         4096
...
Setting kernel parameters as per oracle recommendations...
Altered file /etc/sysctl.conf
Original file backed up at /etc/sysctl.conf.orabackup
Verifying & setting of kernel parameters passed

Setting user limits using /etc/security/limits.d/oracle...
Verifying oracle user OS limits as per Oracle recommendations...
Adding oracle soft nofile 1024
Adding oracle hard nofile 65536
Adding oracle soft nproc 16384
Adding oracle hard nproc 16384
Adding oracle soft stack 10240
Adding oracle hard stack 32768
Setting oracle user OS limits as per oracle recommendations...
Altered file /etc/security/limits.d/oracle-rdbms...
Original file backed up at
Verifying & setting of user limits passed

Verifying kernel boot parameters as per Oracle ...
...
Altered file /boot/grub/grub.conf
Original file backed up at /boot/grub/grub.conf.orabackup
Verifying & setting of boot parameters passed

Taking a backup of old config files under /var/log/oracle...
```

- Notice that the required user and groups are created if necessary.
 - Notice that kernel parameters are set and `/etc/sysctl.conf` is backed up beforehand.
 - Notice that oracle user OS limits are set in `/etc/security/limits.d/oracle-rdbms-server-12cR1-preinstall.conf`.
 - Notice that kernel boot parameters are set and `/boot/grub.grub.conf` is backed up beforehand.
- b. Use the `find` command to list the files backed up before settings were changed, *orabackup.

```
# find / -name "*orabackup"  
/etc/sysctl.conf.orabackup  
/boot/grub/grub.conf.orabackup
```

- c. Use the `diff` command to view the changes made in the `grub.conf` file.

```
# diff /boot/grub/grub.conf /boot/grub/grub.conf.orabackup  
...
```

- Notice that the “`numa=off`” kernel parameter is added to the `grub.conf` file.
- d. Use the `wc -l` command to display the number of lines in the `/etc/sysctl.conf` files.

```
# wc -l /etc/sysctl.conf /etc/sysctl.conf.orabackup  
72 /etc/sysctl.conf  
40 /etc/sysctl.conf.orabackup  
...
```

- Notice that 32 new lines are added to the `sysctl.conf` file.
- e. Use the `diff` command to view the changes made in the `sysctl.conf` file.

```
# diff /etc/sysctl.conf /etc/sysctl.conf.orabackup  
...
```

- f. Use the `cat` command to view the limits set in the `/etc/security/limits.d/oracle-rdbms-server-12cR1-preinstall.conf` file.

```
# cat /etc/security/limits.d/oracle-rdbms-server-12cR1-  
preinstall.conf  
...
```

Practice 18-3: Preparing Disks for ASM Use

Overview

In this practice, you:

- Remove all NFS file systems
- Unmount all file systems on /dev/xvdb and /dev/xvdd
- Delete all existing partitions to use the disks for ASM
- Create one partition using the entire disk on /dev/xvdb
- Create one partition using the entire disk on /dev/xvdd

Assumptions

You are the `root` user on **host03** VM.

Tasks

1. Use the `vi` editor to delete all entries from `/etc/exports`.

```
# vi /etc/exports
Delete all entries and save the file
```

2. Use the `service` command to restart the `nfs` service.

```
# service nfs restart
...
```

3. Use the `showmount -e` command to confirm that there are no exported file systems.

```
# showmount -e
Export list for host03.example.com
```

- No NFS file systems are being exported.
- 4. Unmount all file systems on /dev/xvdb and /dev/xvdd.
 - a. Use the `df` command to list the mounted partitions.

```
# df -h
Filesystem      Size  Used  Avail   Use%  Mounted on
/dev/xvda2      5.7G  4.7G   760M   87%   /
...
/dev/xvdb1      1008M  34M   924M    4%   /Dev
/dev/xvdd1      992M  1.3M   940M    1%   /Test
```

- In this example, the following partitions need to be unmounted (your system may be different).
 - /dev/xvdb1 mounted on /Dev
 - /dev/xvdd1 mounted on /Test

- b. Use the `umount` command to unmount file systems on `/dev/xvdb1` and `/dev/xvdd1`.

```
# umount /Dev
# umount /Test
```

5. Delete all partitions on `/dev/xvdb` and `/dev/xvdd`.

- a. Use the `fdisk` command to display the partition table on `/dev/xvdb` and then delete all the partitions.

```
# fdisk /dev/xvdb
...
Command (m for help): p
Disk /dev/xvdb: 5368 MB, 5368709120 bytes
...
Device Boot Start End Blocks Id System
/dev/xvdb1 ...
/dev/xvdb2 ...

Command (m for help): d
Partition number (1-4): 1

Command (m for help): d
Selected partition 2

Command (m for help): p
...
Device Boot Start End Blocks Id System

Command (m for help): w
...
```

- In this example, two partitions are deleted (your system may be different).
- b. Use the `fdisk` command to display the partition table on `/dev/xvdd` and then delete all the partitions.

```
# fdisk /dev/xvdd
...
Command (m for help): p
Disk /dev/xvdb: 5368 MB, 5368709120 bytes
...
Device Boot Start End Blocks Id System
/dev/xvdd1 ...
/dev/xvdd2 ...

Command (m for help): d
```

```

Partition number (1-4) : 1

Command (m for help) : d
Selected partition 2

Command (m for help) : p
...
      Device Boot Start      End      Blocks   Id  System
Command (m for help) : w
...

```

- In this example, two partitions are deleted (your system may be different).
6. Use the `vi` editor to delete entries pertaining to file systems on `/dev/xvdb` and `/dev/xvdd` from the `/etc/fstab` file.

```

# vi /etc/fstab
LABEL=Dev    /Dev    ... (delete this entry)
LABEL=Test   /Test   ... (delete this entry)

```

- In this example, delete the two entries shown (your system may be different).
7. In preparation for the ASMLib lab, you must disable SELinux completely and reboot.

- ASMLib is not supported with SELinux enabled (it is a known issue).

Use the `vi` editor to edit `/etc/selinux/config` as follows:

```

# vi /etc/selinux/config
SELINUX=enforcing (old value)
SELINUX=disabled  (new value)

```

8. Reboot your system and log back in.

```

# reboot
...
```

- After you reboot your system, your vnc session closes.
- Connect to **host03** by using VNC.
 - Run the `vncviewer&` command.


```
# vncviewer&
```

 - The **VNC Viewer: Connection Details** dialog box is displayed.
 - Enter the command, `localhost:<port_number>`, substituting the correct port number for the **host03** guest. For example, if the port number is 5903, enter the following and click OK:


```
localhost:5903
```
 - Select Oracle Student from the GNOME login window; password is `oracle`.
 - Right-click the GNOME desktop and select **Open in Terminal** from the pop-up menu.

- d. In the terminal window, become the root user by entering the su - command followed by the root password, oracle.

```
$ su -  
Password: oracle
```

9. Create a partition on /dev/xvdb.

- a. Use the fdisk command to partition /dev/xvdb.

```
# fdisk /dev/xvdb  
WARNING: DOS-compatible mode is deprecated. It's strongly  
recommended to switch off the mode (command 'c') and change  
display units to sectors (command 'u').  
Command (m for help) :
```

- b. Add a new primary partition, giving the partition number 1.

```
Command (m for help): n  
Command action  
  e   extended  
  p   primary partition (1-4)  
p  
Partition number (1-4) : 1
```

- c. Continue adding the new partition, using the entire disk as follows:

```
First cylinder (1-652, default 1): ENTER  
Using default value 1  
Last cylinder, +cylinders or +size{K,M,G} (1-652, default 652):  
ENTER
```

- d. Save the new partition table.

```
Command (m for help): w  
The partition table has been altered!  
Calling ioctl() to re-read partition table.  
Syncing disks.
```

10. Create a partition on /dev/xvdd.

- a. Use the fdisk command to partition /dev/xvdd.

```
# fdisk /dev/xvdd  
WARNING: DOS-compatible mode is deprecated. It's strongly  
recommended to switch off the mode (command 'c') and change  
display units to sectors (command 'u').  
Command (m for help) :
```

- b. Add a new **primary** partition, giving the partition number **1**.

```
Command (m for help): n
Command action
  e   extended
  p   primary partition (1-4)
p
Partition number (1-4): 1
```

- c. Continue adding the new partition, using the entire disk as follows:

```
First cylinder (1-652, default 1): ENTER
Using default value 1
Last cylinder, +cylinders or +size{K,M,G} (1-652, default 652):
ENTER
```

- d. Save the new partition table.

```
Command (m for help): w
The partition table has been altered!
Calling ioctl() to re-read partition table.
Syncing disks.
```

Practice 18-4: Installing and Configuring ASMLib

Overview

In this practice, you:

- Determine which package provides the `oracleasm` binary executable
- Configure ASMLib
- Load and initialize the ASMLib driver
- Mark disk partitions for ASM use
- View information about the ASM disk partitions

Assumptions

You are the `root` user on **host03** VM.

Tasks

1. Determine the package name that provides the `oracleasm` binary executable.
 - a. Use the `which` command to display the location of the `oracleasm` binary.

```
# which oracleasm
/usr/sbin/oracleasm
```

- b. Use the `rpm -qf` command to display the package name that provides the `oracleasm` binary.
 - Recall that you installed this optional package during the Oracle Linux installation practice in Lesson 3.

2. Use the `oracleasm configure -i` command to configure ASMLib.

```
# oracleasm configure -i
Configuring the Oracle ASM library driver.
```

This will configure the on-boot properties of the Oracle ASM library driver. The following questions will determine whether the driver is loaded on boot and what permissions it will have. The current values will be shown in brackets ('[]'). Hitting <ENTER> without typing an answer will keep that current value. Ctrl-C will abort.

```
Default user to own the driver interface []: oracle
Default group to own the driver interface []: dba
Start Oracle ASM library driver on boot (y/n) [n]: y
Scan for Oracle ASM disks on boot (y/n) [y]: ENTER
Writing Oracle ASM library driver configuration: done
```

3. Use the `oracleasm init` command to load and initialize the ASMLib driver.

```
# oracleasm init
Creating /dev/oracleasm mount point: /dev/oracleasm
Loading module "oracleasm": oracleasm
Configuring "oracleasm" to use device physical block size
Mounting ASMLib driver filesystem: /dev/oracleasm
```

4. Use the `oracleasm configure` command without the `-i` option.

```
# oracleasm configure
ORACLEASM_ENABLED=true
ORACLEASM_UID=oracle
ORACLEASM_GID=dba
ORACLEASM_SCANBOOT=true
ORACLEASM_SCANORDER=""
ORACLEASM_SCANEXCLUDE=""
ORACLEASM_USE_LOGICAL_BLOCK_SIZE="false"
```

5. Mark disk partitions for ASM use.

- a. Use the `oracleasm createdisk` command to mark `/dev/xvdb1` for ASM use.

```
# oracleasm createdisk VOL1 /dev/xvdb1
Writing disk header: done
Instantiating disk: done
```

- b. Use the `oracleasm createdisk` command to mark `/dev/xvdd1` for ASM use.

```
# oracleasm createdisk VOL2 /dev/xvdd1
Writing disk header: done
Instantiating disk: done
```

6. Use the `ls` command to display a long list of `/dev/oracleasm/disks` directory.

```
# ls -l /dev/oracleasm/disks
brw-rw----. 1 oracle dba ... VOL1
brw-rw----. 1 oracle dba ... VOL2
```

7. Use the `oracleasm listdisks` command to list the disk names of marked ASMLib disks.

```
# oracleasm listdisks
VOL1
VOL2
```

8. Use the `oracleasm scandisks` command to enable cluster nodes to identify which shared disks have been marked as ASMLib disks on another node.

```
# oracleasm scandisks
Reloading disk partitions: done
Cleaning any stale ASM disks...
Scanning system for ASM disks...
```

9. Use the `oracleasm querydisk` command to determine whether a disk name or disk device is being used by ASMLib.

```
# oracleasm querydisk VOL1
Disk "VOL1" is a valid ASM disk
# oracleasm querydisk /dev/xvdd1
Device "/dev/xvdd1" is marked an ASM disk with the label "VOL2"
```

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 19: System Monitoring

Chapter 19

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 19: System Monitoring

Practices Overview

In these practices, you:

- Use the `sosreport` utility to collect system information
- Use standard Linux utilities to monitor system resource usage
- Use OSWatcher Black Box and OSWatcher Black Box Analyzer

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practice 19-1: Using sosreport to Collect System Information

Overview

In this practice, you:

- Use the `sosreport` utility to collect system information
- Extract the compressed TAR file and view the collected information
- View the status of the `sosreport` plug-ins

Assumptions

You are the `root` user on **host03** VM.

Tasks

1. Use the `rpm` command to verify that the `sos` packing is installed.

```
[host03]# rpm -q sos
sos-2.2-47.0.1.el6.noarch
```

- In this example, the package is installed.
2. Run the `sosreport` command.
 - a. Press Enter when prompted to continue.
 - b. Provide your first initial and last name when prompted.
 - c. Enter number 1 as the case number for which you are generating the report.

```
[host03]# sosreport
sosreport (version 2.2)

This command will collect diagnostic and configuration
information from this Oracle Linux system and installed
applications.

An archive containing the collected information will be
generated in /tmp and may be provided to a Oracle USA
support representative.
```

Any information provided to Oracle USA will be treated
in accordance with the published support policies at:

<https://linux.oracle.com/>

The generated archive may contain data considered
sensitive and its content should be reviewed by the
originating organization before being passed to any third
party.

No changes will be made to system configuration.

Press ENTER to continue, or CTRL-C to quit. **ENTER**

Please enter your first initial and last name [host03]: **cmac**

Please enter the case number that you are generating... **1**

```
Running plugins. Please wait ...
Completed...
...
Your sosreport has been generated and saved in:
/tmp/sosreport-cmac.1-...tar.xv
...
```

3. Use the `cd` command to change to the `/tmp` directory.

```
[host03]# cd /tmp
```

4. Use the `ls` command to display a long listing of the `/tmp` directory.

```
[host03]# ls -l
...
-rw----- sosreport-cmac.1-...tar.xz
-rw----- sosreport-cmac.1-...tar.xz.md5
...
```

- Notice the two `sosreport` files, one with the `.xz` extension and one with the `.md5` extension.
 - The `.xz` file is the compressed data file.
 - Note the size of the `.xz` file.
5. Use the `xz -d` command to uncompress the `.xz` file.

```
[host03]# xz -d sosreport-cmac.1-...tar.xz
```

6. Use the `ls` command to display a long listing of the `/tmp` directory.

```
[host03]# ls -l
...
-rw----- sosreport-cmac.1-...tar
-rw----- sosreport-cmac.1-...tar.xz.md5
...
```

- Notice that the `sosreport` file with the `.tar.xz` extension now has a `.tar` extension.
 - Note the size of the `.tar` file—it is considerably larger than the compressed (`.xz`) file.
7. Use the `tar` command to extract the `.tar` file.

```
[host03]# tar xvf sosreport-cmac.1-...tar
host03-...
host03-.../uname -> sos_commands/kernel/uname_-a
host03-.../date -> sos_commands/general/date
host03-.../java -> sos_commands/general/alternatives_--displa...
host03-.../ifconfig -> sos_commands/networking/ifconfig_-a
...
```

- Notice that the `tar` file is extracted in a `host03-...` directory.

8. Use the `cd` command to change to the `host03-...` directory.

```
[host03]# cd host03-*
```

9. Use the `ls` command to display a long listing of the `host03-...` directory.

```
[host03]# ls -l
...
drwxr-xr-x      boot
lrwxrwxrwx      chkconfig -> sos_commands/startup/...
lrwxrwxrwx      date -> sos_commands/general/date
...
drwxr-xr-x      etc
lrwxrwxrwx      free -> sos_commands/memory/free
lrwxrwxrwx      hostname -> sos_commands/general/...
...
drwxr-xr-x      lib
...
```

- Notice that a number of directories that contain data collected from the system exist.
 - Notice that a number of symbolic links that contain the output of several status-related commands exist.
10. Use the `sosreport -l` command to list the plug-ins.

```
[host03]# sosreport -l
sosreport (version 2.2)
The following plugins are currently enabled:
  acpid           acpid related information
  anaconda        Anaconda / Installation information
...
The following plugins are currently disabled:
  amd             Amd automounter information
  cloudforms      Cloudforms related information
...
The following plugin options are available:
  apache.log       off gathers all apache logs
  auditd.syslogsize 15 max size (MiB) to collect per syslog...
...
```

Practice 19-2: Using Standard Linux Performance Monitoring Tools

Overview

In this practice, you use standard Linux system resource monitoring utilities to observe the following:

- CPU statistics
- Memory statistics
- Disk I/O statistics
- Network statistics

Assumptions

- You are the `root` user on **dom0**.
- The **host03** VM is running.
- The output shown in the tasks is only a sample. Your output will be different.

Tasks

1. Open two terminal windows on **dom0** and use the `su -` command to become the `root` user (the `root` user's password on **dom0** is `oracle`) in both windows.
 - You have two terminal windows open:
 - One window to run the Linux system resource monitoring utilities
 - One window to run a command-line utility to generate a system load

```
[dom0] $ su -
Password: oracle
[dom0] #
```

2. Use the `ssh` command to log in to **host03**. The `root` password is `oracle`.
 - Log in to **host03** from both terminal windows.

```
[dom0] # ssh host03
root@host03's password: oracle
[root@host03 ~] #
```

3. Observe the CPU statistics.
 - a. In one window, use the `top` command to display CPU usage and load averages.
 - The `top` command also monitors process statistics and memory usage.

```
# top
```

- b. In the second window, run the following command to generate a system load:

```
# dd if=/dev/zero of=/dev/null bs=1024
```

A sample output of `top` is shown as follows:

```

root@host03:~
File Edit View Search Terminal Help
top - 04:46:26 up 2 days, 12:33, 3 users, load average: 0.22, 0.06, 0.06
Tasks: 136 total, 2 running, 134 sleeping, 0 stopped, 0 zombie
Cpu(s): 19.7%us, 80.3%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 2050816k total, 1964924k used, 85892k free, 97400k buffers
Swap: 3876860k total, 0k used, 3876860k free, 1514156k cached


PID USER PR NI VIRT RES SHR S %CPU %MEM TIME+ COMMAND
15792 root 20 0 102m 512 424 R 99.7 0.0 0:15.19 dd
15747 root 20 0 15088 1124 836 R 0.3 0.1 0:00.15 top
  1 root 20 0 19408 1520 1220 S 0.0 0.1 0:00.74 init
  2 root 20 0 0 0 0 S 0.0 0.0 0:00.00 kthreadd
  3 root 20 0 0 0 0 S 0.0 0.0 0:00.15 ksoftirqd/0
  5 root 0 -20 0 0 0 S 0.0 0.0 0:00.00 kworker/0:0H
  7 root 0 -20 0 0 0 S 0.0 0.0 0:00.00 kworker/u:0H
  8 root RT 0 0 0 0 S 0.0 0.0 0:00.00 migration/0
  9 root 20 0 0 0 0 S 0.0 0.0 0:00.00 rcu_bh
 10 root 20 0 0 0 0 S 0.0 0.0 0:03.88 rcu_sched
 11 root RT 0 0 0 0 S 0.0 0.0 0:02.53 watchdog/0
 12 root 0 -20 0 0 0 S 0.0 0.0 0:00.00 cpuset
 13 root 0 -20 0 0 0 S 0.0 0.0 0:00.00 khelper
 14 root 20 0 0 0 0 S 0.0 0.0 0:00.00 kdevtmpfs
 15 root 0 -20 0 0 0 S 0.0 0.0 0:00.00 netns
 16 root 20 0 0 0 0 S 0.0 0.0 0:00.00 xenwatch
 17 root 20 0 0 0 0 S 0.0 0.0 0:00.01 xenbus

```

- In this example, CPU usage is high as indicated by the 0.0%id (idle) statistic.
- The biggest consumer of the CPU is the `dd` process, which has a PID of 15792.
- Load average is the average number of processes in the run queue or the number of processes waiting to run on the CPU.
 - The load average over the last 1 minute is 0.22, over the last 5 minutes is 0.06, and over the last 15 minutes is 0.06.
 - A high load average is an indication that your system does not have sufficient CPU capacity.
- Press `q` to quit the `top` command.

c. Use the `mpstat -P ALL 1` command to display CPU statistics.

- The `1` argument provides statistics at one-second intervals.
- The `-P ALL` option provides statistics for each individual processor and globally for all processors.



The screenshot shows a terminal window titled "root@host03:~". The command entered is [root@host03 ~]# mpstat -P ALL 1. The output displays CPU usage statistics over time. The columns include time, CPU type (CPU, all, 0), and various percentage metrics: %usr, %nice, %sys, %iowait, %irq, %soft, %steal, %guest, and %idle. The data shows high CPU usage, particularly %iowait, across all processors.

Time	CPU	%usr	%nice	%sys	%iowait	%irq	%soft	%steal	%guest	%idle
04:49:51 AM	CPU									
04:49:51 AM	all	22.77	0.00	77.23	0.00	0.00	0.00	0.00	0.00	0.00
04:49:52 AM	0	22.77	0.00	77.23	0.00	0.00	0.00	0.00	0.00	0.00
04:49:52 AM	CPU									
04:49:52 AM	all	22.22	0.00	77.78	0.00	0.00	0.00	0.00	0.00	0.00
04:49:53 AM	0	22.22	0.00	77.78	0.00	0.00	0.00	0.00	0.00	0.00
04:49:53 AM	CPU									
04:49:53 AM	all	19.80	0.00	80.20	0.00	0.00	0.00	0.00	0.00	0.00
04:49:54 AM	0	19.80	0.00	80.20	0.00	0.00	0.00	0.00	0.00	0.00
04:49:54 AM	CPU									
04:49:54 AM	all	20.20	0.00	79.80	0.00	0.00	0.00	0.00	0.00	0.00
04:49:55 AM	0	20.20	0.00	79.80	0.00	0.00	0.00	0.00	0.00	0.00
04:49:55 AM	CPU									
04:49:55 AM	all	20.00	0.00	80.00	0.00	0.00	0.00	0.00	0.00	0.00
04:49:56 AM	0	20.00	0.00	80.00	0.00	0.00	0.00	0.00	0.00	0.00
04:49:56 AM	CPU									

- In this example, CPU usage is high as indicated by the `0.00 (%idle)` statistic.
- Press `Ctrl + C` after viewing a few intervals.


- d. Use the `sar -u` command to display system-wide CPU usage. Only a partial output is shown as follows:

The screenshot shows a terminal window titled "root@host03:~". The window contains the output of the `sar -u` command. The output is as follows:

Time	CPU	%user	%nice	%system	%iowait	%steal	%idle
12:00:01 AM	CPU	0.01	0.00	0.03	0.01	0.00	99.96
12:10:01 AM	all	0.01	0.00	0.03	0.01	0.00	99.92
12:20:01 AM	all	0.04	0.00	0.03	0.01	0.00	99.97
12:30:01 AM	all	0.01	0.00	0.02	0.01	0.00	99.97
12:40:01 AM	all	0.01	0.00	0.01	0.00	0.00	99.98
12:50:01 AM	all	0.01	0.00	0.02	0.00	0.00	99.98
01:00:01 AM	all	0.01	0.00	0.02	0.00	0.00	99.97
01:10:01 AM	all	0.01	0.00	0.02	0.00	0.00	99.96
01:20:01 AM	all	0.05	0.00	0.03	0.01	0.00	99.91
01:30:01 AM	all	0.01	0.00	0.01	0.00	0.00	99.98
01:40:01 AM	all	0.01	0.00	0.01	0.00	0.00	99.98
01:50:01 AM	all	0.01	0.00	0.01	0.01	0.00	99.97
02:00:01 AM	all	0.01	0.00	0.01	0.00	0.00	99.98
02:10:01 AM	all	0.01	0.00	0.02	0.01	0.00	99.97
02:20:01 AM	all	0.04	0.00	0.03	0.02	0.00	99.90
02:30:01 AM	all	0.01	0.00	0.02	0.00	0.00	99.98
02:40:01 AM	all	0.01	0.00	0.02	0.01	0.00	99.97
02:50:01 AM	all	0.01	0.00	0.02	0.00	0.00	99.98
03:00:01 AM	all	0.01	0.00	0.02	0.01	0.00	99.97
03:10:01 AM	all	0.01	0.00	0.02	0.00	0.00	99.97
03:20:01 AM	all	0.04	0.00	0.04	0.01	0.00	99.91

- In this example, CPU usage is low, with a few exceptions, as indicated by the 99.xx (%idle) statistics.
- An “Average” line is not shown in this example.

- e. Use the `sar -q` command to display run queue length and load averages. Only a partial output is shown as follows:



The screenshot shows a terminal window titled "root@host03:~". The window contains the output of the `sar -q` command. The output includes system information (Linux 3.8.13-16.2.1.el6uek.x86_64, host03.example.com, 12/21/2013, x86_64) and a header row with columns: 12:00:01 AM, runq-sz, plist-sz, ldavg-1, ldavg-5, and ldavg-15. Below the header is a table of data points from 12:10:01 AM to 03:10:01 AM. The runq-sz column consistently shows 0, while the ldavg columns show values ranging from 0.00 to 0.05.

Time	runq-sz	plist-sz	ldavg-1	ldavg-5	ldavg-15
12:10:01 AM	0	223	0.00	0.02	0.05
12:20:01 AM	0	223	0.00	0.01	0.05
12:30:01 AM	0	223	0.00	0.01	0.05
12:40:01 AM	0	223	0.00	0.01	0.05
12:50:01 AM	0	223	0.00	0.01	0.05
01:00:01 AM	0	223	0.00	0.01	0.05
01:10:01 AM	0	223	0.00	0.01	0.05
01:20:01 AM	0	223	0.00	0.01	0.05
01:30:01 AM	0	223	0.00	0.01	0.05
01:40:01 AM	0	223	0.00	0.01	0.05
01:50:01 AM	0	223	0.00	0.01	0.05
02:00:01 AM	0	223	0.00	0.01	0.05
02:10:01 AM	0	223	0.00	0.01	0.05
02:20:01 AM	0	223	0.00	0.01	0.05
02:30:01 AM	0	223	0.00	0.01	0.05
02:40:01 AM	0	223	0.00	0.01	0.05
02:50:01 AM	0	223	0.00	0.01	0.05
03:00:01 AM	0	223	0.00	0.01	0.05
03:10:01 AM	0	224	0.00	0.01	0.05

- In this example, the CPU is not saturated as indicated by the low run queue length (`runq-sz`) and low load averages (last 1 minute: `ldavg-1`, last 5 minutes: `ldavg-5`, and last 15 minutes: `ldavg-15`).
- A run queue size greater than the number of CPUs on your system is usually indicative of a CPU bottleneck.
- An “Average” line is not shown in this example.

- f. Use the `vmstat 1` command to display CPU statistics.


- This command is primarily used for monitoring virtual memory statistics.

The screenshot shows a terminal window titled "root@host03:~". The command entered is [root@host03 ~]# vmstat 1. The output displays system statistics over time intervals. The columns represent various metrics: procs (processes), memory (swap and free memory), io (block input/output), system (context switches, user/system idle, and interrupt activity), and cpu (CPU usage by state: id, sys, user, nice, wait, and steal). The data shows high CPU usage, particularly in the user and nice states, with low idle percentages.

procs		memory				swap		io		system				cpu			
r	b	swpd	free	buff	cache	si	so	bi	bo	in	cs	us	sy	id	wa	st	
1	0	0	85248	97820	1514300	0	0	6	7	36	15	0	1	99	0	0	
1	0	0	85232	97820	1514304	0	0	0	0	1055	49	18	82	0	0	0	
1	0	0	85240	97820	1514304	0	0	0	0	1058	16	19	81	0	0	0	
1	0	0	85240	97820	1514304	0	0	0	0	1054	34	20	80	0	0	0	
1	0	0	85240	97820	1514304	0	0	0	0	1061	24	20	80	0	0	0	
1	0	0	85240	97820	1514304	0	0	0	0	1061	27	20	80	0	0	0	
1	0	0	85272	97824	1514304	0	0	0	0	1045	30	21	79	0	0	0	
1	0	0	85272	97824	1514304	0	0	0	0	1031	28	19	81	0	0	0	
1	0	0	85272	97824	1514304	0	0	0	0	1032	18	19	81	0	0	0	
1	0	0	85272	97824	1514304	0	0	0	0	1052	50	21	79	0	0	0	
1	0	0	85264	97824	1514304	0	0	0	0	1066	16	19	81	0	0	0	
1	0	0	85264	97824	1514304	0	0	0	0	1029	28	17	83	0	0	0	
1	0	0	85264	97832	1514304	0	0	0	0	1037	20	18	82	0	0	0	
1	0	0	85264	97832	1514304	0	0	0	0	1034	33	18	82	0	0	0	
1	0	0	85264	97832	1514304	0	0	0	0	1045	22	20	80	0	0	0	
1	0	0	85264	97832	1514304	0	0	0	0	1029	28	20	80	0	0	0	
1	0	0	85264	97832	1514304	0	0	0	0	1051	16	20	80	0	0	0	
1	0	0	85264	97832	1514304	0	0	0	0	1044	28	20	80	0	0	0	
1	0	0	85272	97840	1514304	0	0	0	0	1080	24	16	84	0	0	0	
1	0	0	85272	97840	1514304	0	0	0	0	1031	36	19	81	0	0	0	
1	0	0	85272	97840	1514304	0	0	0	0	1038	18	19	81	0	0	0	

- The `1` argument provides statistics at one-second intervals.
- In this example, CPU usage is high as indicated by the low `idle` (%idle) statistics.
- In this example, the high run queue length (`r`) statistics are also an indication of CPU saturation.
- Press Ctrl + C after viewing a few intervals.

- g. Use the `iostat` command to display CPU usage.
- This command is primarily used for monitoring system I/O device loads.



The screenshot shows a terminal window titled "root@host03:~". The window displays the output of the `iostat` command. The output includes system information (Linux version, host name, date, architecture) and two sections of statistics: "avg-cpu" and "Device".

avg-cpu:						
	%user	%nice	%system	%iowait	%steal	%idle
	0.27	0.02	0.80	0.11	0.00	98.80

Device:	tps	Blk_read/s	Blk_wrtn/s	Blk_read	Blk_wrtn
xvda	0.70	12.07	13.62	2651004	2993354
xvdb	0.00	0.02	0.00	3800	0
xvdd	0.00	0.02	0.00	3800	0
scd0	0.00	0.00	0.00	532	0

- In this example, the average `%idle` statistic for the CPU is 98.80%.
4. Observe the memory statistics.
- Use the `top` command to display memory usage.

```
# top
```

A sample output of `top` is shown as follows:

```

root@host03:~
File Edit View Search Terminal Help
top - 05:20:03 up 2 days, 13:07, 3 users, load average: 1.00, 0.99, 0.89
Tasks: 136 total, 2 running, 134 sleeping, 0 stopped, 0 zombie
Cpu(s): 19.5%us, 80.5%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 2050816k total, 1964924k used, 85892k free, 97904k buffers
Swap: 3876860k total, 0k used, 3876860k free, 1514284k cached

PID USER      PR  NI    VIRT   RES   SHR S %CPU %MEM TIME+ COMMAND
15792 root      20   0 102m  512  424 R 98.7  0.0 33:45.91 dd
10127 oracle   20   0 401m 11m 8984 S  0.3  0.6 0:01.39 metacity
  1 root      20   0 19408 1520 1220 S  0.0  0.1 0:00.75 init
  2 root      20   0     0   0   0 S  0.0  0.0 0:00.00 kthreadd
  3 root      20   0     0   0   0 S  0.0  0.0 0:00.15 ksoftirqd/0
  5 root      0 -20    0   0   0 S  0.0  0.0 0:00.00 kworker/0:0H
  7 root      0 -20    0   0   0 S  0.0  0.0 0:00.00 kworker/u:0H
  8 root      RT  0     0   0   0 S  0.0  0.0 0:00.00 migration/0
  9 root      20   0     0   0   0 S  0.0  0.0 0:00.00 rcu_bh
 10 root     20   0     0   0   0 S  0.0  0.0 0:03.94 rcu_sched
 11 root     RT  0     0   0   0 S  0.0  0.0 0:02.54 watchdog/0
 12 root     0 -20    0   0   0 S  0.0  0.0 0:00.00 cpuset
 13 root     0 -20    0   0   0 S  0.0  0.0 0:00.00 khelper
 14 root     20   0     0   0   0 S  0.0  0.0 0:00.00 kdevtmpfs
 15 root     0 -20    0   0   0 S  0.0  0.0 0:00.00 netns
 16 root     20   0     0   0   0 S  0.0  0.0 0:00.00 xenwatch
 17 root     20   0     0   0   0 S  0.0  0.0 0:00.01 xenbus

```

- The upper section of `top` displays memory statistics:
 - The `Mem:` line reflects how much physical memory your system has, how much is used, and how much is free.
 - The `Swap:` line reflects how much swap memory your system has, how much is used, and how much is free.
 - When a computer runs out of physical memory and starts using swap space, its performance deteriorates dramatically.
 - If you run out of swap, you will most likely crash your programs or the OS.
- The lower section of `top` displays a list of processes sorted by CPU usage, with the top consumer of CPU listed first.

- b. Sort by memory usage, press F or O to change the sort field, and then select either of the following to sort processes by memory usage:
- n: %MEM = Memory usage (RES)
 - q: RES = Resident size (kb)
- Press the Enter key after selecting a sort field.

The terminal window shows the following output from the top command:

```
top - 05:21:16 up 2 days, 13:08, 3 users, load average: 1.00, 0.99, 0.90
Tasks: 136 total, 2 running, 134 sleeping, 0 stopped, 0 zombie
Cpu(s): 17.7%us, 82.3%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 2050816k total, 1964932k used, 85884k free, 97920k buffers
Swap: 3876860k total, 0k used, 3876860k free, 1514284k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1769	root	20	0	158m	24m	9060	S	0.3	1.2	0:25.32	Xorg
10135	oracle	20	0	521m	20m	13m	S	0.0	1.0	0:36.60	nautilus
10167	oracle	20	0	316m	18m	9236	S	0.0	0.9	0:00.16	python
10300	oracle	20	0	393m	17m	13m	S	0.0	0.9	0:00.08	gnote
10164	oracle	20	0	472m	14m	10m	S	0.0	0.7	0:00.54	gpk-update-icon
10132	oracle	20	0	321m	13m	10m	S	0.0	0.7	0:02.65	gnome-panel
11314	oracle	20	0	287m	13m	9560	S	0.0	0.7	0:03.83	gnome-terminal
10301	oracle	20	0	452m	12m	10m	S	0.0	0.6	0:01.24	clock-applet
10143	oracle	20	0	310m	11m	8348	S	0.0	0.6	0:00.16	nm-applet
10160	oracle	20	0	322m	11m	8976	S	0.0	0.6	0:02.91	wnck-applet
10127	oracle	20	0	401m	11m	8984	S	0.0	0.6	0:01.39	metacity
10299	oracle	20	0	382m	11m	8768	S	0.0	0.6	0:00.07	gdm-user-switch
10165	oracle	20	0	360m	10m	8252	S	0.0	0.5	0:00.23	gnome-volume-co
10102	oracle	20	0	488m	10m	7868	S	0.0	0.5	0:04.47	gnome-settings-
10159	oracle	20	0	308m	9304	7080	S	0.0	0.5	0:00.04	trashapplet
10213	oracle	20	0	260m	8500	4464	S	0.0	0.4	0:03.26	gnome-screensav
10163	oracle	20	0	255m	8092	6036	S	0.0	0.4	0:00.06	bluetooth-apple

- Press q to quit the top command.

- c. Use the `vmstat` command to display memory statistics.

```
[root@host03 ~]# vmstat 1
procs --memory-- --swap-- --io-- --system-- --cpu--
r b swpd free buff cache si so bi bo in cs us sy id wa st
1 0 86220 97924 1514292 0 0 6 7 38 15 0 1 99 0 0
1 0 86256 97924 1514292 0 0 0 0 1056 48 19 81 0 0 0
1 0 86256 97924 1514292 0 0 0 0 1046 24 20 80 0 0 0
1 0 86256 97928 1514288 0 0 0 0 12 1065 58 18 82 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1050 34 20 80 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1083 54 20 80 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1040 22 19 81 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1060 29 19 81 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1044 23 22 78 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1041 34 18 82 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1076 63 19 81 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1072 68 22 78 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1045 29 18 82 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1048 28 18 82 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1039 33 18 82 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1031 32 19 81 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1048 23 19 81 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1046 26 19 81 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1046 22 17 83 0 0 0
1 0 86256 97928 1514292 0 0 0 0 1050 59 20 80 0 0 0
^C
```


- The `1` argument provides statistics at one-second intervals.
- The important memory statistics are:
 - `swpd` – The amount of virtual memory used
 - `free` – The amount of idle memory
 - `si` – The amount of memory swapped in from disk (per second)
 - `so` – The amount of memory swapped out to disk (per second)
- In this example, the system has sufficient free memory and is not swapping.
- Press `Ctrl + C` after viewing a few intervals.

- d. Use the `free` command to display memory statistics.

```
[root@host03 ~]# free -m
total        used        free      shared  buffers   cached
Mem:       2002       1923         79          0        95     1478
-/+ buffers/cache:    348      1654
Swap:      3785          0      3785
[root@host03 ~]#
```

- This example uses the `-m` option to display amounts in megabytes.
- This command displays the total amount of free and used physical and swap memory in your system.

- e. Use the `sar -r` command to display memory usage statistics. Only a partial output is shown as follows:



The screenshot shows a terminal window titled "root@host03:~". The window contains the following text:

```
[root@host03 ~]# sar -r
Linux 3.8.13-16.2.1.el6uek.x86_64 (host03.example.com) 12/21/2013 _x86_64_
1 CPU

12:00:01 AM kbmemfree kbmemused %memused kbbuffers kbcached kbcommit %commit
12:10:01 AM 58468 1992348 97.15 89480 1554664 501988 8.47
12:20:01 AM 58344 1992472 97.16 89604 1554672 501988 8.47
12:30:01 AM 58352 1992464 97.15 89680 1554672 501988 8.47
12:40:01 AM 58196 1992620 97.16 89756 1554676 501988 8.47
12:50:01 AM 58064 1992752 97.17 89836 1554680 501988 8.47
01:00:01 AM 58072 1992744 97.17 89920 1554676 501988 8.47
01:10:01 AM 57940 1992876 97.17 90016 1554688 501988 8.47
01:20:01 AM 57692 1993124 97.19 90148 1554692 501988 8.47
01:30:01 AM 57724 1993092 97.19 90220 1554696 501988 8.47
01:40:01 AM 57716 1993100 97.19 90308 1554704 501988 8.47
01:50:01 AM 57608 1993208 97.19 90372 1554704 501988 8.47
02:00:01 AM 57444 1993372 97.20 90460 1554708 501988 8.47
02:10:01 AM 57320 1993496 97.21 90544 1554712 501988 8.47
02:20:01 AM 57104 1993712 97.22 90672 1554716 501988 8.47
02:30:01 AM 57196 1993620 97.21 90764 1554724 501988 8.47
02:40:01 AM 57080 1993736 97.22 90848 1554724 501988 8.47
02:50:01 AM 56948 1993868 97.22 90928 1554728 501988 8.47
03:00:01 AM 56856 1993960 97.23 91004 1554728 501988 8.47
03:10:01 AM 56460 1994356 97.25 91120 1554736 502612 8.48
03:20:01 AM 55380 1995436 97.30 91236 1554732 503944 8.50
03:30:01 AM 55252 1995564 97.31 91312 1554744 501988 8.47
03:40:01 AM 55112 1995704 97.31 91408 1554744 501988 8.47
```

- Observe the `%memused` value, which is the percentage of used memory.
- An “Average” line is not shown in this example.

- f. Use the `sar -B` command to display memory paging statistics. Only a partial output is shown as follows:

	pgpgin/s	pgpgout/s	fault/s	majflt/s	pgfree/s	pgscank/s	pgscand/s	pgsteal/s	%vmeff
12:00:01 AM	0.01	1.01	11.97	0.00	6.42	0.00	0.00	0.00	0.00
12:10:01 AM	0.00	1.82	20.39	0.00	15.21	0.00	0.00	0.00	0.00
12:20:01 AM	0.00	0.74	3.02	0.00	3.73	0.00	0.00	0.00	0.00
12:30:01 AM	0.00	0.77	3.03	0.00	3.74	0.00	0.00	0.00	0.00
12:40:01 AM	0.00	0.80	3.02	0.00	3.76	0.00	0.00	0.00	0.00
12:50:01 AM	0.00	0.80	4.74	0.00	4.31	0.00	0.00	0.00	0.00
01:00:01 AM	0.00	0.86	11.95	0.00	6.43	0.00	0.00	0.00	0.00
01:10:01 AM	0.00	1.88	20.39	0.00	15.21	0.00	0.00	0.00	0.00
01:20:01 AM	0.00	0.73	3.02	0.00	3.73	0.00	0.00	0.00	0.00
01:30:01 AM	0.00	0.83	3.03	0.00	3.76	0.00	0.00	0.00	0.00
01:40:01 AM	0.00	0.74	3.03	0.00	3.75	0.00	0.00	0.00	0.00
01:50:01 AM	0.00	0.80	3.03	0.00	3.71	0.00	0.00	0.00	0.00
02:00:01 AM	0.00	0.85	11.94	0.00	6.42	0.00	0.00	0.00	0.00
02:10:01 AM	0.00	1.84	20.56	0.00	15.23	0.00	0.00	0.00	0.00
02:20:01 AM	0.00	0.83	3.03	0.00	3.73	0.00	0.00	0.00	0.00
02:30:01 AM	0.00	0.74	4.75	0.00	4.32	0.00	0.00	0.00	0.00
02:40:01 AM	0.00	0.83	3.02	0.00	3.76	0.00	0.00	0.00	0.00
03:00:01 AM	0.00	0.75	3.02	0.00	3.73	0.00	0.00	0.00	0.00
03:10:01 AM	0.00	0.91	11.86	0.00	6.29	0.00	0.00	0.00	0.00
03:20:01 AM	0.01	3.60	24.41	0.00	16.01	0.00	0.00	0.00	0.00
03:30:01 AM	0.00	10.16	147.15	0.00	58.27	0.00	0.00	0.00	0.00
03:40:01 AM	0.00	0.75	3.03	0.00	3.78	0.00	0.00	0.00	0.00
03:50:01 AM	0.00	0.80	3.03	0.00	3.77	0.00	0.00	0.00	0.00

- Observe `pgscank`, which is the number of pages scanned by the `kswapd` daemon per second, and `pgscand`, which is the number of pages scanned directly per second.
- An “Average” line is not shown in this example.

- g. Use the `sar -w` command to display memory swapping statistics. Only a partial output is shown as follows:

```
[root@host03 ~]# sar -w
Linux 3.8.13-16.2.1.el6uek.x86_64 (host03.example.com) 12/21/2013 _x86_6

12:00:01 AM  pswpin/s pswpout/s
12:10:01 AM      0.00      0.00
12:20:01 AM      0.00      0.00
12:30:01 AM      0.00      0.00
12:40:01 AM      0.00      0.00
12:50:01 AM      0.00      0.00
01:00:01 AM      0.00      0.00
01:10:01 AM      0.00      0.00
01:20:01 AM      0.00      0.00
01:30:01 AM      0.00      0.00
01:40:01 AM      0.00      0.00
01:50:01 AM      0.00      0.00
02:00:01 AM      0.00      0.00
02:10:01 AM      0.00      0.00
02:20:01 AM      0.00      0.00
02:30:01 AM      0.00      0.00
02:40:01 AM      0.00      0.00
02:50:01 AM      0.00      0.00
03:00:01 AM      0.00      0.00
03:10:01 AM      0.00      0.00
03:20:01 AM      0.00      0.00
03:30:01 AM      0.00      0.00
03:40:01 AM      0.00      0.00
03:50:01 AM      0.00      0.00
```

- The `pswpin/s` value is the number of swap pages the system swapped in per second, and `pswpout/s` is the number of swap pages the system swapped out per second.
- An “Average” line is not shown in this example.

- h. Use the `cat` command to view the contents of `/proc/meminfo`.

```
[root@host03 ~]# cat /proc/meminfo
MemTotal:       2050816 kB
MemFree:        85756 kB
Buffers:        98244 kB
Cached:         1514344 kB
SwapCached:      0 kB
Active:          621588 kB
Inactive:       1121660 kB
Active(anon):    30036 kB
Inactive(anon): 104536 kB
Active(file):   591552 kB
Inactive(file): 1017124 kB
Unevictable:     0 kB
Mlocked:         0 kB
SwapTotal:      3876860 kB
SwapFree:       3876860 kB
Dirty:            12 kB
Writeback:        0 kB
AnonPages:      130684 kB
Mapped:          60520 kB
Shmem:           3908 kB
Slab:            168484 kB
SReclaimable:   114464 kB
SUnreclaim:     54020 kB
KernelStack:     1800 kB
PageTables:     24048 kB
NFS_Unstable:     0 kB
```

5. Observe the disk I/O statistics.

- a. Use the `iostat -xz` command to display I/O statistics for devices.

```
[root@host03 ~]# iostat -xz
Linux 3.8.13-16.2.1.el6uek.x86_64 (host03.example.com) 12/21/2013 _x86_64_ (1 CPU)

avg-cpu: %user %nice %system %iowait %steal %idle
          0.41   0.02   1.35   0.11   0.00  98.11

Device:    rrqm/s   wrqm/s     r/s     w/s   rsec/s   wsec/s  avgrrq-sz avgqu-sz  await  svctm %util
xvda       0.00     0.22    0.39    0.31    11.98    13.56    36.46     0.03   43.50    1.76   0.12
xvdb       0.00     0.00    0.00    0.00     0.02     0.00     8.39     0.00   1.91    1.79   0.00
xvdd       0.00     0.00    0.00    0.00     0.02     0.00     8.39     0.00   2.59    2.43   0.00
scd0       0.00     0.00    0.00    0.00     0.00     0.00     6.91     0.00   4.06    4.06   0.00

[root@host03 ~]#
```

- Observe the `%util` value, which is the percentage of CPU time during which I/O requests were issued to the device.
- Device saturation occurs when this value is close to 100%.

- Observe the avgqu-sz value, which is the average queue length of the requests that were issued to the device.
 - If average queue length is greater than 1, it is an indication of device I/O saturation.
- b. Use the sar -d command to display I/O statistics for devices. Only a partial output is shown as follows:

```

root@host03:~#
[root@host03 ~]# sar -d
Linux 3.8.13-16.2.1.el6uek.x86_64 (host03.example.com) 12/21/2013      _x86_64_      (1 CPU)

12:00:01 AM      DEV      tps  rd_sec/s  wr_sec/s  avgrrq-sz  avgqu-sz    await   svctm   %util
12:10:01 AM  dev202-0  0.15    0.03    2.03    14.00    0.00    1.11    0.50    0.01
12:10:01 AM  dev202-16  0.00    0.00    0.00    0.00    0.00    0.00    0.00    0.00
12:10:01 AM  dev202-48  0.00    0.00    0.00    0.00    0.00    0.00    0.00    0.00
12:10:01 AM  dev11-0   0.00    0.00    0.00    0.00    0.00    0.00    0.00    0.00
12:20:01 AM  dev202-0  0.23    0.00    3.64   15.94    0.00    1.76    0.36    0.01
12:20:01 AM  dev202-16  0.00    0.00    0.00    0.00    0.00    0.00    0.00    0.00
12:20:01 AM  dev202-48  0.00    0.00    0.00    0.00    0.00    0.00    0.00    0.00
12:20:01 AM  dev11-0   0.00    0.00    0.00    0.00    0.00    0.00    0.00    0.00

12:20:01 AM      DEV      tps  rd_sec/s  wr_sec/s  avgrrq-sz  avgqu-sz    await   svctm   %util
12:30:01 AM  dev202-0  0.11    0.00    1.48   13.66    0.00    3.51    1.57    0.02
12:30:01 AM  dev202-16  0.00    0.00    0.00    0.00    0.00    0.00    0.00    0.00
12:30:01 AM  dev202-48  0.00    0.00    0.00    0.00    0.00    0.00    0.00    0.00
12:30:01 AM  dev11-0   0.00    0.00    0.00    0.00    0.00    0.00    0.00    0.00
12:40:01 AM  dev202-0  0.11    0.00    1.55   13.65    0.00    0.40    0.19    0.00
12:40:01 AM  dev202-16  0.00    0.00    0.00    0.00    0.00    0.00    0.00    0.00
12:40:01 AM  dev202-48  0.00    0.00    0.00    0.00    0.00    0.00    0.00    0.00
12:40:01 AM  dev11-0   0.00    0.00    0.00    0.00    0.00    0.00    0.00    0.00


12:40:01 AM      DEV      tps  rd_sec/s  wr_sec/s  avgrrq-sz  avgqu-sz    await   svctm   %util
12:50:01 AM  dev202-0  0.12    0.00    1.60   13.91    0.00    0.45    0.23    0.00
12:50:01 AM  dev202-16  0.00    0.00    0.00    0.00    0.00    0.00    0.00    0.00
12:50:01 AM  dev202-48  0.00    0.00    0.00    0.00    0.00    0.00    0.00    0.00

```

- This command also provides %util and avgqu-sz statistics.
- An “Average” line is not shown in this example.

6. Observe network statistics.

- a. Use the `ip -s link` command to observe network statistics.



The screenshot shows a terminal window titled "root@host03:~". The window contains the output of the command `ip -s link`. The output lists three network interfaces: lo, eth0, and eth1. For each interface, it provides information about MTU, queueing discipline (qdisc), link layer details, and statistics for RX and TX operations. The statistics include bytes, packets, errors, dropped packets, overrun counters, and collisions. The output is as follows:

```
[root@host03 ~]# ip -s link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    RX: bytes packets errors dropped overrun mcast
        6556      78      0      0      0      0
    TX: bytes packets errors dropped carrier collsns
        6556      78      0      0      0      0
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP qlen 1000
    link/ether 00:16:3e:00:01:03 brd ff:ff:ff:ff:ff:ff
    RX: bytes packets errors dropped overrun mcast
        84986092  186190      0      6      0      0
    TX: bytes packets errors dropped carrier collsns
        6793583   51719      0      0      0      0
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP qlen 1000
    link/ether 00:16:3e:00:02:03 brd ff:ff:ff:ff:ff:ff
    RX: bytes packets errors dropped overrun mcast
        4214922   110919      0      0      0      0
    TX: bytes packets errors dropped carrier collsns
        258       3      0      0      0      0
[root@host03 ~]#
```

- The number of bytes received and bytes transmitted on each interface is provided by RX: bytes and TX: bytes values.
- This command also provides the number of packets transmitted and received, the number of errors, dropped packets, overruns, and collisions.
- Frames are dropped and the overrun counter is incremented when the capacity of the interface is exceeded.

- b. Use the `ifconfig` command to observe network statistics. Only a partial output is shown as follows:



The screenshot shows a terminal window titled "root@host03:~". The window contains the output of the `ifconfig` command. The output details three network interfaces: `eth0`, `eth1`, and `lo`. For each interface, it shows the link layer configuration (HWaddr), IP addresses (inet), IPv6 addresses (inet6), and various statistics (RX/TX packets, errors, dropped, overruns, collisions, frame, carrier, txqueuelen). The `eth0` interface has an IP address of 192.0.2.103 and a MAC address of 00:16:3E:00:01:03. The `eth1` interface has an IP address of fe80::216:3eff:fe00:103 and a MAC address of 00:16:3E:00:02:03. The `lo` interface is a loopback interface with an IP address of 127.0.0.1 and a MAC address of ::1/128.

```
root@host03:~ [root@host03 ~]# ifconfig
eth0      Link encap:Ethernet HWaddr 00:16:3E:00:01:03
          inet addr:192.0.2.103 Bcast:192.0.2.255 Mask:255.255.255.0
          inet6 addr: fe80::216:3eff:fe00:103/64 Scope:Link
            UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
            RX packets:186242 errors:0 dropped:6 overruns:0 frame:0
            TX packets:51719 errors:0 dropped:0 overruns:0 carrier:0
            collisions:0 txqueuelen:1000
            RX bytes:84988068 (81.0 MiB) TX bytes:6793583 (6.4 MiB)

eth1      Link encap:Ethernet HWaddr 00:16:3E:00:02:03
          inet6 addr: fe80::216:3eff:fe00:203/64 Scope:Link
            UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
            RX packets:110971 errors:0 dropped:0 overruns:0 frame:0
            TX packets:3 errors:0 dropped:0 overruns:0 carrier:0
            collisions:0 txqueuelen:1000
            RX bytes:4216898 (4.0 MiB) TX bytes:258 (258.0 b)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1 Mask:255.0.0.0
          inet6 addr: ::1/128 Scope:Host
            UP LOOPBACK RUNNING MTU:65536 Metric:1
            RX packets:78 errors:0 dropped:0 overruns:0 frame:0
            TX packets:78 errors:0 dropped:0 overruns:0 carrier:0
            collisions:0 txqueuelen:0
            RX bytes:6556 (6.4 KiB) TX bytes:6556 (6.4 KiB)
```

- This command provides similar output as the `ip -s link` command.

- c. Use the `netstat -s` command to observe network statistics. Only a partial output is shown as follows:
- If the `ip` and `ifconfig` commands show an excessive amount of errors, more information can be found by examining the `netstat -s` output.

```

root@host03:~#
File Edit View Search Terminal Help
[root@host03 ~]# netstat -s
Ip:
      50309 total packets received
      0 forwarded
      0 incoming packets discarded
      50309 incoming packets delivered
      51533 requests sent out
Icmp:
      31 ICMP messages received
      0 input ICMP message failed.
      ICMP input histogram:
        destination unreachable: 6
        echo requests: 25
      264 ICMP messages sent
      0 ICMP messages failed
      ICMP output histogram:
        destination unreachable: 239
        echo replies: 25
IcmpMsg:
      InType3: 6
      InType8: 25
      OutType0: 25
      OutType3: 239
Tcp:
      23 active connections openings
      10 passive connection openings
      9 failed connection attempts

```

- This command displays summary statistics for each protocol.
- Observe the number of segments retransmitted as an indicator of network interface saturation.
- Many performance problems associated with the network involve retransmission of TCP packets.

- d. Use the `netstat -i` command to observe a table listing of network interfaces.

Iface	MTU	Met	RX-OK	RX-ERR	RX-DRP	RX-OVR	TX-OK	TX-ERR	TX-DRP	TX-OVR	Flg
eth0	1500	0	186341	0	6	0	51719	0	0	0	BMRU
eth1	1500	0	111070	0	0	0	3	0	0	0	BMRU
lo	65536	0	78	0	0	0	78	0	0	0	LRU

- Observe the RX-ERR and TX-ERR values for any receive and transmit errors.

7. Use the System Monitor GUI to display system resource usage.
- This application requires that you access the GNOME desktop.

- a. Use the `exit` command to close the `ssh` connection.

```
# exit  
logout  
Connection to host03 closed.
```

- b. From `dom0`, run the `xm list -l host03 | grep location` command to determine the VNC port number for `host03`.


```
[dom0]# xm list -l host03 | grep location  
(location 0.0.0.0:5900)  
(location 2)
```

- In this example, the VNC port number is 5900. This may not be true in your case.


- c. Run the `vncviewer&` command:

```
[dom0]# vncviewer&
```


- The **VNC Viewer: Connection Details** dialog box is displayed as shown in the following screenshot:




- d. Enter the `localhost:<port_number>` command, substituting the VNC port number displayed from the previous `xm list` command. For example, if the port number is 5900, enter `localhost:5900` as shown and click **OK**.




- e. The GNOME desktop login window appears as shown in the following screenshot.
Select **Oracle Student** from the list of users.



- f. You are prompted for the **Oracle Student** password as shown in the following screenshot. Enter the password `oracle` and click **Log In**.



- g. The GNOME desktop appears. Right-click the desktop to display the pop-up menu as shown in the following screenshot. Click **Open in Terminal** to open a terminal window.




- h. From the terminal window, use the `su -` command to become the `root` user. The root user's password is `oracle`.

```
$ su -  
Password: oracle  
#
```

- i. From the terminal window, enter the `gnome-system-monitor` command to display the System Monitor GUI.

```
# gnome-system-monitor
```

The System Monitor is displayed as shown in the following screenshot:



- j. Click the other tabs and observe the information displayed on each tab:
 - 1) System tab
 - 2) Processes tab
 - 3) File Systems tab
- k. Click **Monitor > Quit** to exit the System Monitor GUI.
 - Do not exit the GNOME desktop.

Practice 19-3: Installing and Using OSWatcher

Overview

In this practice, you install and run the OSWatcher Black Box (OSWbb) product and view the collected data. Refer to Doc ID 580513.1 for “How To Start OSWatcher Black Box Every System Boot.”

Assumptions

- You are the `root` user on the **host03** VM.
- You are the `root` user on **dom0**.

Tasks

1. Use the `sftp` to copy the `oswbb703.tar` file from **dom0** to **host03**.
 - a. From **host03**, use the `service` command to verify that `sshd` is running. Start the service if necessary.

```
[host03]# service sshd status
opensh-demon (pid ...) is running...
```

- In this example, the service is running.

- b. From **dom0**, use the `cd` command to change to the `/OVS/seed_pool/sfws` directory.

```
[dom0]# cd /OVS/seed_pool/oswbb
```

- c. Use the `ls` command to view the directory for the OSWbb TAR file.

```
[dom0]# ls *tar
oswbb703.tar
```

- d. Use the `sftp` command to connect to **host03** as `root`. Password is `oracle`.

```
[dom0]# sftp root@host03
root@host03's password: oracle
sftp>
```

- e. Use the `put` command to copy the OSWbb TAR file to **host03**.

```
sftp> put oswbb703.tar
Uploading oswbb473.tar to /root/oswbb703.tar
oswbb703.tar ...
```

- f. Use the `quit` command to exit `sftp`.

```
sftp> quit
```

Perform all remaining steps in this practice from **host03**.

2. From **host03**, install, start, and stop OSWbb.
 - a. Use the `cd` command to change to the `root` user's home directory.

```
# cd
```

- b. Use the `tar` command to extract the `oswbb703.tar` file.

```
# tar xvf oswbb703.tar
oswbb/
...
```

- c. Use the `cd` command to change to the `oswbb` directory, and then use `ls` to view the contents of the directory.

```
# cd oswbb
# ls
analysis/      iosub.sh      oswnet.sh      tarupfiles.sh
...
extras.txt     oswbba.jar    startOSWbb.sh
...
```

- Notice the `startOSWbb.sh` file, which is the script used to start OSWbb.

- d. Use the `startOSWbb.sh` command to start OSWbb.

```
# ./startOSWbb.sh
Info...You did not enter a value for snapshotInterval.
Info...Using default value = 30
Info...You did not enter a value for archiveInterval.
Info...Using default value = 48
Setting the archive log directory to /root/oswbb/archive

Testing for discover of OS Utilities
VMSTAT found on your system
IOSTAT found on your system
MPSTAT found on your system
NETSTAT found on your system
TOP found on your system

Testing for discover of OS CPU COUNT
...
CPU COUNT = 1
Discovery completed.
Starting OSWatcher Black Box v7.0 ...
With SnapshotInterval = 30
With ArchiveInterval = 48
...
Starting Data Collection...
oswbb heartbeat:...
oswbb heartbeat:...
oswbb heartbeat:...
...
```

- OSWbb started successfully, the discovery process completed and data collection begins.
 - The default intervals (`snapshotInterval = 30` and `archiveInterval = 48`) are used.
- e. After a few data collection events (`oswbb heartbeat`) have completed, use the `stopOSWbb.sh` command to stop OSWbb.

```
# ./stopOSWbb.sh
```

3. View the data collection directories.

Use the `cd` command to change to the `archive` directory, and then use `ls` to view the contents of the directory.

```
# cd archive
# ls
oswiostat/    oswmpstat/    oswprvtnet/    oswslabinfo/
oswvmstat/    oswmeminfo/    oswnetstat/    oswps/        oswtop/
```

- The `archive` directory is created when OSWbb is started for the first time.
- The directory contains nine subdirectories, one for each data collector.

4. View the `oswiostat` directory.

- a. Use the `cd` command to change to the `oswiostat` directory, and then use `ls` to view the contents of the directory.

```
# cd oswiostat
# ls
host03.example.com_iostat...
```

- b. Use the `less` command to view the file.

```
# less host03...
Linux OSW v7.0
zzz ***...
...
```

- Notice that this file contains the output of the `iostat` command.
- The `iostat` command ran at 30-second intervals (the value of `snapshotInterval`).
- Each interval begins with `zzz ***` characters followed by a time stamp.

5. View the `oswmpstat` directory.

- a. Use the `cd` command to change to the `oswmpstat` directory, and then use `ls` to view the contents of the directory.

```
# cd ../oswmpstat
# pwd
/root/oswbb/archive/oswmpstat
# ls
host03.example.com_mpstat...
```

- b. Use the `less` command to view the file.

```
# less host03...
Linux OSW v7.0
zzz ***...
...
```

- Notice that this file contains the output of the `mpstat` command.
- The `mpstat` command ran at 30-second intervals (the value of `snapshotInterval`).
- Each interval begins with `zzz ***` characters followed by a time stamp.

6. View the `oswprvtnet` directory.

Use the `cd` command to change to the `oswprvtnet` directory, and then use `ls` to view the contents of the directory.

```
# cd ../oswprvtnet
# pwd
/root/oswbb/archive/oswprvtnet
# ls
```

- Notice that this directory is empty.
- This directory contains the status of RAC private networks.
- You need to manually create the `private.net` file to run `traceroute` commands.

7. View the `oswslabinfo` directory.

- a. Use the `cd` command to change to the `oswslabinfo` directory, and then use `ls` to view the contents of the directory.

```
# cd ../oswslabinfo
# pwd
/root/oswbb/archive/oswslabinfo
# ls
host03.example.com_slabinfo...
```

- b. Use the `less` command to view the file.

```
# less host03...
zzz ***...
slabinfo - version: 2.1
...
```

- Notice that this file contains the contents of the `/proc/slabinfo` file.
- The `/proc/slabinfo` file is read at 30-second intervals (the value of `snapshotInterval`).
- Each interval begins with `zzz ***` characters followed by a time stamp.

8. View the `oswvmstat` directory.

- a. Use the `cd` command to change to the `oswvmstat` directory, and then use `ls` to view the contents of the directory.

```
# cd ../oswvmstat  
# pwd  
/root/oswbb/archive/oswvmstat  
# ls  
host03.example.com_vmstat...
```

- b. Use the `less` command to view the file.

```
# less host03...  
Linux OSW v7.0 ...  
SNAP_INTERVAL 30  
CPU_COUNT 1  
zzz ***...  
...
```

- Notice that this file contains the output of the `vmstat` command.
- The `vmstat` command ran at 30-second intervals (the value of `snapshotInterval`).
- Each interval begins with `zzz ***` characters followed by a time stamp.

9. View the `oswmeminfo` directory.

- a. Use the `cd` command to change to the `oswmeminfo` directory, and then use `ls` to view the contents of the directory.

```
# cd ../oswmeminfo  
# pwd  
/root/oswbb/archive/oswmeminfo  
# ls  
host03.example.com_meminfo...
```

- b. Use the `less` command to view the file.

```
# less host03...  
zzz ***...  
MemTotal: ...  
MemFree: ...  
...
```

- Notice that this file contains the contents of the `/proc/meminfo` file.
- The `/proc/meminfo` file is read at 30-second intervals (the value of `snapshotInterval`).
- Each interval begins with `zzz ***` characters followed by a time stamp.

10. View the `oswnetstat` directory.

- Use the `cd` command to change to the `oswnetstat` directory, and then use `ls` to view the contents of the directory.

```
# cd ../oswnetstat  
# pwd  
/root/oswbb/archive/oswnetstat  
# ls  
host03.example.com_netstat...
```

- Use the `less` command to view the file.

```
# less host03...  
Linux OSW v7.0  
zzz ***...  
...
```

- Notice that this file contains the output of the `netstat` command.
- The `netstat` command ran at 30-second intervals (the value of `snapshotInterval`).
- Each interval begins with `zzz ***` characters followed by a time stamp.

11. View the `oswps` directory.

- Use the `cd` command to change to the `oswps` directory, and then use `ls` to view the contents of the directory.

```
# cd ../oswps  
# pwd  
/root/oswbb/archive/oswps  
# ls  
host03.example.com_ps...
```

- Use the `less` command to view the file.

```
# less host03...  
Linux OSW v7.0  
zzz ***...  
...
```

- Notice that this file contains the output of the `ps` command.
- The `ps` command ran at 30-second intervals (the value of `snapshotInterval`).
- Each interval begins with `zzz ***` characters followed by a time stamp.

12. View the `ostop` directory.

- Use the `cd` command to change to the `ostop` directory, and then use `ls` to view the contents of the directory.

```
# cd ../oswtop  
# pwd  
/root/oswbb/archive/oswtop  
# ls  
host03.example.com_top...
```

- Use the `less` command to view the file.

```
# less host03...  
Linux OSW v7.0  
zzz ***...  
...
```

- Notice that this file contains the output of the `top` command.
- The `top` command ran at 30-second intervals (the value of `snapshotInterval`).
- Each interval begins with `zzz ***` characters followed by a time stamp.

Practice 19-4: Using OSWatcher Black Box Analyzer

Overview

In this practice, you perform the following:

- Start OSWatcher Black Box Analyzer (OSWbba) on **host03**
- View CPU and Memory Graphs from OSWbba
- Use OSWbba to analyze the data
- View the analysis report

Assumptions

- You are the `root` user on **host03**.
- You completed Practice 19-3: Installing and Using OSWatcher.

Tasks

1. Start OSWatcher Black Box Analyzer (OSWbba) on **host03**.

- a. Use the `java -version` command to display the Java version number.

```
# java -version
java version "1.7.0_45"
...
```

- In this example, version `1.7.0_45` is installed.
- The minimum version is `1.4.2`.

- b. Use the `java -jar` command to display the OSWbba menu. First, ensure that you are in the directory where OSWbba is installed (`~/oswbb`).

```
# cd ~/oswbb
# java -jar oswbba.jar -i ~/oswbb/archive
...
Enter 1 to Display CPU Process Queue Graphs
Enter 2 to Display CPU Utilization Graphs
Enter 3 to Display CPU Other Graphs
Enter 4 to Display Memory Graphs
Enter 5 to Display Disk IO Graphs

Enter 6 to Generate All CPU Gif Files
Enter 7 to Generate All Memory Gif Files
Enter 8 to Generate All Disk Gif Files

Enter L to Specify Alternate Location of Gif Directory
Enter T to Specify Different Time Scale
Enter D to Return to Default Time scale
Enter R to Remove Currently Displayed Graphs
```


```
Enter A to Analyze Data
Enter S to Analyze Subset of Data(Changes analysis dataset
including graph time scale)

Enter P to Generate A Profile
Enter X to Export Parsed Data to File
Enter Q to Quit Program

Please Select an Option:
```


2. View CPU and Memory Graphs. A sample output is shown as follows. Your graphs will be different.

- a. From the OSWbba menu, select 1 to display the CPU Process Queue graph.




- After viewing the graph, select R from the menu to remove the graph.

- b. From the OSWbba menu, select 2 to display the CPU Utilization graph.




- After viewing the graph, select R from the menu to remove the graph.

- c. From the OSWbba menu, select 3 to display the CPU Other graph.



- After viewing the graph, select R from the menu to remove the graph.

- d. From the OSWbba menu, select 4 to display the Memory Graph.



- After viewing the graph, select R from the menu to remove the graph.

3. Run the OSWbba analyzer and view the report.

- a. From the OSWbba menu, select A to analyze the collected data and produce a report. Provide a file name or press Enter to accept the default name.

```
...
Enter A to Analyze Data
...
Please Select an Option:a
Enter a fully qualified analysis file name or enter <CR> to
accept default name: ENTER
A new analysis file analysis/host03...txt has been created.
```

- This message “A new analysis file analysis/host03...txt has been created” appears and the menu is then re-displayed.

- You could alternatively run the analyzer from the command line by using the following command (do not run this command, this is information only):

```
# java -jar oswbba.jar -i ~/oswbb/archive -A
```

- b. Exit the OSWbba menu by selecting option Q to quit.
- c. From the **host03** command line, use the **ls** command to display the analysis file in the **~/oswbb/analysis** directory.

```
# ls ~/oswbb/analysis  
host03.example.com...txt
```

- Your analysis file name is different from the sample file name shown.
- d. Use the **less** command to view the analysis file.
- Use your analysis file name observed in the previous step.

```
# less ~/oswbb/analysis/host03*  
...
```

- Notice the file has system information at the top followed by sections such as:
 - Section 1: System Status
 - Section 2: System Slowdown
 - Section 2.1: System Slowdown RCA Ordered By Impact
 - Section 3: System General Findings
 - Section 4: CPU Detailed Findings
 - Section 4.1: CPU Run Queue:
 - Section 4.2: CPU Utilization: Percent Busy
 - Section 4.3: CPU Utilization: Percent Sys
 - Section 5: Memory Detailed Findings
 - Section 5.1: Memory: Process Swap Queue
 - Section 5.2: Memory: Scan Rate
 - Section 5.3: Memory: Page In:
 - Section 5.4: Memory: Page Tables (Linux only):
 - Section 5.5: Top 5 Memory Consuming Processes Beginning
 - Section 5.6: Top 5 Memory Consuming Processes Ending
 - Section 6: Disk Detailed Findings
 - Section 6.1: Device Percent Busy Findings
 - Section 6.2: Device Service Times Findings
 - Section 6.3: Device Wait Queue Times Findings
 - Section 7: Network Detailed Findings
 - Section 7.1: Network Data Link Findings
 - Section 7.2: Network IP Findings
 - Section 7.3: Network UDP Findings
 - Section 7.4: Network TCP Findings
 - Section 8: Process Detailed Findings

- Section 8.1: PS Process Summary Ordered By Time
 - Section 8.2: PS for Processes With Status = D or T Ordered By Time
 - Section 8.3: PS for (Processes with CPU > 0) When System Idle CPU < 30% Ordered By Time
 - Section 8.4: Top VSZ Processes Increasing Memory Per Snapshot
 - Section 8.5: Top RSS Processes Increasing Memory Per Snapshot
- e. After viewing the analysis file, press **q** to exit the **less** command.
4. Terminate the **dd** command started in Practice 19-2 (3b).

In the second terminal window on **host03** in which you initiated the following **dd** command, press **Ctrl + C** to terminate the command.

```
# dd if=/dev/zero of=/dev/null bs=1024  
CTRL-C
```

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 20: System Logging

Chapter 20

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 20: System Logging

Practices Overview

In these practices, you configure system logging, use `rsyslog` templates to format message logs, and install and run `logwatch`.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practice 20-1: Viewing the System Log File Configuration

Overview

In this practice, you view the system logging configuration file, modify the file, and observe the impact of the modifications. You also configure log file rotation.

Assumptions

You are the root user on the **host03** VM.

Tasks

1. Explore the main configuration file for system logging, /etc/rsyslog.conf.
 - a. Use the less command to view the system logging configuration file.

```
# less /etc/rsyslog.conf
# rsyslog v5 configuration file
...
#####
#MODULES #####
$ModLoad imuxsock.so # provides support for local system...
$ModLoad imklog.so    # provides kernel logging support...
...
#####
#GLOBAL DIRECTIVES #####
# use default timestamp format
$ActionFileDefaultTemplate RSYSLOG_TraditionalFileFormat
...
#####
#RULES #####
...
# Log all kernel messages to the console.
# Logging much else clutters up the screen
#kern.*                                     /dev/console

# Log anything (except mail) of level info or higher.
# Don't log private authentication messages!
*.info;mail.none;authpriv.none;cron.none    /var/log/messages

# The authpriv file has restricted access.
authpriv.*                                    /var/log/secure

# Log all the mail messages in one place.
mail.*                                         -/var/log/maillog

# Log cron stuff
cron.*                                         /var/log/cron
...
```

- b. Use the up-arrow and down-arrow keys to view the various sections of the file.
 - The **MODULES** section uses the \$ModLoad command to load the modules.
 - The **GLOBAL DIRECTIVES** section specifies configuration options.
 - The **RULES** section defines a filter (facility.priority) and an action.
 - c. Press the letter **q** to exit the “less” command.
2. Change the action for cron logging.
 - a. Use the `ls` command to list the `cron*` files in the `/var/log` directory.
 - Because this is a freshly installed system, you might not have any rotated log files (files with a date stamp).

```
# ls /var/log/cron*
cron    cron-20131222
```

- b. Use the `vi` editor to modify the system logging configuration file. Change the action for cron logging to log to a different log file: `/var/log/cron_new`.

```
# vi /etc/rsyslog.conf
...
# Log cron stuff
cron.*          /var/log/cron          (old entry)
cron.*          /var/log/cron_new        (new entry)
```

- c. Use the `service` command to restart the `rsyslog` service.

```
# service rsyslog restart
Shutting down system logger:      [  OK  ]
Starting system logger:           [  OK  ]
```

- d. Use the `crontab -e` command to create a cron job that runs the `ls` command every minute.

```
# crontab -e
* * * * * ls
```

- e. Use the `ls` command to list the `cron*` files in the `/var/log` directory.

```
# ls /var/log/cron*
cron    cron-20131222    cron_new
```

- Notice the new log file, `cron_new`.

- f. Use the `tail` command to view the last two lines in the `cron` log file.

```
# tail -2 /var/log/cron
Dec 13 11:22:51 host03 crontab[8129]: (root) LIST (root)
Dec 13 11:23:01 host03 CROND[8129]: (root) CMD (root)
```

- This is sample output only.

- g. Use the `head` command to view the first entries in the `cron_new` log file.

```
# head /var/log/cron_new
Dec 13 11:24:01 host03 CROND[8178]: (root) CMD (ls)
Dec 13 11:25:01 host03 CROND[8196]: (root) CMD (ls)
...
```

- Notice from the date_time stamps that the new log entries are being written to `cron_new`.

- h. Use the `vi` editor to modify the system logging configuration file. Change the action for `cron` logging back to the original log file.

```
# vi /etc/rsyslog.conf
...
# Log cron stuff
cron.*          /var/log/cron_new      (old entry)
cron.*          /var/log/cron          (new entry)
```

- i. Use the `service` command to restart the `rsyslog` service.

```
# service rsyslog restart
Shutting down system logger: [ OK ]
Starting system logger: [ OK ]
```

- j. Use the `tail` command to ensure that `cron` is now logging to `/var/log/cron`.

```
# tail /var/log/cron
...
Dec 13 11:44:01 host03 CROND[8404]: (root) CMD (ls)
# tail /var/log/cron_new
...
Dec 13 11:43:01 host03 CROND[8376]: (root) CMD (ls)
```

- Notice that the entry in `cron` has a later time stamp than the entry in `cron_new`.
- You might need to wait a minute for the `cron` job to run.

- k. Use the `rm` command to delete the `cron_new` log file.

```
# rm /var/log/cron_new
rm: remove regular file 'cron_new'? y
```

- l. Use `crontab -r` to remove the crontab.

```
# crontab -r
```

3. Configure `rsyslog` to log debug messages.

- a. Use the `vi` editor to modify the system logging configuration file. Add an entry at the bottom of the file to log all debug messages to `/var/log/debug`.

```
# vi /etc/rsyslog.conf
...
* .debug          /var/log/debug
```

- b. Use the `service` command to restart the `rsyslog` service.

```
# service rsyslog restart
Shutting down system logger: [ OK ]
Starting system logger: [ OK ]
```

- c. Use the `logger` command to generate an **informational** log message.

```
# logger -p info "This is an info-priority message"
```

- d. Use the `tail` command to view the log files.

```
# tail /var/log/messages
...
Dec 13 13:36:12 host03 oracle: This is an info-priority message
# tail /var/log/debug
...
Dec 13 13:36:12 host03 oracle: This is an info-priority message
```

- Notice that the message was written to both log files.

- e. Use the `logger` command to generate a **debug-level** log message.

```
# logger -p debug "This is a debug-priority message"
```

- f. Use the `tail` command to view the log files.

```
# tail /var/log/messages
...
Dec 13 13:36:12 host03 oracle: This is an info-priority message
# tail /var/log/debug
...
Dec 13 13:42:16 host03 oracle: This is a debug-priority message
```

- Notice that the debug-level message was written only to `/var/log/debug`.

- g. Use the `vi` editor to modify `/etc/rsyslog.conf` and remove the entry at the bottom of the file to log all debug messages.

- Remove this entry: `*.debug /var/log/debug`.

- h. Use the `service` command to restart the `rsyslog` service.

```
# service rsyslog restart
Shutting down system logger: [ OK ]
Starting system logger: [ OK ]
```

4. Configure log file rotation.

- a. Use the `ls` command to view the contents of the `/var/log` directory.

```
# ls /var/log/messages*
/var/log/messages  /var/log/messages-20131222
# ls /var/log/maillog*
/var/log/maillog  /var/log/maillog-20111211
# ls /var/log/cron*
/var/log/cron  /var/log/cron-20111211
```

- Notice that some files in `/var/log` have numbers at the end of the file name.
- These numbers represent a rotated log with the time stamp added to the log file name.
- You may not have log files with a time stamp appended to the file name. It depends on how long your system has been running.

- b. Use the `vi` editor to modify the `/etc/logrotate.conf` configuration file. Change the frequency of the default log file rotation from **weekly** to **daily**.

```
# vi /etc/logrotate.conf
...
# rotate log files weekly
weekly          (old entry)
daily           (new entry)
```

- Your log files now rotate daily after making this change.
- Because this is a freshly installed system, there will be no rotated log files.

Practice 20-2: Using rsyslog Templates

Overview

In this practice, you use `rsyslog` templates to format `rsyslog` output.

Assumptions

You are the `root` user on the **host03** VM.

Tasks

1. Use the `vi` editor to modify `/etc/rsyslog.conf` and define a template.

Add a new line at the bottom of the file as follows:

```
$template class, "Message: %msg%\n"
```

- This entry creates a template named `class`.

2. Continuing editing `/etc/rsyslog.conf` and create a log file that uses the template.

Add a new line at the bottom of the file as follows:

```
*.* /var/log/class.log;class
```

- This entry must be after the entry that defined the template.
- This entry writes all messages to the `/var/log/class.log` file and formats the entries using the template `class`.

3. After saving the changes to `/etc/rsyslog.conf`, use the `service` command to restart the `rsyslog` service.

```
# service rsyslog restart
Shutting down system logger: [ OK ]
Starting system logger: [ OK ]
```

4. Use the `cat` command to view the `/var/log/class.log` file.

```
# cat /var/log/class.log
Message: ...
...
```

- Notice that all entries are preceded by the text “Message:” followed by the actual message, as defined in the template `class`.

5. Use the `vi` editor to modify `/etc/rsyslog.conf` and modify the `class` template.

Change the template definition as follows:

```
$template class, "Time: %timestamp%, Facility: %syslogfacility-text%, Priority: %syslogpriority-text%, Hostname: %hostname%, Message: %msg%\n"
```

6. After saving the changes to /etc/rsyslog.conf, use the service command to restart the rsyslog service.

```
# service rsyslog restart
Shutting down system logger: [ OK ]
Starting system logger: [ OK ]
```

7. Use the cat command to view the /var/log/class.log file.

```
# cat /var/log/class.log
Message: ...
...
Time: ...
...
```

- Notice that the newest entries now include the Time, Facility, Priority, Hostname, and Message properties, as defined in the template class.
8. Use the vi editor to modify /etc/rsyslog.conf and delete the following entries you added in this practice.
- \$template class,"Time: %timestamp%, Facility: %syslogfacility-text%, Priority: %syslogpriority-text%, Hostname: %hostname%, Message: %msg%\n"
 - *.* /var/log/class.log;class
9. After saving the changes to /etc/rsyslog.conf, use the service command to restart the rsyslog service.

```
# service rsyslog restart
Shutting down system logger: [ OK ]
Starting system logger: [ OK ]
```

10. Use the rm command to remove the /var/log/class.log file.

```
# rm /var/log/class.log
rm: remove regular file '/var/log/class.log'? y
```

Practice 20-3: Using logwatch

Overview

In this practice, you install the `logwatch` package, view the main configuration file (the `cron` file), and run the `logwatch` utility from the command line.

Assumptions

You are the `root` user on the **host03** VM.

Tasks

1. Install the `logwatch` package.
 - a. Use the `rpm` command to determine whether the `logwatch` package is already installed.
 - In this example, the `logwatch` package is not installed.
 - b. Use the `yum` command to install the `logwatch` package.

```
# rpm -q logwatch
• In this example, the logwatch package is not installed.

# yum install logwatch
...
Setting up Install Process
Resolving Dependencies
--> Running transaction check
...
--> Finished Dependency Resolution
Dependencies Resolved
=====
Package
=====
Installing:
logwatch
Installing for dependencies:
perl-Date-Manip
perl-YAML-Syck
Transaction Summary
=====
Install      3 Package(s)
Total download size: 1.7 M
Installed size: 11 M
Is this ok [y/N]: y
Downloading packages
...
Running rpm_check debug
```

```
Running Transaction Test
...
Installing : perl-YAML-Syck-1.07-4.el6.x86_64
Installing : perl-Date-Manip-6.24-1.el6.noarch
Installing : logwatch-7.3.6-49.el6.noarch
...
Installed:
logwatch.noarch 0:7.3.6-52.el6
Dependency Installed:
perl-Date-Manip.noarch 0:6.24-1.el6      perl-YAML-Syck.x86_64...
Complete!
```

- c. Use the `rpm` command to verify that `logwatch` has successfully installed.

```
# rpm -q logwatch
logwatch-7.3.6-52.el6.noarch
```

2. View `logwatch` files.

- a. Use the `find` command to list all `logwatch` files.

```
# find / -name "*logwatch*"
/var/cache/logwatch
/var/lib/yum/yumdb...logwatch-7.3.6-52.el6.noarch
/usr/sbin/logwatch
...
/usr/share/logwatch/default.conf/logwatch.conf
/usr/share/logwatch/scripts/logwatch/pl
...
/etc/cron.daily/0logwatch
/etc/logwatch
/etc/logwatch/conf/logwatch.conf
...
```

- b. Use the `less` command to view the main `logwatch` configuration file.

```
# less /usr/share/logwatch/default.conf/logwatch.conf
...
```

- Notice various configurable items such as the following:

- LogDir
- TmpDir
- MailTo
- Print
- Save
- Range
- Detail

- Service
- DailyReport

- c. Use the less command to view the logwatch cron file.

```
# less /etc/cron.daily/0logwatch  
...
```

3. Run logwatch from the command line.

- a. Run the logwatch --help command to view the logwatch help.

```
# logwatch --help  
Usage: /usr/sbin/logwatch [--detail <level>] [--logfile <name>]  
[--print] [--mailto <addr>] ...
```

- b. Run logwatch and display output to your screen.

```
# logwatch --print  
##### logwatch 7.3.6 (05/19/07) #####  
Processing Initiated: Fri Jan...  
Date Range Processed: yesterday  
...
```

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 21: Troubleshooting

Chapter 21

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practices for Lesson 21: Troubleshooting

Practices Overview

In these practices, you perform a variety of troubleshooting exercises and fix some common problems. Specific problems you encounter include the following:

- System boots into single-user mode by default.
- Status utilities are not producing expected output.
- A cron job fails to run.
- A user cannot log in.
- File system does not mount.
- Logical volume space is exhausted.
- You cannot ping remote hosts.
- There are NFS permission problems.
- You cannot log in to remote hosts using ssh.
- Log file is not getting updated.

Each practice:

- Tells you what problem the activity is designed to simulate
- Tells you how to set up the problem
- Gives you some hints or things to check
- Has an associated “Solution” that provides steps to diagnose and resolve the problem

Each solution:

- Includes only the steps to resolve the problem
- Does not include steps to set up the problem

Appendix:

- An appendix exists, which provides the source code for the executables that cause the problems that you need to troubleshoot and fix.

Practice 21-1: Transferring Utilities from dom0

Overview

In this practice, you transfer files hosted on **dom0** to **host01**. In subsequent practices, you are instructed to run one of these executables that causes the problem that you need to troubleshoot and fix.

Assumptions

You are the `root` user on the **host01** VM.

Tasks

1. Use the `sftp` command to transfer files from **dom0** to **host01**.
 - a. From the root user's home directory on **host01**, use the `sftp` command to connect to **dom0**.
 - Answer "yes" when prompted, "Are you sure..."
 - Provide `root` password of `oracle` when prompted.

```
# cd
# pwd
/root
# sftp dom0
The authenticity of host 'dom0 (192.0.2.1)' can't be
established. RSA key fingerprint is ...
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '192.0.2.1' (RSA) to the list of
known hosts.
root@dom0's password: oracle
sftp>
```

- b. From the `sftp` prompt, use the `cd` command to change to the `/OVS/seed_pool/new_stuff/ts_scripts` directory on **dom0**. Use the `ls` command to list the directory contents.

```
sftp> cd /OVS/seed_pool/ts_scripts
sftp> ls
two.x
three.x
four.x
five.x
six.x
eight.x
nine.x
ten.x
eleven.x
sftp>
```

- c. From the `sftp` prompt, use the `mget *` command to upload all the files in the `/OVS/seed_pool/ts_scripts` directory from **dom0** to **host01**.

```
sftp> mget *
Fetching .../two.x to two.x
...
sftp>
```

- d. Use the `quit` command to exit `sftp`.

```
sftp> quit
```

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practice 21-2: System Boots into Single-User Mode

Overview

In this practice, your system boots into single-user mode by default, which is not what you want.

Assumptions

You are the `root` user on the **host01** VM.

Tasks

1. Display the current run level.

Use the `who -r` command to view the current run level.

```
# who -r
run-level 3 <date_time>
```

- Notice the run level is 3.

2. Execute the `two.x` program from the `root` user's home directory.

```
# cd
# pwd
/root
# ./two.x
```

3. Reboot the system.

- a. Use the `reboot` command to reboot your system:

```
# reboot
...
```

- After you reboot your system, your VNC session closes.

- b. Connect to **host01** by using VNC.

- 1) Run the `vncviewer&` command.

```
# vncviewer&
```

- The VNC Viewer: Connection Details dialog box is displayed.

- 2) Enter the command `localhost :<port_number>`, substituting the correct port number for the **host01** guest. For example, if the port number is 5900, enter the following and click **OK**:

```
localhost:5900
```

- Notice you did not need to log in. You were automatically logged in as `root`.

- c. Use the `who -r` command to view the current run level.

```
# who -r
run-level S <date_time>
```

- Notice the run level is `S` or single-user mode.

4. Diagnose and fix the problem so that the system does not boot into single-user mode.

- Review Lesson 4: Linux Boot Process
 - Which file specifies the default run level? Is the default run level set to s?
 - Which file allows you to view kernel boot parameters? Is single a kernel boot parameter?
 - Which file allows you to specify kernel boot parameters? Is single specified as a kernel boot parameter?
 - How do you change kernel boot parameters at boot time?
 - How do you permanently change kernel boot parameters?
5. Reboot the system after problem is fixed to confirm it does not boot into single-user mode.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Solution 21-2: System Boots into Single-User Mode

Steps

1. Determine why the system boots into single-user mode.

- Use the `cat` command to view the default run level set in the `/etc/inittab` file.

```
# cat /etc/inittab
...
id:3:initdefault:
```

- Notice the default run level is 3, not S. Therefore, this is not causing the problem.

- Use the `cat` command to view the kernel boot parameters in the `/proc/cmdline` file.

```
# cat /proc/cmdline
ro root=UUID=... rd_NO_LUKS rd_NO_LVM LANG=en_US.UTF-8 rd_NO_MD
SYSFONT=latacyrheb-sun16 KEYBOARDTYPE=pc KEYTABLE=us rd_NO_DM
rhgb quiet single
```

- Notice the word `single` is a kernel boot parameter that is causing your system to boot into single-user mode.

- Use the `grep` command and search for the word “`single`” in the `/boot/grub/grub.conf` file.

```
# grep single /boot/grub/grub.conf
kernel /vmlinuz-3.8.13-16... single
```

- Notice the word `single` appears on the kernel line for the `3.8.13-16` kernel.
- This would cause the `3.8.13-16` kernel to boot into single-user mode by default.

- Use the `uname -r` command to determine which kernel is running.

```
# uname -r
3.8.13-16.2.1.el6uek.x86_64
```

- The `3.8.13-16` kernel is running. The word `single` needs to be removed from the associated kernel line in the `/boot/grub/grub.conf` file.

2. Fix the system so it does not boot into single-user mode.

- Use the `vi` editor to modify `/boot/grub/grub.conf` and remove the word `single` at the end of the `3.8.13-16` kernel line.

```
# vi /boot/grub/grub.conf
...
default=0
...
title Oracle Linux Server Unbreakable Kernel (3.8.13-16...
    root (hd0,0)
    kernel /vmlinuz-3.8.13-16... single
...
```

- Press `Ctrl + D` to continue the boot process.

- The login prompt appears after the boot process completes.

```
# CTRL-d  
...  
host01 login:
```

- Log in as root. Password is oracle.

```
host01 login: root  
Password: oracle
```

- Use the `who -r` command to view the current run level.

```
# who -r  
run-level 3 <date_time> last=S
```

- Notice the run level is 3. The last run level was S.

- Confirm the system does not boot into single-user mode.

- Use the `reboot` command to reboot your system:

```
# reboot  
...
```

- After you reboot your system, your VNC session closes.

- Connect to **host01** by using VNC.

- Run the `vncviewer&` command.

```
# vncviewer&  
• The VNC Viewer: Connection Details dialog box is displayed.
```

- Enter the command `localhost:<port_number>`, substituting the correct port number for the **host01** guest. For example, if the port number is 5900, enter the following and click **OK**:

```
localhost:5900
```

- Notice you do need to log in this time.

- Log in as root. The password is oracle.

```
host01 login: root  
Password: oracle
```

- Use the `who -r` command to view the current run level.

```
# who -r  
run-level 3 <date_time>
```

- Notice the run level is now 3 and not single-user.

Practice 21-3: Status Commands Fail

Overview

In this practice, you notice that some of the “status” utilities are not producing the expected output. You diagnose and fix this problem.

Assumptions

You are the `root` user on the **host01** VM.

Tasks

1. Execute the `three.x` program from the `root` user's home directory.

```
# cd
# pwd
/root
# ./three.x
```

2. Run some “status” tools and notice the errors.

- a. Run the `mpstat` command.

```
# mpstat
Linux 3.8.13-16...
Cannot open /proc/stat: No such file or directory
```

- b. Run the `iostat` command.

```
# iostat
Linux 3.8.13-16...
Cannot open /proc/stat: No such file or directory
```

- c. Run the `netstat` command.

```
# netstat
Active Internet connections (w/o servers)
Proto Recv-Q Send-Q Local Address Foreign Address State
Active UNIX domain sockets (w/o servers)
Proto RefCnt Flags      Type      State      I-Node Path
```

- d. Run the `route` command.

```
# route
/proc/net/route: No such file or directory
INET (IPv4) not configured in this system.
```

- e. Run the `ifconfig` command.

```
# ifconfig
Warning: cannot open /proc/net/dev (No such file or directory).
Limited output.
...
```

- Notice that none of these tools produce the expected output.

3. Diagnose and fix the problem with the “status” tools.
 - Review Lesson 5: System Configuration
 - Commands that are failing are commands that provide information about the current state of the kernel.
 - Which file system contains a hierarchy of special files that represent the current state of the kernel?
 - How do you check if this file system is mounted?
 - If this file system is not mounted, how do you mount it?
4. Re-run some of the previous commands that failed to ensure that the problem was fixed.
 - a. Run the `mpstat` command.
 - b. Run the `iostat` command.
 - c. Run the `netstat` command.
 - d. Run the `route` command.
 - e. Run the `ifconfig` command.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Solution 21-3: Status Commands Fail

Steps

1. Some of the output indicates a problem with the `proc` file system. Diagnose the problem.

- a. Use the `ls` command to display the contents of `/proc`.

```
# ls /proc
```

- No output is not normal and suggests the `proc` file system is not mounted.

- b. Run the `mount` command to display mounted file systems.

```
# mount
```

```
/dev/xvda2 on / type ext4 (rw)
sysfs on /sys type sysfs (rw)
...
none on /proc/sys/fs/binfmt_misc type binfmt_misc (rw)
```

- Notice that `proc` is not mounted.

- c. Run the `df -h` command to display mounted file systems.

```
# df -h
```

```
df: `/proc/sys/fs/binfmt_misc': No such file or directory
Filesystem      Size  Used  Avail   Use%  Mounted on
/dev/xvda2      5.8G  2.2G  3.3G   40%   /
...

```

- Notice this output reports that `/proc/sys/fs/binfmt_misc` does not exist.

2. Fix the problem.

- a. View the `/etc/fstab` file to ensure the entry to mount `proc` exists.

```
# cat /etc/fstab
```

```
...
```

```
proc          /proc      proc      defaults 0 0
```

- Notice there is an entry to mount the `proc` file system in the `/etc/fstab` file.

- b. Run the `mount -a` command to re-mount entries in the `/etc/fstab` file.

```
# mount -a
```

3. Re-run some of the previous “status” commands to ensure that the problem is fixed.

- a. Run the `mpstat` command.
- b. Run the `iostat` command.
- c. Run the `netstat` command.
- d. Run the `route` command.
- e. Run the `ifconfig` command.

Practice 21-4: cron Job Fails to Run

Overview

In this practice, you diagnose and fix a problem that is preventing a cron job from running.

Assumptions

You are the root user on the host01 VM.

Tasks

1. Create a crontab for the root user.

- a. Use the crontab -e command to create a crontab job that runs the vmstat command every minute.

```
# crontab -e
* * * * * vmstat
```

- The crontab -e command uses the vi editor. Save your changes and exit vi.

- b. After one minute, use the tail command to view the last few lines in the /var/log/cron log file.

```
# tail /var/log/cron
...
<date_time> host01 CROND[...] (root) CMD (vmstat)
```

- Notice there is an entry in this file that states the date and time that the cron daemon ran the vmstat command.
- You also have mail, because the output from cron jobs is sent to the user's mailbox.

- c. Use the mail command to view the results of your job.

```
# mail
...
N # Cron Daemon <date_time> ... "Cron <root@host01> vm"
&
```

- Notice you have mail that contains the output of the vmstat command.

- d. View the details of mailbox entries by entering the associated number and then press **Enter**.

- To re-display the header, press **h** and then press **Enter**.
- To quit the mail program, press **q** and then press **Enter**.

- e. Press **q** and then **Enter** to quit the mail program.

```
& q
```

2. Execute the `four.x` program from the `root` user's home directory.

```
# cd  
# pwd  
/root  
# ./four.x
```

3. Edit the `crontab` for the `root` user.

- a. Use the `crontab -e` command to edit the `crontab` job and replace `vmstat` with `iostat`.

```
# crontab -e  
* * * * * vmstat  
* * * * * iostat
```

- b. After one minute, use the `tail` command to view the last few lines in the `/var/log/cron` log file.

```
# tail /var/log/cron  
...
```

- Notice there is NO entry in this file that states that the `iostat` command ran.

- c. Use the `mail` command to view your incoming mailbox.

```
# mail  
...  
&
```

- Notice there is NO mail entry containing the output of the `iostat` command.

- d. Press `q` and then **Enter** to quit the `mail` program.

```
& q
```

4. Diagnose and fix the problem of the `cron` job failing to run.

- Review Lesson 8: Automating Tasks

- Does the `root` user have a `crontab` entry?
 - Is the `root` user's `crontab` entry valid?
 - Are the permissions correct on the `root` user's `crontab` file?
 - Is the `cron` daemon running?

5. Verify that the `cron` job is running every minute.

- a. Run the `mail` command to view the output of the `iostat` command.
 - b. Use the `tail` command to view `/var/log/cron` and ensure the `cron` job is running every minute.

6. Remove the `root` user's `crontab` before continuing to the next practice.

Solution 21-4: cron Job Fails to Run

Steps

1. Diagnose the problem of the cron job failing to run.

- a. Use the crontab -l command to list the crontab entry.

```
# crontab -l
* * * * * iostat
```

- Notice the crontab entry exists and the format is valid.

- b. Use the ls -l command to view the root user's crontab entry in the /var/spool/cron directory.

```
# ls -l /var/spool/cron
-rw-----. root root ... root
```

- The root file exists with the proper permissions (read-write for the owner).

- c. Use the cat command to view the /var/spool/cron/root file.

```
# cat /var/spool/cron/root
* * * * * iostat
```

- This confirms the crontab entry is valid.

- d. Use the service command to view the status of the crond daemon.

```
# service crond status
crond is stopped
```

- Notice the crond service is not running. This is the cause of the problem.

2. Fix the problem of the cron job failing to run.

Use the service command to start the crond daemon.

```
# service crond start
Starting crond: [ OK ]
```

3. Verify that the cron job is running every minute.

- a. After one minute, use the tail command to view the last few lines in the /var/log/cron log file.

```
# tail /var/log/cron
...
<date_time> host01 CROND[...] (root) CMD (iostat)
```

- Notice there is now an entry in this file that states the date and time that the cron daemon ran the iostat command.

- b. Use the `mail` command to view the results of your job.

```
# mail  
...  
N # Cron Daemon <date_time> ... "Cron <root@host01> io"
```

- Notice you have mail that contains the output of the `iostat` command.

- c. Press `q` and then **Enter** to quit the mail program.

```
& q
```

4. Remove the `root` user's `crontab` before continuing to the next practice.

```
# crontab -r
```

Practice 21-5: User Cannot Log In

Overview

In this practice, you diagnose and fix a problem which is preventing a user from being able to log in.

Assumptions

You are the `root` user on the **host01** VM.

Tasks

1. Add a new user.
 - a. Use the `useradd john` command to add user `john`.

```
# useradd john
```

- b. Use the `passwd john` command to assign a password (`password`) to user `john`.
 - Disregard the “BAD PASSWORD” warning.

```
# passwd john
Changing password for user john.
New password: password
BAD PASSWORD: it is based on a dictionary word
Retype new password: password
passwd: all authentication tokens updated successfully.
```

2. Verify you can log in as user `john`.
 - a. Use the `exit` command to log out as user `root`.

```
# exit
logout
```

 - The “host01 login” prompt appears.
 - b. Log in as user `john`. Password is `password`. Verify you can successfully log in as user `john`.

```
host01 login: john
Password: password
$ pwd
/home/john
$ whoami
john
```

- You can successfully log in as user `john`.
- c. Use the `exit` command to log out as user `john`.

```
$ exit
logout
```

- The “host01 login” prompt appears.

3. Log in as root and execute the five.x program from the root user's home directory.

- Password is oracle.

```
host01 login: root  
Password: oracle  
# ./five.x
```

4. Attempt to log in as user john from the login prompt.

- a. Use the exit command to log out as user root.

```
# exit  
logout
```

- The "host01 login" prompt appears.

- b. Log in as the john user. Password is oracle.

```
host01 login: john  
Password: password  
Login incorrect  
login:
```

- Notice you cannot log in as user john.

5. Diagnose and fix the problem of user john not being able to log in.

- Review Lesson 10, "User and Group Administration".
 - Which configuration file contains usernames? Is the user john present in this file?
 - Which configuration file contains user passwords? Is the user john present in this file?
 - If entries are missing in these files, how do you re-create the entries?
 - If passwords are corrupted or forgotten, how do you re-create passwords?

6. Verify you can successfully log in as user john.

Solution 21-5: User Cannot Log In

Steps

1. Diagnose and fix the problem of user john not being able to log in.

- a. Log in as the root user. Password is oracle.

```
host01 login: root  
Password: oracle
```

- b. Use the grep command to search for john in the /etc/passwd file.

```
# grep john /etc/passwd  
john:x:501:501::/home/john:/bin/bash
```

- Notice the john entry is present.

- c. Use the grep command to search for john in the /etc/shadow file.

```
# grep john /etc/shadow  
#john:$6$...:0:99999:7:::
```

- Notice the john entry is commented out.

- d. Use the vi editor to modify /etc/shadow. Remove the # character from the beginning of the john line.

- Because /etc/shadow is a read-only file, use :wq! to write, save, and exit vi after removing the # character.

```
# vi /etc/shadow  
...  
#john:$6$...:0:99999:7::: (old entry)  
john:$6$...:0:99999:7::: (new entry)
```

- Alternatively you could run the following command to re-create a password entry in the /etc/shadow file for user john.

```
# passwd john  
Changing password for user john.  
...
```

- e. Use the exit command to log out as user root.

```
# exit  
logout
```

2. Verify you can log in as user `john`.

- a. Attempt to log in as user `john` from the log in prompt. Password is `password`.

```
host01 login: john  
Password: password  
$ whoami  
john
```

- Notice you can now log in as user `john`.

- b. Use the `exit` command to log out as user `john`.

```
$ exit  
logout
```

Practice 21-6: File System Troubleshooting

Overview

In this practice, you diagnose and fix an unmountable file system.

Assumptions

You are the `root` user on the **host01** VM.

Tasks

1. Display the partition table on `/dev/xvdb`.

Use the `fdisk -l` command to display the partition table on `/dev/xvdb`.

```
# fdisk -l /dev/xvdb
Disk /dev/xvdb: 5368 MB, 5368709120 bytes
...
Device Boot Start End Blocks Id System
/dev/xvdb1      1 652 5237158+ 83 Linux
```

- Notice there is one partition on `/dev/xvdb` that consists of the entire device.

2. Make a file system on `/dev/xvdb1`.

Use the `mkfs` command to make an **ext4** file system on `/dev/xvdb1`.

```
# mkfs -t ext4 /dev/xvdb1
mke2fs 1.43-WIP (20-Jun-2013)
Filesystem label=
OS type: Linux
...
Writing inode tables: done
Creating journal (32768 blocks): done
...
```

3. Mount the file system.

- a. Use the `mkdir` command to create `/Test`.

```
# mkdir /Test
```

- b. Use the `mount` command to mount `/dev/xvdb1` on `/Test`.

```
# mount /dev/xvdb1 /Test
```

4. Copy files to `/Test`.

Use the `cp` command to copy `/boot/init*` to `/Test`. Use the `ls` command to display the contents of the `/Test` directory.

```
# cp /boot/init* /Test
# ls /Test
initramfs-2.6.32.431.el6.x86_64.img          lost+found
initramfs-3.8.13-16.2.1.el6uek.x86_64.img
```

5. Unmount the file system on /dev/xvdb1.

Use the `umount` command to unmount /Test.

```
# umount /Test
```

6. Execute the `six.x` program from the `root` user's home directory.

```
# cd  
# pwd  
/root  
# ./six.x
```

7. Attempt to mount the file system on /dev/xvdb1 on /Test.

Use the `mount` command as shown in Task 3b.

```
# mount /dev/xvdb1 /Test  
mount: you must specify the filesystem type
```

- The `mount` command fails.

8. Diagnose and fix the problem with the file system.

- Review Lesson 11, “File Systems”.
 - How do you specify the file system type with the `mount` command?
 - What commands are available to check and repair file systems?

9. Ensure you can mount the file system on /Test.

- Ensure the files (`initramfs*`) exists on /Test.

Solution 21-6: File System Troubleshooting

Steps

1. Diagnose the problem.

- a. Run the `mount` command but include the `-t ext4` option to specify the file system type.

```
# mount -t ext4 /dev/xvdb1 /Test
mount: wrong fs type, bad option, bad superblock on /dev/xvdb1,
      missing codepage or helper program, or other error
      in some cases useful info is found in syslog - try
      dmesg | tail or so
```

- b. Per the preceding suggestion, run the `dmesg | tail` command.

```
# dmesg | tail
...
EXT4-fs (xvdb1): VFS: Can't find ext4 filesystem
```

- Notice the VFS error: “Can’t find ext4 filesystem” on xvdb1.

- c. Use the `tail` command to view the last lines in the `/var/log/messages` file.

```
# tail /var/log/messages
...
<date_time> host01 kernel: EXT4-fs (xvdb1): VFS: Can't find ext4
filesystem
```

- Notice the error message written to the `/var/log/messages` file is similar to the error message given by the `dmesg` command.

- d. Use the `tune2fs -l /dev/xvdb1` command to list the contents of the file system superblock.

```
# tune2fs -l /dev/xvdb1
tune2fs 1.43-WIP (20-Jun-2013)
tune2fs: Bad magic number in super-block while trying to open
/dev/xvdb1
Couldn't find valid filesystem superblock
```

- Notice the “superblock” error.

- e. Use the `dumpe2fs /dev/xvdb1` command to dump file system information.

```
# dumpe2fs /dev/xvdb1
dumpe2fs 1.43-WIP (20-Jun-2013)
dumpe2fs: Bad magic number in super-block while trying to open
/dev/xvdb1
Couldn't find valid filesystem superblock
```

- Notice this command reports the same “superblock” error on `/dev/xvdb1`.

2. Fix the file system.

- a. Use the `fsck /dev/xvdb1` command to check (and repair) the file system.

- Press **Enter** to accept the default **y** (yes) answer.

```
# fsck /dev/xvdb1
fsck from util-linux-ng 2.17.2
e2fsck 1.43-WIP (20-Jun-2013)
ext2fs_open2: Bad magic number in super-block
fsck.ext2: Superblock invalid, trying backup blocks...
/dev/xvdb1 was not cleanly unmounted, check forced.
Resize inode not valid. Recreate<y>? ENTER
Pass 1: Checking inodes, blocks, and sizes
...
Free blocks count wrong for group #0 ...
Fix<y>? ENTER
Free blocks count wrong for group #1 ...
Fix<y>? ENTER
Free blocks count wrong ...
Fix<y>? ENTER
Free inodes count wrong for group #0 ...
Fix<y>? ENTER
Free inodes count wrong ...
Fix<y>? ENTER
/dev/xvdb1: ***** FILE SYSTEM WAS MODIFIED *****
/dev/xvdb1: ...
```

- Run the `fsck /dev/xvdb1` command a second time to check the file system.

```
# fsck /dev/xvdb1
fsck from util-linux-ng 2.17.2
e2fsck 1.43-WIP (20-Jun-2013)
/dev/xvdb1: clean ...
```

- Notice that the file system is fixed (clean).

- Mount the file system.

- Use the `mount` command to mount `/dev/xvdb1` on `/Test`.

```
# mount /dev/xvdb1 /Test
```

- The `mount` command is successful. The `fsck` command fixed the corrupt superblock.

- Use the `ls` command to list the contents of the `/Test` directory.

```
# ls /Test
initramfs-2.6.32.431.el6.x86_64.img          lost+found
initramfs-3.8.13-16.2.1.el6uek.x86_64.img
```

- All original files are present in the directory.

Practice 21-7: Logical Volume Space is Exhausted

Overview

In this practice, you exhaust space in a logical volume and expand the file system to fix the problem.

Assumptions

You are the root user on the **host01** VM.

Tasks

1. Create three new partitions on /dev/xvdd.
 - a. Change the system ID on each new partition to 8e (Linux LVM).

a. Use the `fdisk` command to create a 1 GB **primary** partition on /dev/xvdd by using the following parameters:

```
# fdisk /dev/xvdd
...
Command (m for help): n
Command action
  e   extended
  p   primary partition (1-4)
p
Partition number (1-4): 1
First cylinder (1-652, default 1): ENTER
Using default value 1
Last cylinder, +cylinders or +size{K,M,G} (1-652, default 652):
+1G
```

b. Use the “t” command to change the system ID on the new partition.

```
Command (m for help): t
Selected partition 1
Hex code (type L to list codes): 8e
Changed system type of partition 1 to 8e (Linux LVM)
```

c. Continue using `fdisk` to create a second 1 GB **primary** partition on /dev/xvdd by using the following parameters:

```
Command (m for help): n
Command action
  e   extended
  p   primary partition (1-4)
p
Partition number (1-4): 2
First cylinder (133-652, default 133): ENTER
Using default value 133
```

```
Last cylinder, +cylinders or +size{K,M,G} (133-652, default  
652) : +1G
```

- d. Use the “t” command to change the system ID on the new partition.

```
Command (m for help): t  
Partition number (1-4): 2  
Hex code (type L to list codes): 8e  
Changed system type of partition 2 to 8e (Linux LVM)
```

- e. Continue using fdisk to create a third 1 GB **primary** partition on /dev/xvdd by using the following parameters:

```
Command (m for help): n  
Command action  
    e   extended  
    p   primary partition (1-4)  
p  
Partition number (1-4): 3  
First cylinder (265-652, default 265): ENTER  
Using default value 265  
Last cylinder, +cylinders or +size{K,M,G} (265-652, default  
652) : +1G
```

- f. Use the “t” command to change the system ID on the new partition.

```
Command (m for help): t  
Partition number (1-4): 3  
Hex code (type L to list codes): 8e  
Changed system type of partition 3 to 8e (Linux LVM)
```

- g. Print the new partition table.

```
Command (m for help): p  
Disk /dev/xvdb: 5368 MB, 5368709120 bytes  
...  
  Device Boot Start End Blocks Id System  
  /dev/xvdd1          1 132 1060258+ 8e Linux LVM  
  /dev/xvdd2         133 264 1060290 8e Linux LVM  
  /dev/xvdd3         265 396 1060290 8e Linux LVM
```

- h. Save the new partition table.

```
Command (m for help): w  
The partition table has been altered!  
Calling ioctl() to re-read partition table.  
Syncing disks.
```

2. Initialize the new partitions for use by LVM (create physical volumes).

- a. List the partitions with the Linux LVM (**8e**) system ID.

```
# fdisk -l | grep 8e
/dev/xvdd1            1    132   1060258+    8e    Linux LVM
/dev/xvdd2            133   264   1060290    8e    Linux LVM
/dev/xvdd3            265   396   1060290    8e    Linux LVM
```

- b. Use the **pvccreate** command to create physical volumes on all three partitions.

```
# pvccreate /dev/xvdd[123]
Physical volume "/dev/xvdd1" successfully created
Physical volume "/dev/xvdd2" successfully created
Physical volume "/dev/xvdd3" successfully created
```

3. Create a volume group.

- a. Use the **vgcreate** command to create a volume group named **myvolg** from the three physical volumes.

```
# vgcreate myvolg /dev/xvdd[123]
Volume group "myvolg" successfully created
```

- b. Use the **vgs** command to display attributes of the volume group.

```
# vgs
VG      #PV #LV #SN Attr   VSize   VFree
myvolg     3   0   0 wz--n- 3.02g  3.02g
```

- c. Use the **pvs** command to display information about the physical volumes.

```
# pvs
PV      VG     Fmt   Attr  PSize   PFree
/dev/xvdd1  myvolg lvm2  a--  1.01g  1.01g
/dev/xvdd2  myvolg lvm2  a--  1.01g  1.01g
/dev/xvdd3  myvolg lvm2  a--  1.01g  1.01g
```

4. Create a logical volume.

- a. Use the **lvcreate** command to create a 4 MB logical volume named **myvol** from the **myvolg** volume group.

```
# lvcreate -L 4m -n myvol myvolg
Logical volume "myvol" created
```

- b. Use the **lvs** command to display attributes of the logical volume.

```
# lvs
LV      VG     Attr   LSize Pool Origin Data% Move Log Cpy%...
myvol  myvolg -wi-a- 4.00m
```

5. Create and mount a file system on the logical volume.

- a. List the /dev entries for the **myvol** logical volume.

```
# ls -l /dev/myvolg/myvol
lrwxrwxrwx.  /dev/myvolg/myvol -> ../dm-0
# ls -l /dev/mapper/myvolg-myvol
lrwxrwxrwx.  /dev/mapper/myvolg-myvol -> ../dm-0
```

- b. Create an **ext4** file system on the **myvol** logical volume.

```
# mkfs.ext4 /dev/mapper/myvolg-myvol
mke2fs 1.43-WIP (20-Jun-2013)
Filesystem label=
OS type: Linux
...
Writing inode tables: done
Creating journal (1024 blocks): done
...
```

- c. Create a /myvol mount point.

```
# mkdir /myvol
```

- d. Mount the file system.

```
# mount /dev/mapper/myvolg-myvol /myvol
```

- e. Run the **df -h** command to display mounted file systems.

```
# df -h
Filesystem      Size  Used  Avail   Use%  Mounted on
...
/dev/mapper/myvolg-myvol
        2.9M    29K    2.7M     2%    /myvol
```

- f. Use the **vi** editor to add the following line to /etc/fstab.

```
/dev/mapper/myvolg-myvol /myvol ext4 defaults 0 0
```

6. Exhaust space on logical volume.

- a. Use the **cd** command to change to the /myvol directory.

```
# cd /myvol
```

- b. Use the **dd if=/dev/zero of=bigfile bs=1024M count=5120** command to attempt to create a 5 MB file on /myvol.

```
# dd if=/dev/zero of=bigfile bs=1024 count=5120
dd: writing `bigfile': No space left on device
2910+0 records in
2909+0 records in
2978816 bytes (3.0 MB) copied...
```

- Notice there is no space left on the device.

- c. Use the `ls -l` command to display the contents of the `/myvol` directory.

```
# ls -l /myvol
-rw-r--r--. 1 root root 2978816 <date_time> bigfile
...
```

- d. Run the `df -h` command to display mounted file systems.

```
# df -h
Filesystem      Size  Used  Avail   Use%  Mounted on
...
/dev/mapper/myvolg-myvol
        2.9M    2.9M     0M  100%  /myvol
```

- Notice the logical volume mounted on `/myvol` is full.

7. Determine what can be done to allocate more space to `/myvol`.

- Review Lesson 12, “Storage Administration”.
 - What is the total size of the logical volume?
 - What volume group does the logical volume belong to?
 - What is the total size of the volume group?
 - Can you expand the logical volume by using available space in the volume group?
 - Do you need to add more space to the volume group before you can expand the logical volume?
 - What is the command that allows you to allocate more space from a volume group to a logical volume?

8. Allocate an additional 1 GB to `/myvol`

- After allocating more space, ensure you can successfully create a 5 MB file on `/myvol`.

Solution 21-7: Logical Volume Space is Exhausted

Steps

- Determine what is needed to increase the size of the `myvolg/myvol` logical volume and the file system.

- Use the `lvs` command to display attributes of the logical volume.

```
# lvs
  LV      VG      Attr      LSize   Pool Origin Data% Move Log Cpy%...
  myvol  myvolg -wi-ao-  4.00m
```

- Notice the total size of the `myvol` logical volume is 4 MB.
- Notice the `myvol` logical volume belongs to the `myvolg` volume group.

- Use the `vgs` command to display attributes of the volume group.

```
# vgs
  VG      #PV #LV #SN Attr      VSize   VFree
  myvolg     3   1   0 wz--n-  3.02g  3.02g
```

- Notice there is 3.02 GB available in the `myvolg` volume group.
- Therefore, you can expand the logical volume by using available space in the volume group.
- You do not need to add more space to the volume group.

- Allocate an additional 1 GB to `/myvol`

- Use the `lvextend` command to increase the size of the `myvolg/myvol` logical volume and the file system by 1 GB.
- Using the `-r` option resizes the underlying file system together with the logical volume.

```
# lvextend -L +1G -r myvolg/myvol
Extending logical volume myvol to 1.00 GiB
Logical volume myvol successfully resized
resize2fs 1.43-WIP (20-Jun-2013)
Filesystem at /dev/mapper/myvolg-myvol is mounted on /myvol; online resizing required
old_desc_blocks = 1, new_desc_blocks = 5
The file system on /dev/mapper/myvolg-myvol is now 1052672
blocks long.
```

- Use the `lvs` command to display attributes of the logical volume.

```
# lvs
  LV      VG      Attr      LSize   Pool Origin Data% Move Log Cpy%...
  myvol  myvolg -wi-ao-  1.00g
```

- Notice the `myvol` logical volume is now 1 GB in size.

- c. Run the `df -h` command to display mounted file systems.

```
# df -h
Filesystem      Size   Used   Avail   Use%   Mounted on
...
/dev/mapper/myvolg-myvol
    1011M   3.0M   987M    1%   /myvol
```

- Notice there is now 987 MB of available space on /myvol.

3. Attempt to create the `bigfile` again.

- a. Ensure the current directory is /myvol before running the `dd` command.

```
# cd /myvol
# dd if=/dev/zero of=bigfile bs=1024 count=5120
5120+0 records in
5120+0 records in
5242880 bytes (5.2 MB) copied...
```

- Notice you can now successfully create a 5 MB file on /myvol.

- b. Use the `ls -l` command to display the contents of the /myvol directory.

```
# ls -l /myvol
-rw-r--r--. 1 root root 5242880 <date_time> bigfile
...
```

- c. Run the `df -h` command to display mounted file systems.

```
# df -h
Filesystem      Size   Used   Avail   Use%   Mounted on
...
/dev/mapper/myvolg-myvol
    1012M   5.2M   985M    1%   /myvol
```

- Notice you still have 985 MB of space available on /myvol.

Practice 21-8: Network Connectivity Problem

Overview

In this practice, you diagnose and fix a network connectivity problem.

Assumptions

You are the `root` user on the **host01** VM.

Tasks

1. Verify network interface is configured properly.

Use the `ping` command to contact **dom0** and **host03**. Use CTRL-C to abort the commands after connectivity is confirmed.

```
# ping dom0
64 bytes from example.com (192.0.2.1): icmp_seq=1 ttl=64 ...
64 bytes from example.com (192.0.2.1): icmp_seq=1 ttl=64 ...
64 bytes from example.com (192.0.2.1): icmp_seq=1 ttl=64 ...
CTRL-C
# ping host03
64 bytes from example.com (192.0.2.103): icmp_seq=1 ttl=64 ...
64 bytes from example.com (192.0.2.103): icmp_seq=1 ttl=64 ...
64 bytes from example.com (192.0.2.103): icmp_seq=1 ttl=64 ...
CTRL-C
```

2. Execute the `eight.x` program from the `root` user's home directory.

```
# cd
# pwd
/root
# ./eight.x
```

3. Repeat Task 1 to test network interface configuration.

Use the `ping` command to contact **dom0** and **host03**.

```
# ping dom0
connect: Network is unreachable
# ping host03
connect: Network is unreachable
```

- The remote systems are no longer reachable.

4. Diagnose and fix the network connectivity problem.

- Review Lesson 13, “Network Configuration”
 - How do you display the configuration of your network?
 - Are the network interfaces up?
 - What are the IP addresses of **dom0** and **host03**?
 - Are **dom0** and **host03** on the same network as **host01**?

- Do you have a route to **dom0** and **host03**?
 - How do you view the route table?
 - In which files do you configure network interfaces?
 - Are these network interface configuration files configured properly?
 - Is the network service running?
5. Verify network connectivity to the remote hosts is working.
- a. Run the `ping dom0` command.
 - b. Run the `ping host03` command.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Solution 21-8: Network Connectivity Problem

Steps

1. Diagnose the network connectivity problem.

- a. Use the `ifconfig` command to display the configuration of your network interfaces.

```
# ifconfig
eth0      Link encap:Ethernet HWaddr 00:16:3E:00:01:01
          inet addr:192.0.3.101 ...
          inet6 addr: ...
          UP ...
          ...
          ...
lo       Link encap:Local Loopback
          inet addr:127.0.0.1
          ...
```

- Output shows that the `eth0` interface is `UP` and has an IP address of `192.0.3.101`.

- b. Use the `route` command to display the route table.

```
# route
Destination     Gateway         Genmask        Flags ...   Iface
link-local      *              255.255.0.0    U            eth0
192.0.3.0       *              255.255.255.0  U            eth0
```

- The route table indicates the route to the `192.0.3.0` network is through `eth0`.

- c. Use the `cat` command to view the contents of the `/etc/hosts` file.

```
# cat /etc/hosts
127.0.0.1      localhost.localdomain localhost
192.0.2.1       example.com                  dom0
192.0.2.101     host01.example.com        host01
192.0.2.102     host02.example.com        host02
192.0.2.103     host03.example.com        host03
```

- The contents of the `/etc/hosts` file indicate the IP address should be configured with `192.0.2` network addresses and not `192.0.3`.

2. Fix the network connectivity problem.

- a. Use the `vi` editor to modify `/etc/sysconfig/network-scripts/ifcfg-eth0` and change the `IPADDR` directive as shown.

```
# vi /etc/sysconfig/network-scripts/ifcfg-eth0
...
IPADDR=192.0.3.101          (old value)
IPADDR=192.0.2.101          (new value)
...
```

- b. Use the `service` command to restart the network service.
- Always restart the service whenever a configuration file is changed.

```
# service network restart
Shutting down interface eth0: [ OK ]
Shutting down loopback interface: [ OK ]
Bringing up loopback interface: [ OK ]
Bringing up interface eth0: [ OK ]
```

3. Verify network connectivity to the remote hosts is working.

Use the `ping` command to contact **dom0** and **host03**. Use CTRL-C to abort the commands after connectivity is confirmed.

```
# ping dom0
64 bytes from example.com (192.0.2.1): icmp_seq=1 ttl=64 ...
64 bytes from example.com (192.0.2.1): icmp_seq=1 ttl=64 ...
64 bytes from example.com (192.0.2.1): icmp_seq=1 ttl=64 ...
CTRL-C
# ping host03
64 bytes from example.com (192.0.2.103): icmp_seq=1 ttl=64 ...
64 bytes from example.com (192.0.2.103): icmp_seq=1 ttl=64 ...
64 bytes from example.com (192.0.2.103): icmp_seq=1 ttl=64 ...
CTRL-C
```

- Network connectivity to the remote hosts works.

Practice 21-9: NFS Permission Problem

Overview

In this practice, you diagnose and fix an NFS permission problem.

Assumptions

- You are the `root` user on the **host01** VM.
- You are the `root` user on the **host03** VM.

Tasks

Perform Task 1 from **host03**.

1. Configure **host03** as an NFS server.

- a. Use the `rpm` command to verify that the `nfs-utils` packages are installed.

```
[host03]# rpm -qa | grep nfs-utils
nfs-utils-lib-...
nfs-utils-...
```

- In this example, the `nfs-utils` packages are installed.
- b. Use the `vi` editor and edit `/etc/exports` to add an entry to export `/home/oracle` to all client systems. Include the `rw` option.

```
[host03]# vi /etc/exports
/home/oracle *(rw)
```

- c. Use the `service` command to verify that the `rpcbind` service is started.

```
[host03]# service rpcbind status
rpcbind (pid ...) is running...
```

- In this example, the `rpcbind` service is running.

- d. Verify that the `nfs` service is started.

```
[host03]# service nfs status
rpc.svcgssd is stopped
rpc.mountd is stopped
nfsd is stopped
rpc.rquotad is stopped
```

- In this example, the services are not running and must be started.

- e. Start the `nfs` service and associated services.

```
[host03]# service nfs start
Starting NFS services: [ OK ]
Starting NFS quotas: [ OK ]
Starting NFS mountd: [ OK ]
Starting NFS daemon: [ OK ]
Starting RPC idmapd: [ OK ]
```

- f. Verify that the `nfslock` service is started.

```
[host03]# service nfslock status
rpc.statd (pid ...) is running...
```

- In this example, the `nfslock` service is running.

- g. Use the `showmount` command to display exported file systems.

```
[host03]# showmount -e
Export list for host03.example.com:
/home/oracle *
```

- h. Disable the firewall by running the following command.

- An `iptables` rule could be written to allow NFS connectivity.
- But for purposes of this practice, you simply stop the `iptables` service.

```
[host03]# service iptables stop
iptables: Flushing firewall rules: [ OK ]
iptables: Setting chains to policy ACCEPT: filter [ OK ]
iptables: Unloading modules: [ OK ]
```

- i. Use the `ls` command to list the contents of `/home/oracle` on **host03**.

```
[host03]# ls /home/oracle
Desktop Documents Downloads Music Pictures Public
Templates Videos
```

Perform Tasks 2, 3, and 4 from **host01**.

2. Create a mount point in the `oracle` user's home directory on **host01**.

- a. Use the `su - oracle` command to become the `oracle` user.

```
[host01]# su - oracle
```

- b. As the `oracle` user, use the `mkdir` command to create a directory named `remote_home`.

```
[host01]$ mkdir remote_home
```

- c. Use the `exit` command to log off as the `oracle` user.

```
[host01]$ exit
logout
```

- You are now the `root` user on **host01**.

3. Mount the exported `/home/oracle` file system on **host01**.

- a. As the `root` user, use the `mount` command to mount the exported file system—`host03:/home/oracle`—on the local mount point, `/home/oracle/remote_home` with read-write permissions (`-o rw`).

```
[host01]# mount -t nfs -o rw host03:/home/oracle
/home/oracle/remote_home
```

- b. Use the `df` command to verify that the NFS file system is mounted.

```
[host01]# df -h
Filesystem           Size  Used  Avail  Use%  Mounted on
...
host03:/home/oracle   1.9G  19M  1.8G   1%  /home/oracle...
```

4. Access the remote file system as the `oracle` user.

- a. From **host01**, use the `su - oracle` command to become the `oracle` user.

- Use the `pwd` command to display the current directory.
- Use the `ls -l` command to view the content of the current directory.

```
[host01]# su - oracle
[host01]$ pwd
/home/oracle
[host01]$ ls -l
drwx----- ... remote_home
```

- Notice you are in the `oracle` user's home directory and that the `remote_home` directory exists.

- b. Use the `ls` command to list the contents of `remote_home`.

```
[host01]$ ls remote_home
Desktop  Documents  Downloads  Music  Pictures  Public
Templates  Videos
```

- Notice the contents of `/home/oracle/remote_home` on **host01** are the same as `/home/oracle` on **host03** because they are the same directories.

Perform Task 5 from both **host01** and **host03**. Note the host name in the prompt.

5. Confirm you have read-write permission on the NFS file system.

- a. As the `oracle` user on **host01**, use the `vi` editor to create a "test" file on `remote_home` to verify you have read-write permission.

```
[host01]$ cd ~/remote_home
[host01]$ pwd
/home/oracle/remote_home
[host01]$ vi test
created from host01
```

- b. From **host03**, use the `su - oracle` command to become the `oracle` user.

- Use the `pwd` command to display the current directory.
- Use the `ls` command to view the content of the current directory.

```
[host03]# su - oracle
[host03]$ pwd
/home/oracle
[host01]$ ls
```

```
Desktop Documents Downloads Music Pictures Public
Templates test Videos
```

- Notice the presence of the “test” file that was created from **host01**.
- c. As the **oracle** user on **host03**, use the **vi** editor to edit the “test” file.

```
[host03]$ vi test
created from host01
edited from host03
```

- d. From **host01**, use the **cat** command to view the contents of the edited “test” file in **/home/oracle/remote_home**.

```
[host01]$ cat ~/remote_home/test
created from host01
edited from host03
```

- This confirms you have read-write permission.

Perform Tasks 6, 7, 8, and 9 from **host01**.

6. Unmount the NFS file system.

- a. From **host01**, use the **exit** command to log out as the **oracle** user.

```
[host01]$ exit
logout
```

- You are now the **root** user on **host01**.

- b. Use the **umount** command to unmount **/home/oracle/remote_home**.

```
[host01]# umount /home/oracle/remote_home
```

7. On **host01**, execute the **nine.x** program from the **root** user’s home directory.

```
[host01]# cd
[host01]# pwd
/root
[host01]# ./nine.x
```

8. Mount the exported **/home/oracle** file system on **host01**.

- a. As the **root** user, use the **mount** command to mount the exported file system—**host03:/home/oracle**—on the local mount point, **/home/oracle/remote_home** with read-write permissions (**-o rw**).

```
[host01]# mount -t nfs -o rw host03:/home/oracle
/home/oracle/remote_home
```

- b. Use the **df** command to verify that the NFS file system is mounted.

```
[host01]# df -h
Filesystem           Size   Used   Avail   Use%   Mounted on
...
host03:/home/oracle    1.9G   18M   1.8G    1%   /home/oracle...
```

9. Access the remote file system as the **oracle** user.

- a. Use the `su - oracle` command to become the `oracle` user.
 - Use the `pwd` command to display the current directory.
 - Use the `ls -l` command to view the content of the current directory.

```
[host01]# su - oracle
[host01]$ pwd
/home/oracle
[host01]$ ls -l
drwx----- . . . remote_home
```

- Notice you are in the `oracle` user's home directory and that the `remote_home` directory exists.

- b. Use the `ls` command to list the contents of `remote_home`.

```
[host01]$ ls remote_home
ls: cannot open directory remote_home/: Permission denied
```

- Notice you get a permission error when attempting to view the contents of the NFS file system.

- c. Use the `exit` command to log off as the `oracle` user.

```
[host01]$ exit
logout
```

- You are now the root user on `host01`.

10. Diagnose and fix the cause of the permission error.

- Review Lesson 14, “File Sharing”.
 - Since the `mount` command succeeded, you know that:
 - Network connectivity from **host01** to **host03** exists
 - The file system is being shared by the server
 - Host name resolution (DNS) is not a problem
 - The firewall is not blocking NFS packets
 - Check the share options on the server.
 - Check the permissions on the NFS share.
 - Check the mount options on the client.
 - Check the UIDs and the GIDs of the `oracle` user on the client and the server.
 - View the `/var/log/secure` log file for messages that could indicate the cause of the problem.

11. After fixing the problem, ensure the `oracle` user from **host01** has read-permission on the `/home/oracle/remote_home/test` file.

Note: Task 12 is a clean-up task. Only perform this task after the NFS permission problem is fixed.

12. After ensuring the problem is fixed, unmount the NFS share and restore the hosts to their beginning state.

- a. As the root user on **host01**, use the **umount** command to unmount **/home/oracle/remote_home**.

```
[host01]# umount /home/oracle/remote_home
```

- b. As the root user on **host03**, use the **vi** editor and delete all lines in **/etc/exports**.

```
[host03]# vi /etc/exports  
/home/oracle *(rw) delete this entry
```

- c. As the root user on **host03**, use the **service** command to stop the **nfs** service.

```
[host03]# service nfs stop  
Shutting down NFS daemon: [ OK ]  
...
```

- d. As the root user on **host03**, use the **service** command to start the **iptables** service.

```
[host03]# service iptables start  
iptables: Applying firewall rules [ OK ]
```

Solution 21-9: NFS Permission Problem

Steps

1. Diagnose the cause of the permission error.

- a. As the `root` user on **host03**, view the share options on the NFS file system.

```
[host03]# cat /var/lib/nfs/etab
/home/oracle
*(rw,sync,wdelay,hide,nocrossmnt,secure,root_squash,...)
```

- The NFS share has read-write (`rw`) permission, which is correct.

- b. View the permissions on the NFS file system.

```
[host03]# ls -ld /home/oracle
drwx----- oracle oracle ... /home/oracle/
[host03]# ls -l /home/oracle/test
-rw-rw-r---- oracle oracle ... /home/oracle/test
```

- The `/home/oracle` directory has `rwx` permission for the owner (`oracle`), which is correct.
- The `/home/oracle/test` file has `rw` permission for the owner (`oracle`) and the group (`oracle`), which is correct.

- c. As the `root` user on **host01**, view the mount options on the NFS file system.

```
[host01]# mount
...
host03:/home/oracle on /home/oracle/remote_home type nfs
(rw,vers=4,addr=192.0.2.103,clientaddr=192.0.2.101)
```

- The NFS share is mounted with read-write (`rw`) permission, which is correct.

- d. From **host03**, use the `grep` command to view the UID and GID for `oracle` in the `/etc/passwd` and the `/etc/group` files.

```
[host03]# grep oracle /etc/passwd
oracle:x:500:500:Oracle Student:/home/oracle:/bin/bash
[host03]# grep oracle /etc/group
oracle:x:500:
```

- Notice the `oracle` user's UID and GID are 500 on **host03**.

- e. From **host01**, use the `grep` command to view the UID and GID for `oracle` in the `/etc/passwd` and the `/etc/group` files.

```
[host01]# grep oracle /etc/passwd
oracle:x:555:555:Oracle Student:/home/oracle:/bin/bash
[host01]# grep oracle /etc/group
oracle:x:555:
```

- Notice the `oracle` user's UID and GID are 555 on **host01**.
- This difference in UIDs and GIDs is the cause of the permission error.

- f. On **host01**, use the `tail` command to view the latest messages written to the `/var/log/secure` file.

```
[host01]# tail /var/log/secure
...
<date_time> host01 userdel[...]: delete user 'oracle'
<date_time> host01 userdel[...]: removed group 'oracle' ...
<date_time> host01 useradd[...]: new group: name=oracle, ...
<date_time> host01 useradd[...]: new user: name=oracle, ...
...
```

- Notice the `oracle` user was deleted on **host01** and re-added with UID and GID of 555.
2. Fix the permission problem by changing `oracle` UID and GID for `oracle` on **host01** from 555 to 500.
- The `oracle` UID and GID on **host01** needs to be same as `oracle` user and group (500) on **host03**.
 - If possible, the easiest way is simply to delete the `oracle` user on **host01** and then re-add with the proper UID and GID.
 - The following steps are provided to make the changes without deleting and re-adding the user.
- a. Before changing UIDs and GIDs, use the `umount` command to unmount the NFS file system on **host01**.

```
[host01]# umount /home/oracle/remote_home
```

- b. Use the `usermod -u` command to change the UID for the `oracle` user.

```
[host01]# usermod -u 500 oracle
```

- c. Use the `groupmod -g` command to change the GID for the `oracle` group.

```
[host01]# groupmod -g 500 oracle
```

- d. Use the following `find` command to change the owner permissions to the new UID.

- Redirect error messages to `/dev/null`.

```
[host01]# find / -user 555 -exec chown -h 500 {} \; 2> /dev/null
```

- e. Use the following `find` command to change the group permissions to the new GID.

- Redirect error messages to `/dev/null`.

```
[host01]# find / -group 555 -exec chgrp -h 500 {} \; 2>
/dev/null
```

- f. Use the `usermod -g` command to change the GID for the `oracle` user.

- This command may return a “`usermod: no changes`” message, which is of no concern.

```
[host01]# usermod -g 500 oracle
```

3. On **host01**, mount the exported /home/oracle file system on **host01**.
 - a. As the root user, use the `mount` command to mount the exported file system—`host03:/home/oracle`—on the local mount point, `/home/oracle/remote_home` with read-write permissions (`-o rw`).

```
[host01]# mount -t nfs -o rw host03:/home/oracle  
/home/oracle/remote_home
```

- b. Use the `df` command to verify that the NFS file system is mounted.

```
[host01]# df -h  
Filesystem           Size   Used   Avail   Use%   Mounted on  
...  
host03:/home/oracle     1.9G   5.0M   1.8G    1%   /home/oracle...
```

4. Attempt to access the remote file system as the `oracle` user.

- a. From **host01**, use the `su - oracle` command to become the `oracle` user.

- Use the `pwd` command to display the current directory.
 - Use the `ls -l` command to view the content of the current directory.

```
[host01]# su - oracle  
[host01]$ pwd  
/home/oracle  
[host01]$ ls -l  
drwx----- . . . remote_home
```

- Notice you are in the `oracle` user's home directory and that the `remote_home` directory exists.

- b. Use the `ls` command to list the contents of `remote_home`.

```
[host01]$ ls remote_home  
Desktop  Documents  Downloads  Music  Pictures  Public  
Templates  test  Videos
```

- You can now view the contents of the NFS file system.
- c. Ensure you have read-permission by using the `vi` command to edit the "test" file.

```
[host01]$ cd ~/remote_home  
[host01]$ vi test  
<make some changes to the content>
```

- You are able to successfully read and write to files on the NFS file system.

5. After ensuring the problem is fixed, unmount the NFS share and restore the hosts to their beginning state.

- a. Use the `exit` command to log off as the `oracle` user.

```
[host01]$ exit  
logout
```

- You are now the `root` user on **host01**.

- b. As the root user on **host01**, use the **umount** command to un-mount **/home/oracle/remote_home**.

```
[host01]# umount /home/oracle/remote_home
```

- c. As the root user on **host03**, use the **vi** editor and delete all lines in **/etc/exports**.

```
[host03]# vi /etc/exports  
/home/oracle *(rw)                                     delete this entry
```

- d. As the root user on **host03**, use the **service** command to stop the **nfs** service.

```
[host03]# service nfs stop  
Shutting down NFS daemon:                                [ OK ]  
...
```

- e. As the root user on **host03**, use the **service** command to start the **iptables** service.

```
[host03]# service iptables start  
iptables: Applying firewall rules                         [ OK ]
```

Practice 21-10: Remote Access Problem

Overview

In this practice, you diagnose and fix a remote access problem.

Assumptions

- You are the `root` user on the **host01** VM.
- You are the `root` user on the **host03** VM.

Tasks

1. Ensure remote connectivity is working from **host03** to **host01**.

- As the `root` user on **host03**, confirm that you can use the `ssh` command to connect to **host01**.
 - Password is `oracle`.
 - Use the `hostname` command to confirm that you did connect.

```
[host03]# ssh host01
root@host01's password: oracle
Last login: ...
[host01]# hostname
host01.example.com
```

- Use the `logout` command to log off.

```
[host01]# logout
Connection to host01 closed.
```

2. Ensure remote connectivity is working from **host01** to **host03**.

- As the `root` user on **host01**, confirm that you can use the `ssh` command to connect to **host03**.
 - Password is `oracle`.
 - Use the `hostname` command to confirm that you did connect.

```
[host01]# ssh host03
root@host03's password: oracle
Last login: ...
[host03]# hostname
host03.example.com
```

- Use the `logout` command to log off.

```
[host03]# logout
Connection to host03 closed.
```

3. From **host01**, execute the `ten.x` program from the `root` user's home directory.

```
[host01]# cd
[host01]# pwd
```

```
/root  
[host01]# ./ten.x
```

4. Ensure remote connectivity is working from **host01** to **host03**.
 - a. As the **root** user on **host01**, confirm that you can use the **ssh** command to connect to **host03**.

```
[host01]# ssh host03  
root@host03's password: oracle  
Last login: ...  
[host03]# hostname  
host03.example.com
```

- b. Use the **logout** command to log off.
5. Ensure remote connectivity is working from **host03** to **host01**.

As the **root** user on **host03**, confirm that you can use the **ssh** command to connect to **host01**.

```
[host03]# ssh host01  
ssh: connect to host host01 port 22: No route to host
```

- Notice the **ssh** command fails.
6. Diagnose and fix the remote connectivity problem.
 - Think about what you can and what you cannot do:
 - You can **ssh** from **host01** to **host03**
 - You cannot **ssh** from **host03** to **host01**
 - Therefore, the problem is most likely a configuration issue on **host01**.
 - Can you **ping** **host01**? If so, then there is nothing wrong with the network interface configuration.
 - Review Lesson 15, “OpenSSH”.
 - Ensure the **sshd** service is configured properly and that the service is running.
 - Notice the error message, “**ssh: connect to host host01 port 22: No route to host**.”
 - View the **/var/log/messages** log file for entries related to **ssh** or port 22.
 - View the **/var/log/secure** log file for entries related to **ssh** or port 22.
 - Review Lesson 16, “Pluggable Authentication Modules (PAM)”.
 - Are there any PAM modules that are causing the problem with **sshd**?
 - Review Lesson 17, “Security Administration”.
 - Is the firewall prohibiting a connection on port 22?
 - Is there a TCP wrapper configured that is causing the problem?
 7. After fixing the problem, ensure you can **ssh** from **host03** to **host01**.

Solution 21-10: Remote Access Problem

Steps

1. Attempt to diagnose the problem by testing network connectivity.

From **host03**, use the `ping host01` command to test network connectivity to **host01**.

```
[host03]# ping host01
64 bytes from example.com (192.0.2.101): icmp_seq=1 ttl=64 ...
64 bytes from example.com (192.0.2.101): icmp_seq=1 ttl=64 ...
64 bytes from example.com (192.0.2.101): icmp_seq=1 ttl=64 ...
CTRL-C
```

- Notice the `ping` command verifies network connectivity to **host01** exists.

2. Attempt to diagnose the problem by checking `sshd` configuration.

a. From **host01**, use the `service` command to view the status of the `sshd` service.

```
[host01]# service sshd status
openSSH-daemon (pid ...) is running...
```

- The `sshd` service is running.

b. From **host03**, use the `service` command to view the status of the `sshd` service.

```
[host03]# service sshd status
openSSH-daemon (pid ...) is running...
```

- The `sshd` service is running.

c. Use the `grep` command to search for "ssh" in the `/var/log/messages*` files.

- Repeat the command but search for "port 22".

```
[host01]# grep ssh /var/log/messages*
[host01]# grep "port 22" /var/log/messages*
```

- Notice that neither command returned any output.

d. Use the `grep` command to search for "ssh" in the `/var/log/secure*` files.

- Repeat the command but search for "port 22".

```
[host01]# grep ssh /var/log/secure*
...
[host01]# grep "port 22" /var/log/secure*
...
```

- Both commands returned output but the `<date_time>` stamp on the log file entries does not correspond to the `<date_time>` of the problem.

3. Attempt to diagnose the problem by checking PAM configuration.

- It is unlikely the problem is caused by PAM.
- If PAM were the cause of the problem, a log file entry would usually be written to the `/var/log/secure` file.
- However, some of the steps used to investigate PAM configuration follow.

- a. Use the `ls` command to determine if there are any PAM modules related to ssh.

```
[host01]# ls /etc/pam.d/*ssh*
/etc/pam.d/sshd           /etc/pam.d/ssh-keycat
```

- Notice there are two PAM configuration files that are related to ssh.

- b. Use the `cat` command to view the contents of the `/etc/pam.d/sshd` file.

```
[host01]# cat /etc/pam.d/sshd
#%PAM-1.0
auth      required    pam_sepermit.so
auth      include     password-auth
...
```

- Notice the first PAM module is `pam_sepermit.so`.

- c. View the man pages for the `pam_sepermit` authentication module.

```
[host01]# man pam_sepermit
...
pam_sepermit - PAM module to allow/deny login depending
On SELinux enforcement state
...
```

- Notice that when SELinux is in enforcing mode, the `pam_sepermit` authentication module allows or denies login.

- d. SELinux stands for “Security-Enhanced Linux” and is covered in a later course but for the purposes of this practice, use the `sestatus` command to display information about SELinux.

```
[host01]# sestatus
SELinux status:      enabled
...
Current mode:        enforcing
...
```

- Notice that SELinux is enabled and is in enforcing mode.

- e. Use the `setenforce 0` command to temporarily change SELinux to “permissive” mode.

- Temporarily setting SELinux to “permissive” mode and re-testing is a good practice to perform to ensure SELinux is not the cause of a problem with a service.
- Use the `sestatus` command to display the change in the SELinux mode.

```
[host01]# setenforce 0
[host01]# sestatus
SELinux status:      enabled
...
Current mode:        permissive
...
```

- Notice the “Current mode” is now set to “permissive.”

- f. Test if setting SELinux to permissive mode allows the ssh connection from **host03** to **host01**.

```
[host03]# ssh host01
ssh: connect to host host01 port 22: No route to host
```

- Notice the ssh command still fails. Therefore SELinux is not the cause of the problem.

- g. Use the **setenforce 1** command to change SELinux back to “enforcing” mode.

- Use the **sestatus** command to display the change in the SELinux mode.

```
[host01]# setenforce 1
[host01]# sestatus
SELinux status:      enabled
...
Current mode:        enforcing
...
```

- Notice the “Current mode” is now set to “enforcing.”

You could continue viewing man pages for the other PAM modules listed in the /etc/pam.d/sshd and /etc/pam.d/ssh-keycat files to determine if any of these are the cause of the problem. However, as was previously stated, entries are usually written to the /var/log/secure log file when PAM is denying access.

4. Attempt to diagnose the problem by checking firewall (iptables) configuration.

- a. Use the **service** command to view the status of the **iptables** service.

```
[host01]# service iptables status
...
```

- Notice that the **iptables** service is running.

- b. The **ssh** service uses port 22, which can be confirmed by using the **grep ssh /etc/services** command.

```
[host01]# grep ssh /etc/services
ssh          22/tcp      # The Secure Shell (SSH) Protocol
ssh          22/udp      # The Secure Shell (SSH) Protocol
...
```

- c. View the status of the **iptables** service but pipe the output to **grep** and search for 22.

```
[host01]# services iptables status | grep 22
```

- No output indicates there is no **iptables** rule for this port number.

5. Stop the **iptables** service on **host01** and re-test remote connectivity.

- a. Use the **service** command to stop the **iptables** service.

```
[host01]# service iptables stop
...
```

- b. From **host03**, use the **ssh** command to attempt to connect to **host01**.

- Password is oracle.
- Use the hostname command to confirm that you did connect.

```
[host03]# ssh host01
root@host01's password: oracle
Last login: ...
[host01]# hostname
host01.example.com
```

- Connection is successful with the iptables service stopped on host01.

- Use the logout command to log off.

```
[host01]# logout
Connection to host01 closed.
```

- Fix the remote connectivity problem by creating an iptables rule on host01 to trust ssh.

- Stopping iptables on host01 allows the ssh connection.
- However, to keep this service running you need to create a rule that trusts the ssh service on port # 22.

 - Use the service command to start the iptables service.

```
[host01]# service iptables start
iptables: Applying firewall rules
```

- Use the iptables command to list rules in the INPUT chain and include line numbers.

```
[host01]# iptables -L INPUT --line-numbers
Chain INPUT (policy ACCEPT)
num  target     prot opt    source        destination
 1    ACCEPT     all  --    anywhere     anywhere    state ...
 2    ACCEPT     icmp --    anywhere     anywhere
 3    ACCEPT     all  --    anywhere     anywhere
 4    REJECT     all  --    anywhere     anywhere    reject-with...
```

- You want to insert the new rule before the REJECT entry on line 4.
- Use the iptables command to insert the “ssh” rule before line 4 with the following characteristics from the command line:
 - Chain = INPUT
 - Protocol = tcp
 - State = NEW
 - Destination port = ssh
 - Target = ACCEPT

```
[host01]# iptables -I INPUT 4 -p tcp -m state --state NEW --
dport ssh -j ACCEPT
```

- This rule accepts incoming tcp traffic for ssh.

- d. Use the `iptables` command to list rules in the INPUT chain and include line numbers.

```
[host01]# iptables -L INPUT --line-numbers
Chain INPUT (policy ACCEPT)
num  target     prot opt source      destination
1    ACCEPT     all  --  anywhere   anywhere   state ...
2    ACCEPT     icmp --  anywhere   anywhere
3    ACCEPT     all  --  anywhere   anywhere
4    ACCEPT     tcp  --  anywhere   anywhere   state NEW tcp
dpt:ssh
5    REJECT     all  --  anywhere   anywhere   reject-with...
```

- Notice the new line 4 rule that accepts `tcp` traffic when the destination port is `ssh`.

- e. Use the `service` command to save the `iptables` rules.

- After making any rule changes, always save the `iptables` rules.
- Saving `iptables` rules updates the `/etc/sysconfig/iptables` file.

```
[host01]# service iptables save
iptables: Saving firewall rules to /etc/sysconfig/iptables...
```

- f. Use the `cat` command to view the `/etc/sysconfig/iptables` file.

```
[host01]# cat /etc/sysconfig/iptables
...
-A INPUT -p tcp -m state --state NEW -m tcp --dport 22 -j ACCEPT
...
```

- Notice that in the configuration file for `iptables`, the new rule has the actual port number, 22, for `ssh`.

- g. Use the `service` command to restart the `iptables` service.

- After making any rule changes, restart the `iptables` service.

```
[host03]# service iptables restart
iptables: Setting chains to policy ACCEPT: filter      [  OK  ]
iptables: Flushing firewall rules:                      [  OK  ]
iptables: Unloading modules:                            [  OK  ]
iptables: Applying firewall rules:                     [  OK  ]
```

7. Ensure remote connectivity is working.

- a. From `host03`, confirm that you can use the `ssh` command to connect to `host01`.

- Password is `oracle`.
- Use the `hostname` command to confirm that you did connect.

```
[host03]# ssh host01
root@host01's password: oracle
Last login: ...
[host01]# hostname
host01.example.com
```

- b. Use the `logout` command to log off.

```
[host01]# logout  
Connection to host01 closed.
```

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Practice 21-11: Log File Is Not Getting Updated

Overview

In this practice, you diagnose and fix a problem with a log file not getting updated.

Assumptions

- You are the `root` user on the **host01** VM.
- You are the `root` user on the **host03** VM.

Tasks

1. From **host01**, set the `oracle` user's password to `oracle`.

Use the `passwd oracle` command to change the password.

- Ignore the "BAD PASSWORD" messages.

```
[host01]# passwd oracle
Changing password for user oracle.
New password: oracle
BAD PASSWORD: it is based on a dictionary word
BAD PASSWORD: is too simple
Retype new password: oracle
passwd: all authentication tokens updated successfully.
```

2. From **host03**, confirm you can use the `ssh` command to log in to **host01** as the `oracle` user.

- a. Use the `exit` command to log out as the `root` user.

- Use the `whoami` command to confirm you are the `oracle` user.

```
[host03]# exit
logout
[host03]$ whoami
oracle
```

- b. Use the `ssh` command to log in to **host01**.

- Password is `oracle`.
- Use the `hostname` command to verify you were able to log in.

```
[host03]$ ssh host01
oracle@host01's password: oracle
Last login: ...
[host01]$ hostname
host01.example.com
```

- c. Use the `exit` command to log off **host01**.

```
[host01]$ exit  
logout  
Connection to host01 closed.
```

3. From **host01**, save a copy of the `/var/log/secure` file.

Use the `cp` command to copy `/var/log/secure` to `~/secure_before`.

```
[host01]# cp /var/log/secure ~/secure_before
```

- You review the differences in these copies later in this practice.

4. From **host01**, configure the `pam_nologin` authentication module to prevent non-root login.

- This exercise was performed in the PAM lesson, practice 16-2.

Use the `vi` editor and create the `/etc/nologin` file with the following contents:

```
[host01]# vi /etc/nologin  
No logins allowed at this time.
```

5. As the `oracle` user on **host03**, attempt to log in to **host01**.

Use the `ssh` command to connect to **host01**.

- Password is `oracle`.

```
[host03]$ ssh host01  
oracle@host01's password: oracle  
No logins allowed at this time.  
Connection closed by 192.0.2.101
```

- Notice that the connection is denied by the PAM authentication module.

6. From **host01**, view the new entries in the `/var/log/secure` log file.

Use the `diff` command to view the differences in the `/var/log/secure` log file and the copy of the log file you made in Task 3.

```
[host01]# diff /var/log/secure ~/secure_before  
< <date_time> host01 [sshd[...]: Failed password for oracle from  
192.0.2.103 port ... ssh2  
< <date_time> host01 [sshd[...]: fatal: Access denied for user  
oracle by PAM account configuration
```

- Notice that two entries were written to this file when the `oracle` user attempted to ssh from **host03** to **host01**.
- Also notice that one entry specifically references PAM account configuration.

7. From **host01**, execute the `eleven.x` program from the `root` user's home directory.

```
[host01]# cd  
[host01]# pwd  
/root  
[host01]# ./eleven.x
```

8. From **host01**, save a copy of the `/var/log/secure` file.

Use the `cp` command to copy `/var/log/secure` to `~/secure_before`.

- Answer `y` to overwrite the file.

```
[host01]# cp /var/log/secure ~/secure_before  
cp: overwrite '/root/secure_before'? y
```

- You review the differences in these copies later in this practice.

9. As the **oracle** user on **host03**, attempt to log in to **host01**.

Use the `ssh` command to connect to **host01**. Password is **oracle**.

```
[host03]$ ssh host01  
oracle@host01's password: oracle  
No logins allowed at this time.  
Connection closed by 192.0.2.101
```

- Notice that the connection is denied by the PAM authentication module.

10. From **host01**, view the new entries in the `/var/log/secure` log file.

Use the `diff` command to view the differences in the `/var/log/secure` log file and the copy of the log file you made in Task 8.

```
[host01]# diff /var/log/secure ~/secure_before
```

- Notice that there are no differences in the files.
- The `/var/log/secure` log file is not getting updated as expected.

11. Diagnose and fix the problem of the log file not getting updated.

- Review Lesson 20, “System Logging”.
- Is the logging daemon running?
- Is logging configured for the `/var/log/secure` log file?

12. After fixing the problem, ensure the `/var/log/secure` log file is getting updated when you attempt to `ssh` from **host03** to **host01** as the **oracle** user.

13. Return **host01** back to the original state.

Use the `rm` command to remove the `/etc/nologin` file.

```
# rm /etc/nologin  
rm: remove regular file '/etc/nologin'? y
```

Solution 21-11: Log File Is Not Getting Updated

Steps

1. Diagnose the cause of the logging problem.

- As the root user on **host01**, view the status of the `rsyslog` service.

```
[host01]# service rsyslog status  
rsyslogd (pid ...) is running...
```

- The `rsyslogd` service is running.

- Use the `grep` command to search for `/var/log/secure` in the `rsyslog` configuration file.

```
[host01]# grep /var/log/secure /etc/rsyslog.conf  
#authpriv.*          /var/log/secure
```

- Notice the configuration file contains an entry for `/var/log/secure`; however, the entry is commented out (preceded by a `#` sign).

2. Fix the problem.

- Use the `vi` editor to remove the `#` from the beginning of the line containing `/var/log/secure`.

```
[host01]# vi /etc/rsyslog.conf  
...  
#authpriv.*          /var/log/secure          (old entry)  
authpriv.*           /var/log/secure          (new entry)  
...
```

- Use the `service` command to restart the `rsyslog` service.

```
[host01]# service rsyslog restart  
Shutting down system logger: [ OK ]  
Starting system logger: [ OK ]
```

3. Verify the `/var/log/secure` log file is getting updated.

- From **host01**, use the `cp` command to copy `/var/log/secure` to `~/secure_before`.

- Answer `y` to overwrite the file.

```
[host01]# cp /var/log/secure ~/secure_before  
cp: overwrite '/root/secure_before'? y
```

- As the `oracle` user on **host03**, use the `ssh` command to connect to **host01**.

- Password is `oracle`.

```
[host03]$ ssh host01  
oracle@host01's password: oracle  
No logins allowed at this time.  
Connection closed by 192.0.2.101
```

- Notice that the connection is denied by the PAM authentication module.
- c. From **host01**, use the `diff` command to view the differences in the `/var/log/secure` log file and the copy, `~/secure_before`.

```
[host01]# diff /var/log/secure ~/secure_before
< <date_time> host01 [sshd[...]: Failed password for oracle from
192.0.2.103 port ... ssh2
< <date_time> host01 [sshd[...]: fatal: Access denied for user
oracle by PAM account configuration
```

- Notice that entries are now being written to the `/var/log/secure` file as expected.
- 4. Return **host01** back to the original state.

Use the `rm` command to remove the `/etc/nologin` file.

```
# rm /etc/nologin
rm: remove regular file '/etc/nologin'? y
```

Appendix A: Source Code for Problem-Causing Executables

Overview

In this appendix, the source code for the executables that cause the problems are given.

Practice 21-2: System Boots into Single-User Mode

The “two.x” script appends the word `single` to the end of the kernel line for the default kernel using the following commands:

```
#!/bin/bash
cp /boot/grub/grub.conf /boot/grub/grub.conf.save
sed -i -e '/vmlinuz-3.8.13/s/$/ single/' /boot/grub/grub.conf
```

This result of this program follows:

```
# cat /boot/grub/grub.conf
...
default=0
...
title Oracle Linux Server Unbreakable Kernel (3.8.13-16...)
    root (hd0,0)
    kernel /vmlinuz-3.8.13-16... single
...
```

Practice 21-3: Status Commands Fail

The “three.x” script runs the following `umount` command to unmount the `proc` file system.

```
#!/bin/bash
umount -l /proc
```

Practice 21-4: cron Job Fails to Run

The “four.x” script stops the `crond` service by running the following command:

```
#!/bin/bash
service crond stop > /dev/null
```

Practice 21-5: User Cannot Log In

The “five.x” script inserts the “#” sign at the beginning of the `john` line in the `/etc/shadow` file by running the following commands:

```
#!/bin/bash
cp /etc/shadow /etc/shadow.save
sed -i -e '/john/s/^/#/' /etc/shadow
```

The result of this program follows:

```
# cat /etc/shadow
...
#john:$6$...:0:99999:7:::
```

Practice 21-6: File System Troubleshooting

The “six.x” script runs the following dd command to corrupt the file system superblock.

```
#!/bin/bash
dd if=/dev/zero of=/dev/xvdb bs=1024 skip=1000 count=300 2>
/dev/null
```

Practice 21-8: Network Connectivity Problem

The “eight.x” script changes the IP address for **host01** from 192.0.2.101 to 192.0.3.101 by running the following commands:

```
#!/bin/bash
cd /etc/sysconfig/network-scripts
/bin/cp ifcfg-eth0 ~/ifcfg-eth0.save
sed -e 's/2.101/3.101/' ifcfg-eth0 > ifcfg-eth0.new
/bin/mv ifcfg-eth0.new ifcfg-eth0
service network restart > /dev/null
cd
```

The result of this program follows:

```
# cat /etc/sysconfig/network-scripts/ifcfg-eth0
...
IPADDR=192.0.3.101
...
# ifconfig eth0
eth0      Link encap: Ethernet  HWaddr: 00:16:3E:00:01:01
          inet addr: 192.0.3.101 ...
...
```

Practice 21-9: NFS Permission Problem

The “nine.x” script deletes the oracle user, re-adds the oracle user with UID=555, and recreates the /home/oracle/remote_home directory by running the following commands:

```
#!/bin/bash
userdel -r oracle
useradd -u 555 oracle
mkdir /home/oracle/remote_home
chown -R oracle.oracle /home/oracle
```

The result of this program follows:

```
# grep oracle /etc/passwd
```

```
oracle:x:555:555::/home/oracle:/bin/bash
# ls -l /home/oracle
drwxr-xr-x. 2 oracle oracle 4096 <date_time> remote_home
```

Practice 21-10: Remote Access Problem

The “ten.x” script removes the iptables rule on **host01**, which trusts the sshd service and restarts the iptables service by running the following commands:

```
#!/bin/bash
cp /etc/sysconfig/iptables /etc/sysconfig/iptables.save
sed -i '/22/d' /etc/sysconfig/iptables
service iptables restart > /dev/null
```

Practice 21-11: Log File Is Not Getting Updated

The “eleven.x” script inserts the “#” sign at the beginning of the secure line in the /etc/rsyslog.conf file and restarts the rsyslog service by running the following commands:

```
#!/bin/bash
cp /etc/rsyslog.conf /etc/rsyslog.conf.save
sed -i -e '/secure/s/^/#/' /etc/rsyslog.conf
service rsyslog restart > /dev/null
```

The result of this program follows:

```
# cat /etc/rsyslog.conf
...
#authpriv.*                                     /var/log/secure
```

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Appendices: Remote Access Options

Chapter 22

Unauthorized reproduction or distribution prohibited. Copyright© 2017, Oracle and/or its affiliates.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.

Appendix

Appendices Overview

The three appendixes show various options for accessing your student PC remotely:

- Appendix A: Using an NX Client to Connect to **dom0**
- Appendix B: Using an NX Player to Connect to **dom0**
- Appendix C: Using VNC (TightVNC) to Connect Directly to VM Guests

Appendix A: Using an NX Client to Connect to dom0

Overview

This appendix discusses accessing your student PC (**dom0**) remotely by using NX Client. The NX Client in this appendix is NX Client for Windows, Version 3.5.0-9.


Steps

1. Install **NX Client** (if necessary) from <http://www.nomachine.com/download.php>.
2. Run **NX Client** (for example, select **NX Client for Windows** from the Windows Start menu).
 - An NX Connection Wizard steps you through creating the initial session.
 - The following Welcome window appears.




- a. Click **Next**.

- The following Session window appears.




- Enter anything you like for **Session** (for example, OracleLinux).
- Enter the IP address (provided by your instructor) for **Host**.
- Accept the remaining defaults and click **Next**.
- The Desktop window appears. Change KDE to **GNOME** by selecting from the drop-down list.

- Your window should look like the following:




- f. Accept all other defaults and click **Next**.
- The following Configuration completed window appears.



- g. Click **Finish**.

- The NX Login window appears.



- For **Login**, enter vncuser.
- For **Password**, enter vnctech.
- Your **Session** defaults to the session that you just created. In this example, the **Session** is OracleLinux. Your session name may be different.
- Click **Login**.
 - The **dom0** GNOME virtual desktop window appears.
 - Future connections will bypass the configuration wizard and only bring up the NX Login window.

Mohammed Ameen (mohammedameen@oracle.com) has a
non-transferable license to use this Student Guide.


Appendix B: Using an NX Player to Connect to dom0

Overview


This appendix discusses accessing your student PC (**dom0**) remotely using NX Player. The NX Player in this appendix is NX Player for Windows, Preview 5, version 4.0.132.

Steps

1. Install **NX Player** (if necessary) from <http://www.nomachine.com/download.php>.
2. Run **NX Player** (for example, select **NX Player for Windows** from the Windows Start menu).
 - a. Ensure that the **Favorites** tab is selected.
 - The Connect window appears.




- b. Click **New connection** to display the following window.



- c. Enter the IP address (provided by your instructor) for **Host**.
- d. Accept the defaults:
- 1) Port 22
 - 2) Use the SSH service
 - 3) Use the NX login
- e. Note the connection name. In this example it is **Unnamed connection 2**. Yours is most likely **Unnamed connect 1**.

f. Press **Enter**. The following window appears.




g. Click the connection that you just created (**Unnamed connection 1**, for example).

- The Login window appears.

3. Log in.


- The window shown in the following screenshot appears, prompting for login authentication.



- Ensure that **Login as a named user** is selected.
- For **Username**, enter vncuser.
- For **Password**, enter vnctech.
- Click **Ok**.

4. Create a new session.


- The window shown in the following screenshot appears.




- Click **Create a new session**.

5. Create a new GNOME virtual desktop.

- The window shown in the following screenshot appears.



- Click **Create a new GNOME virtual desktop**.
- Click the **X** in the **Information** message box to close the box.



- The **dom0** GNOME virtual desktop window appears.


Appendix C: Using VNC (TightVNC) to Connect Directly to VM Guests

Overview

This appendix discusses accessing the VM guest systems that directly uses VNC (TightVNC). It is not recommended to connect to **dom0** or to the **host03** VM by using VNC. Both **dom0** and **host03** have the GNOME user interface, which causes various problems when connecting using VNC.

Steps


1. Install **tightvnc** (if necessary) from <http://www.tightvnc.com/>.
2. Run **TightVNC Viewer** (for example, select **TightVNC Viewer** from the Windows Start menu).
 - The following New TightVNC Connection window appears.




3. Connect directly to your **host01** virtual machine.
 - The following assumptions are made:
 - The **host01** VM was created first (has a port number of 5900).
 - The **host02** VM was created next (has a port number of 5902).
 - The **host03** VM was created last (has a port number of 5903).
 - The output of the following commands (from **dom0** as **root**) indicates that this assumption is true.

```
# xm list -l host01 | grep location
      (location 0.0.0.0:5900)
# xm list -l host02 | grep location
      (location 0.0.0.0:5902)
# xm list -l host03 | grep location
      (location 0.0.0.0:5903)
```

- Enter the IP address (provided by your instructor), followed by the port number to connect directly a VM guest.
 - a. To connect directly to the **host01** VM, enter the following.
 - In this example, the IP address of your student PC is 10.150.30.68. Your IP address is different.



- 1) Click **Connect**.
 - A terminal window appears.



- 2) Log in as root password with 0racle (leading zero, not letter O).
- 3) Enter the hostname command to confirm that you are logged in to **host01**.

```
# hostname  
host01.example.com
```

- 4) Log off by entering the exit command.
- 5) Close the VNC window by clicking the X in the top-right corner of the window.

Mohammed Ameen (mohammed.a.ameen@oracle.com) has a
non-transferable license to use this Student Guide.